

Kubernetes Architecture



Nigel Poulton

@nigelpoulton www.nigelpoulton.com



Overview



Big Picture View

Masters

Nodes

Pods

Services

Deployments

The API and API Server

Recap

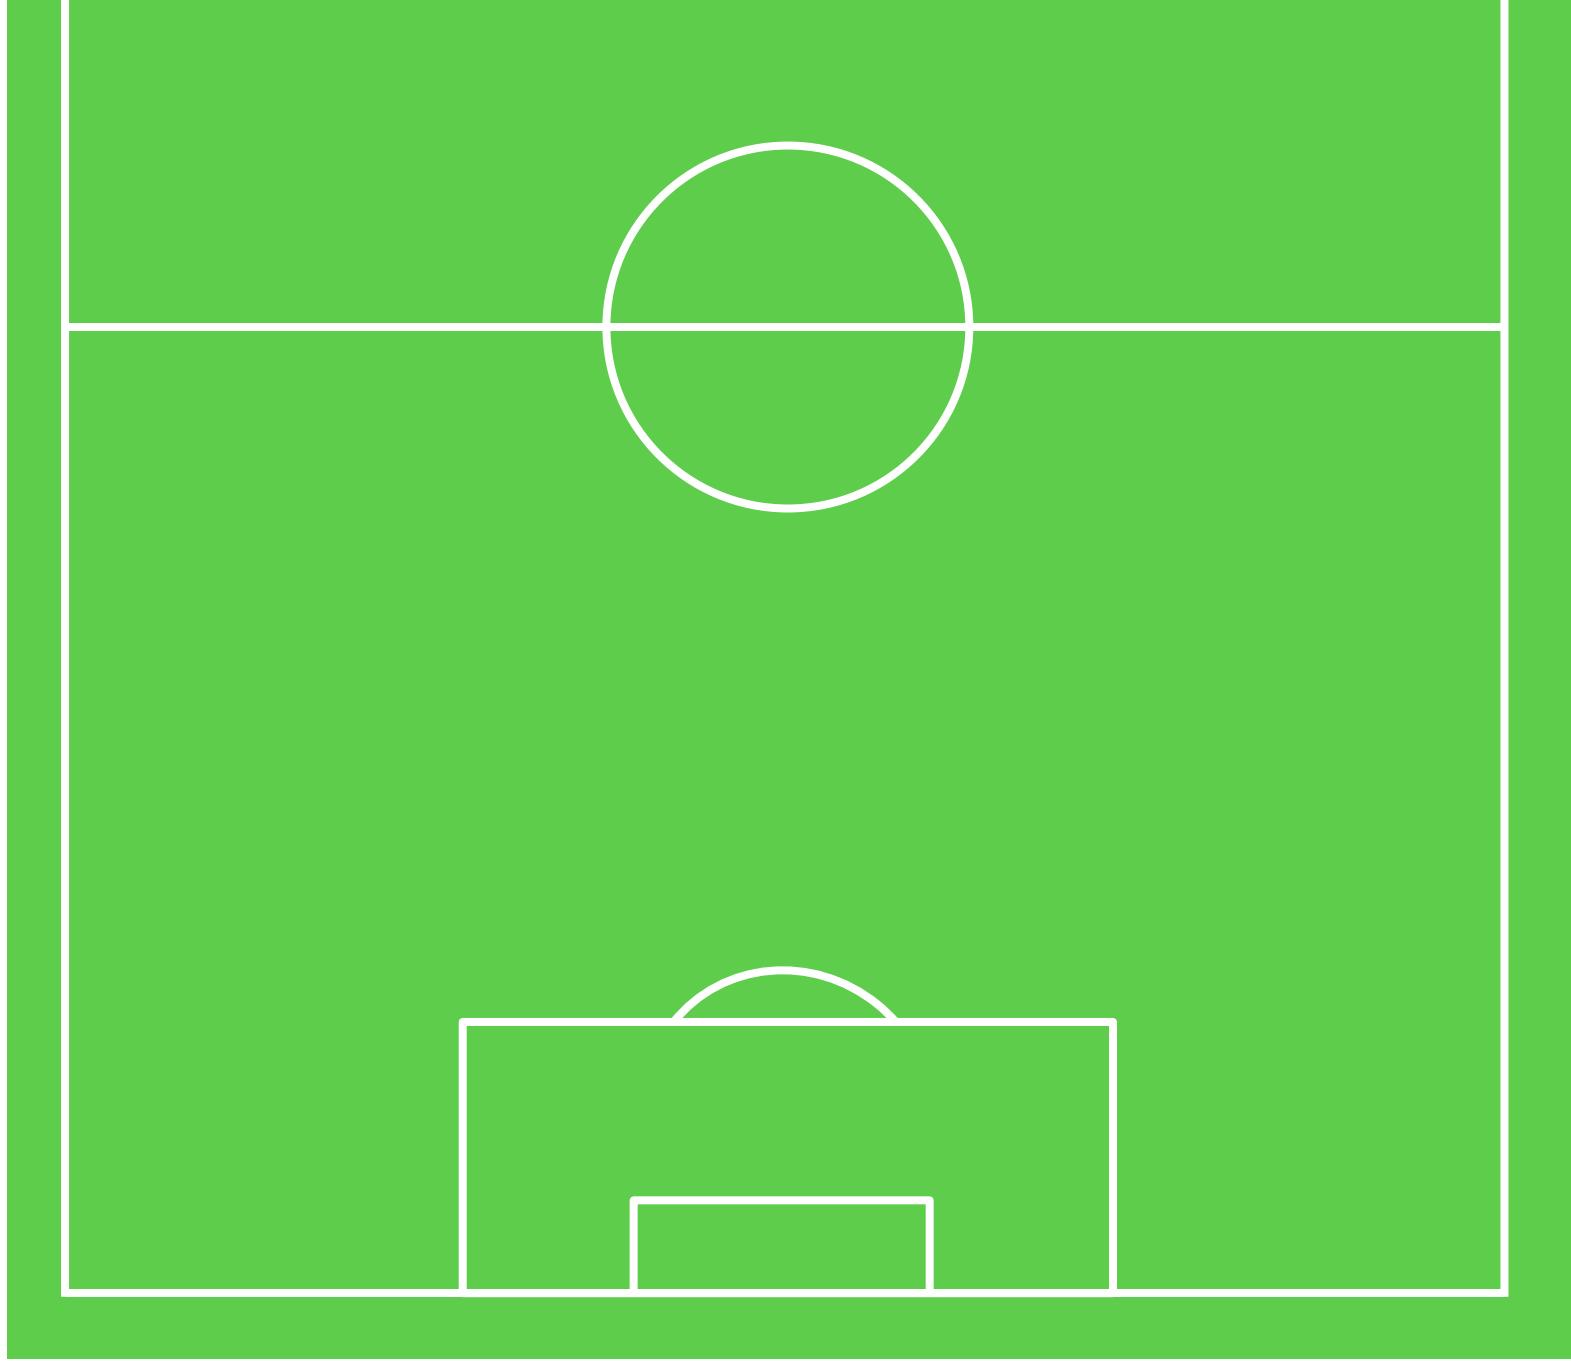


Up Next:
Kubernetes: Big Picture Overview

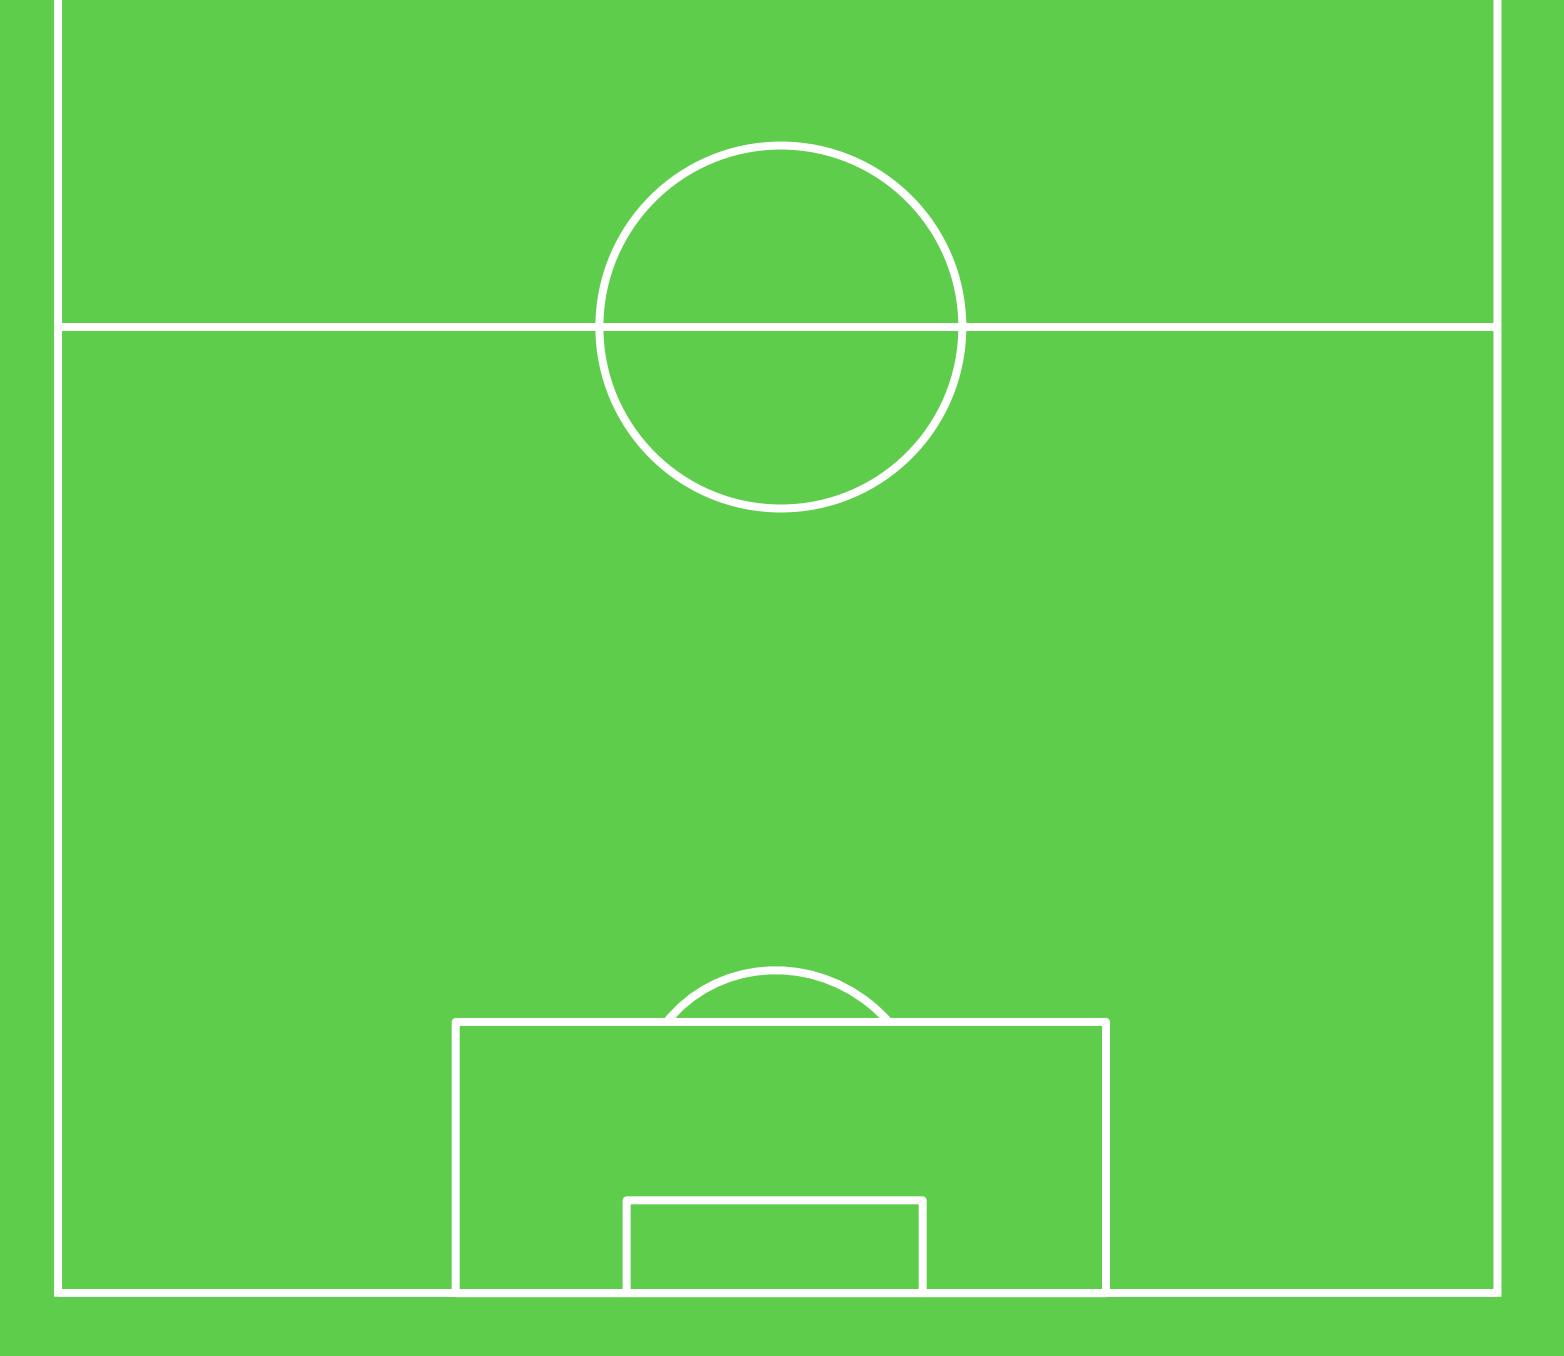
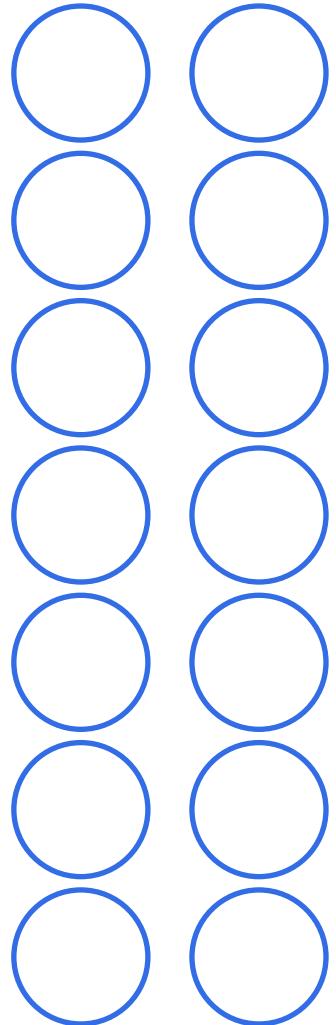


Kubernetes: Big Picture Overview

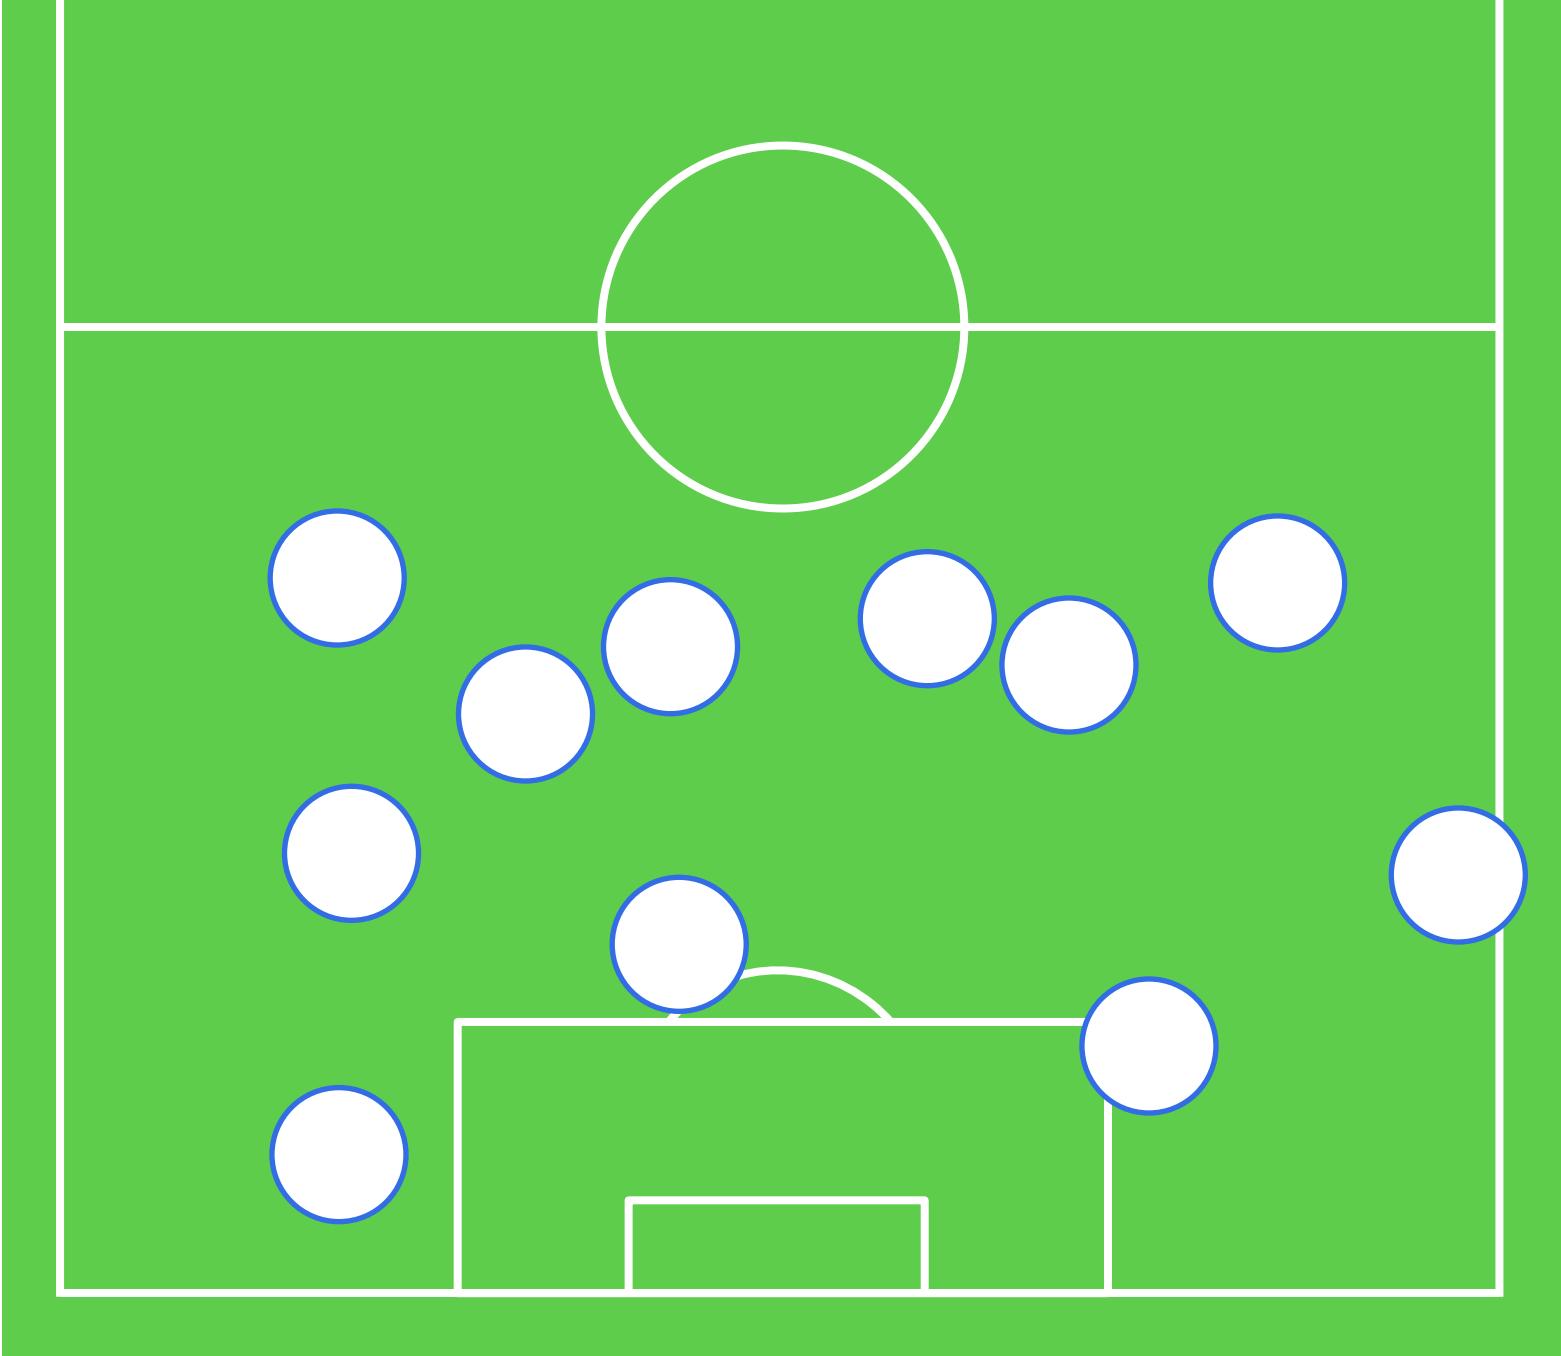
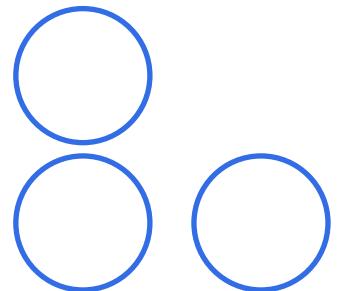




Team



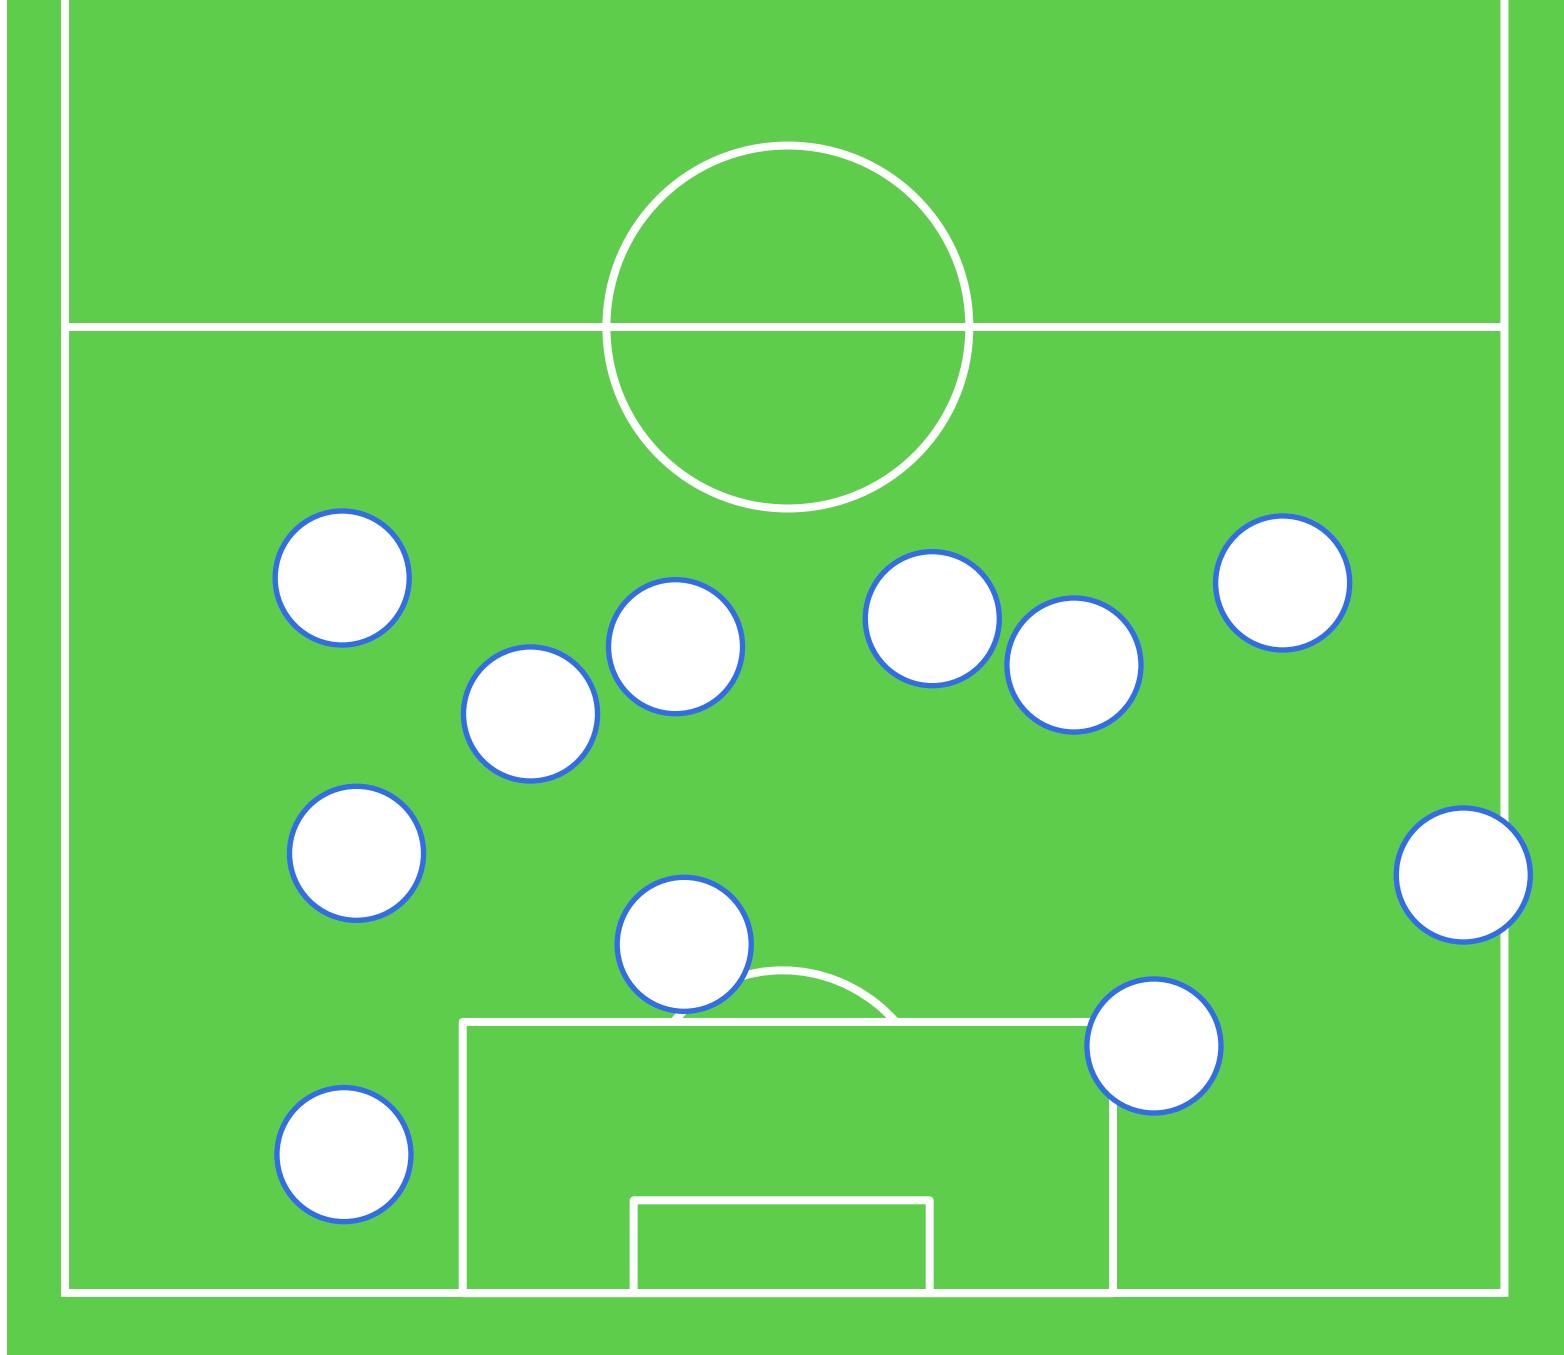
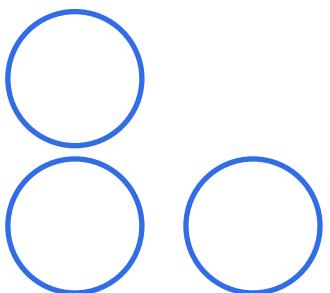
Team



Team



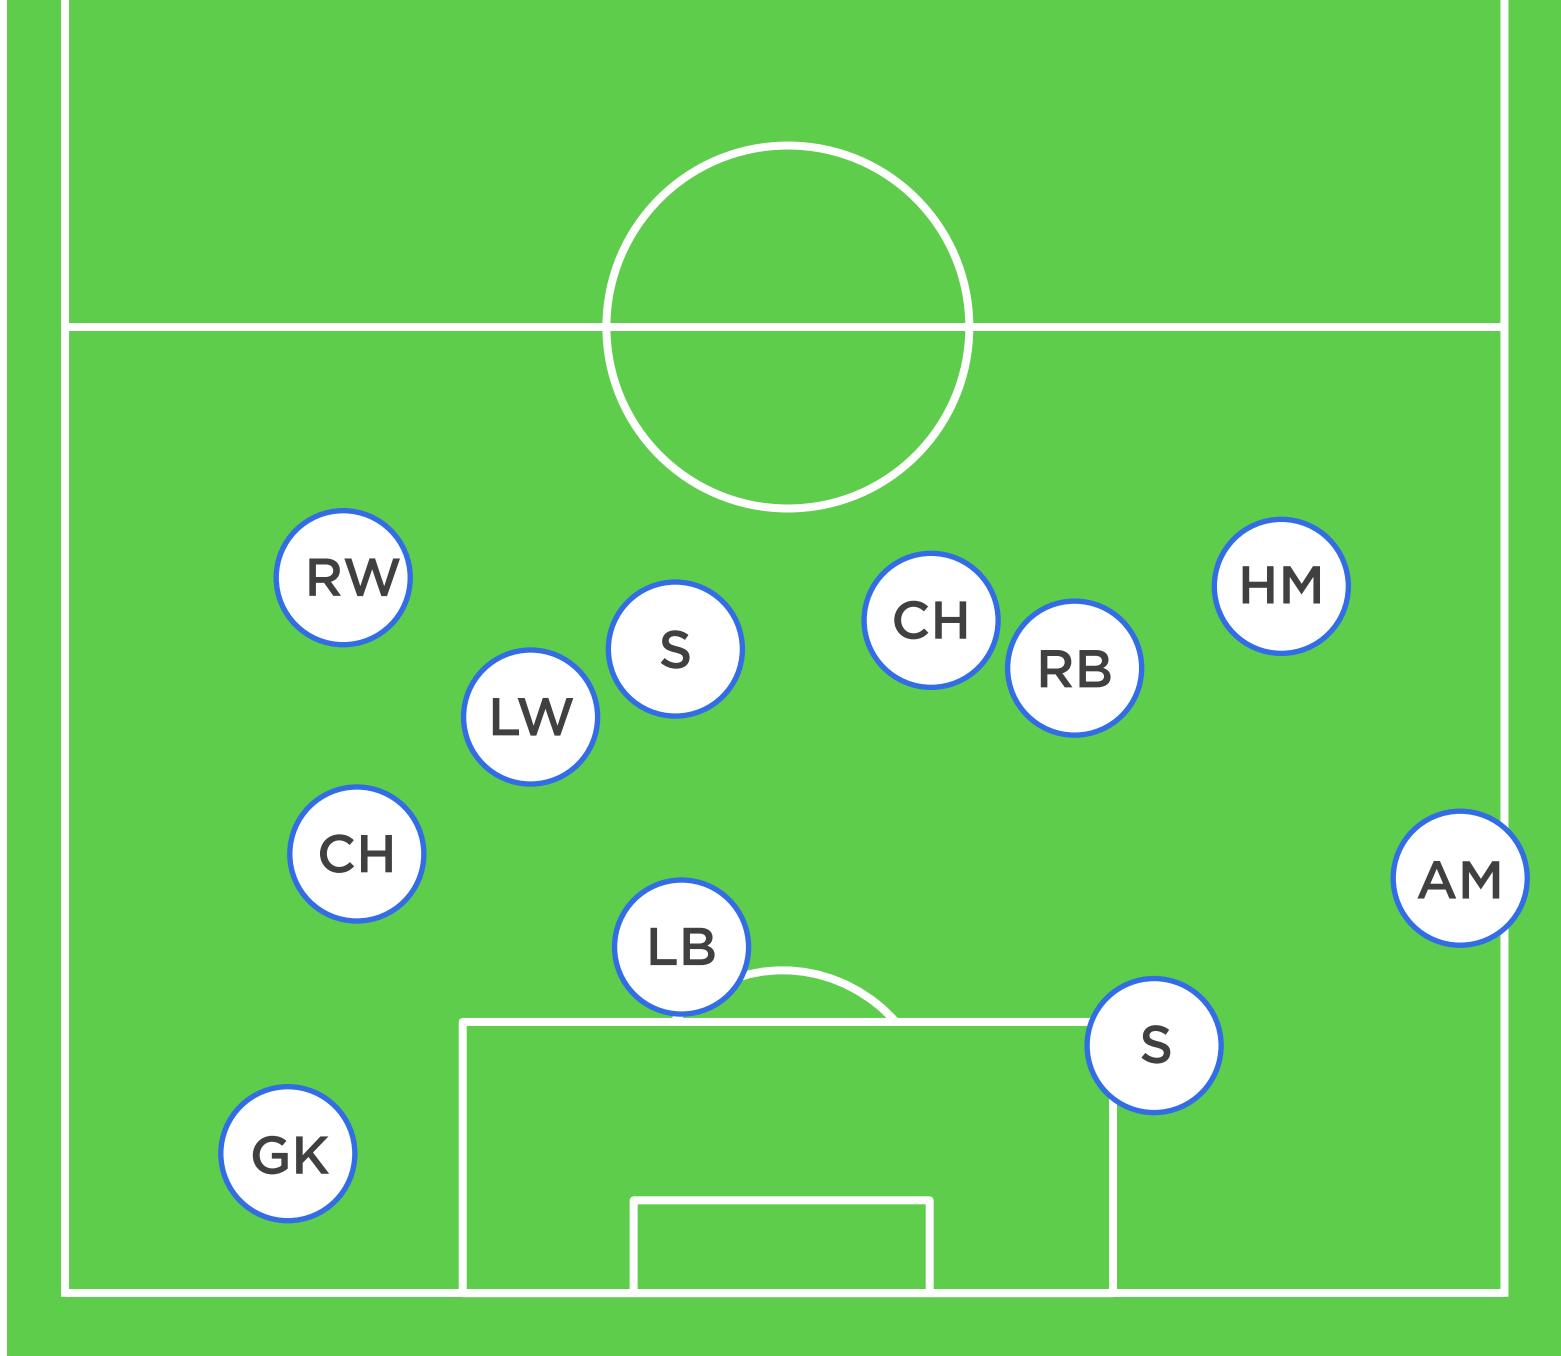
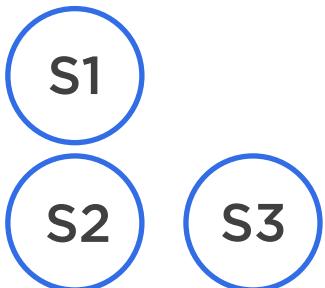
Manager
(coach)

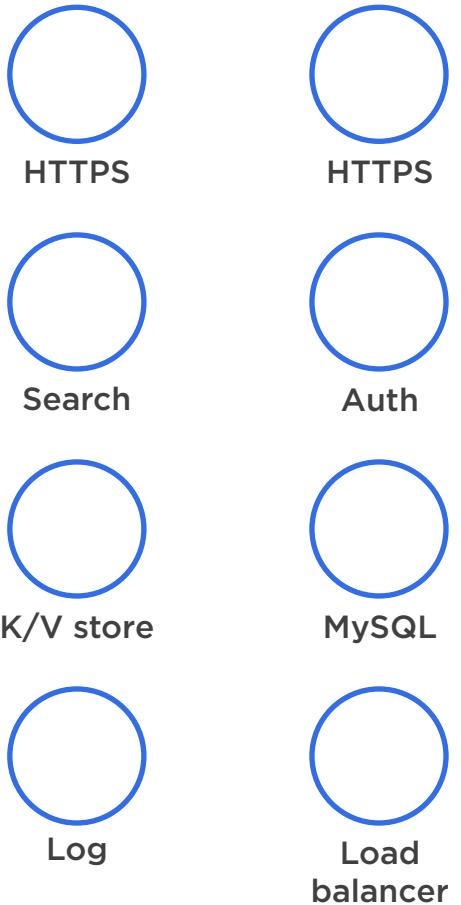


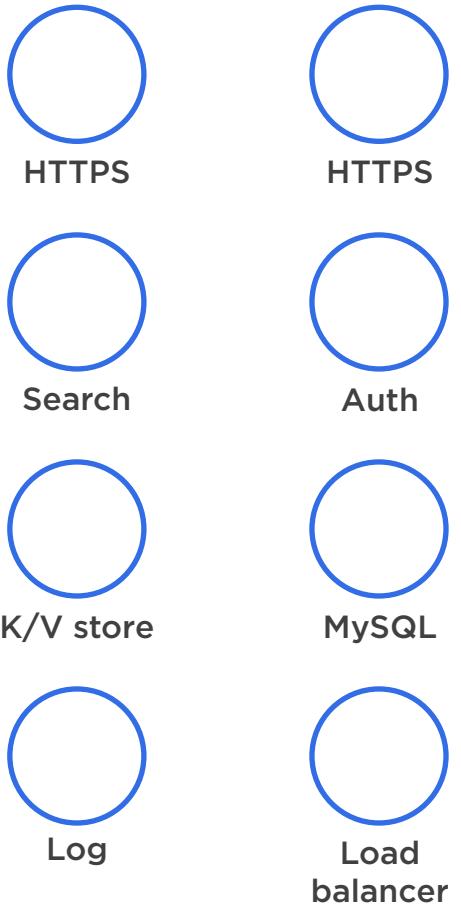
Team

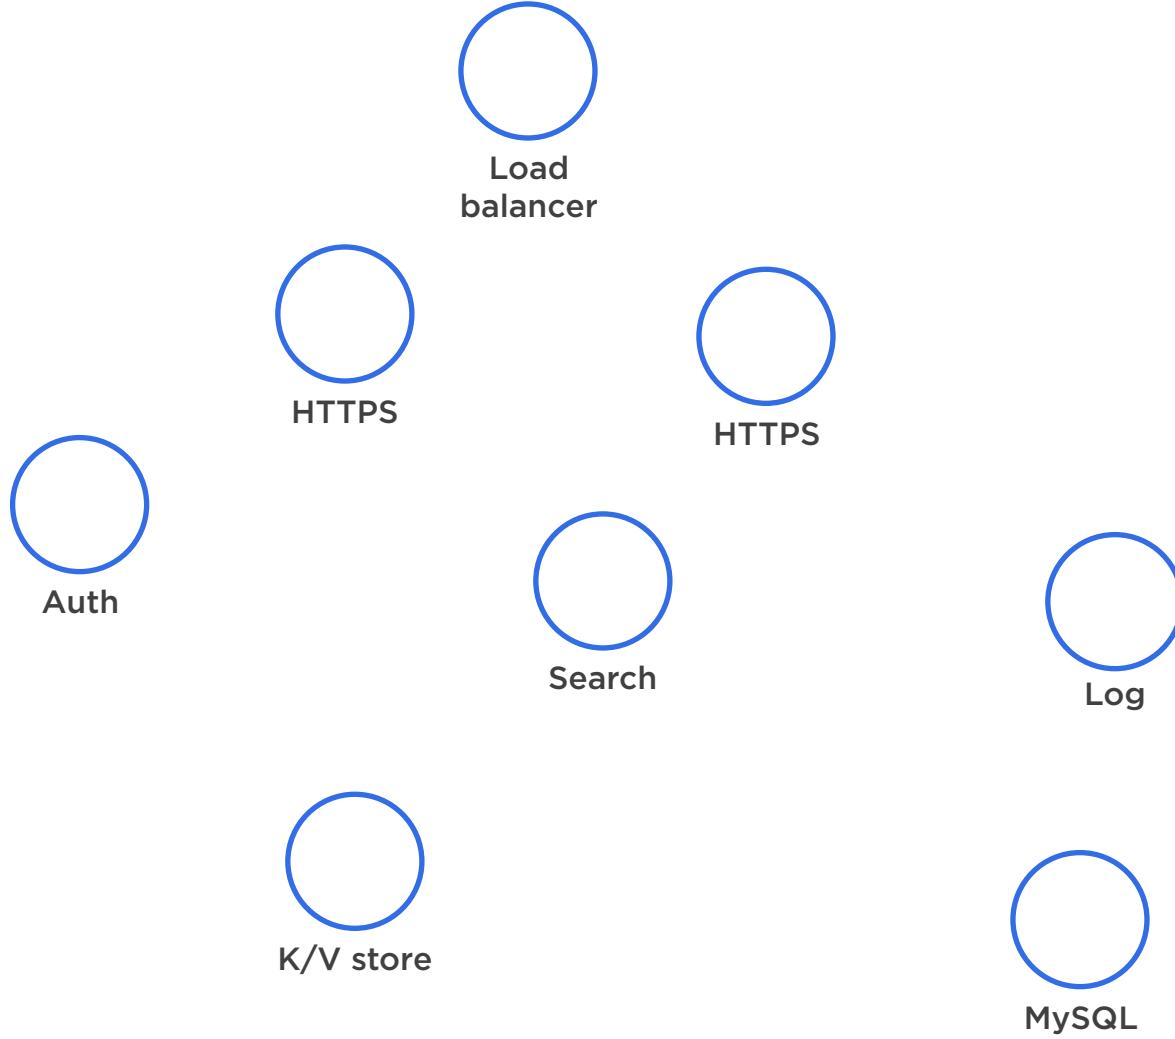


Manager
(coach)



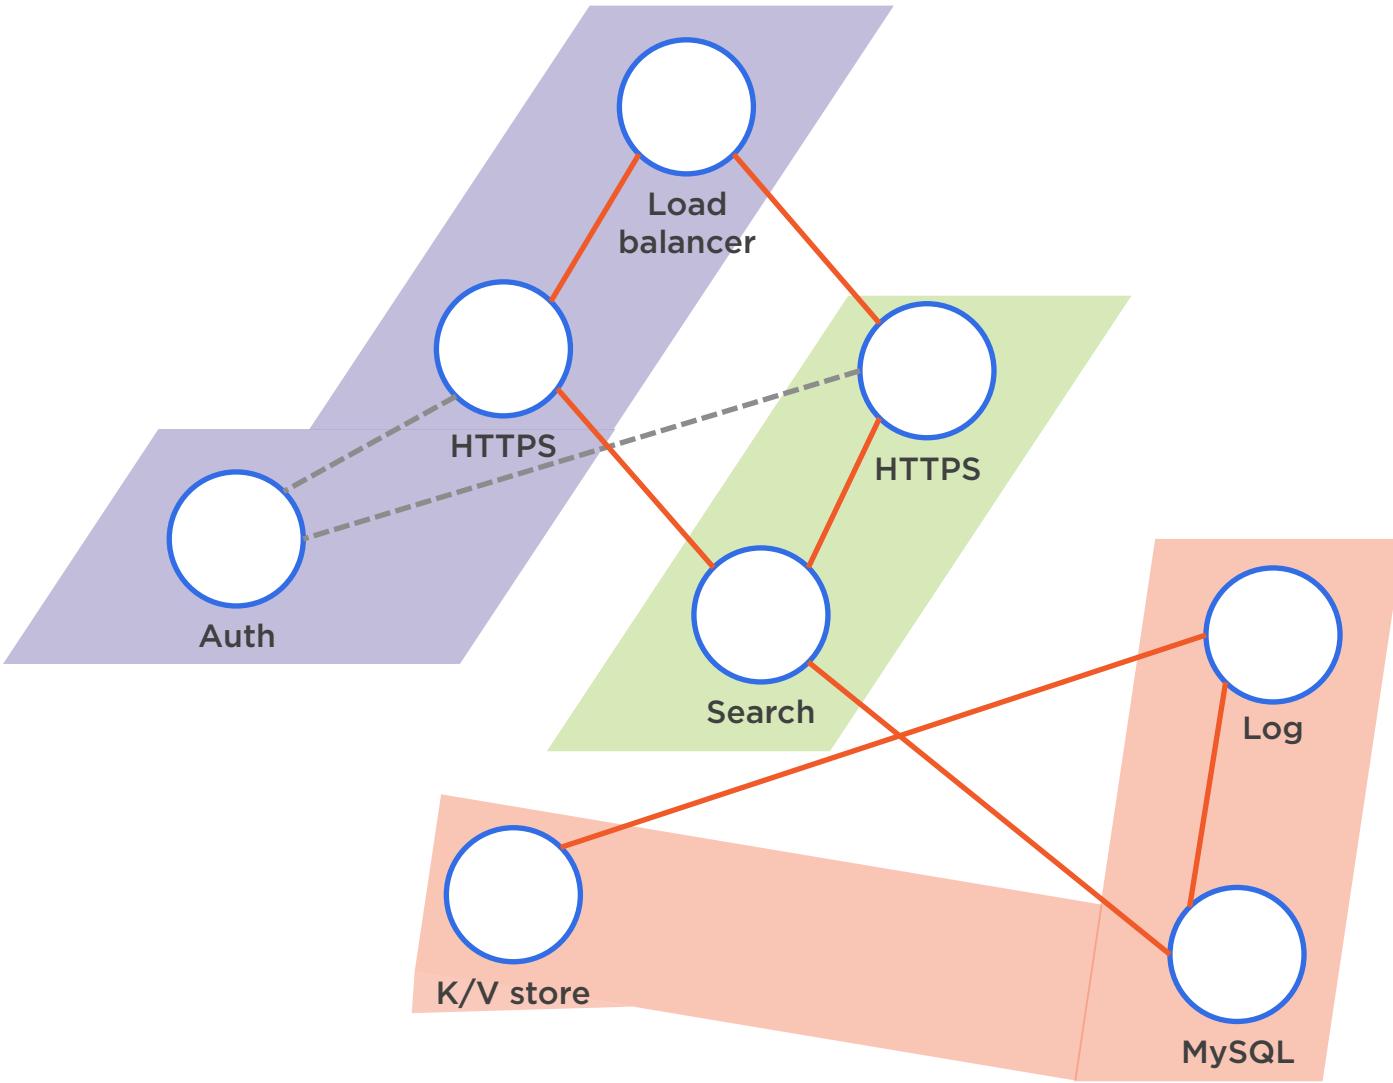


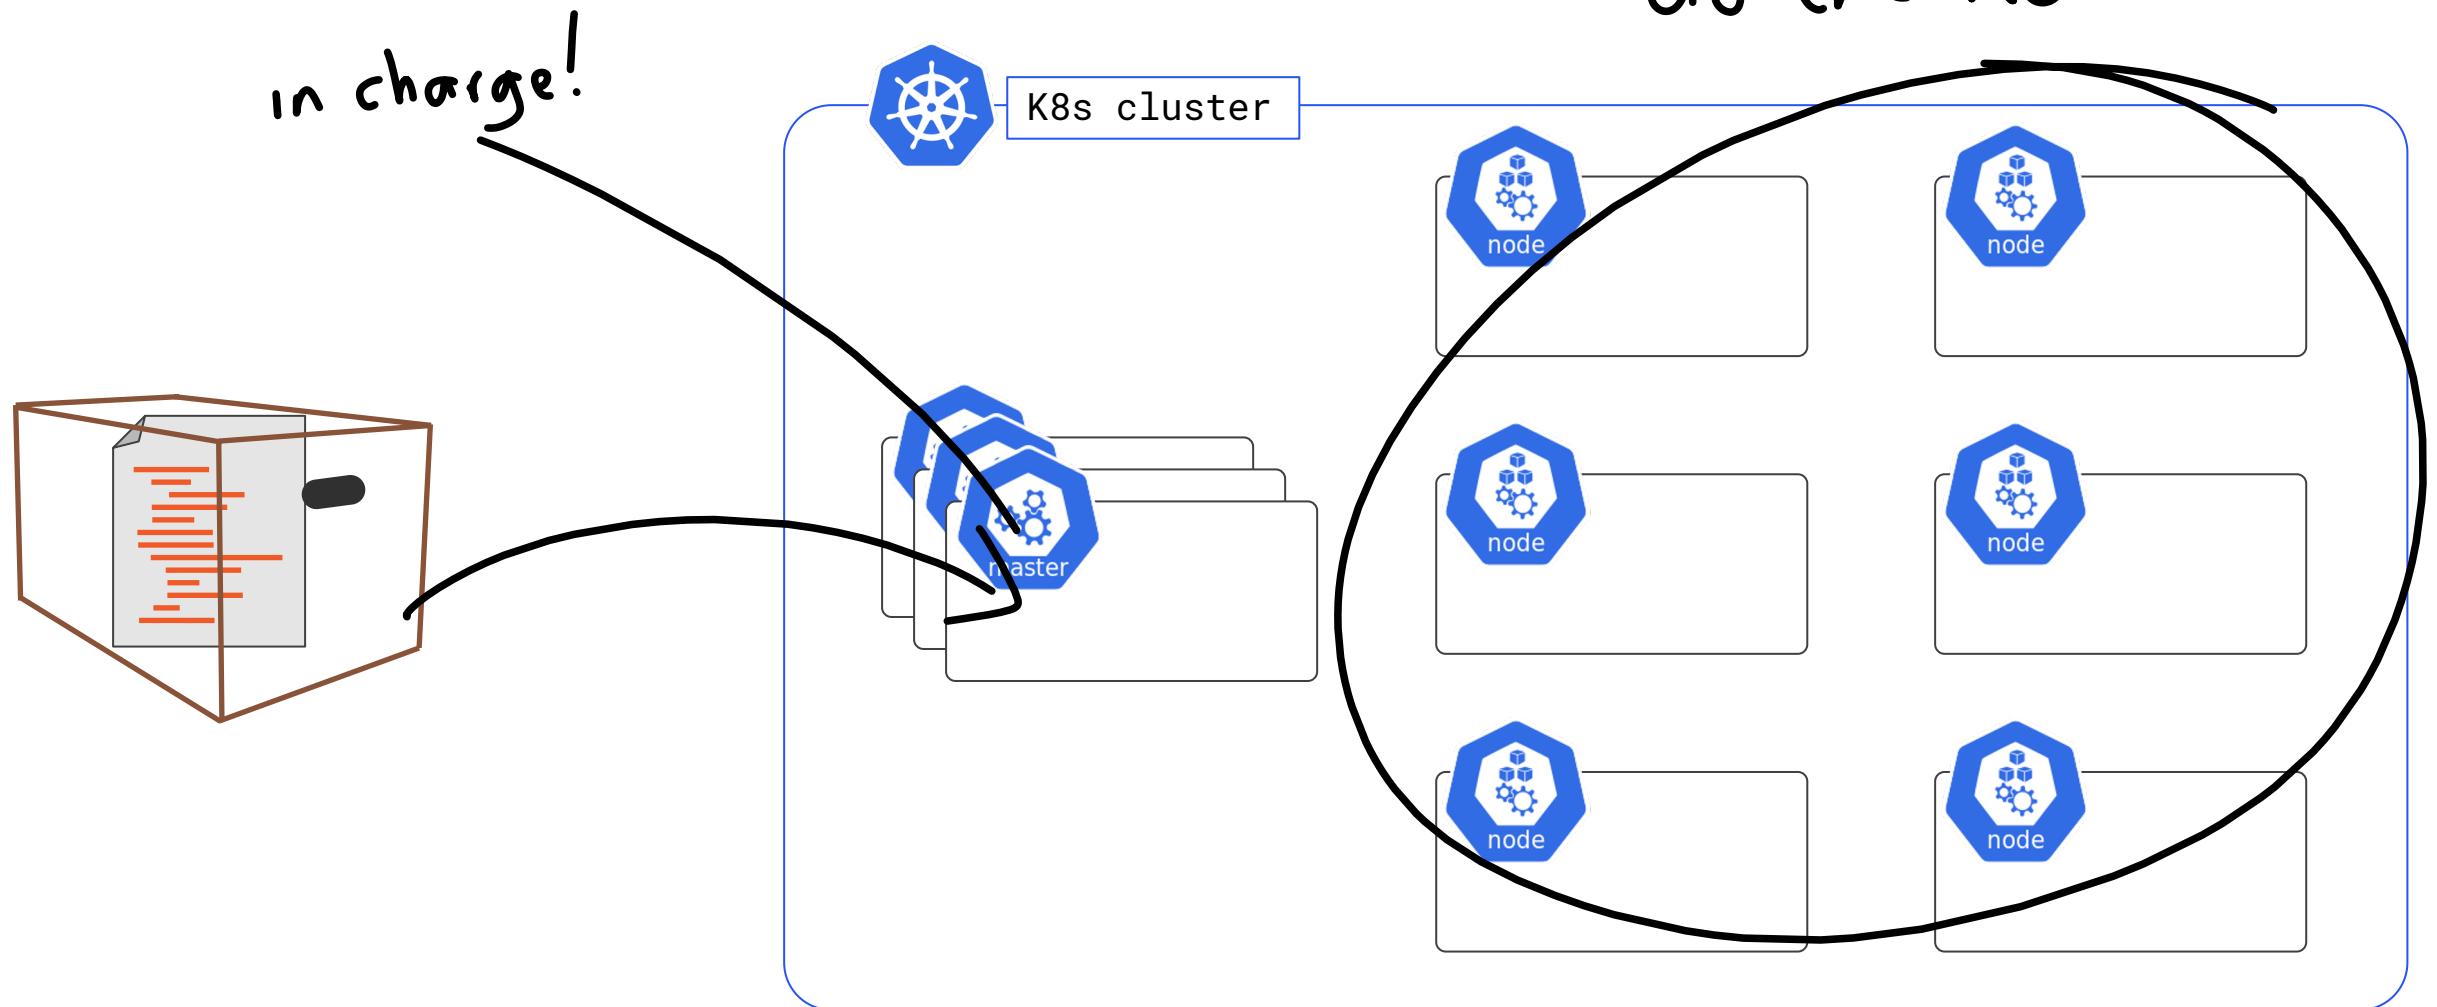


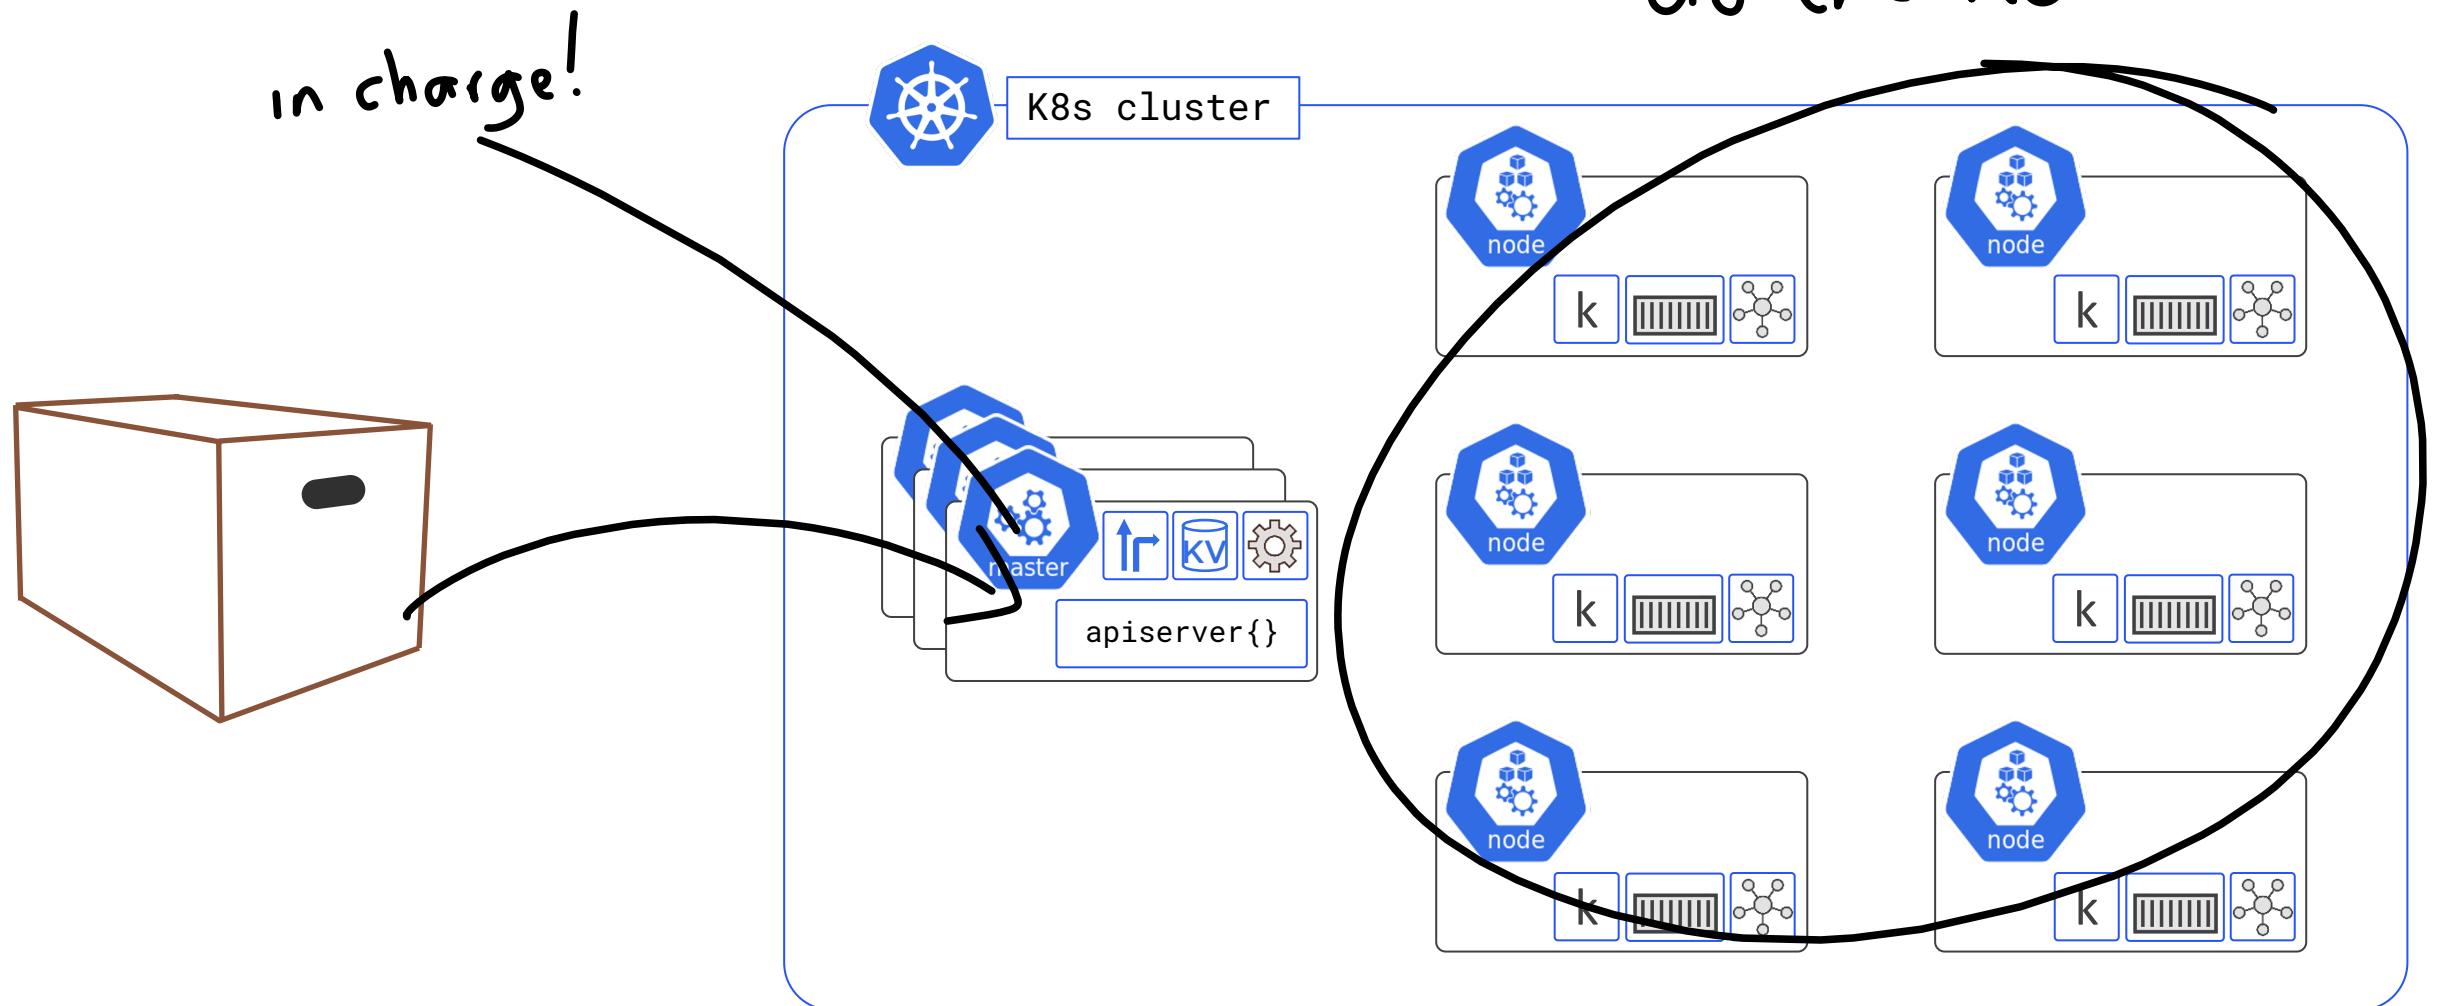


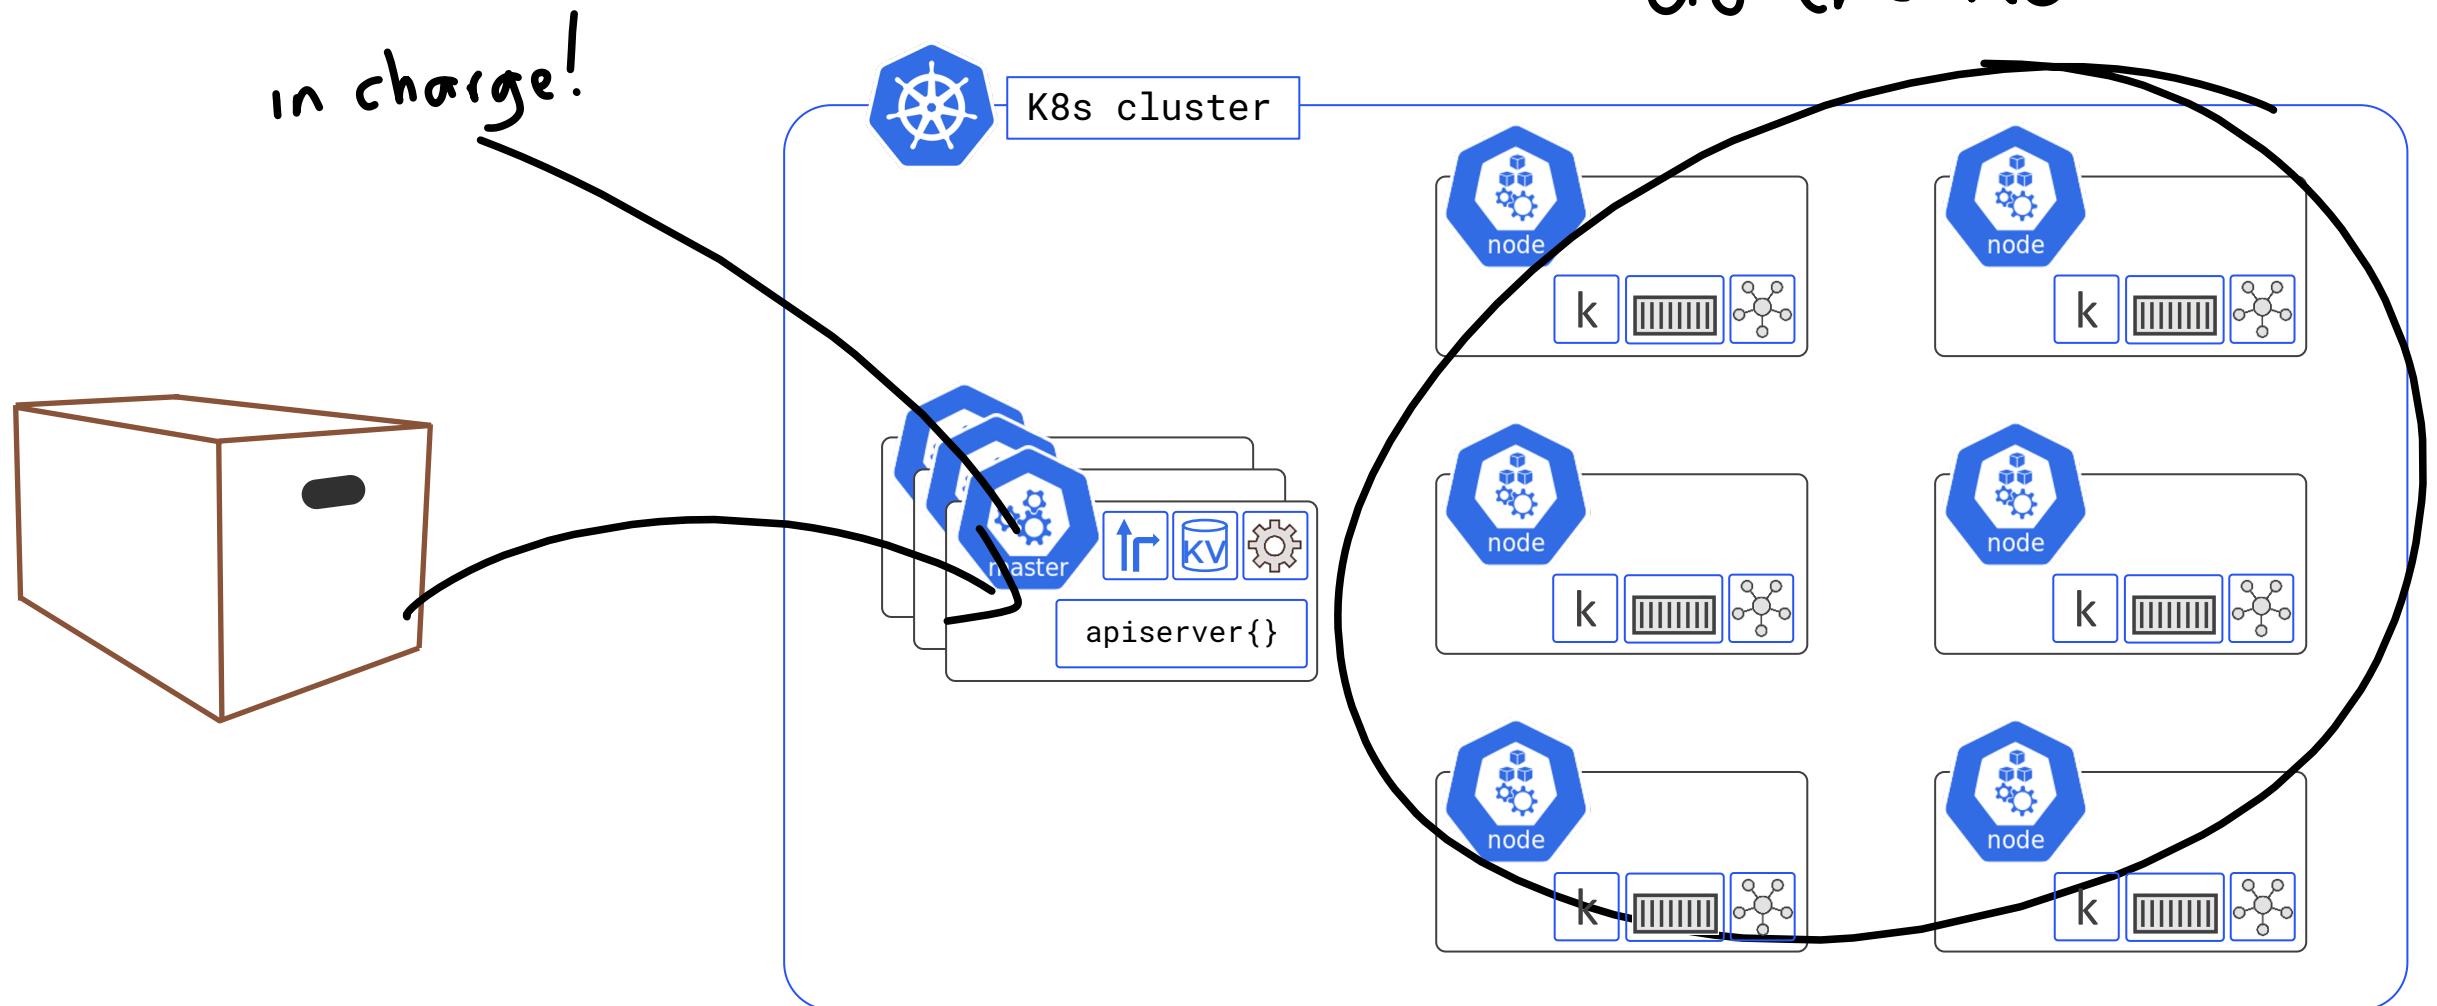


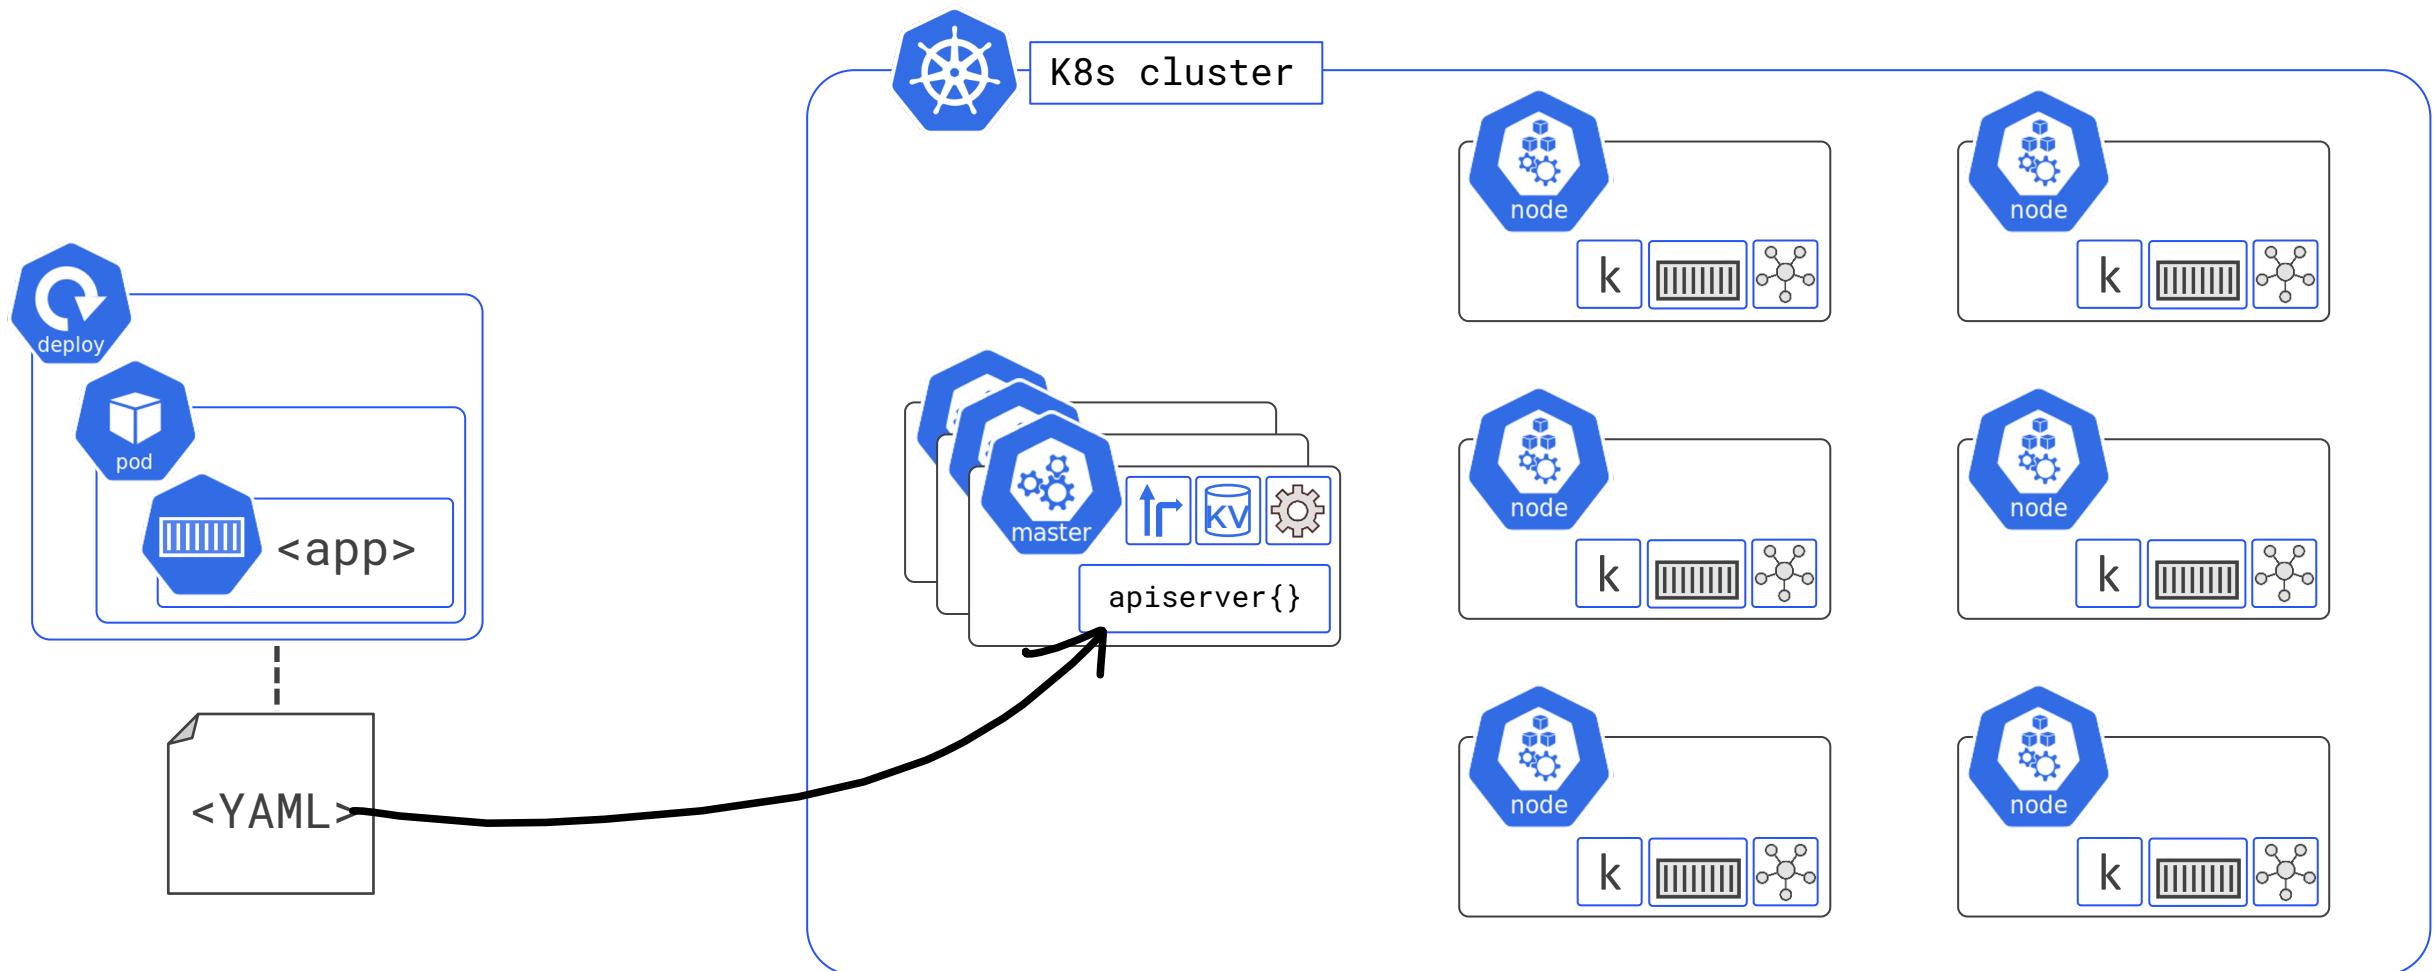
- Node 1
- Node 2
- Node 3











Up Next:
Masters



Masters

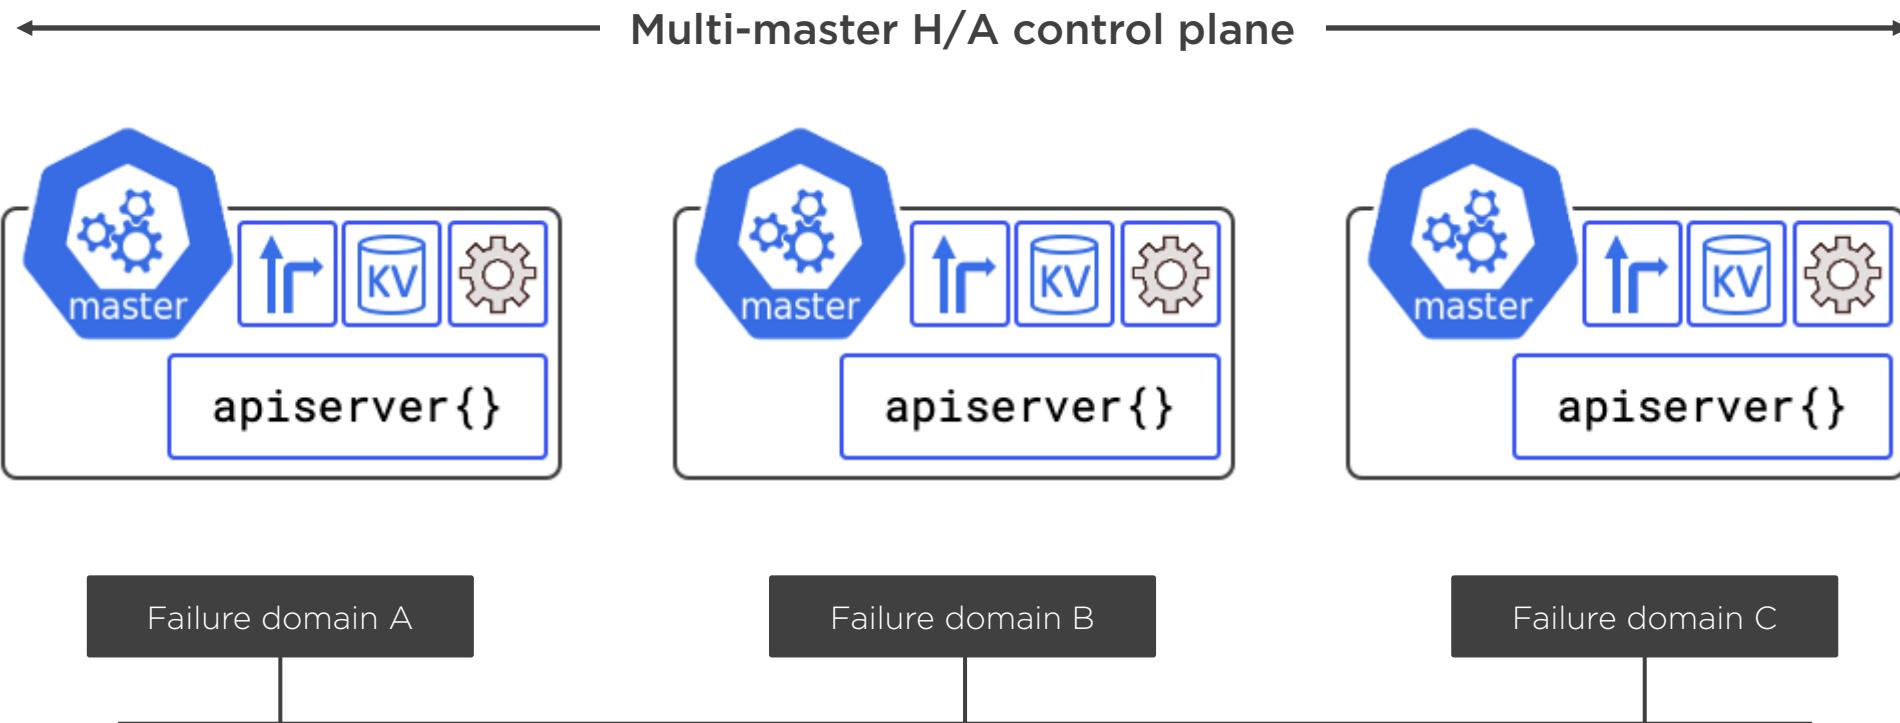


Masters

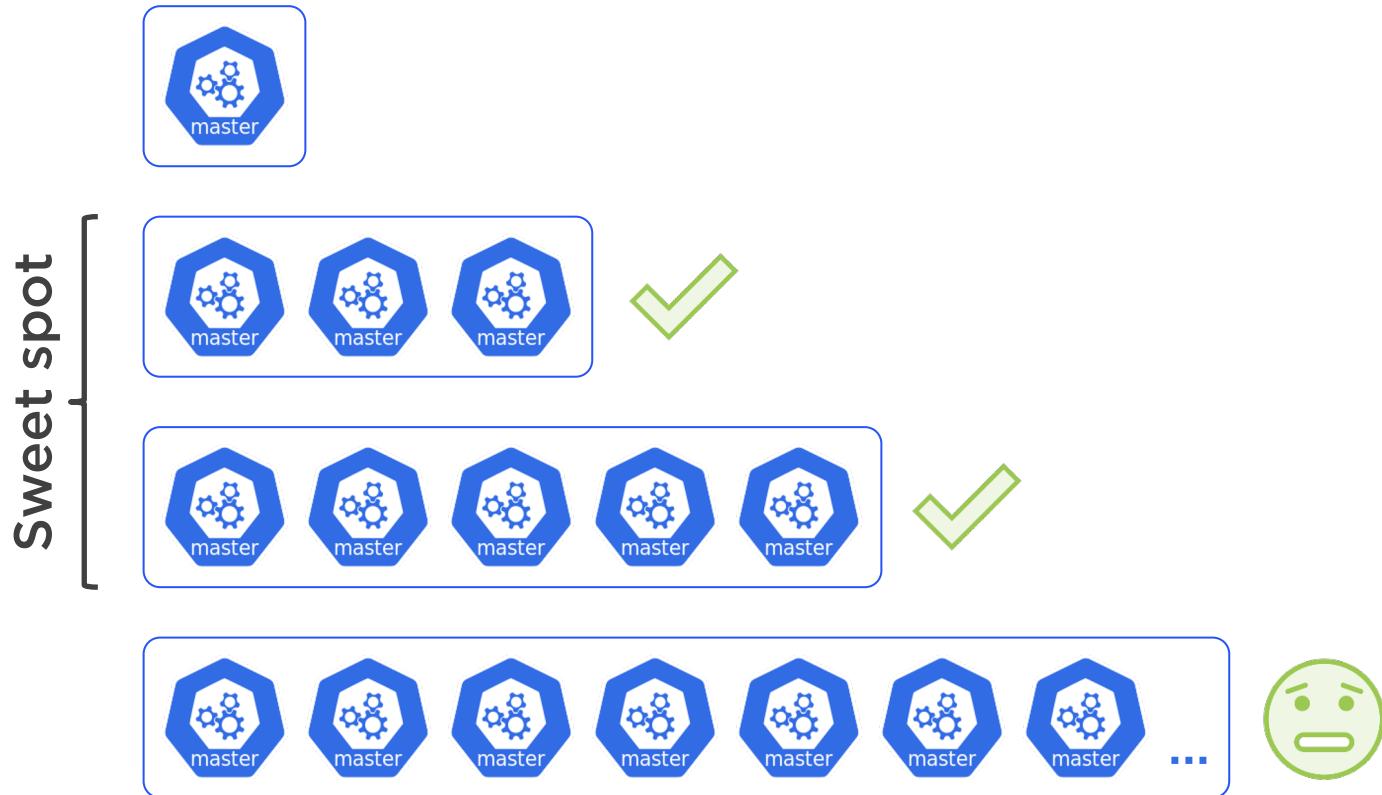
Also known as *head nodes* or the *control plane*







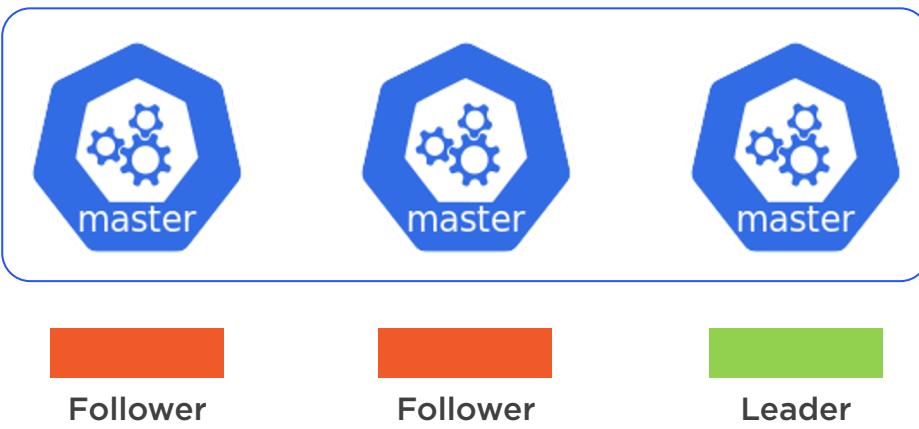
H/A Design

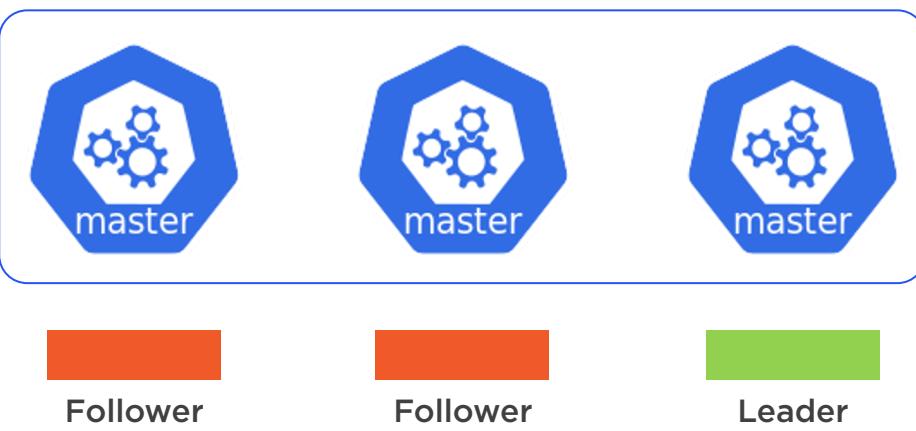




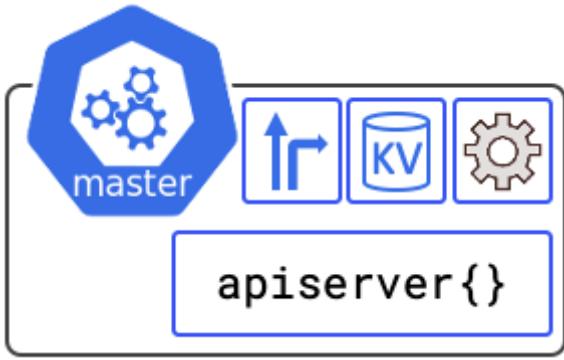
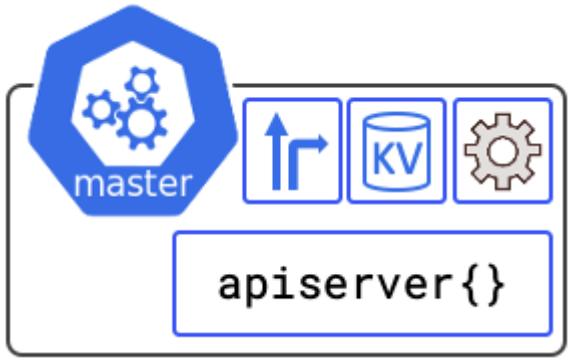


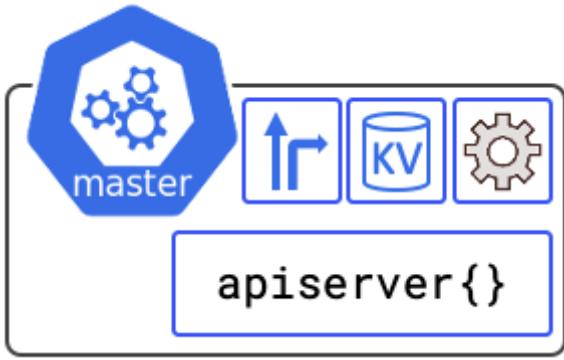
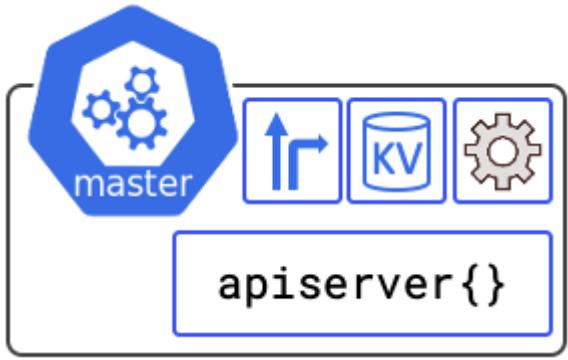


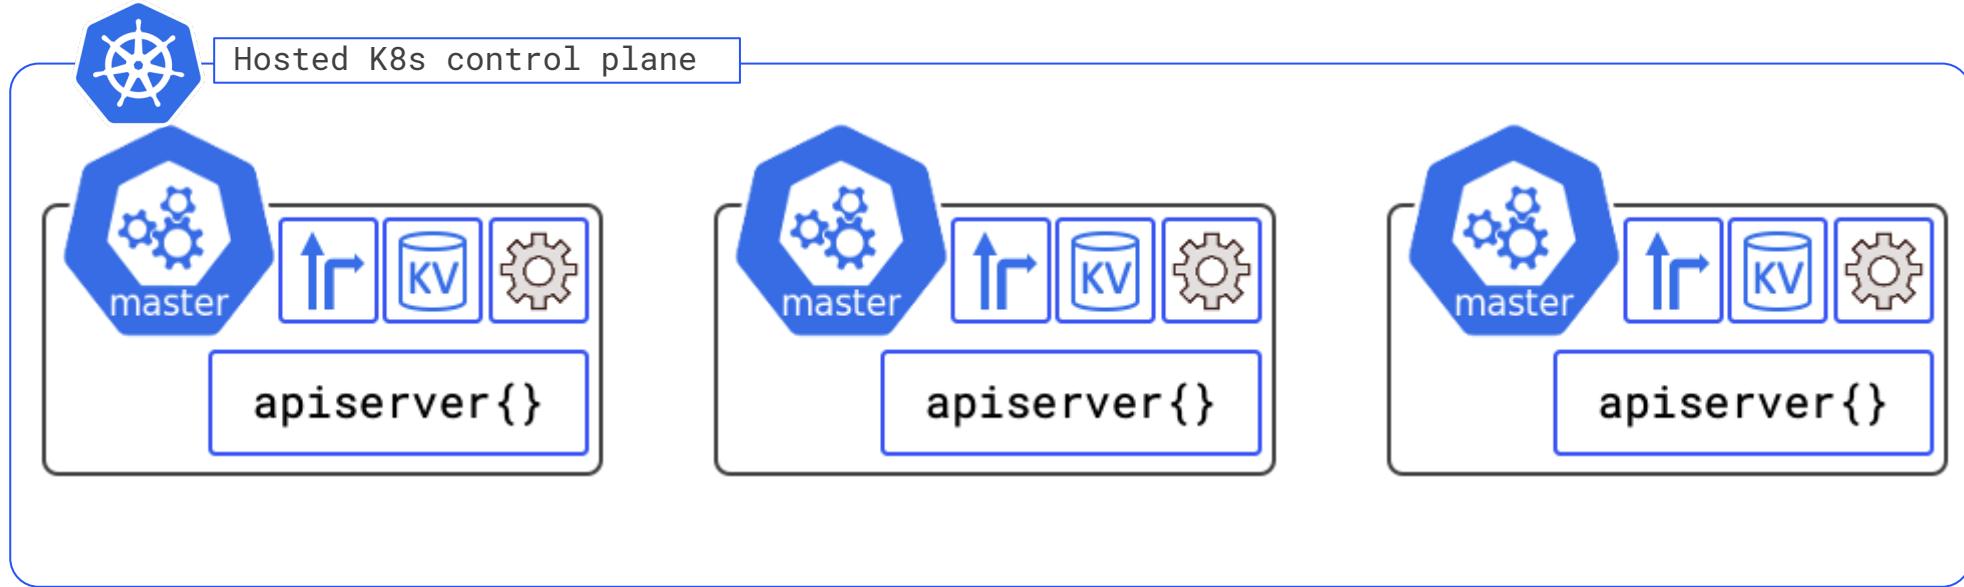


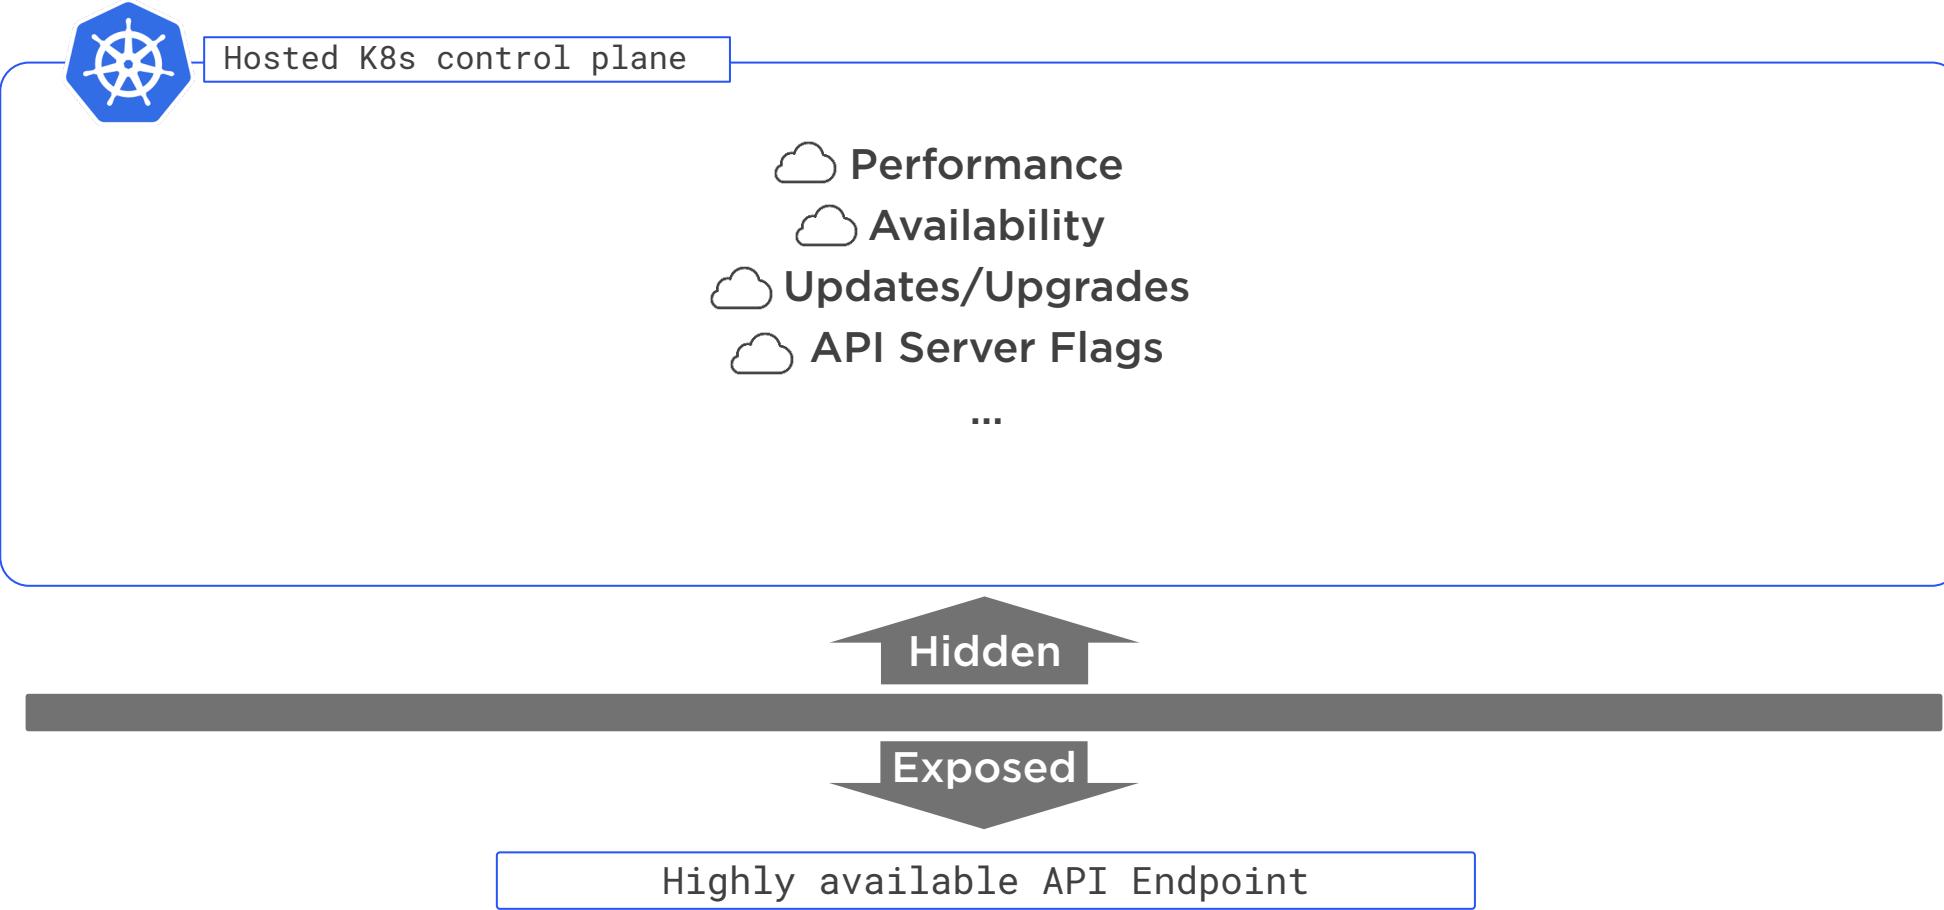


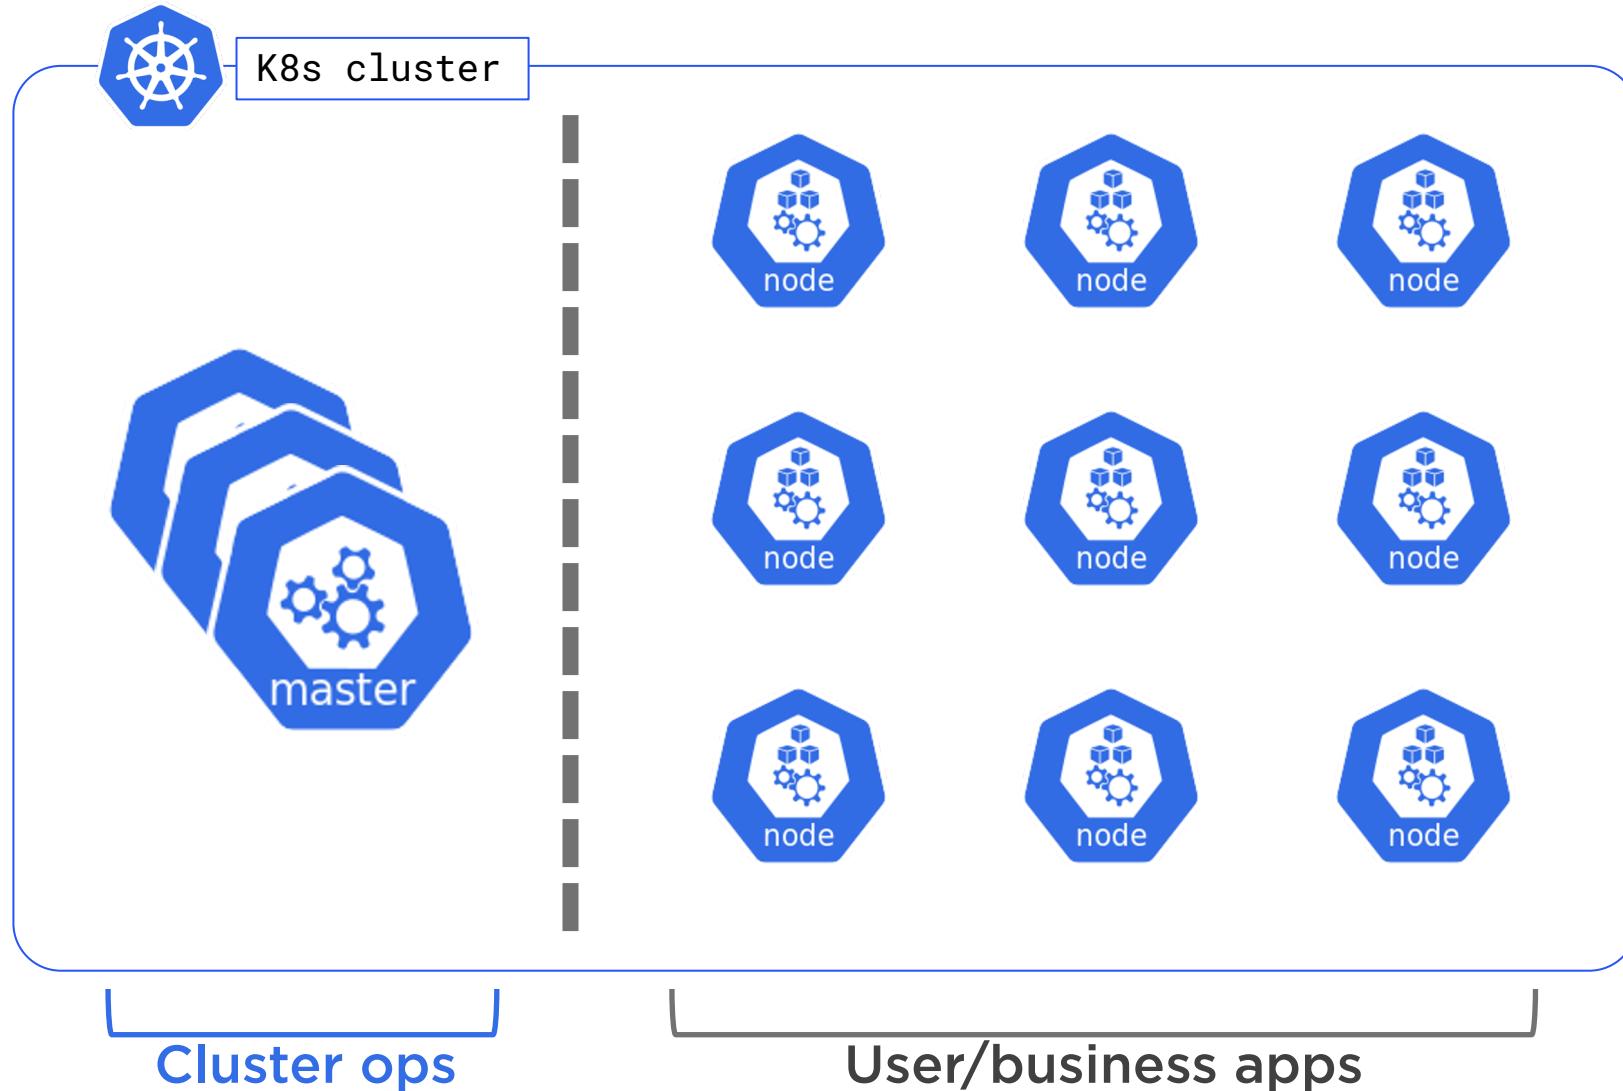




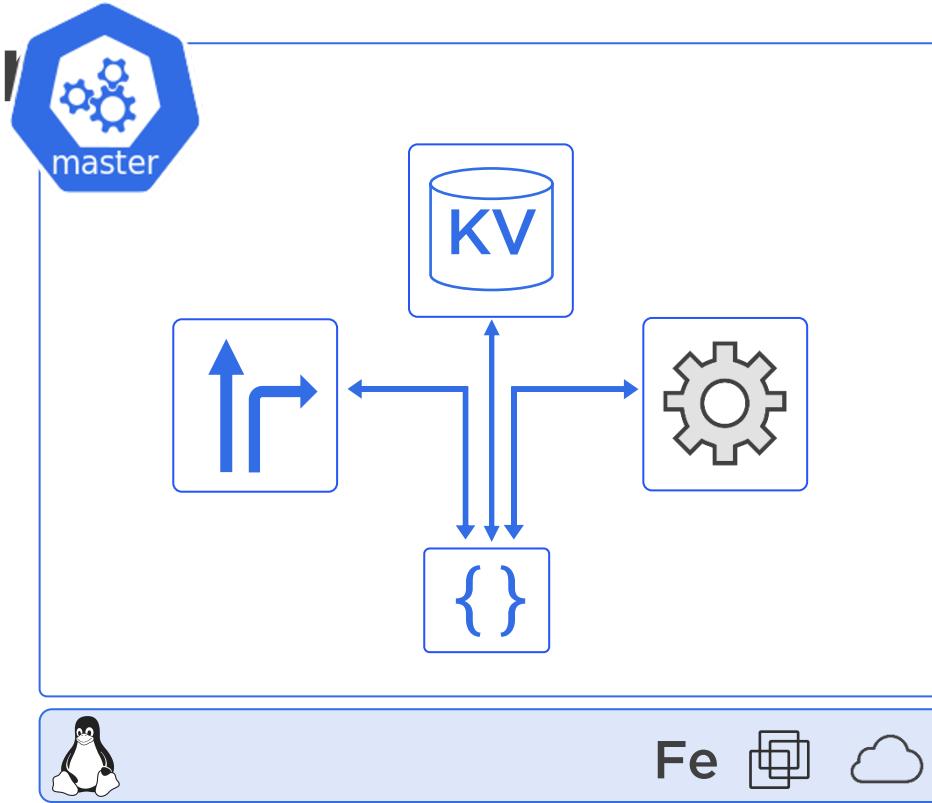








kube-apiserver

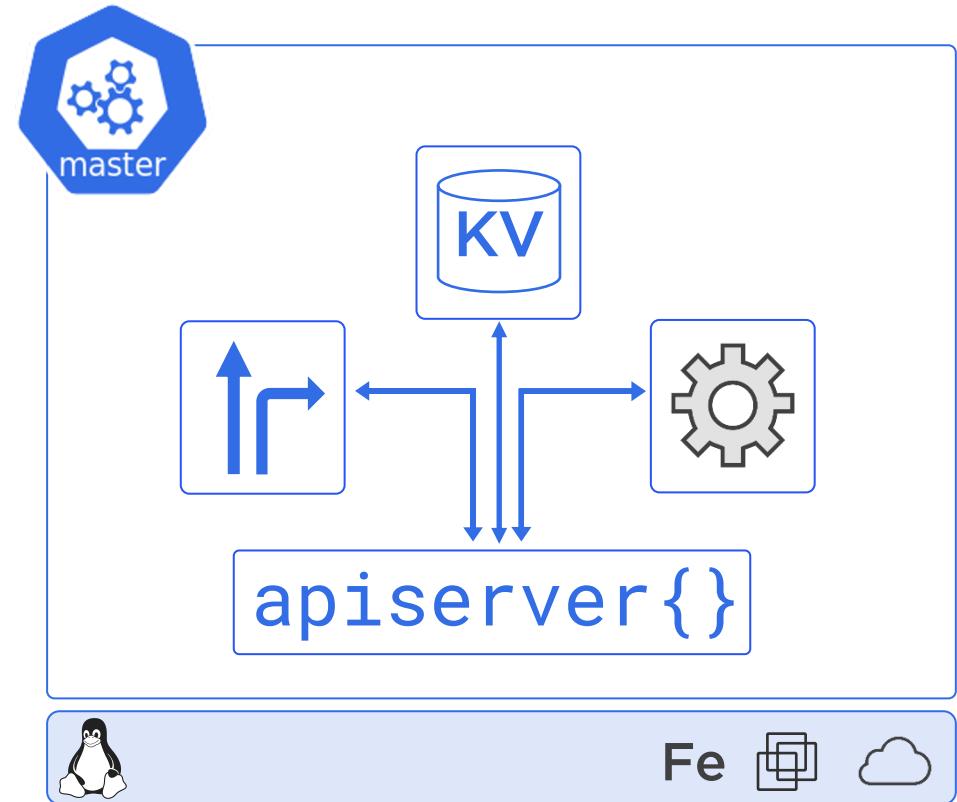


Fe



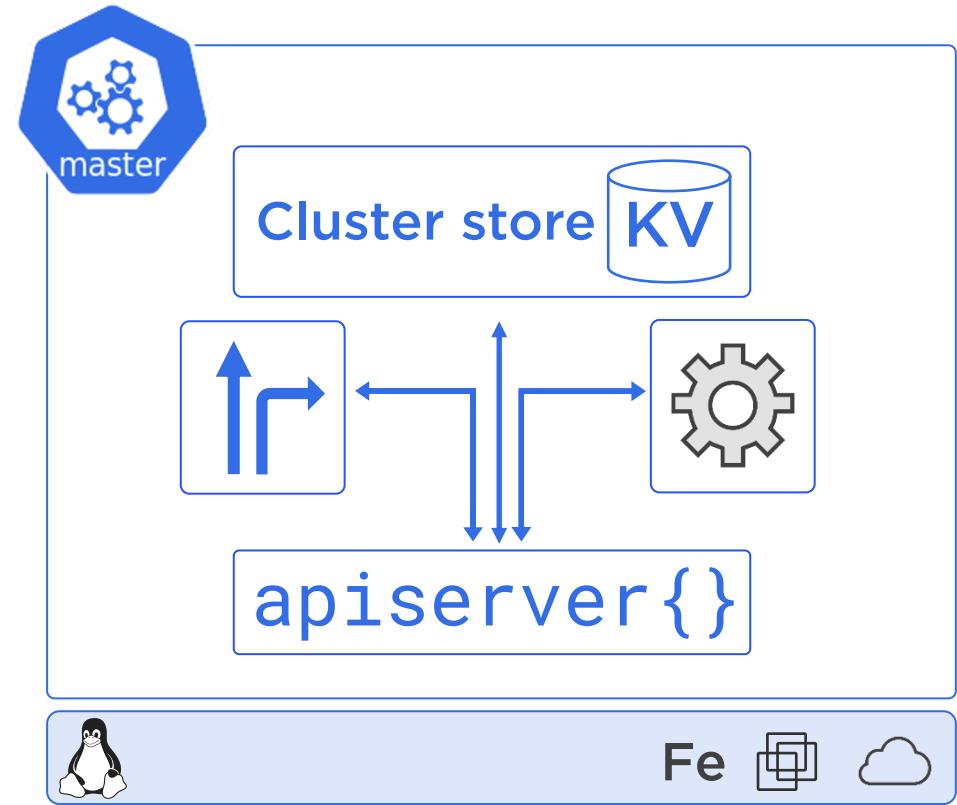
kube-apiserver

- Front-end to the control plane
- Exposes the API (REST)
- Consumes JSON/YAML



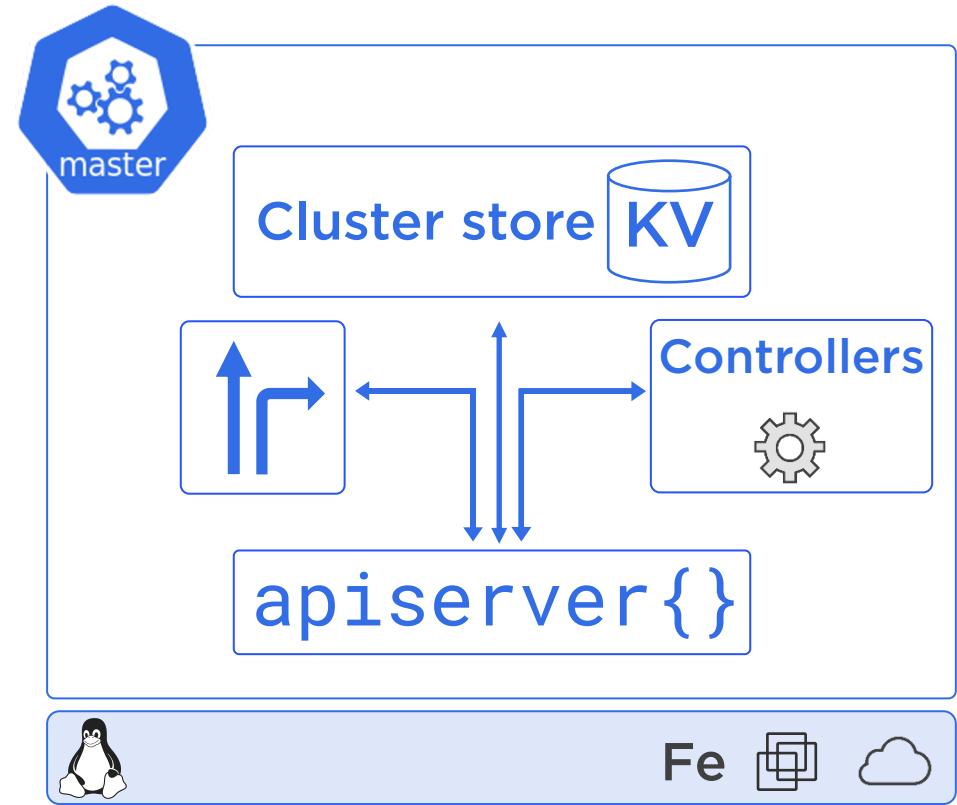
Cluster Store

- Persists cluster state and config
- Based on **etcd**
- Performance is critical
- Have recovery plans in place



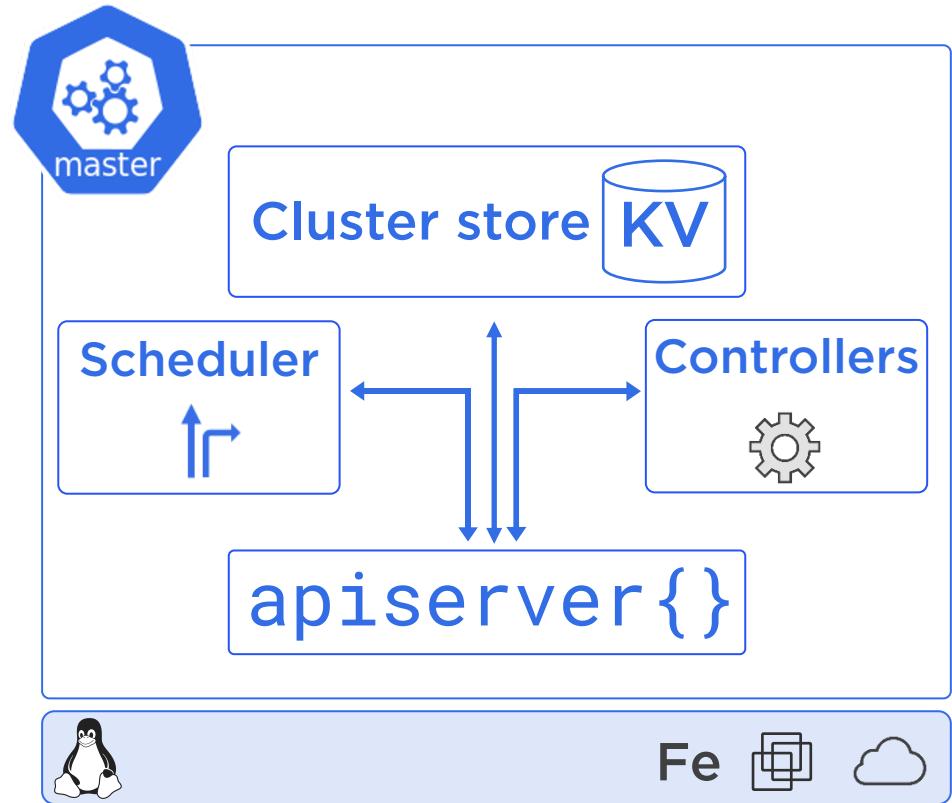
Kube-controller-manager

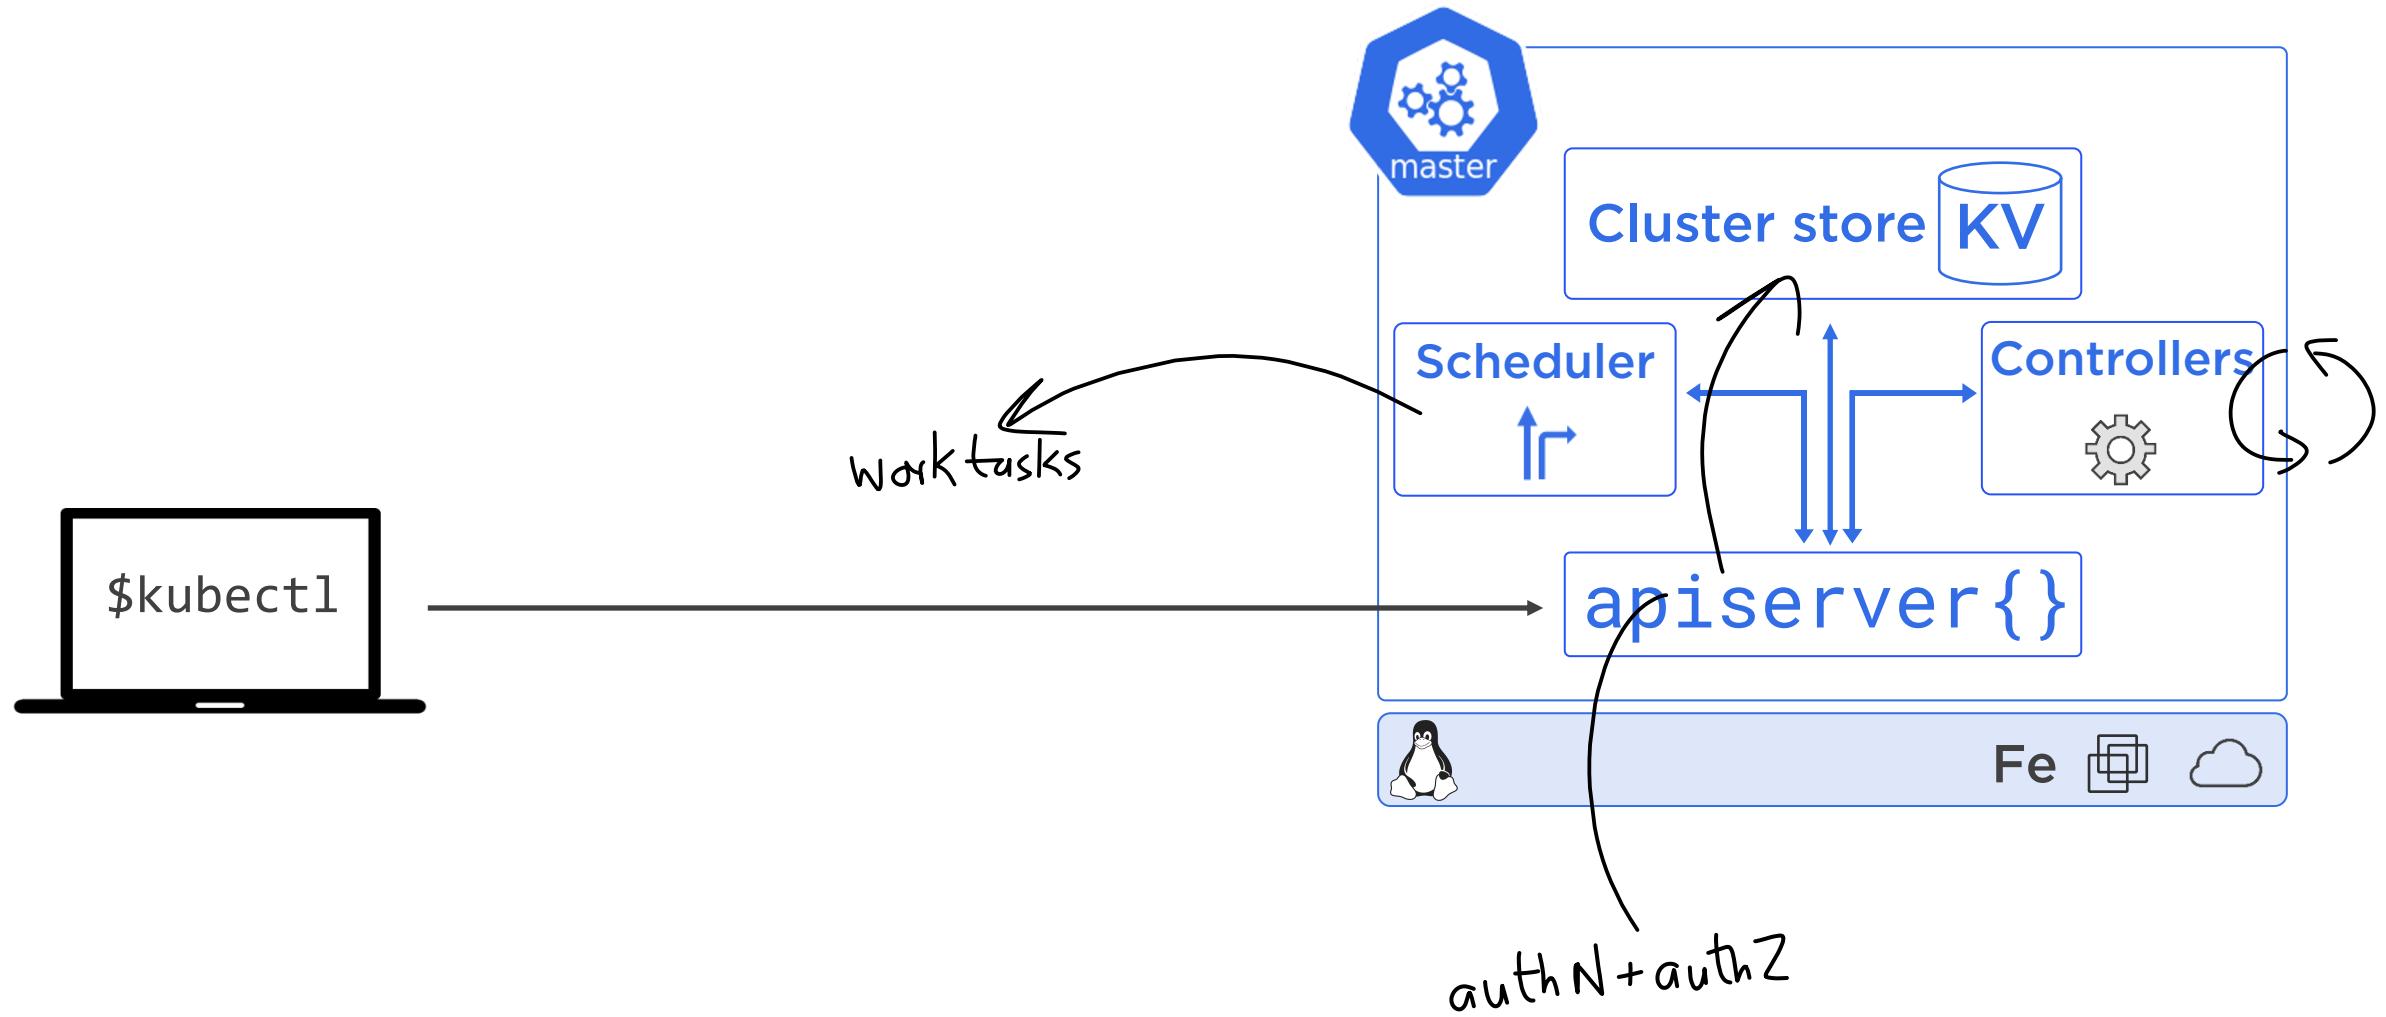
- Controller of controllers
 - Node controller
 - Deployment controller
 - Endpoints/EndpointSlice controller...
- Watch loops
- Reconciles observed state with desired state



Kube-scheduler

- Watches API Server for new work tasks
- Assigns work to cluster nodes
 - Affinity/Anti-affinity
 - Constraints
 - Taints
 - Resources...



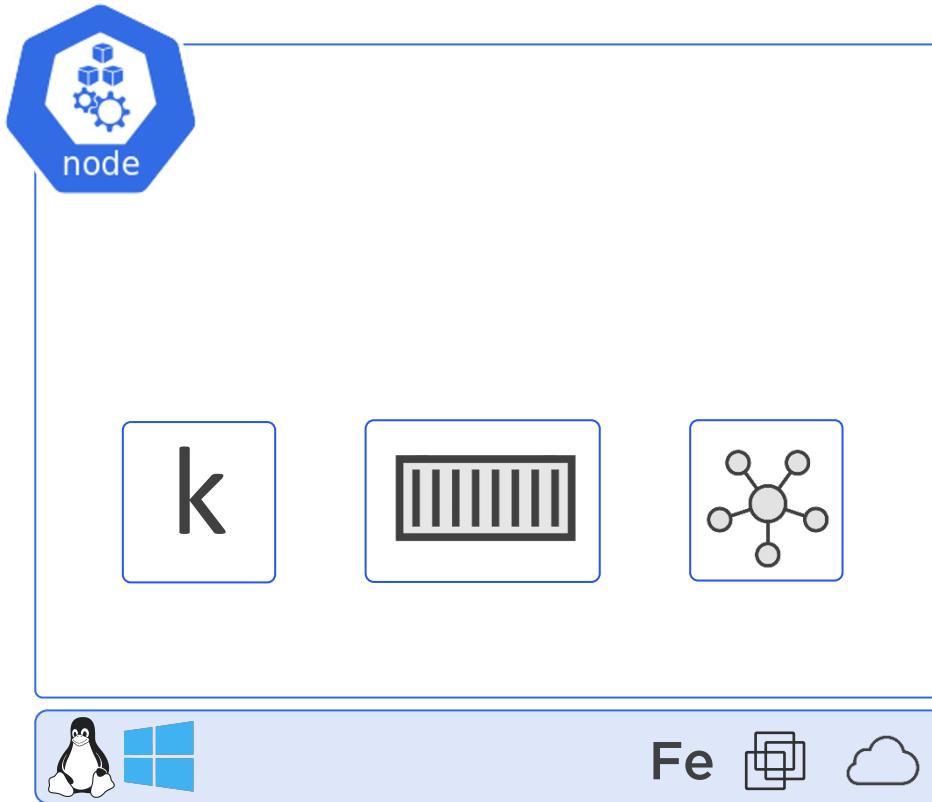


Up Next:
Nodes



Nodes

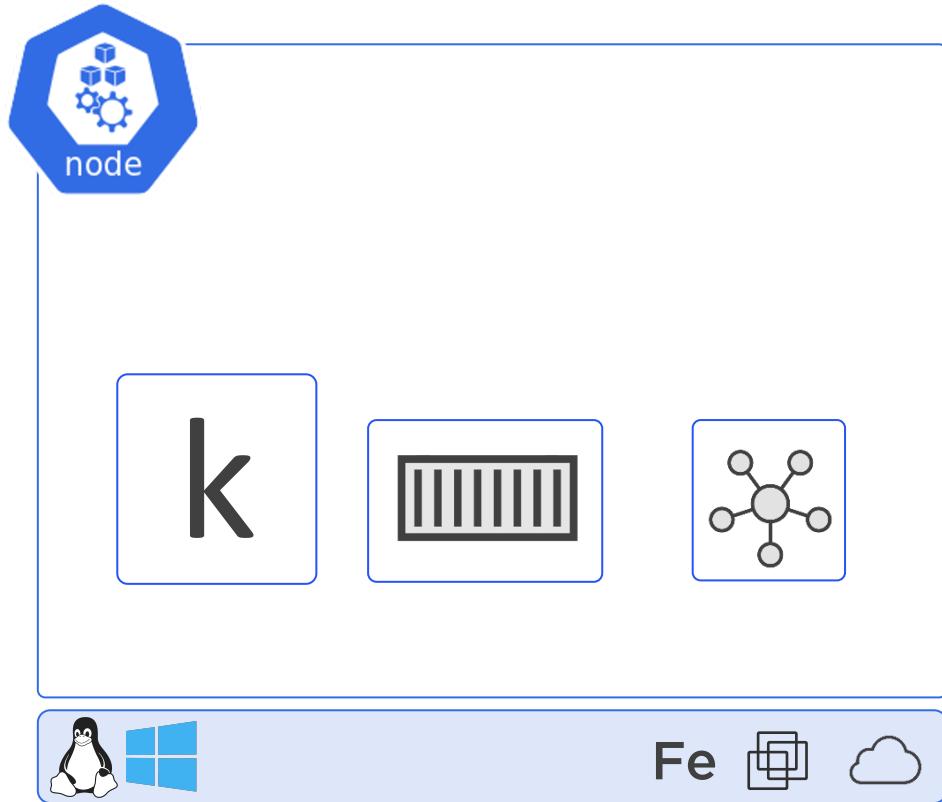






Kubelet

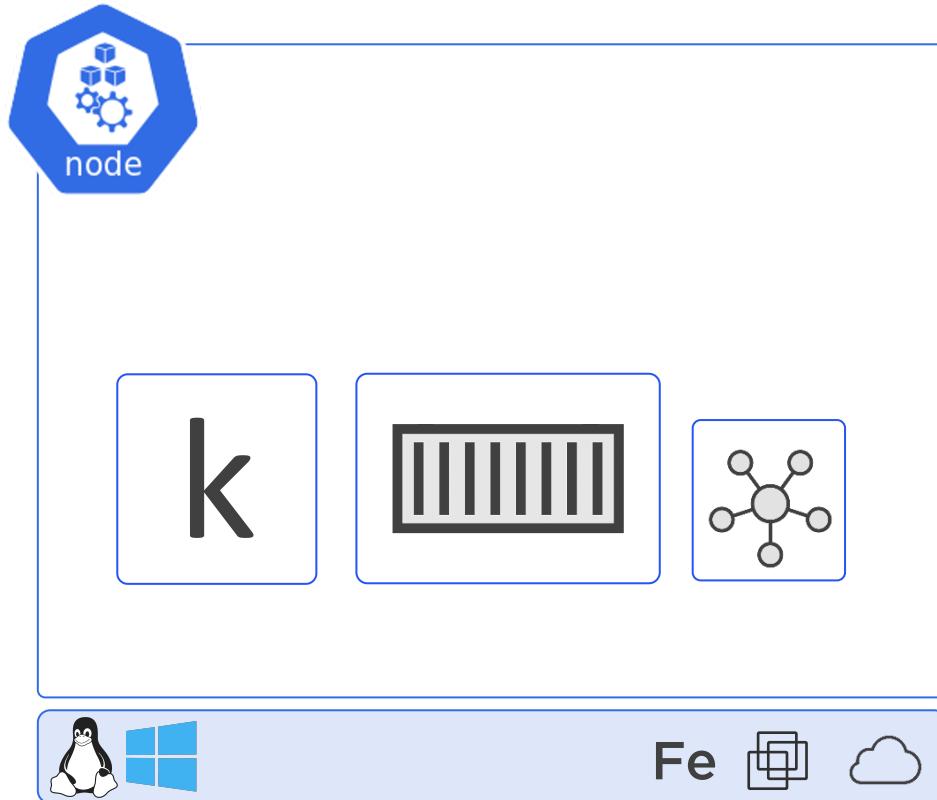
- Main Kubernetes agent
- Registers node with cluster
- Watches API Server for work tasks (Pods)
- Executes Pods
- Reports back to Masters

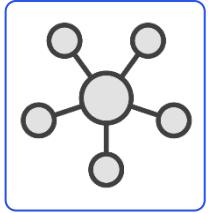




Container runtime

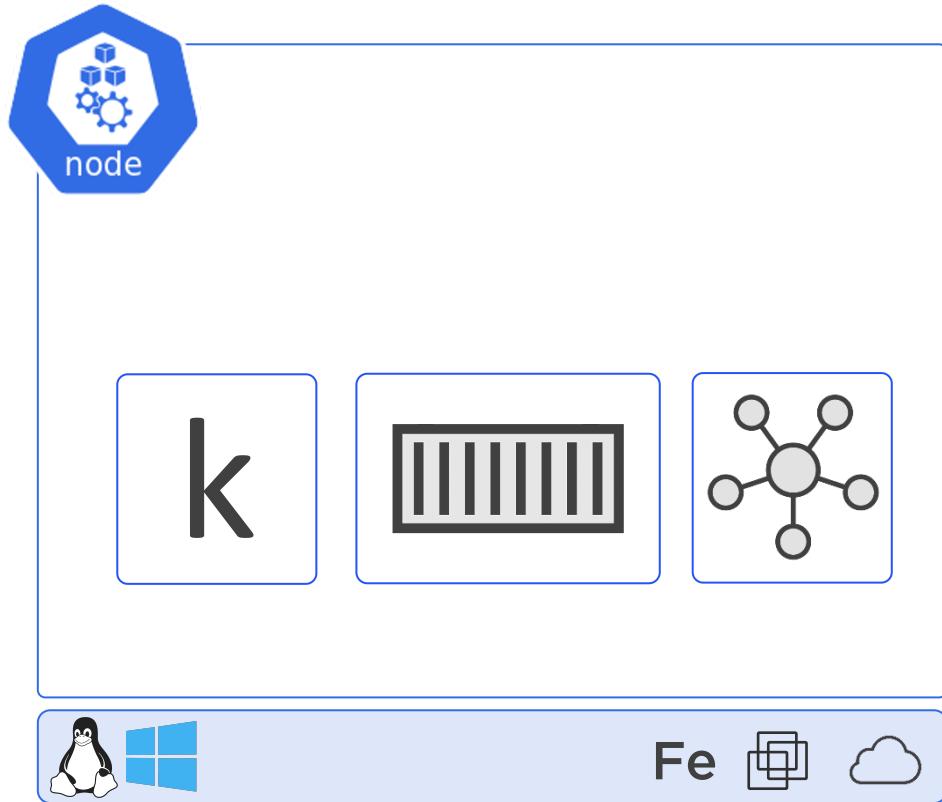
- Can be Docker
- Pluggable: Container Runtime Interface (CRI)
 - Docker, containerd, CRI-O, Kata...
- Low-level container intelligence

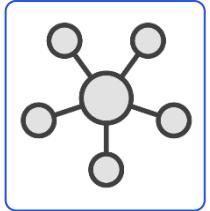




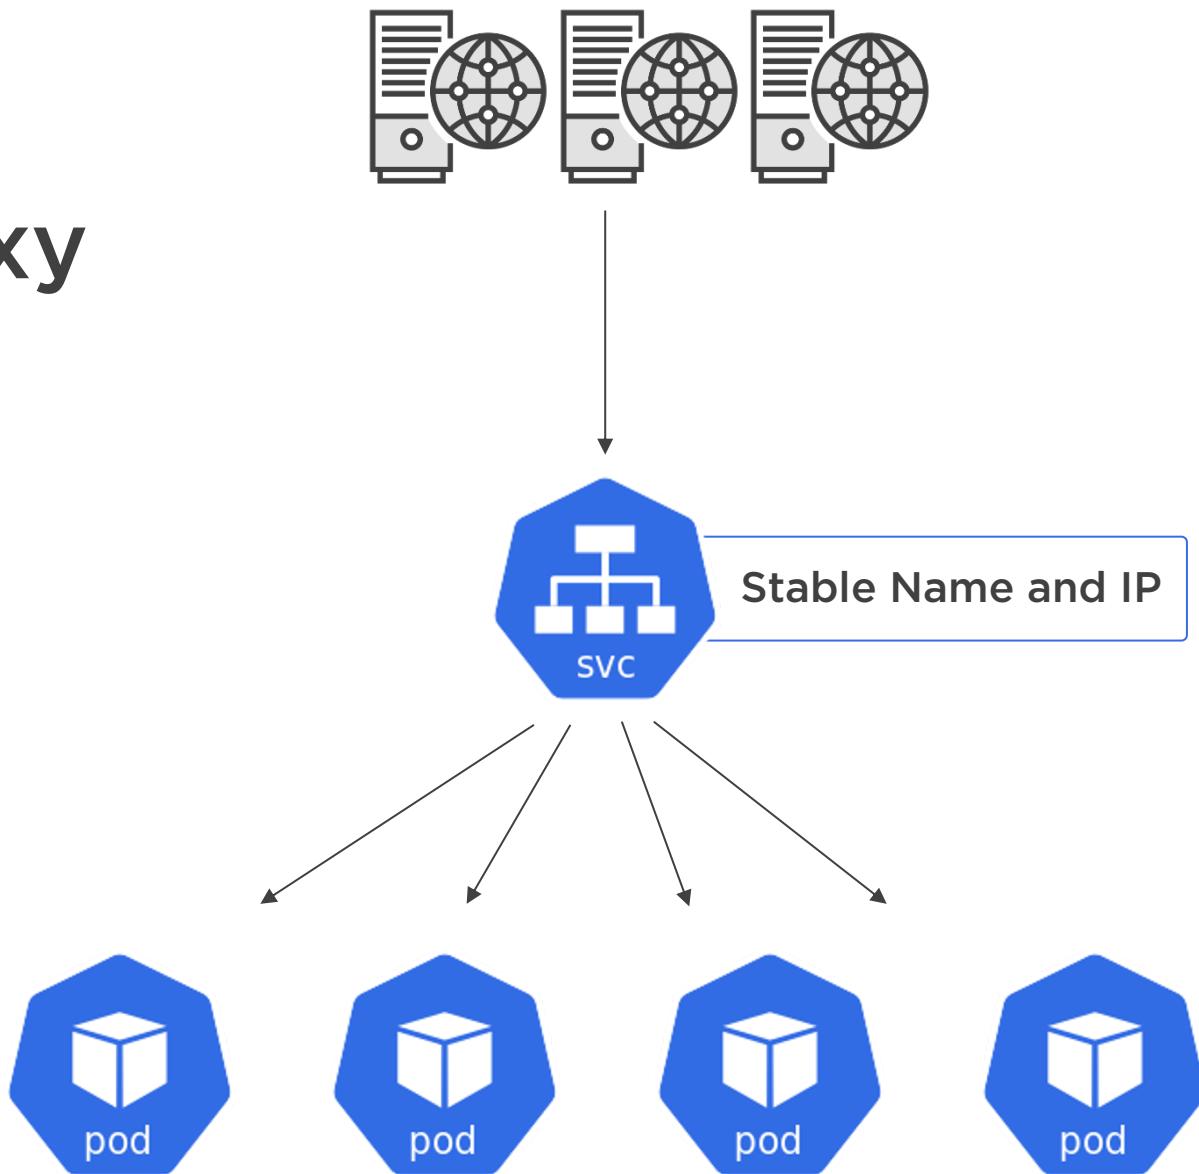
Kube-proxy

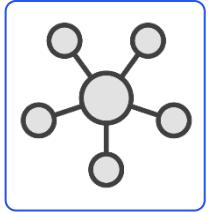
- Networking component
- Pod IP addresses





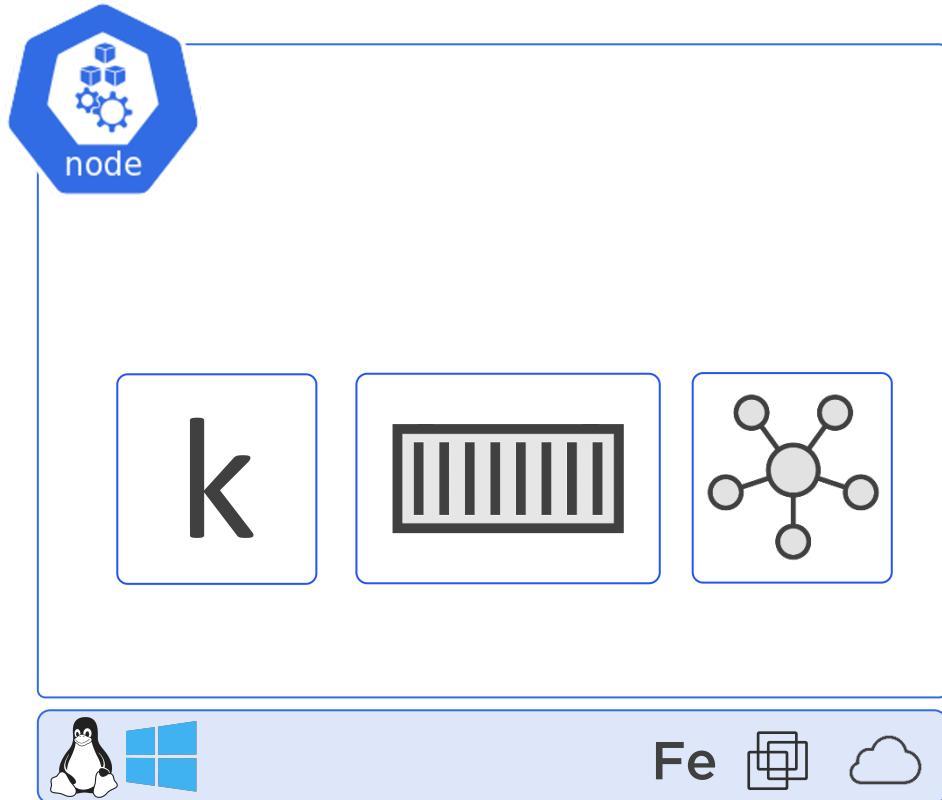
Kube-proxy

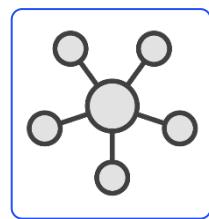
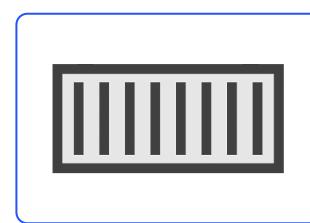
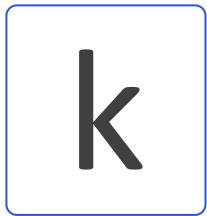




Kube-proxy

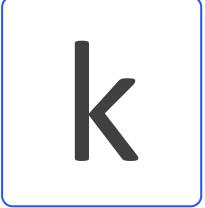
- Networking component
- Pod IP addresses
- Basic load-balancing





Fe





k

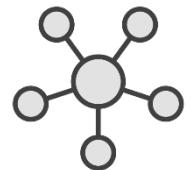
Kubelet

Main K8s agent



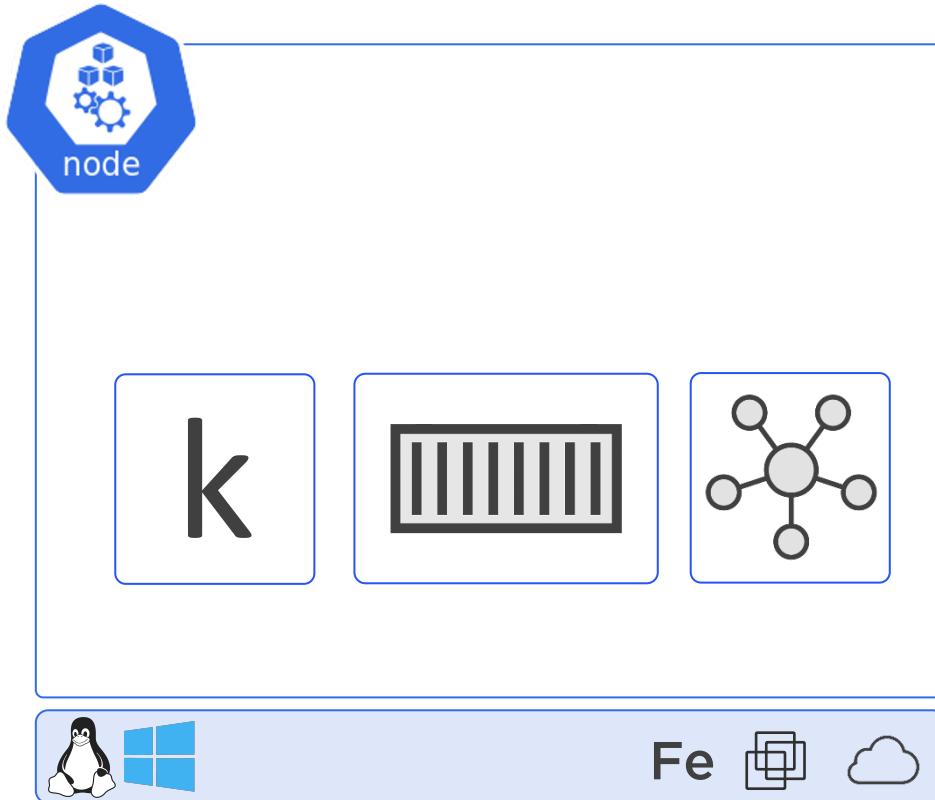
Container runtime

Docker, containerd, CRI-O, more...



Kube-proxy

Vital role in networking





Virtual Kubelet



Nodeless Kubernetes





Virtual Kubelet



Nodeless Kubernetes





Virtual Kubelet

Pods run on cloud's hosted container back-end



Up Next:
The Declarative Model & Desired State



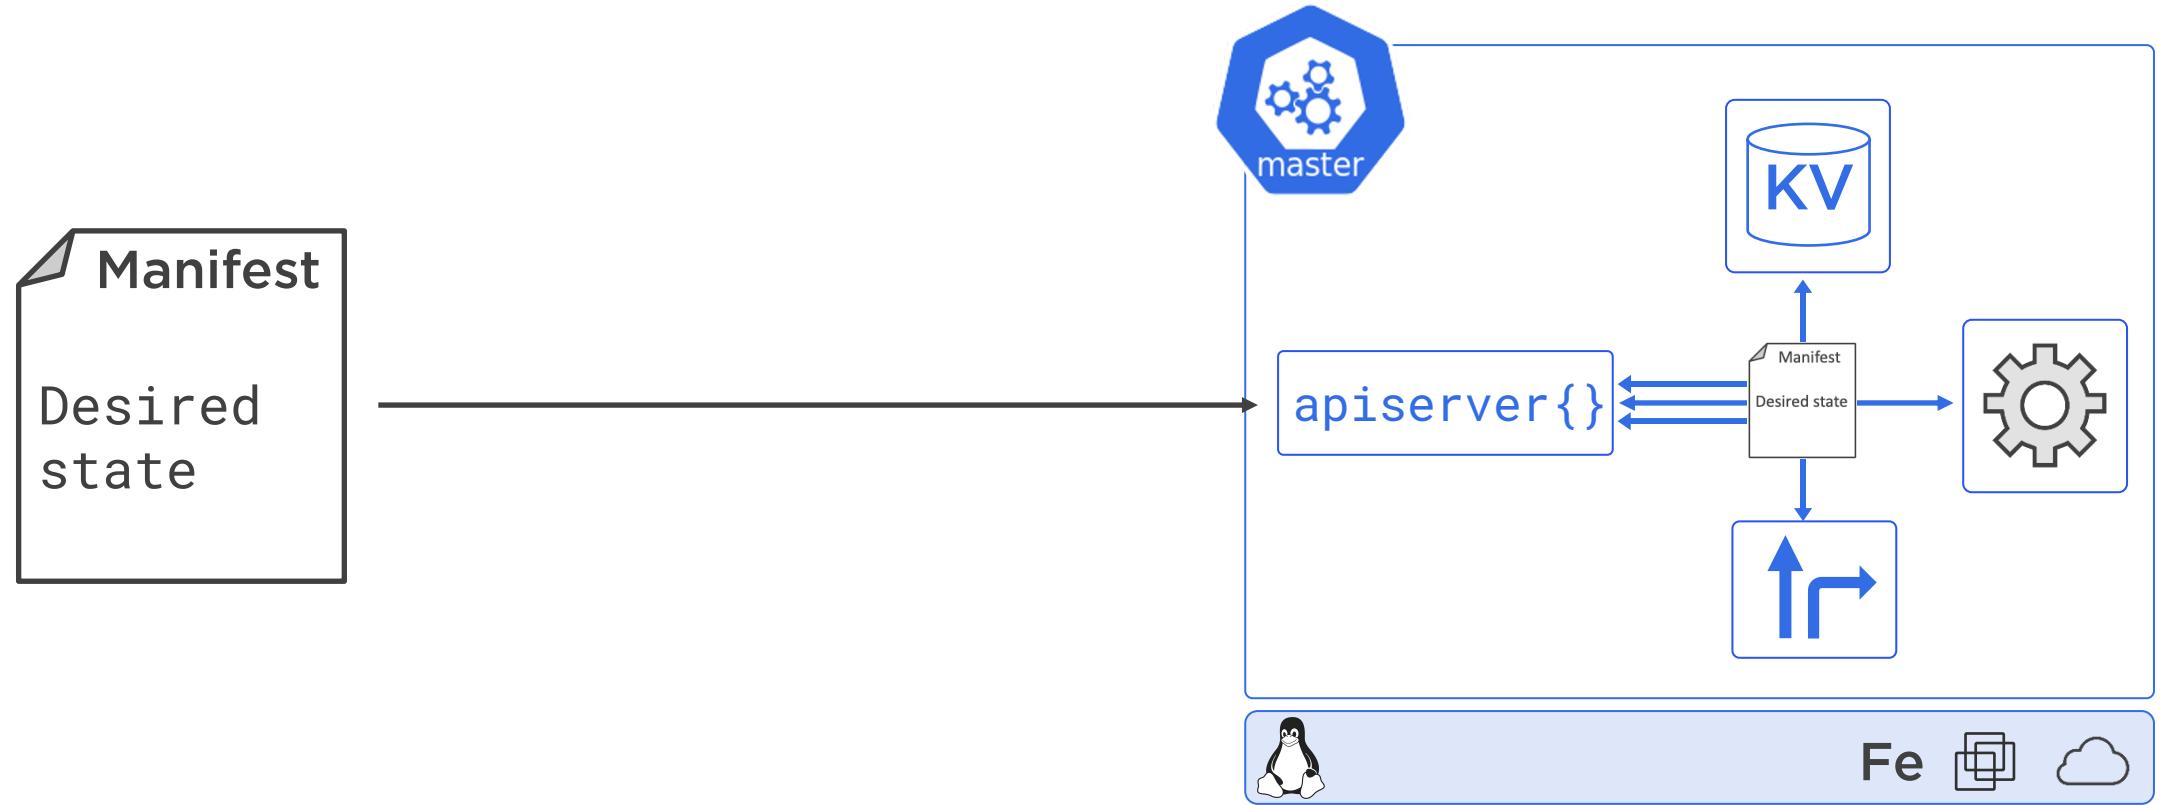
The Declarative Model and Desired State

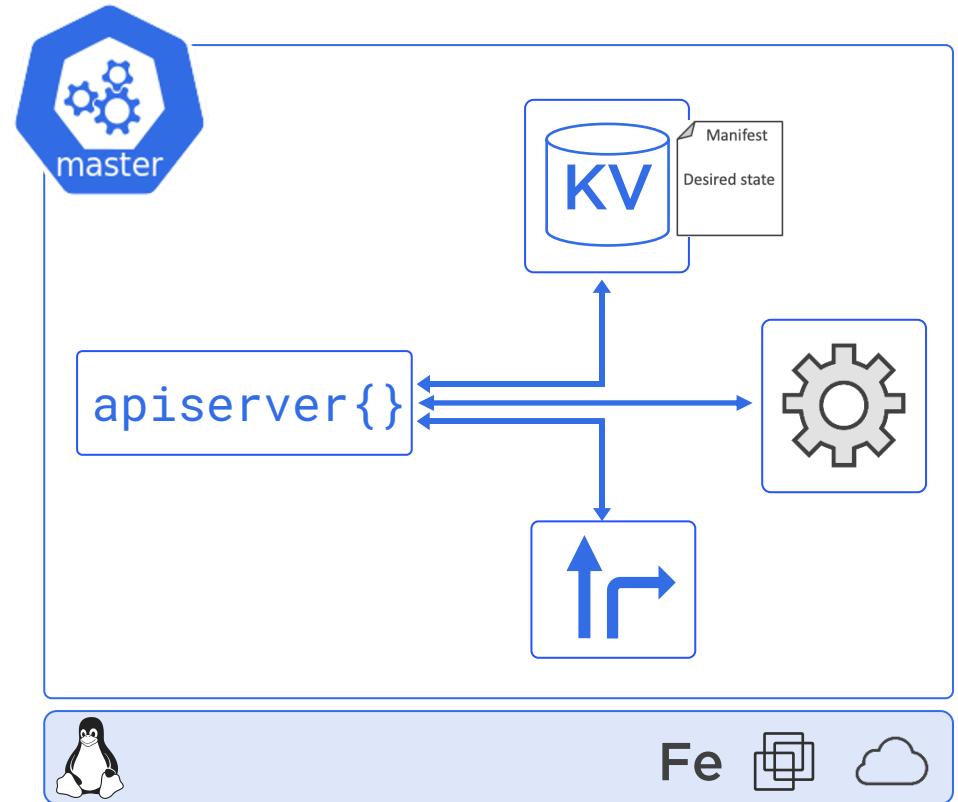


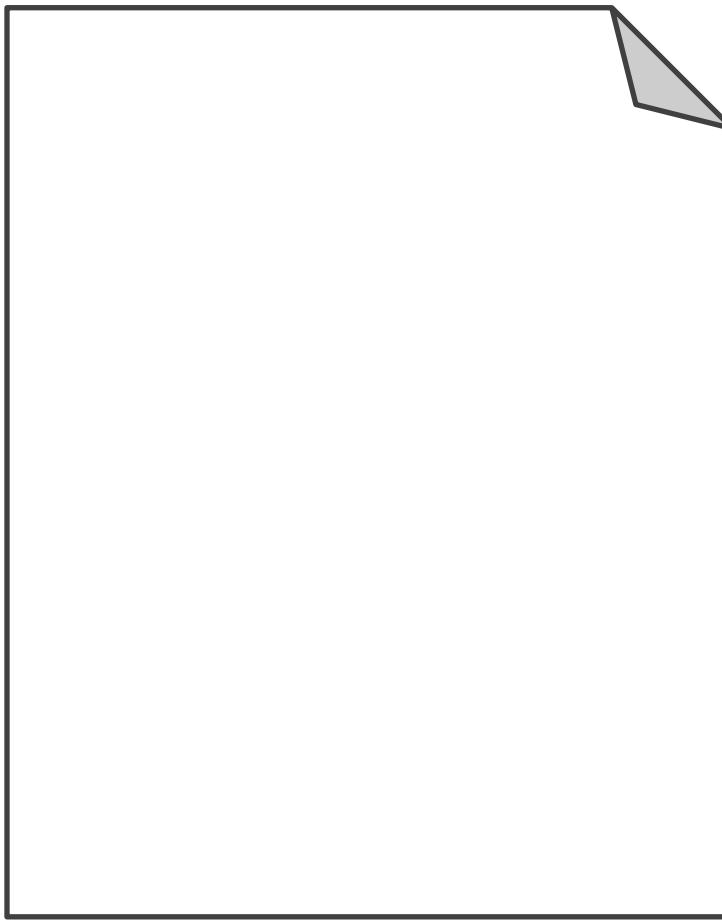
Declarative model

Describe what you *want (desired state)* in a manifest file







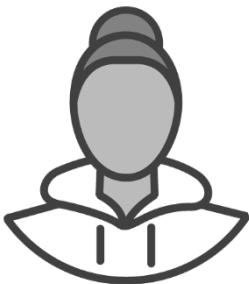


```
kind: Kitchen
spec:
  type: New
  location: Rear
  style: OpenPlan
  heating:
    type: Underfloor
    medium: Water
  windows:
    type: FloorToCeiling
    aspect: South
  doors:
    type: FireDoor
    accessTo: Garage
  island: Yes
  roofGarden: Yes
  ...
```



```
kind: Kitchen
spec:
  type: New
  location: Rear
  style: OpenPlan
  heating:
    type: Underfloor
    medium: Water
  windows:
    type: FloorToCeiling
    aspect: South
  doors:
    type: FireDoor
    accessTo: Garage
  island: Yes
  roofGarden: Yes
  ...
```





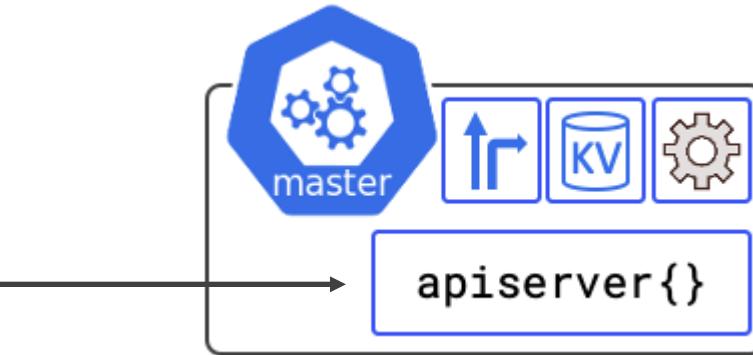
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: test
spec:
  replicas: 6
  selector:
    matchLabels:
      app: ps-test
  template:
    spec:
      containers:
        - name: c1
          image: web1:1.3
          ports: 8080
            - containerPort
              ...

```

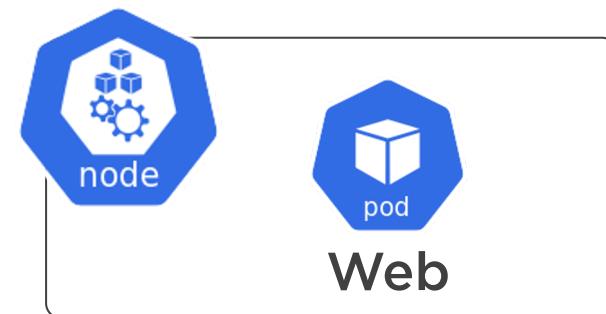
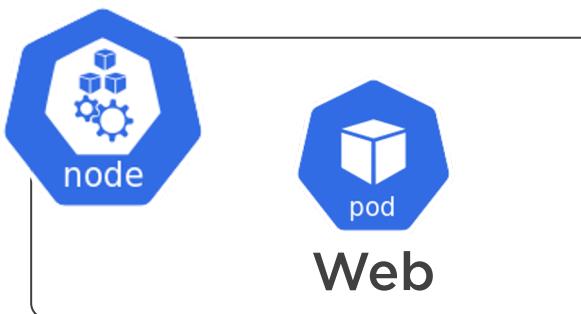
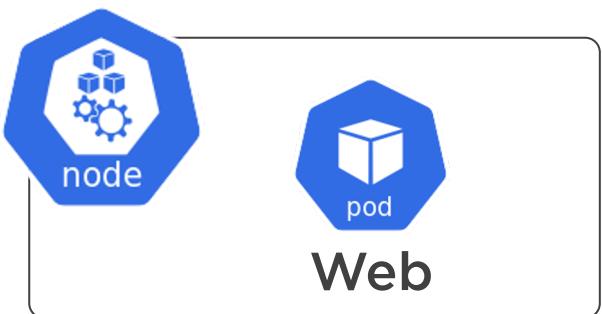
Declarative
(Declaring what you want)



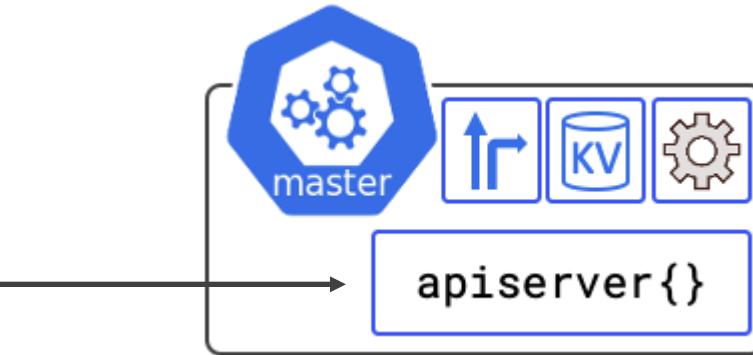
```
name: web  
replicas: 3  
...
```



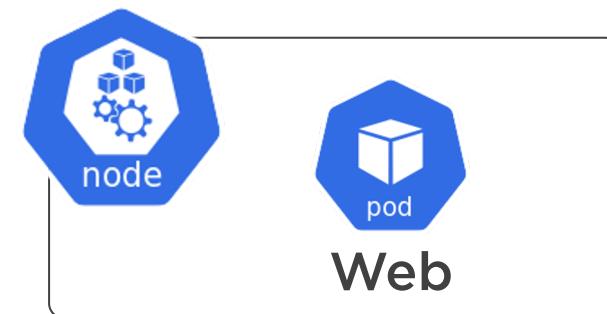
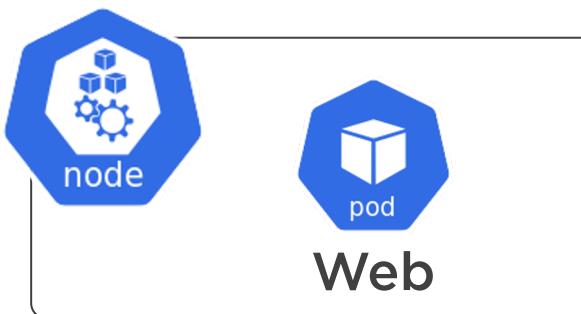
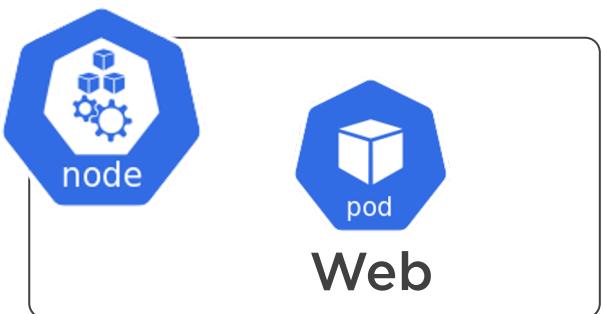
Desired state: 3
Observed state: 3



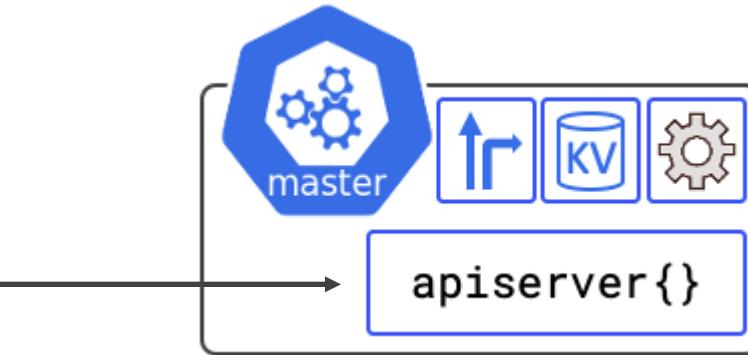
```
name: web  
replicas: 3  
...
```



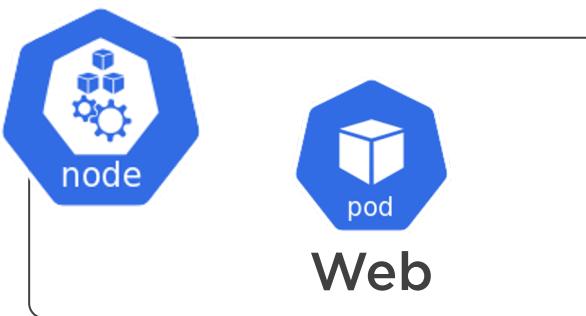
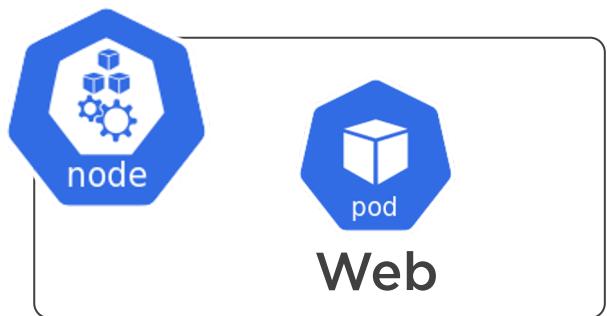
Desired state: 3
Observed state: 3



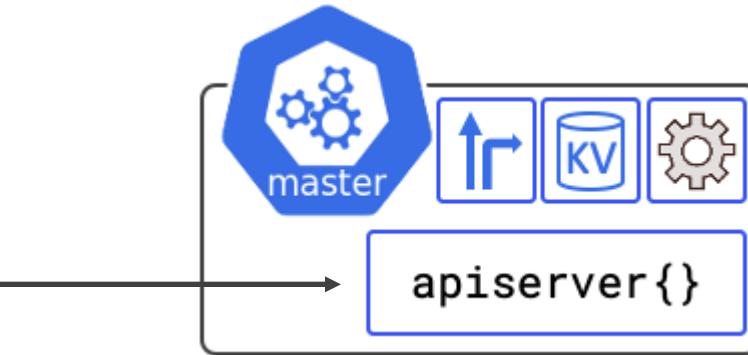
```
name: web  
replicas: 3  
...
```



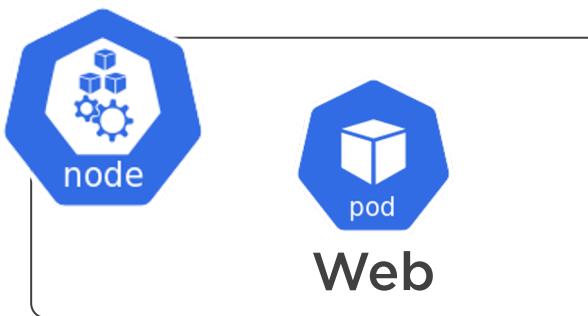
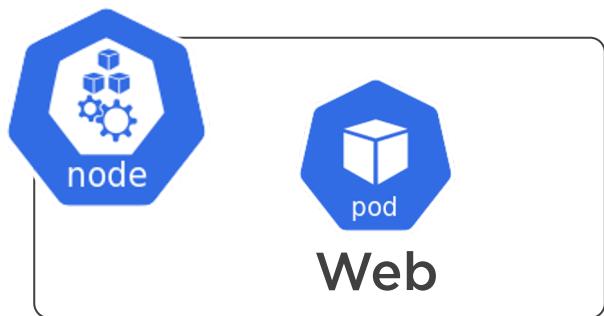
Desired state: 3
Observed state: 3



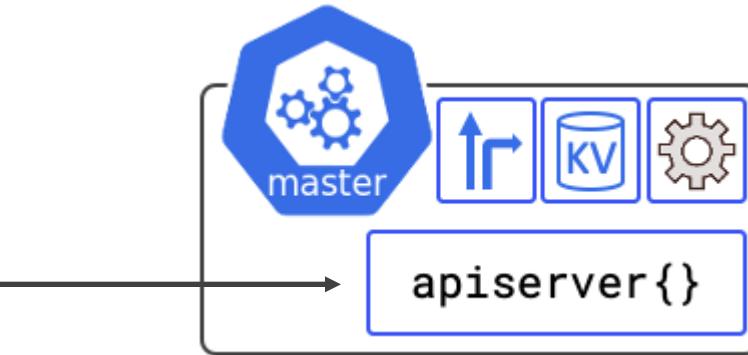
```
name: web  
replicas: 3  
...
```



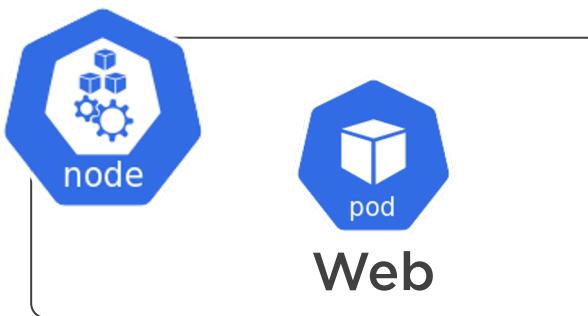
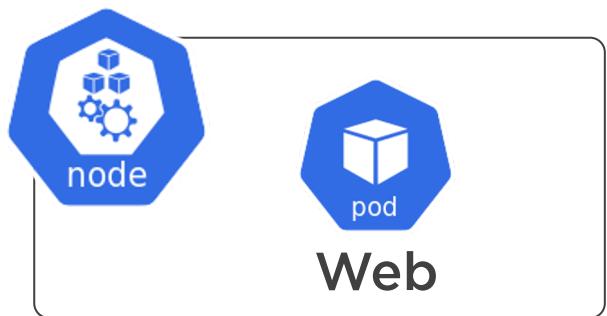
Desired state: 3
Observed state: 2



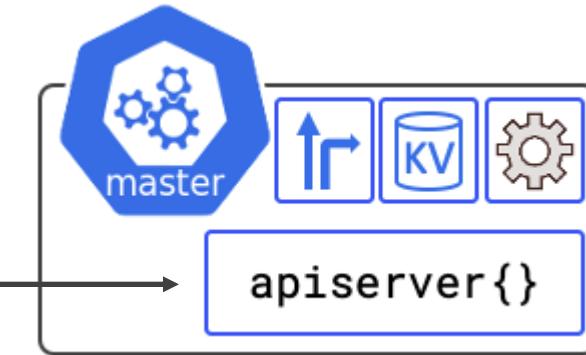
```
name: web  
replicas: 3  
...
```



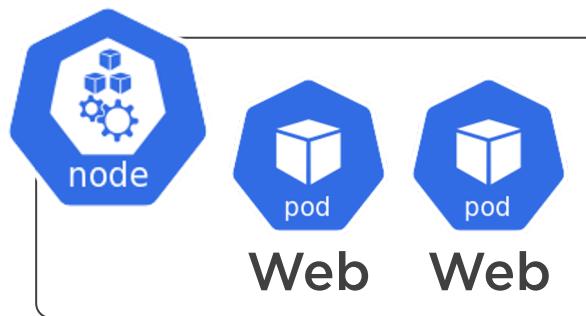
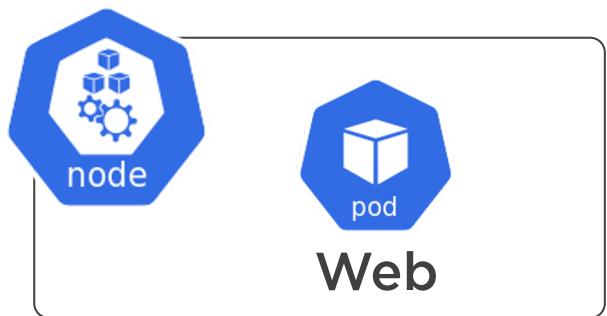
Desired state: 3
Observed state: 2



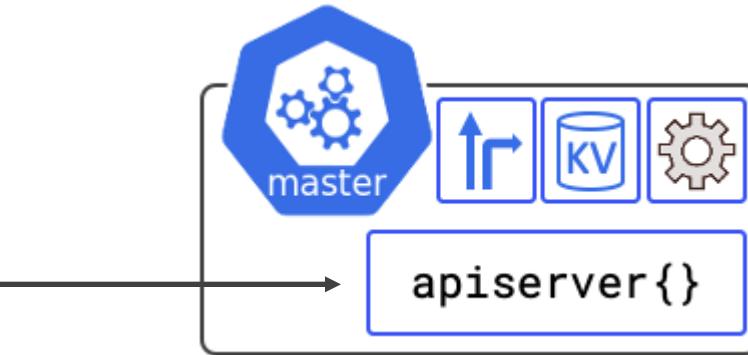
```
name: web  
replicas: 3  
...
```



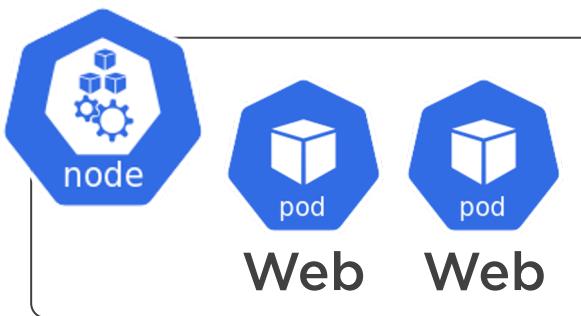
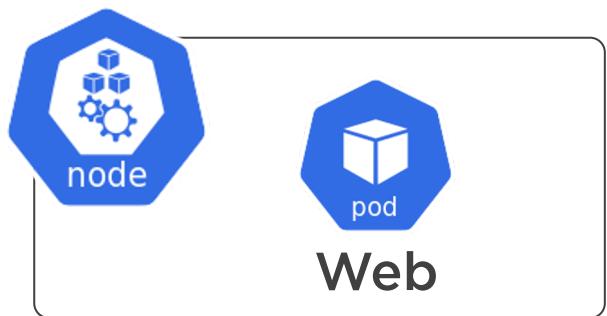
Desired state: 3
Observed state: 2

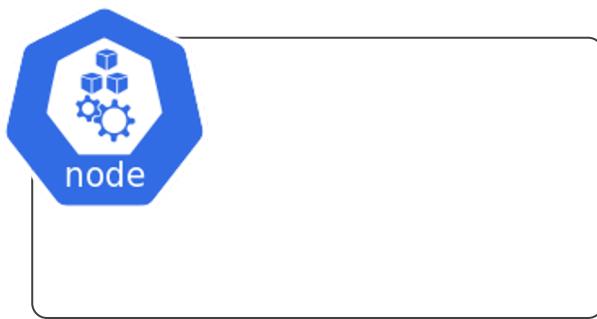
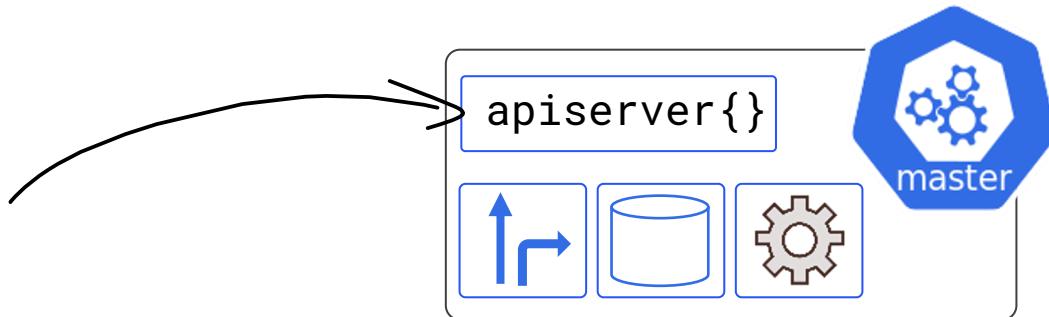


```
name: web  
replicas: 3  
...
```

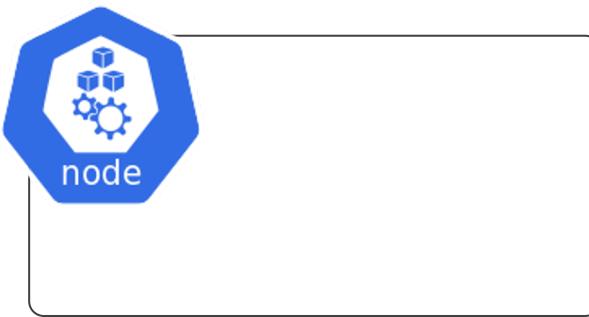
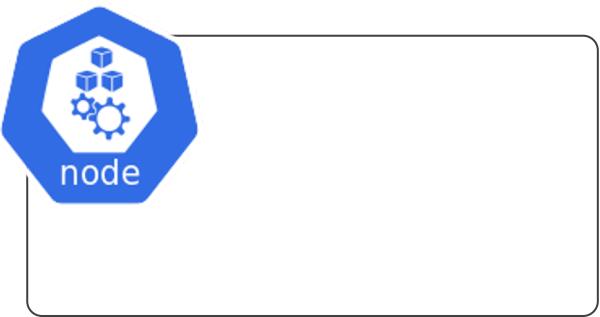
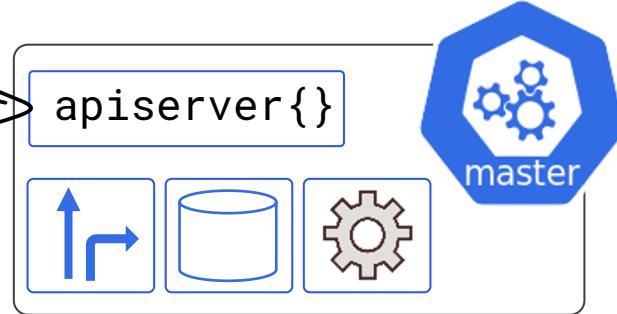


Desired state: 3
Observed state: 3

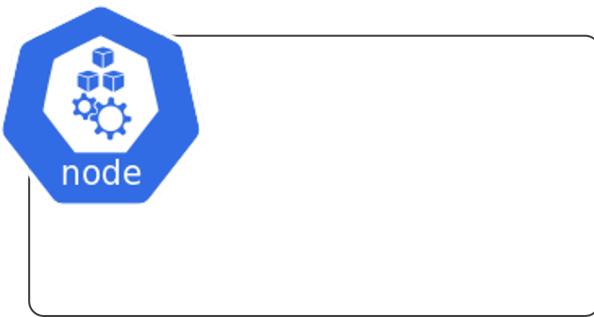
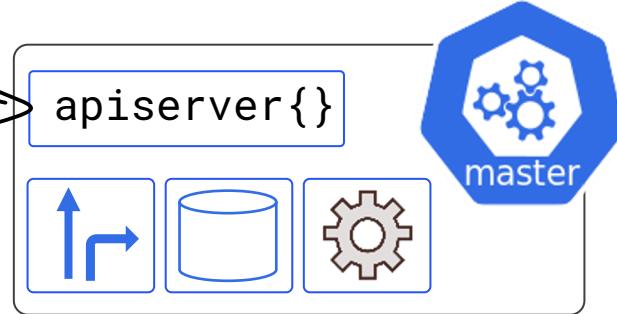


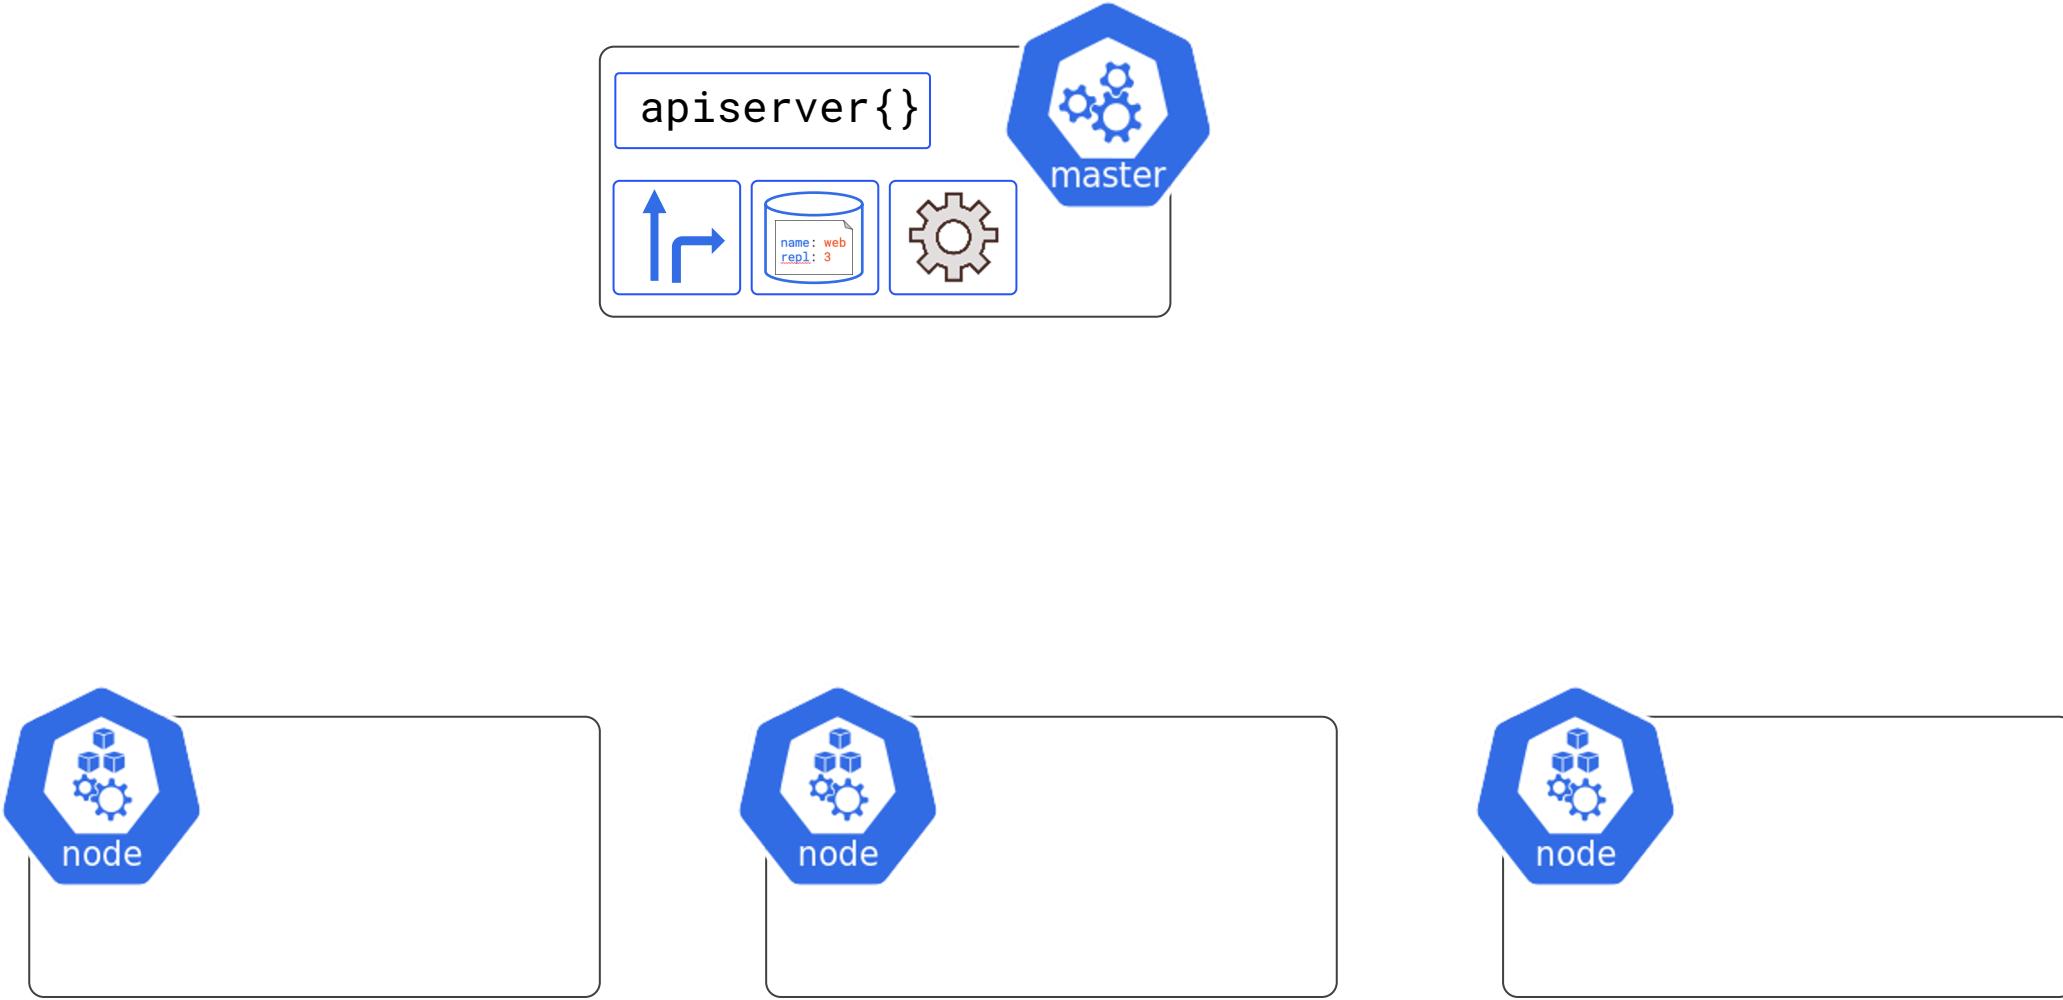


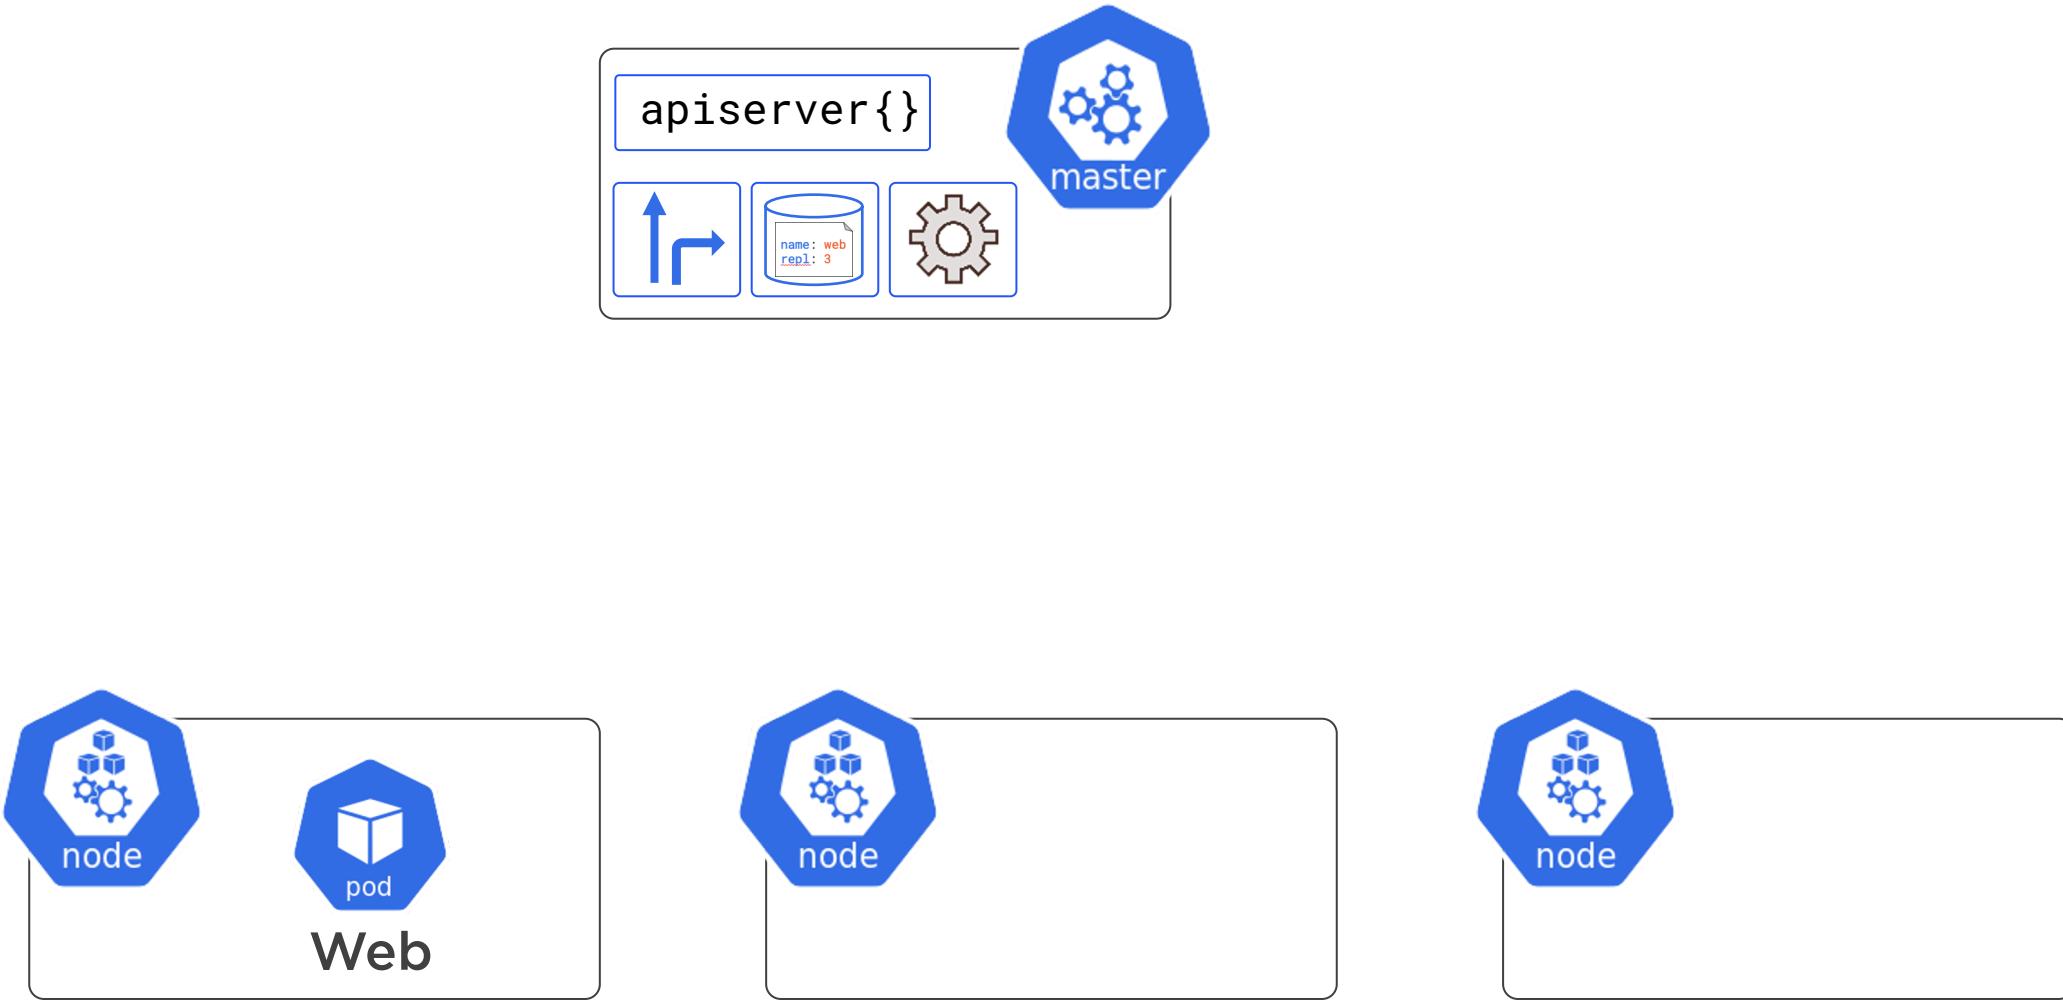
```
name: web  
replicas: 3  
...
```

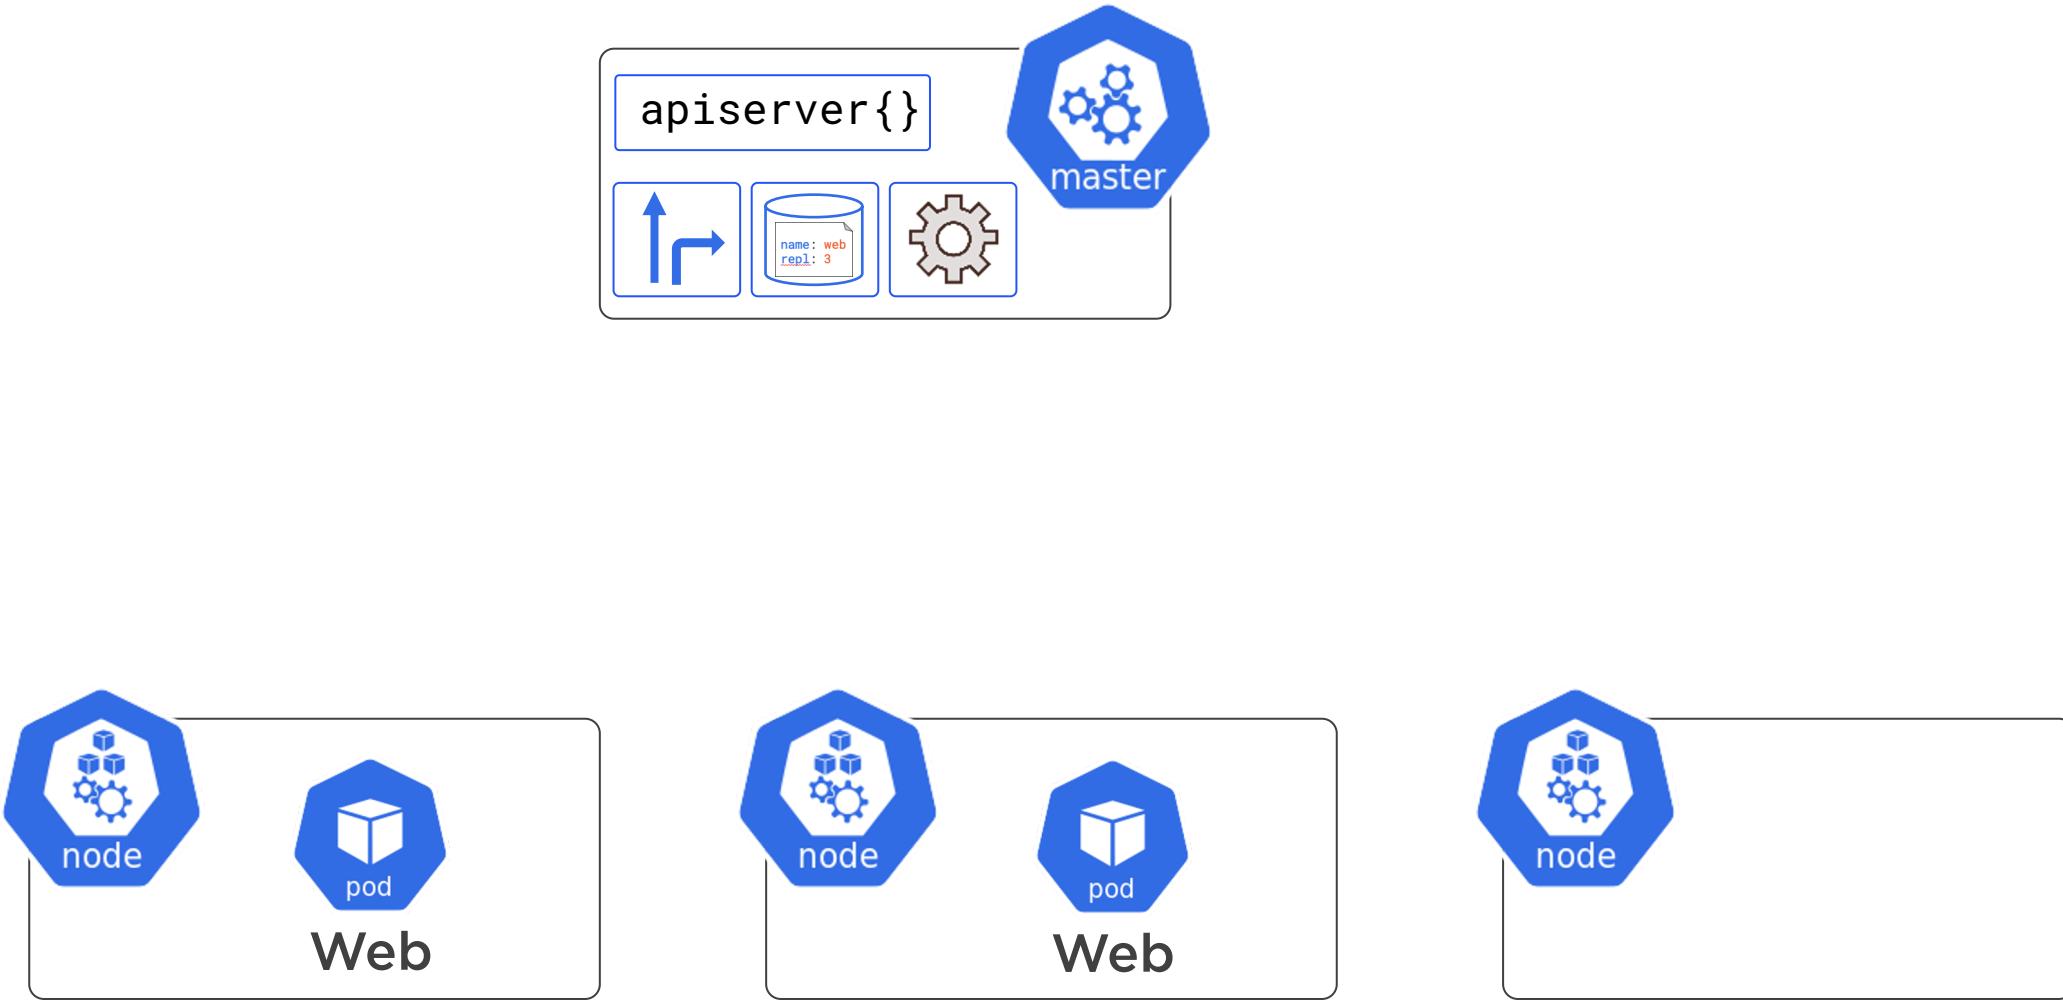


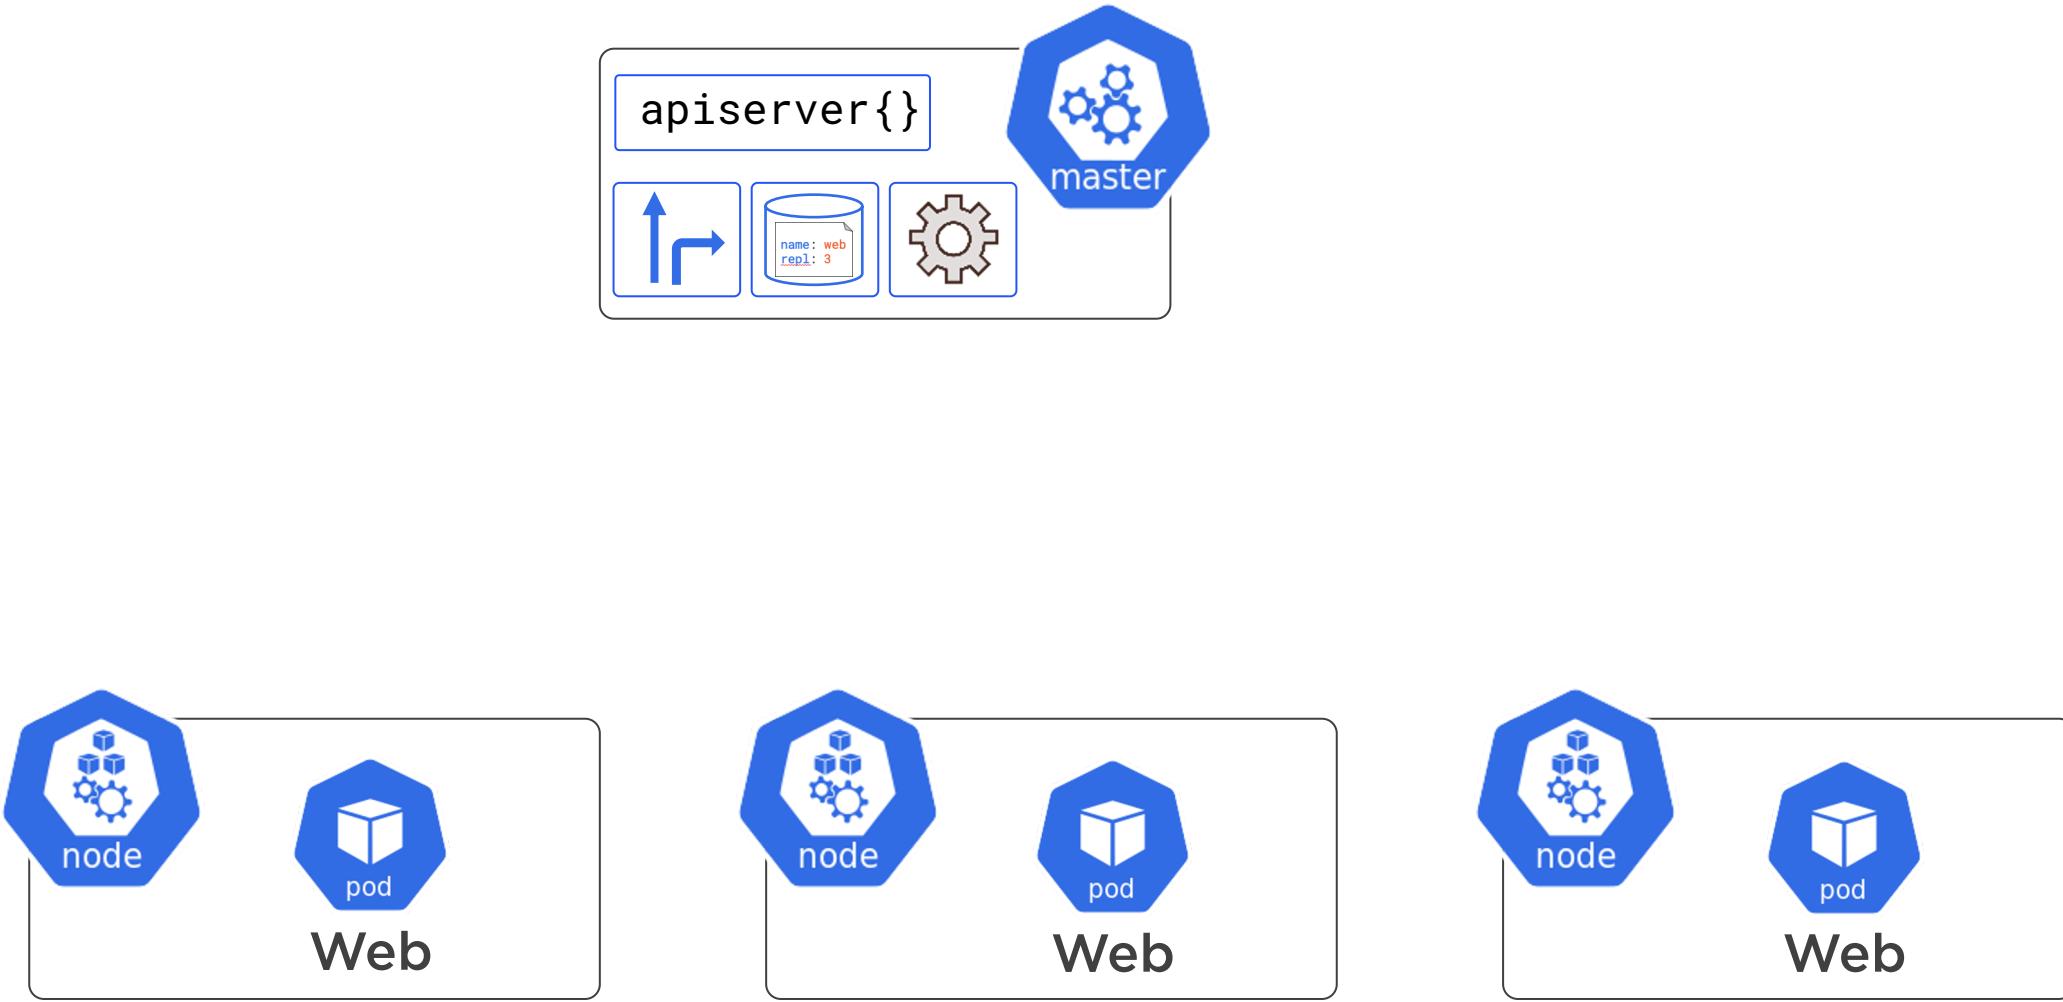
```
name: web  
replicas: 3  
...
```

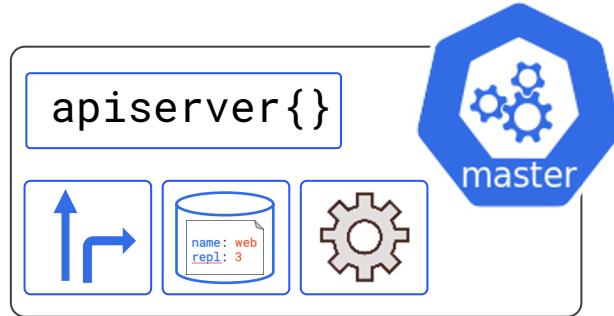




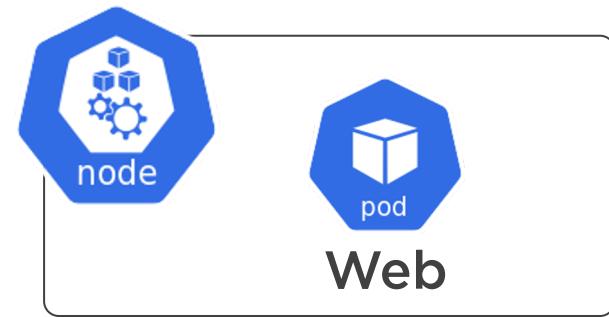
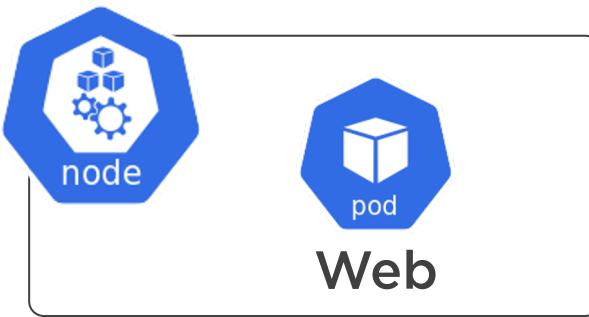
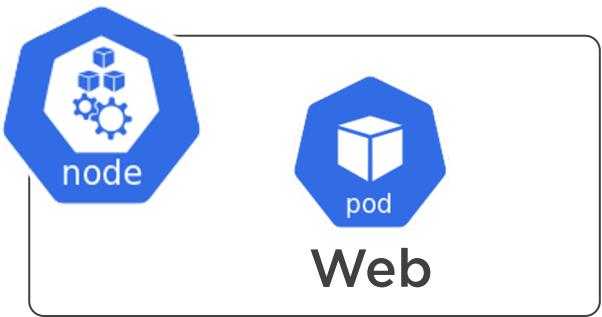


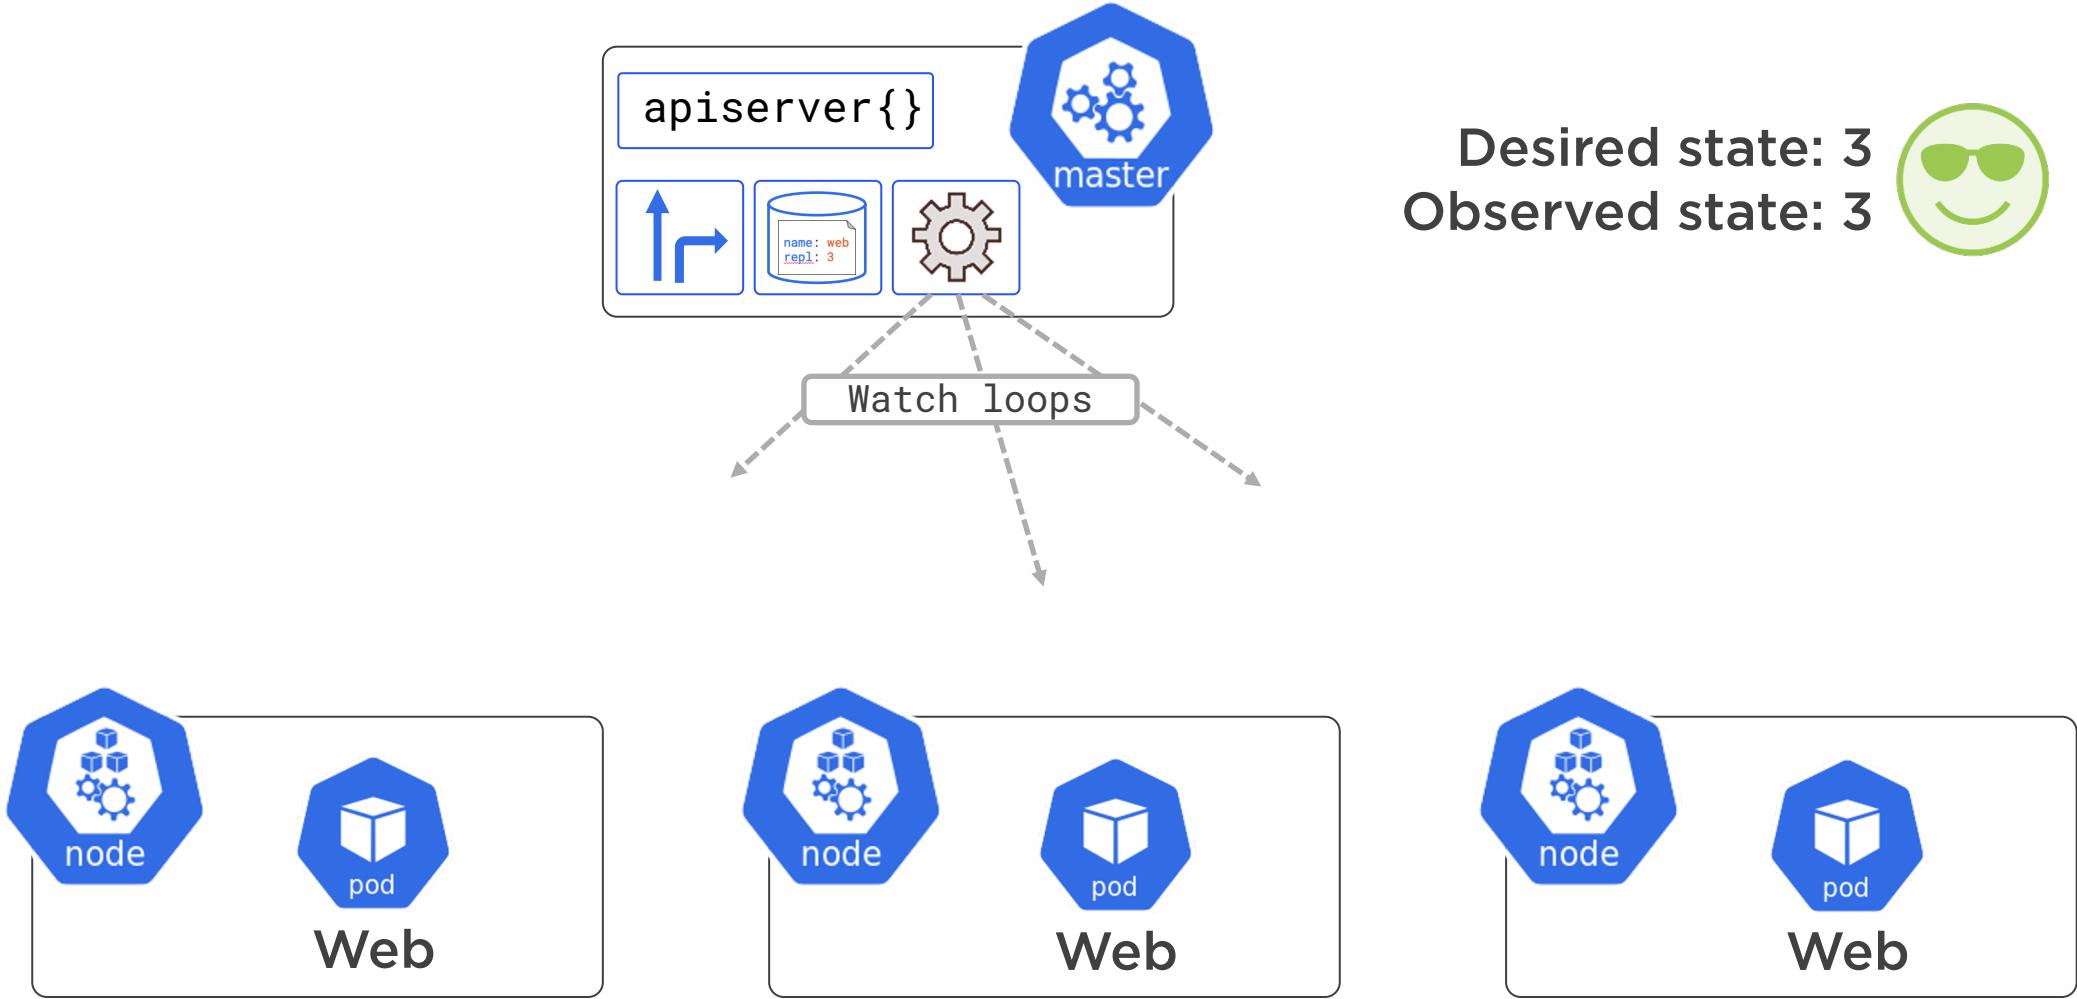






Desired state: 3
Observed state: 3



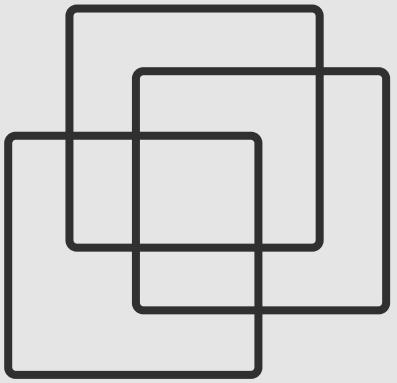


Up Next:
The Mighty Pod

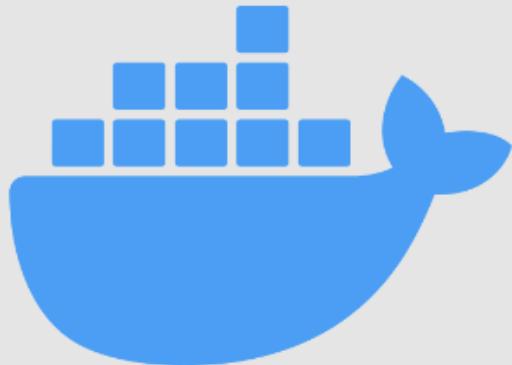


The Mighty Pod





Virtual Machine

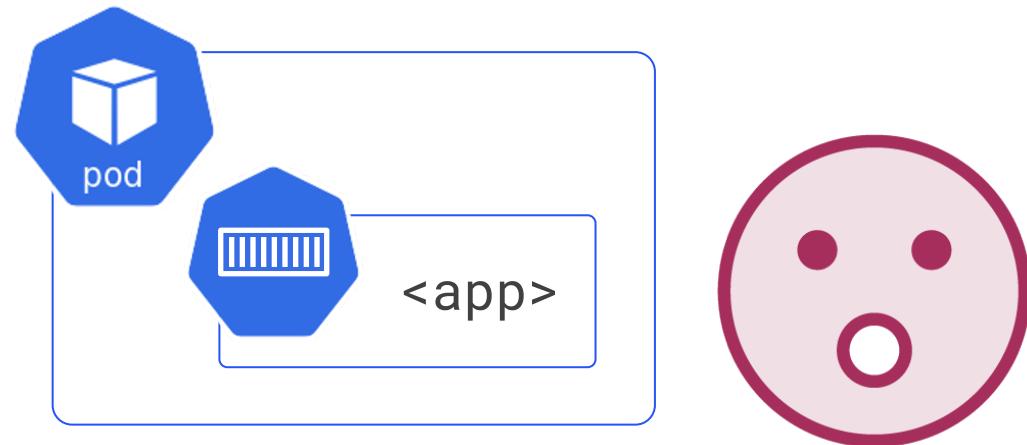


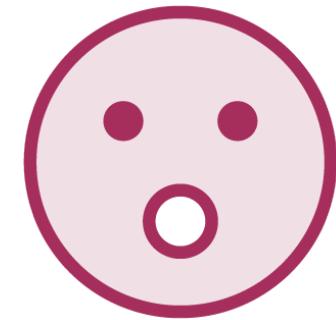
Container



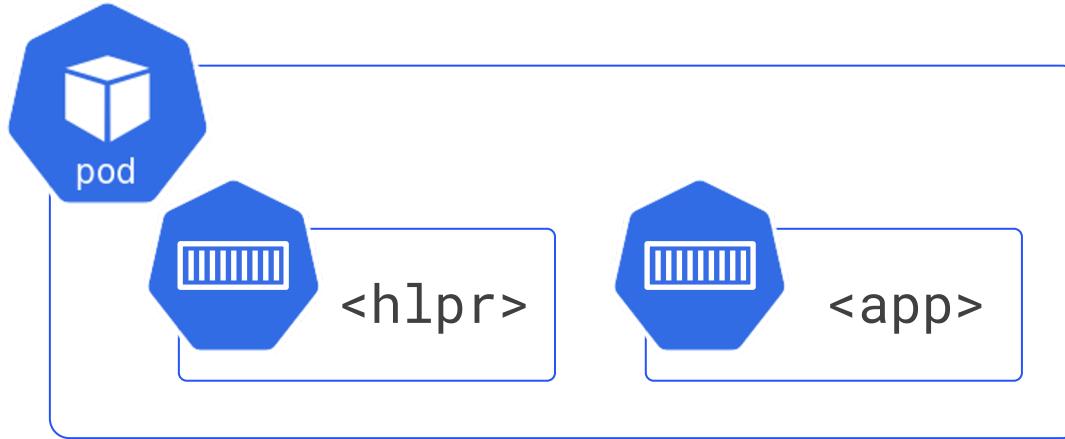
Pod





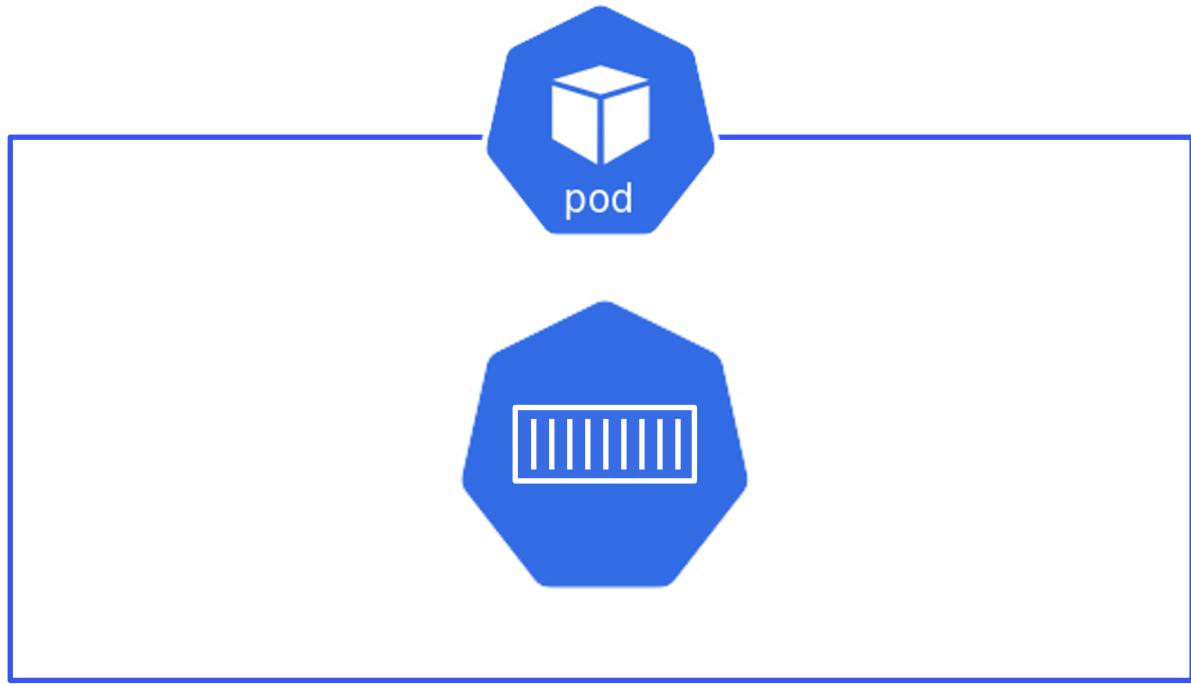


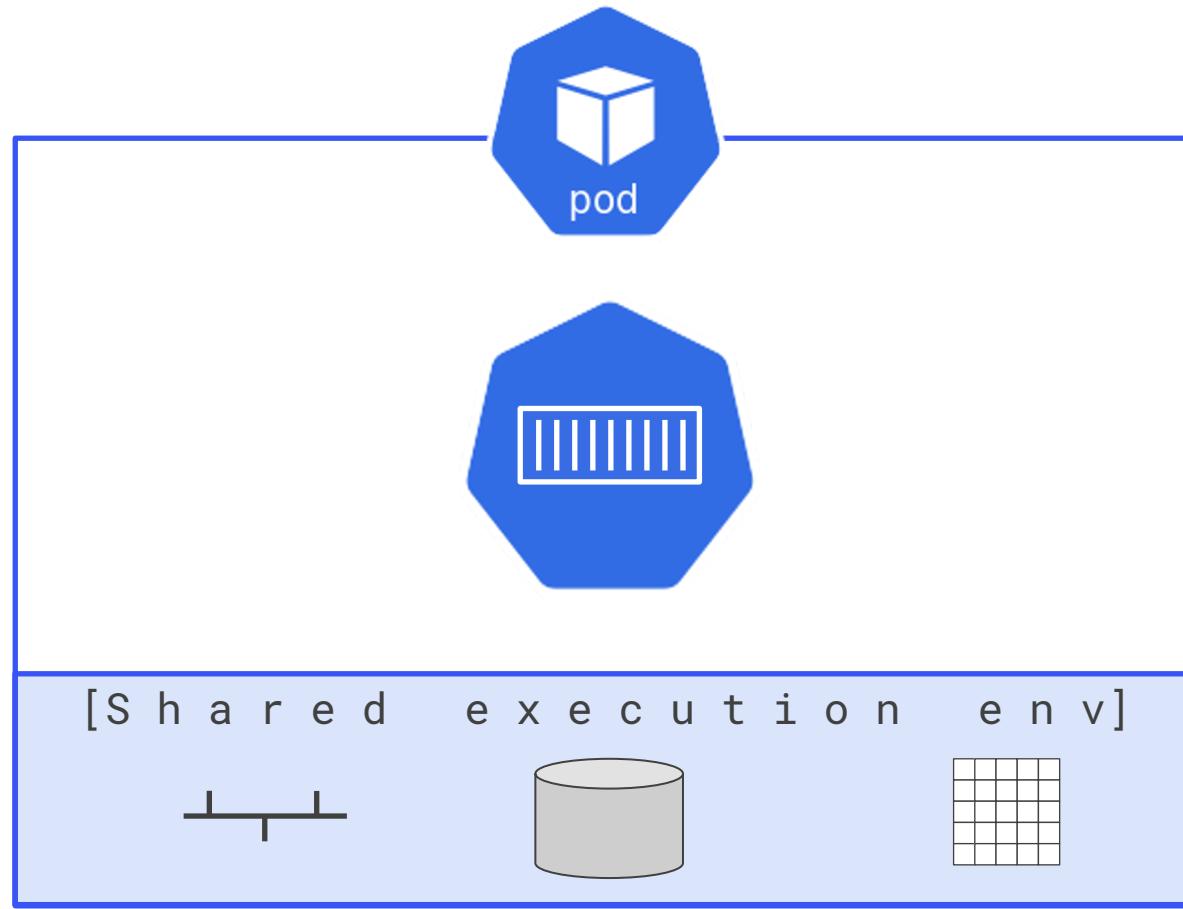


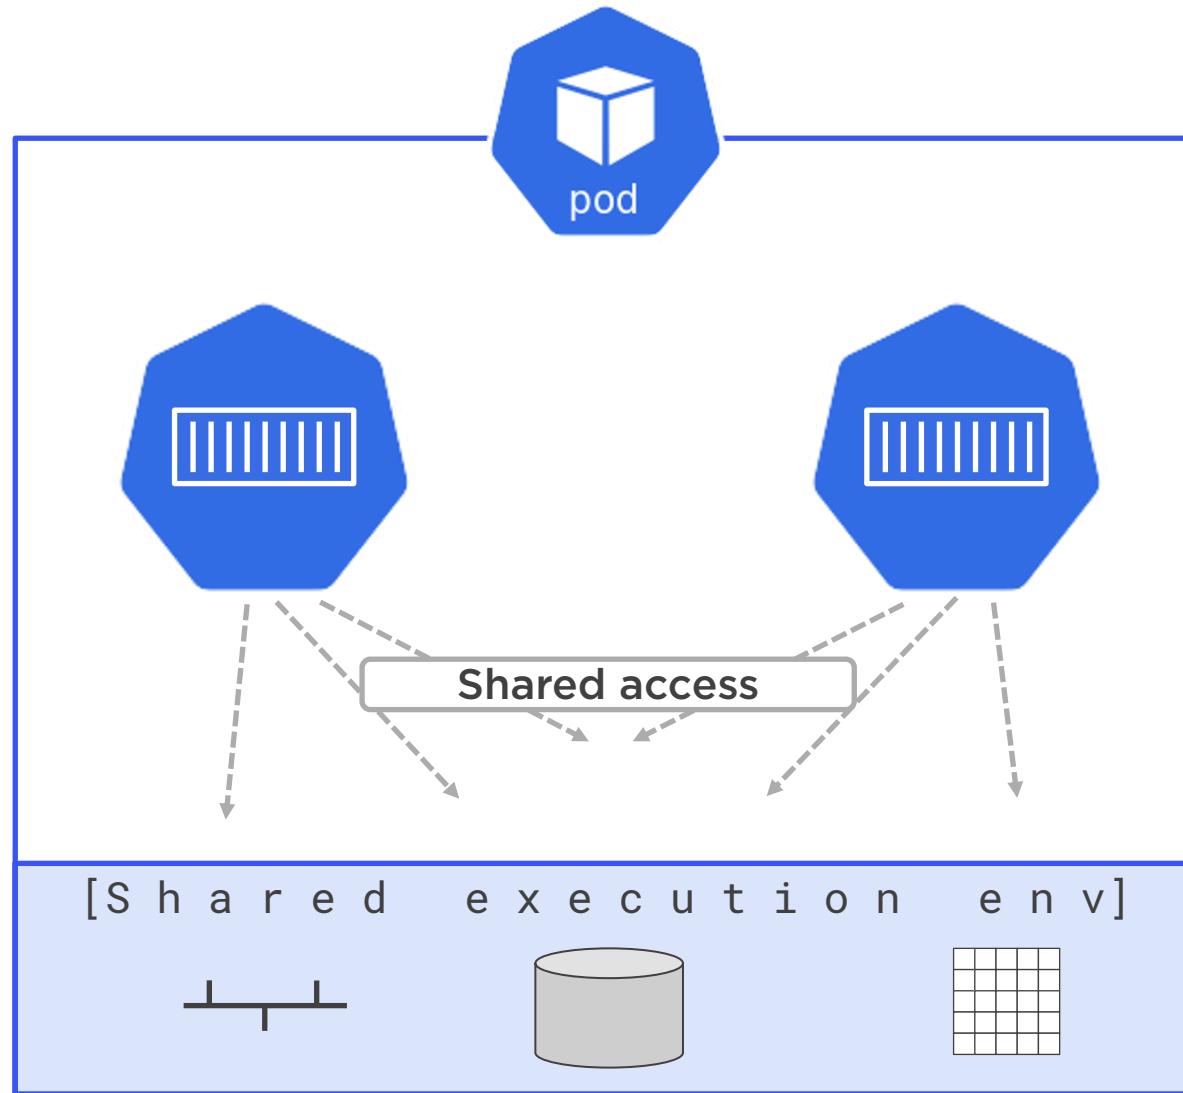


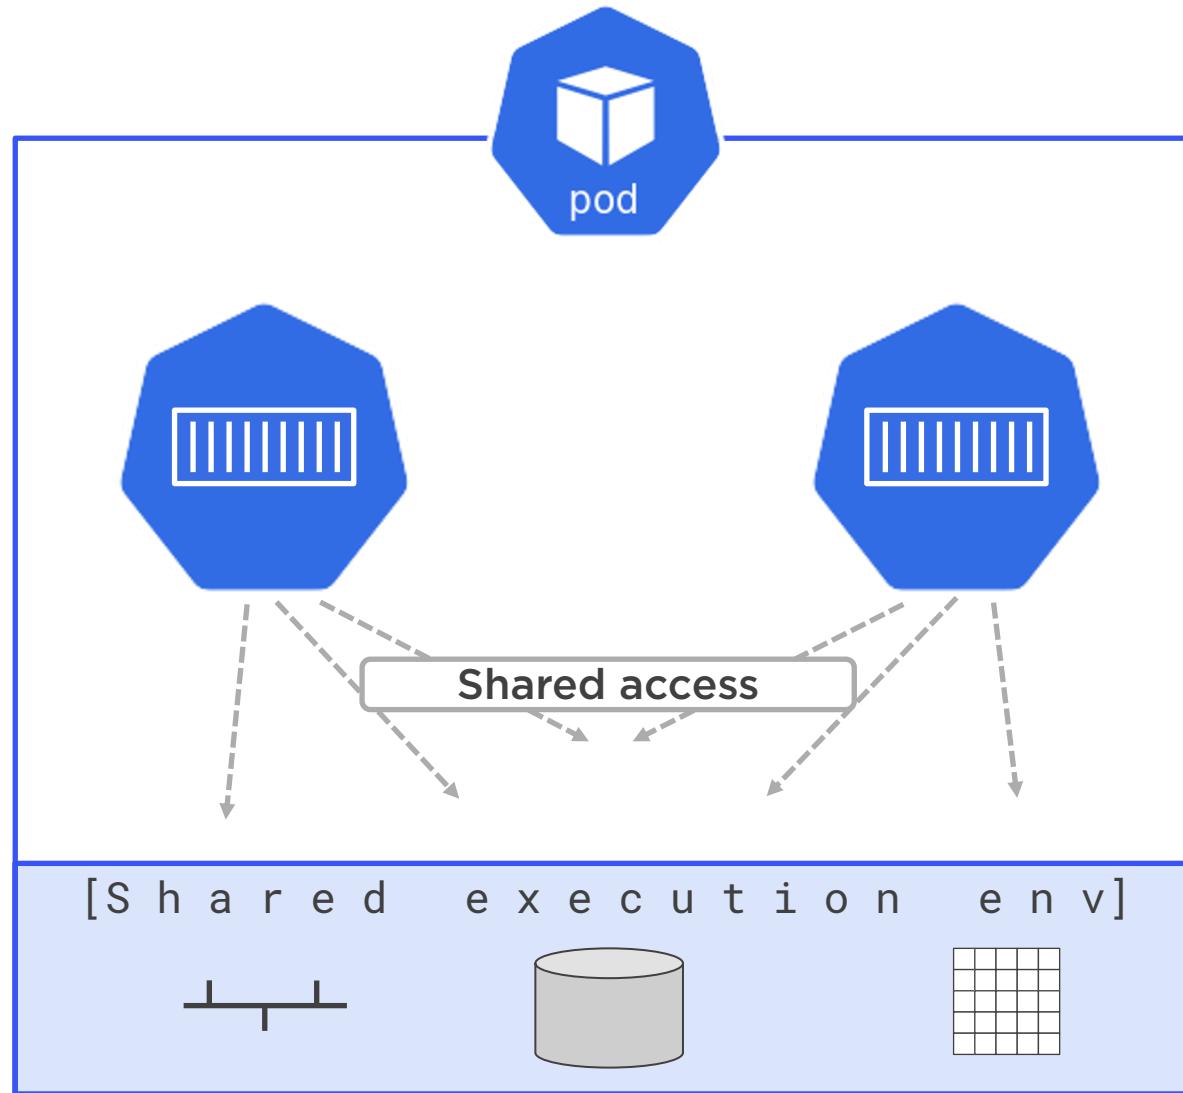
Kubernetes for Developers: Integrating Volumes and Using Multi-container Pods

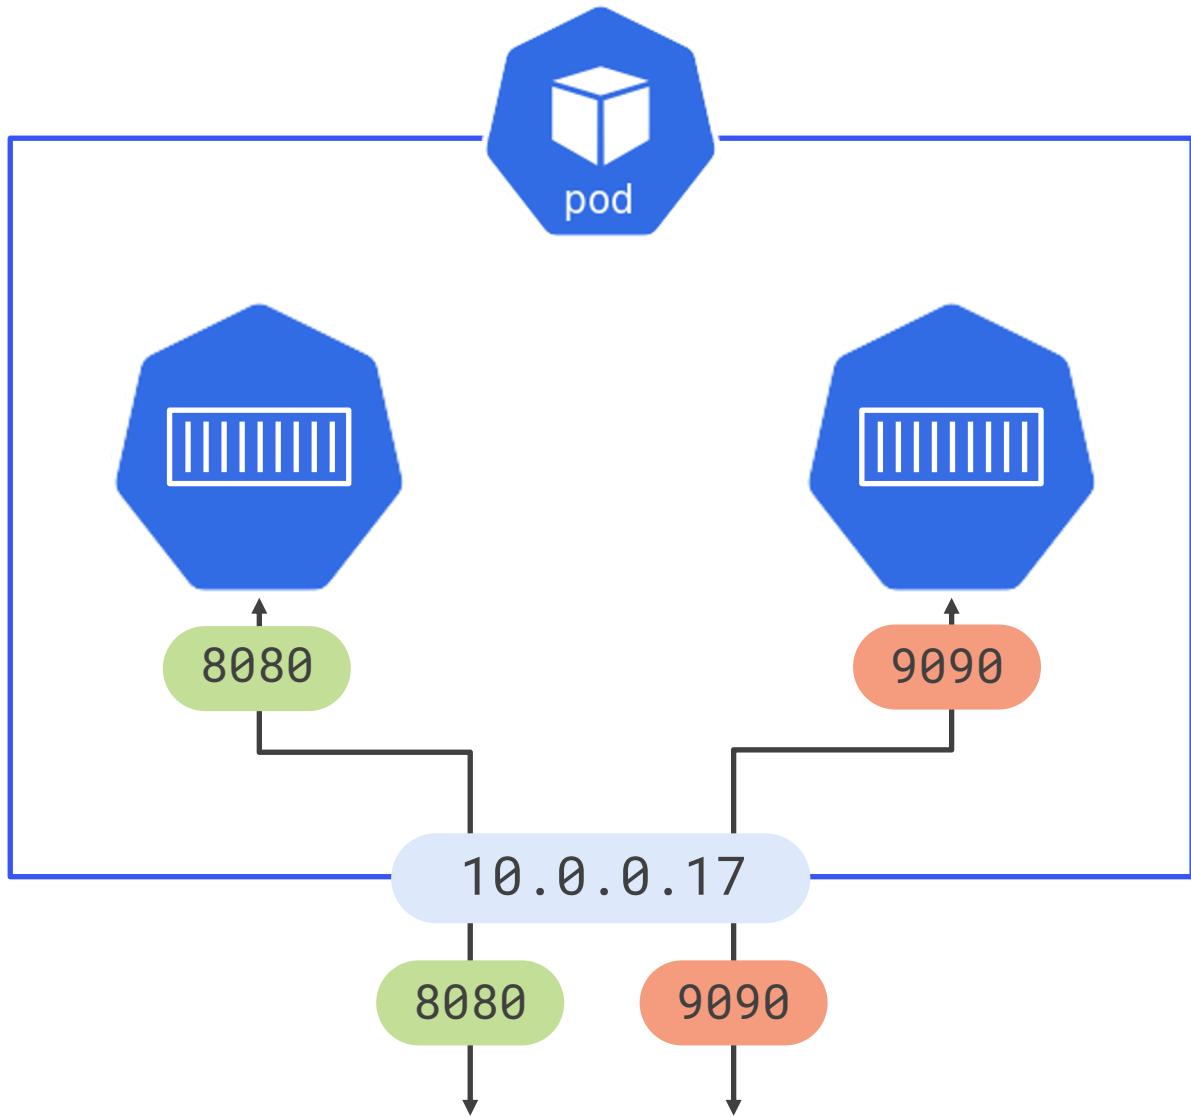


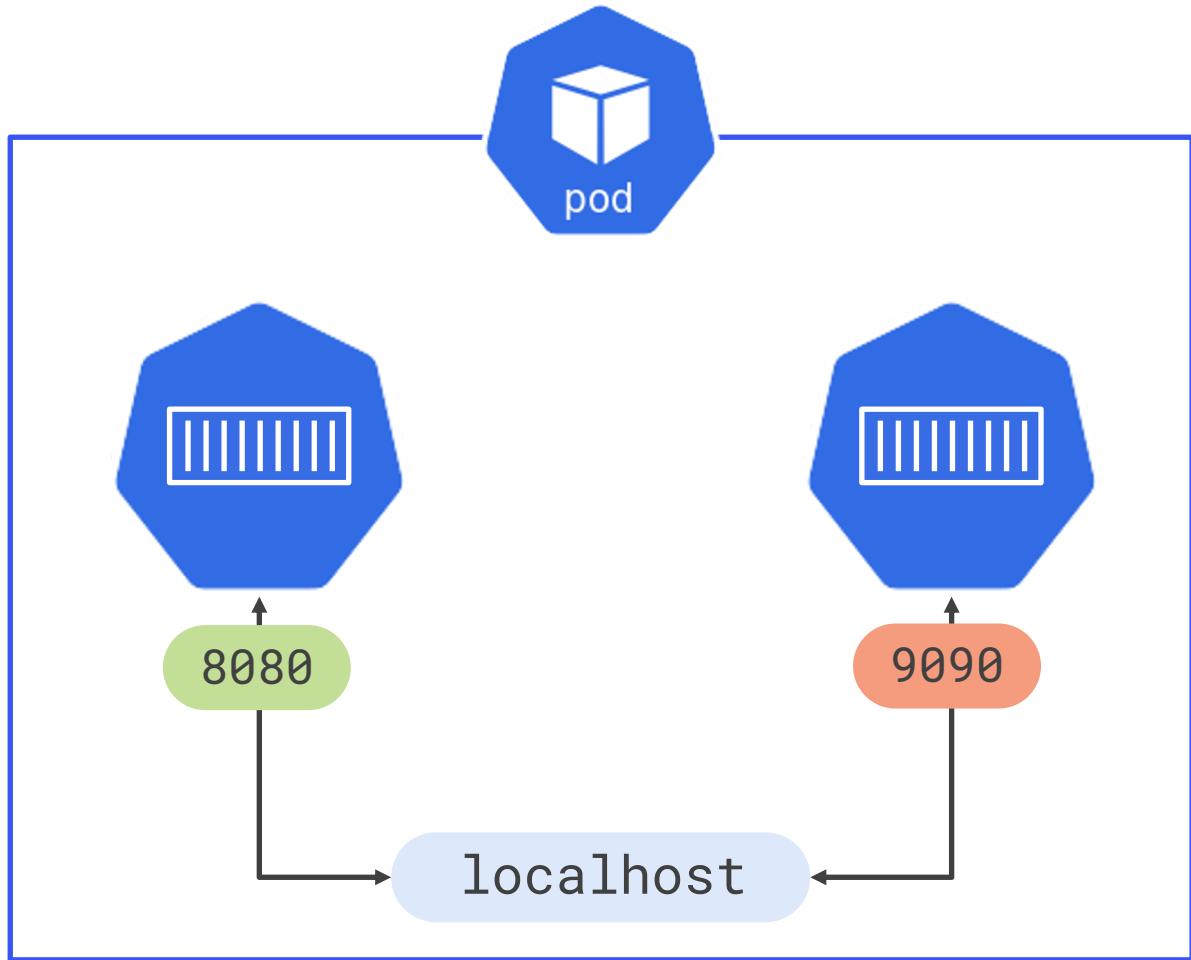


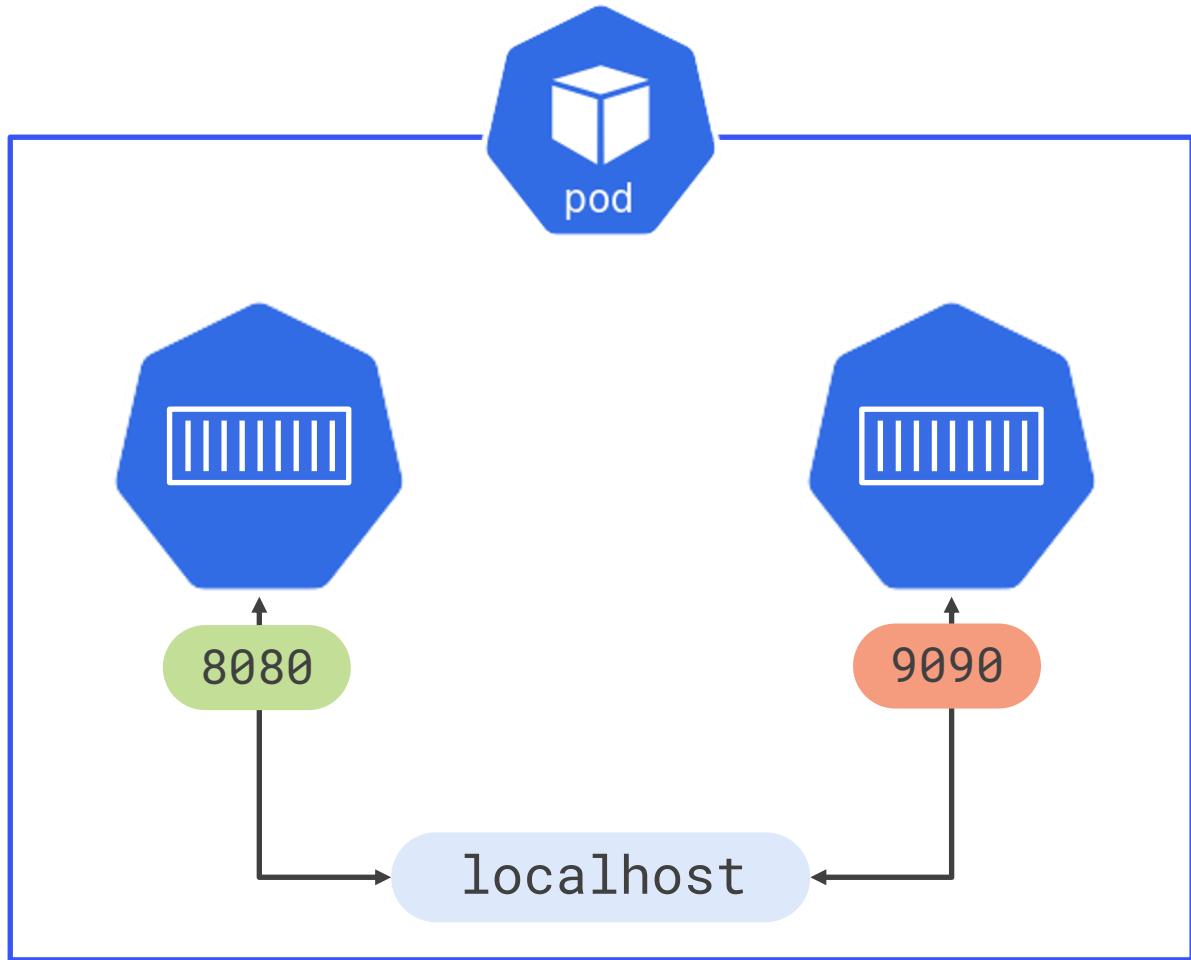


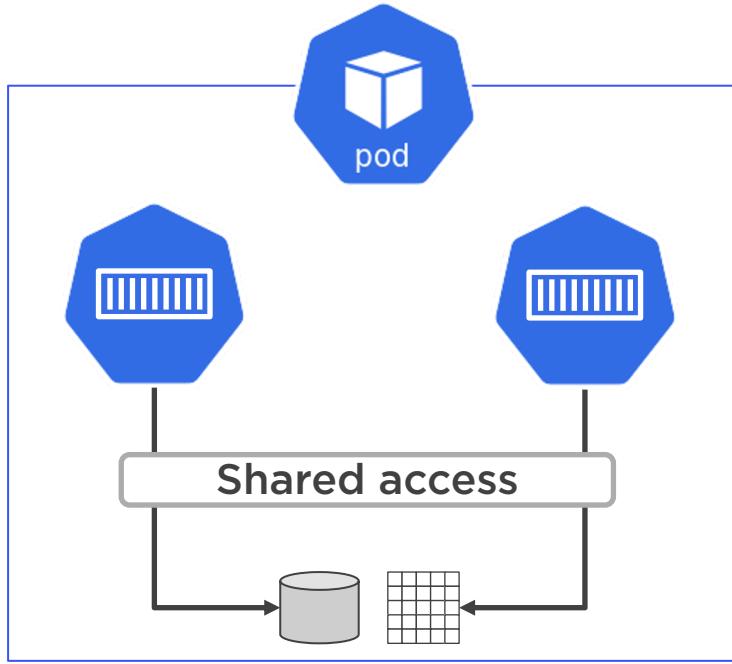






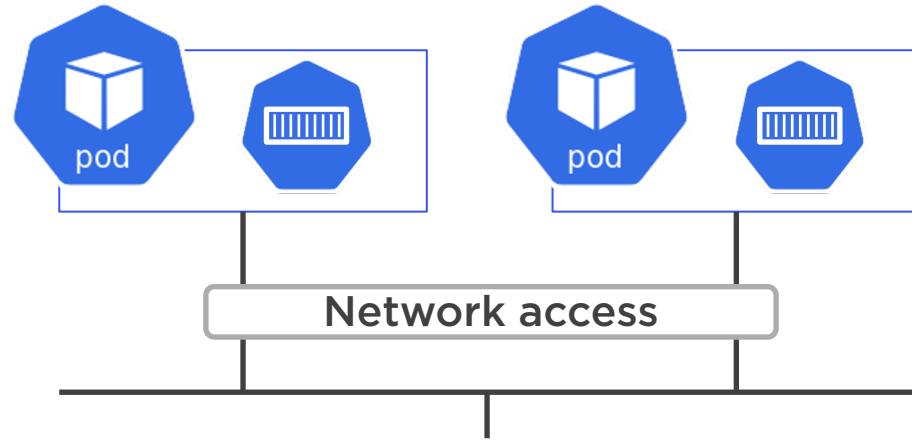






Tightly coupled

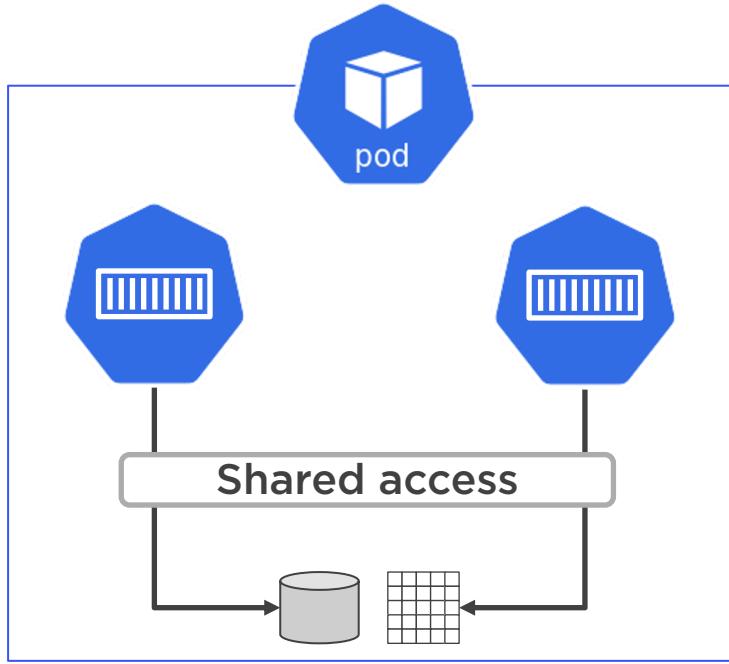
When two containers (app services) **absolutely need** to share vols, memory etc.



Loosely coupled

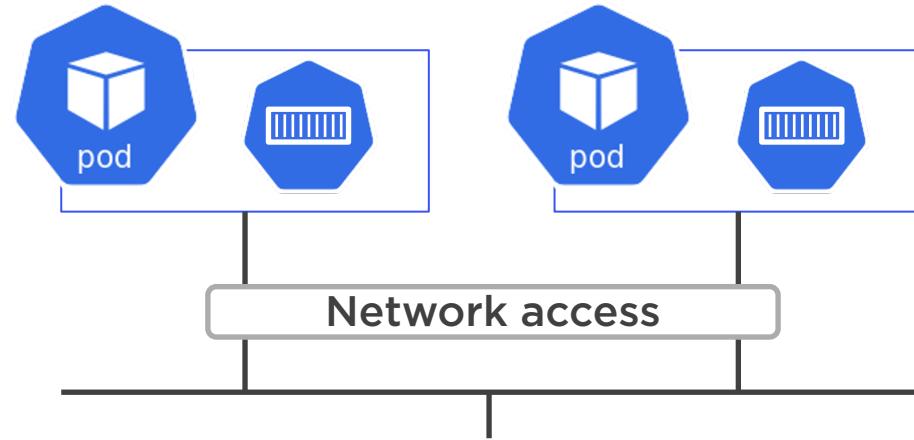
When two containers (app services) **don't absolutely need** to share resources





Tightly coupled

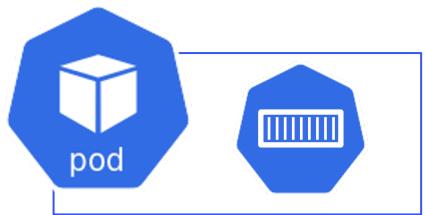
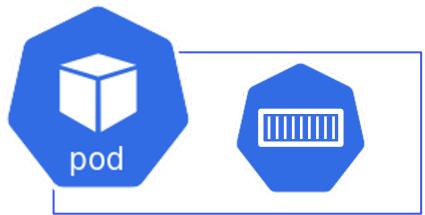
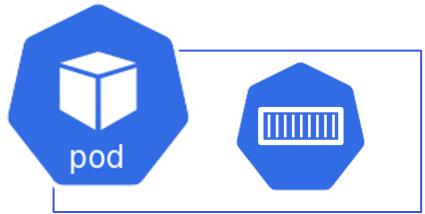
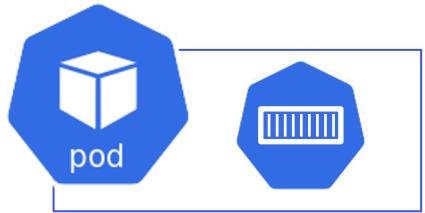
When two containers (app services) absolutely need to share vols, memory etc.



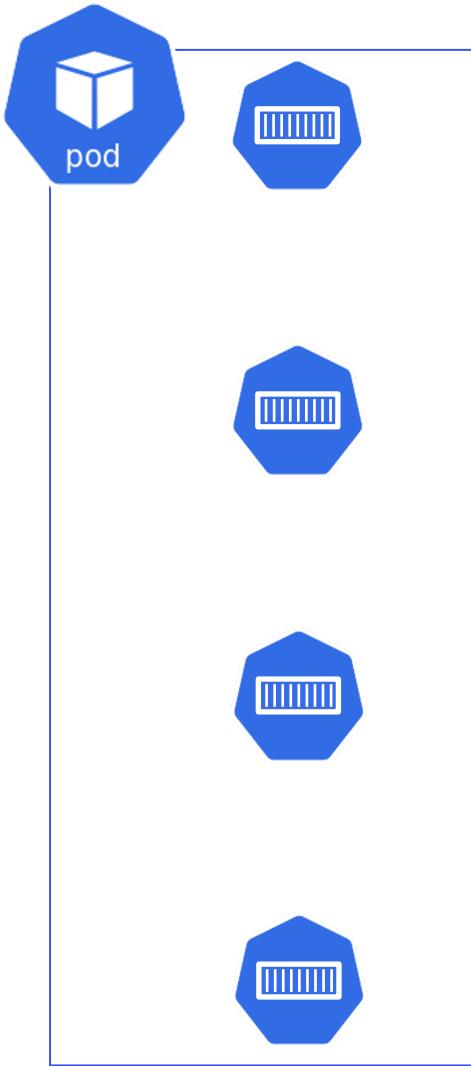
Loosely coupled

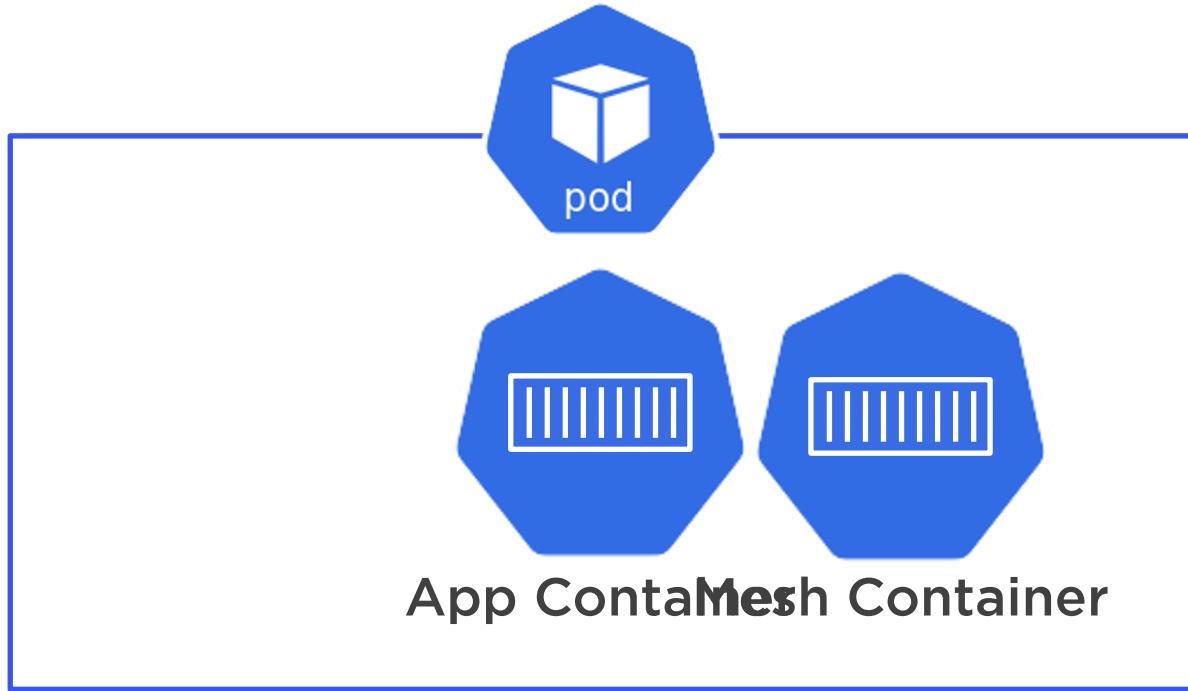
When two containers (app services) don't absolutely need to share resources

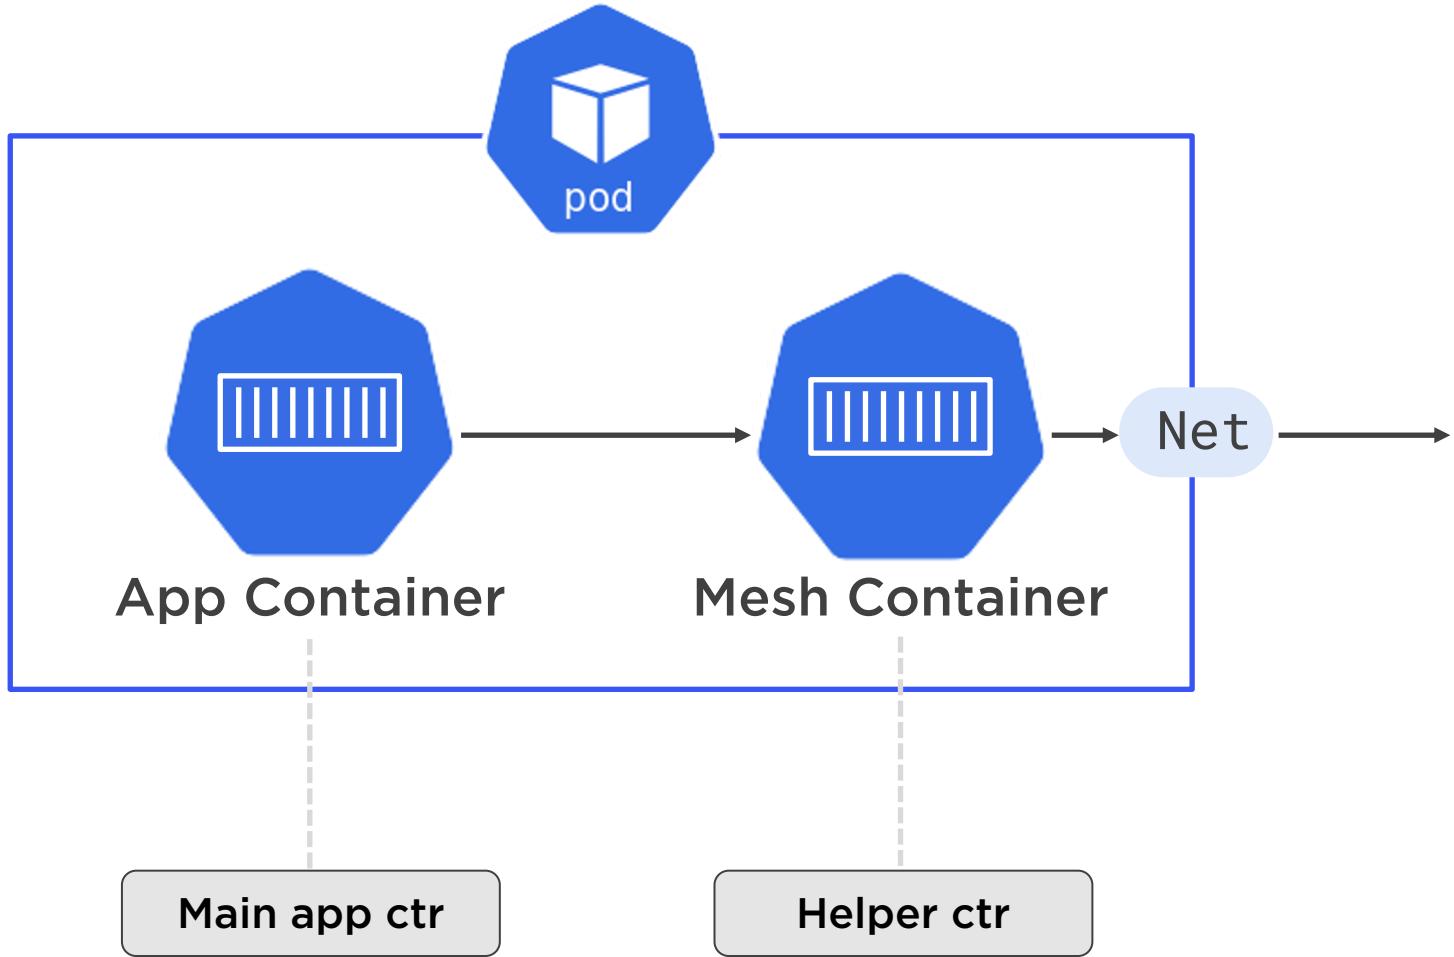


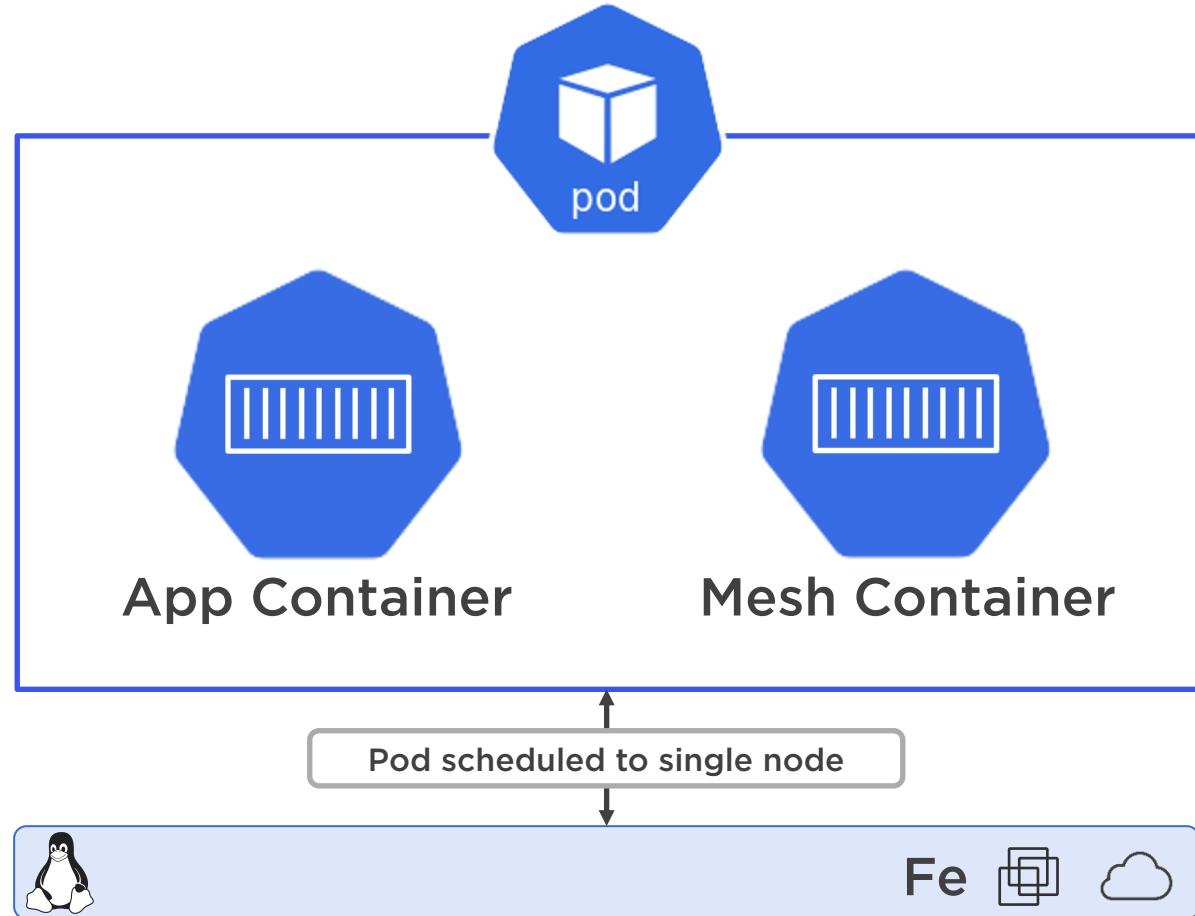


Scaling

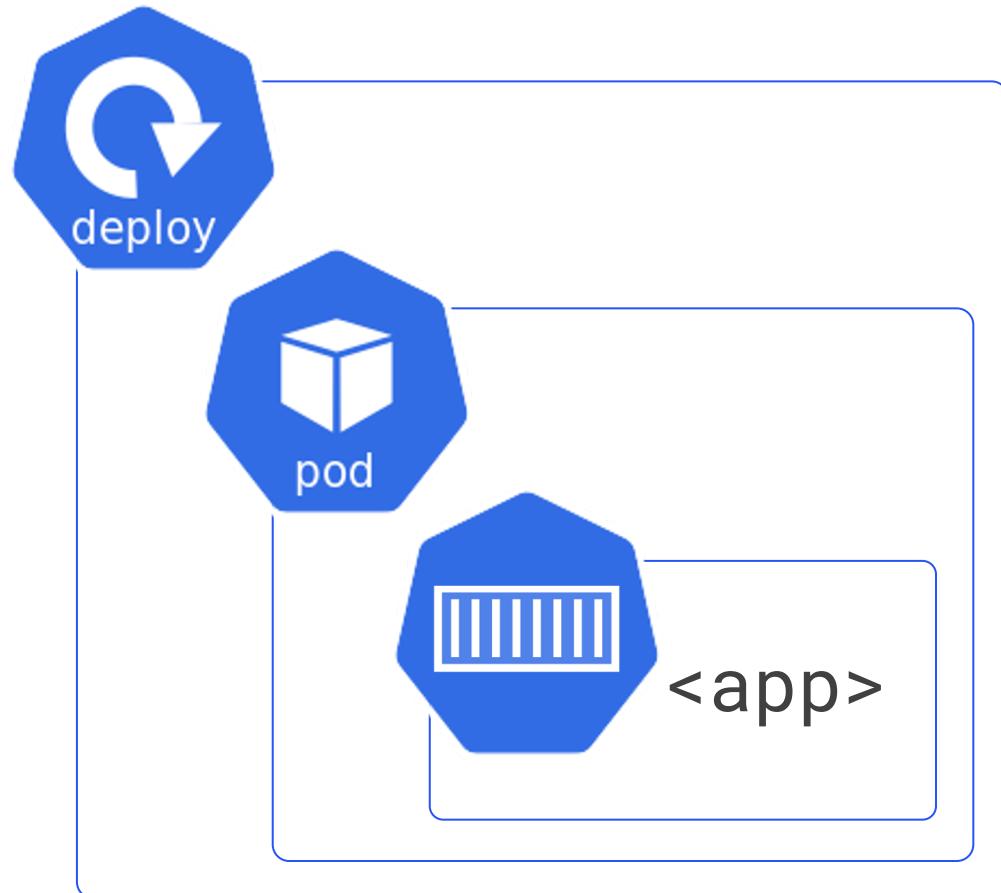


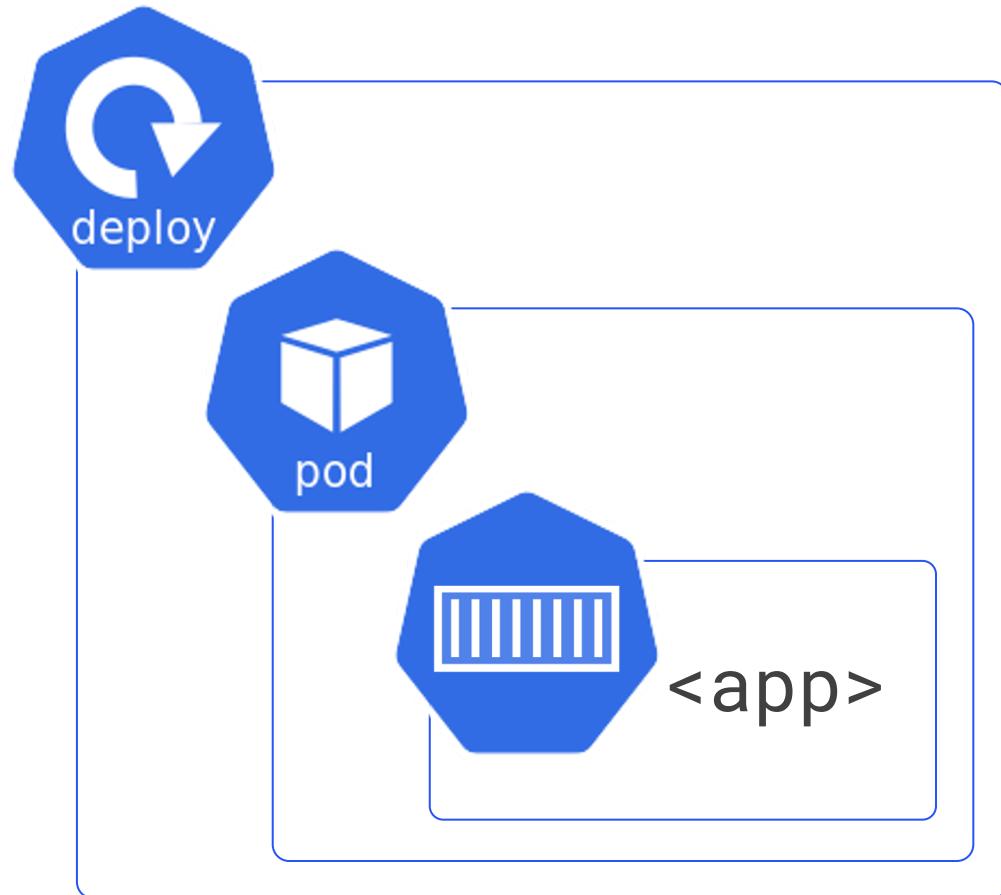














Annotations

Labels

Policies

Resources

Co-scheduling containers

...

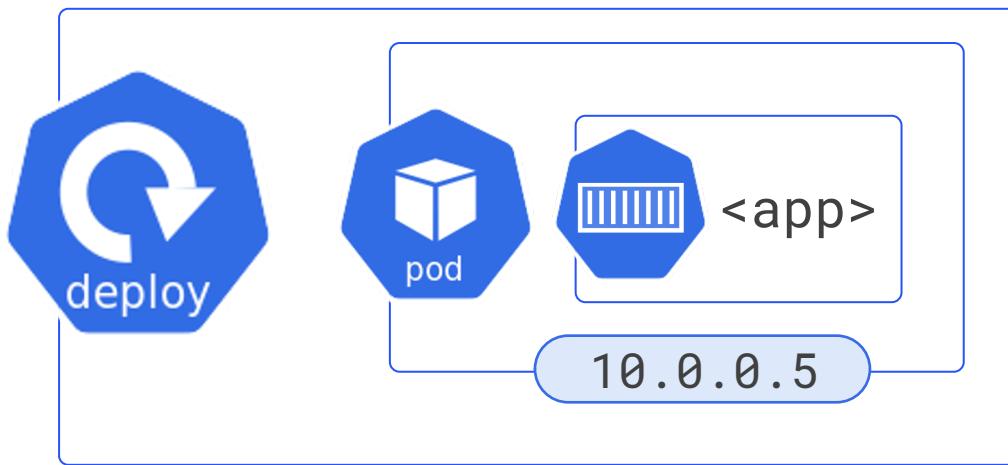


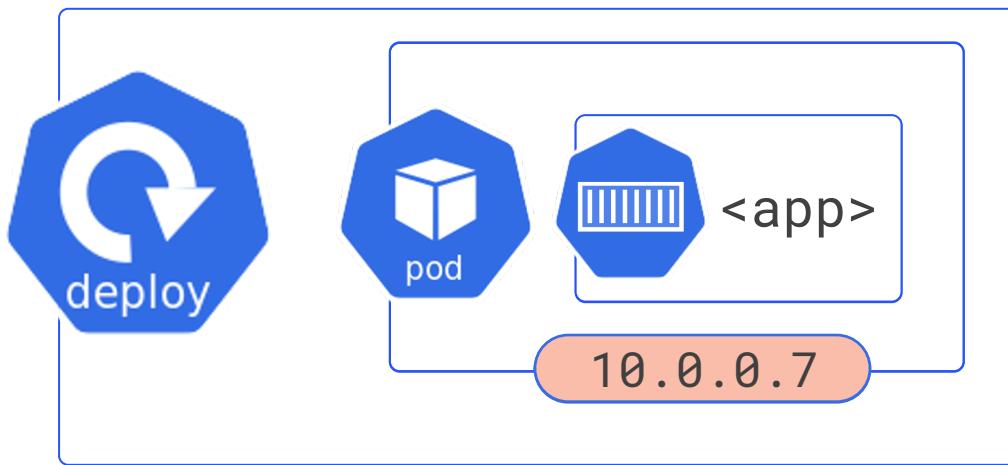
Up Next:
Stable Networking with Services

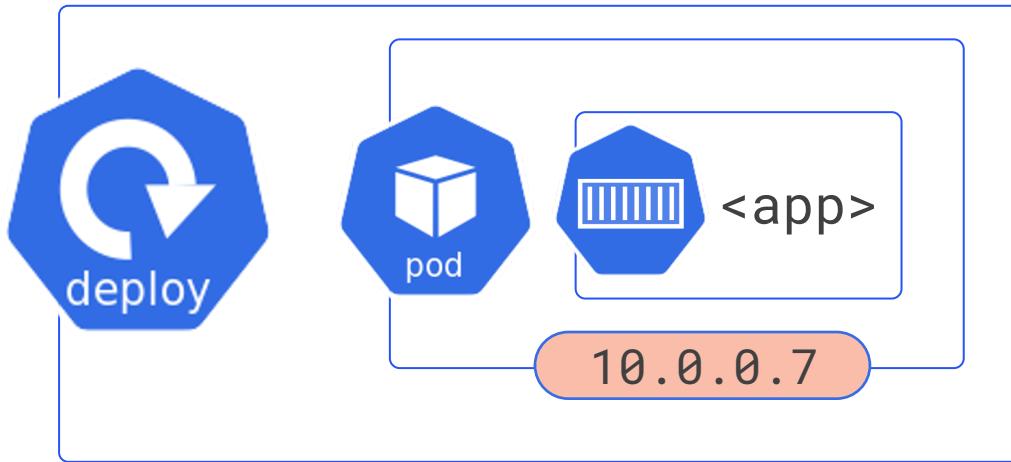


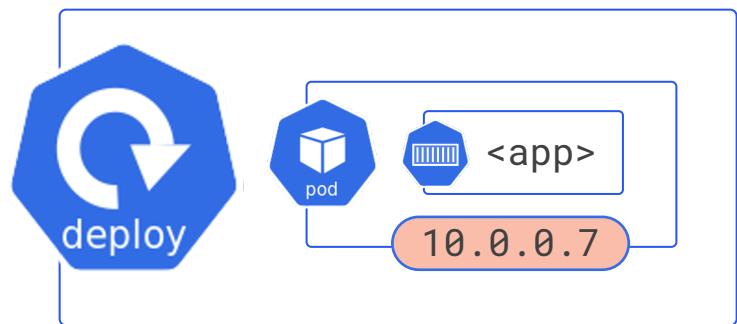
Stable Networking with Services

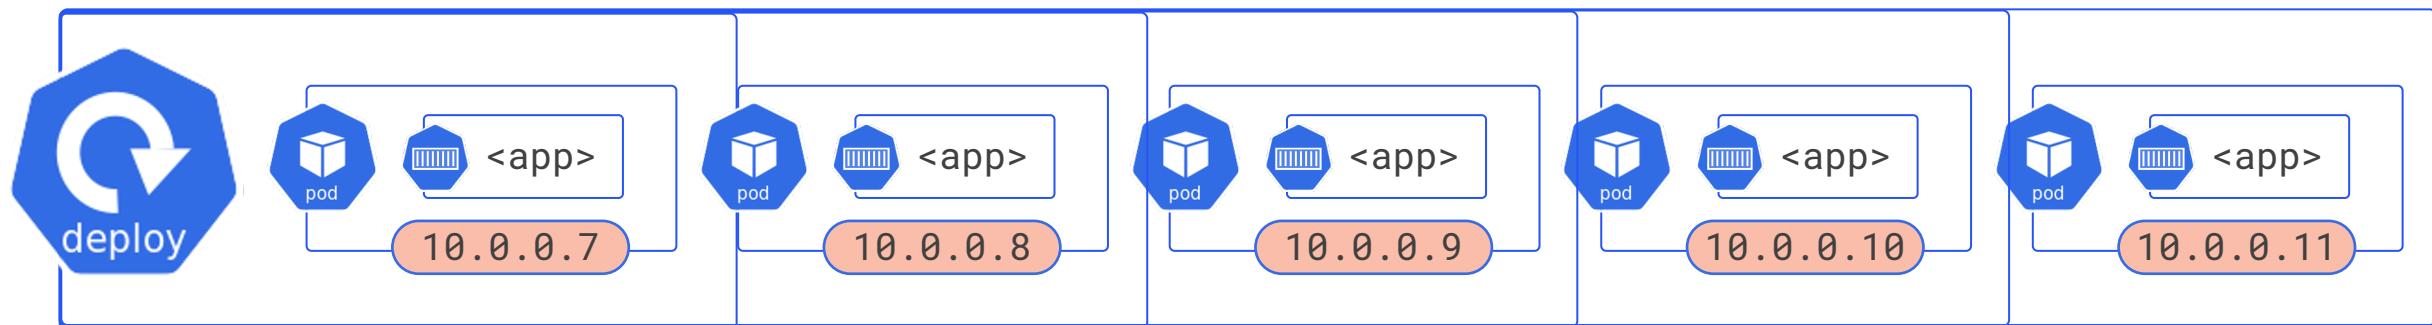


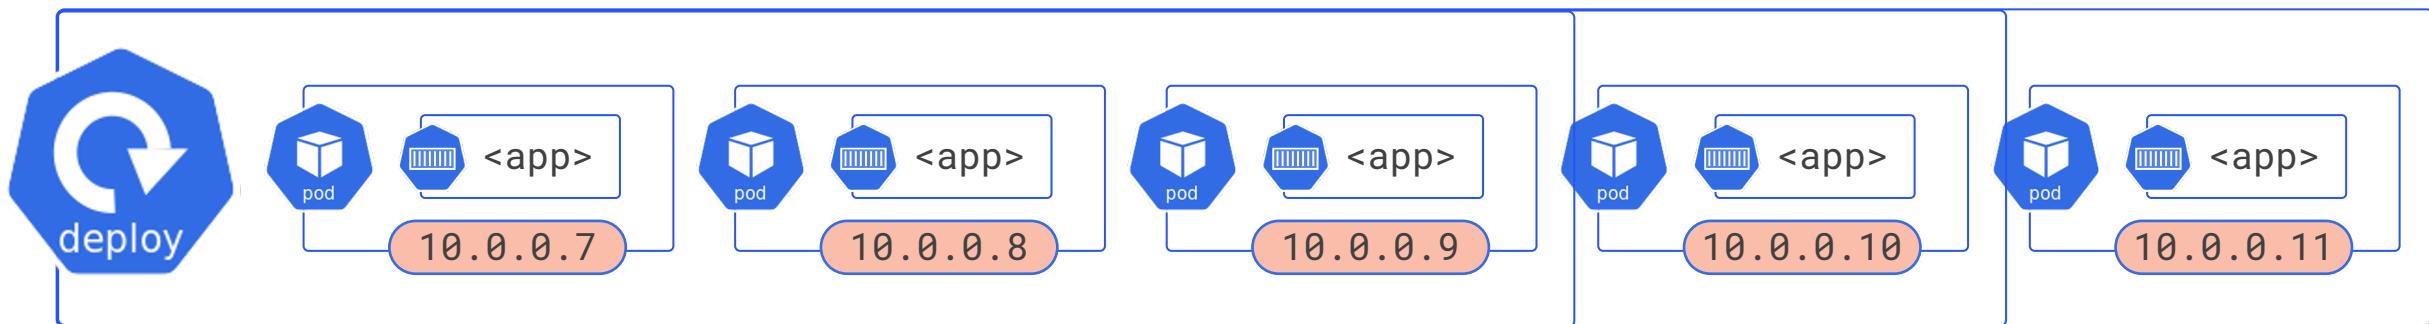


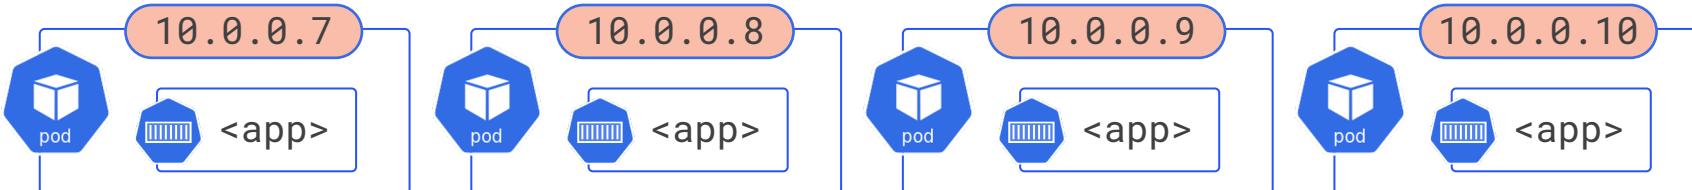
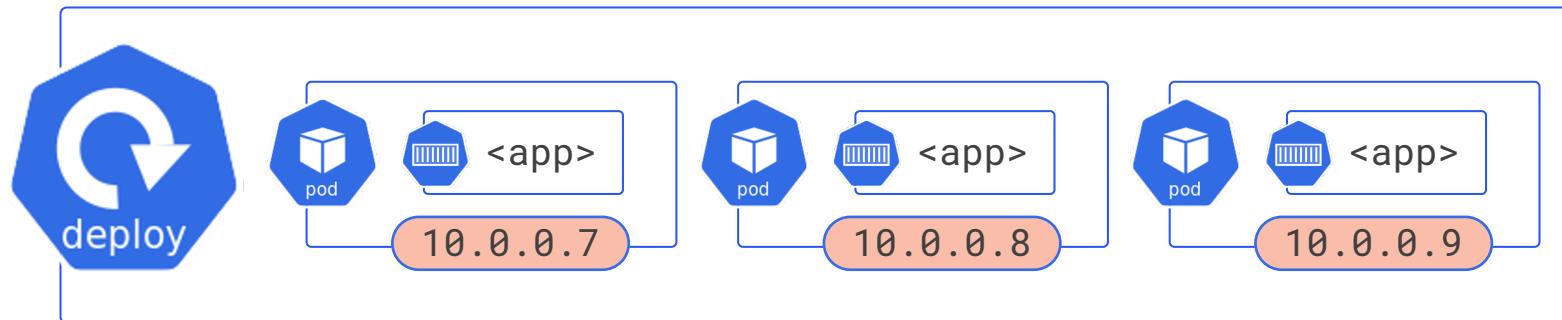


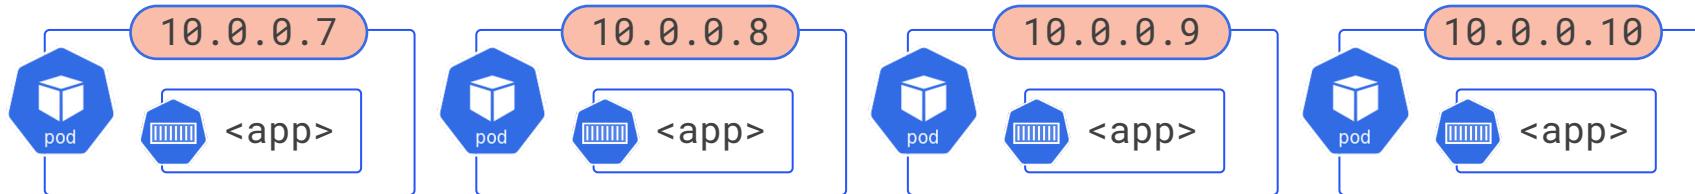


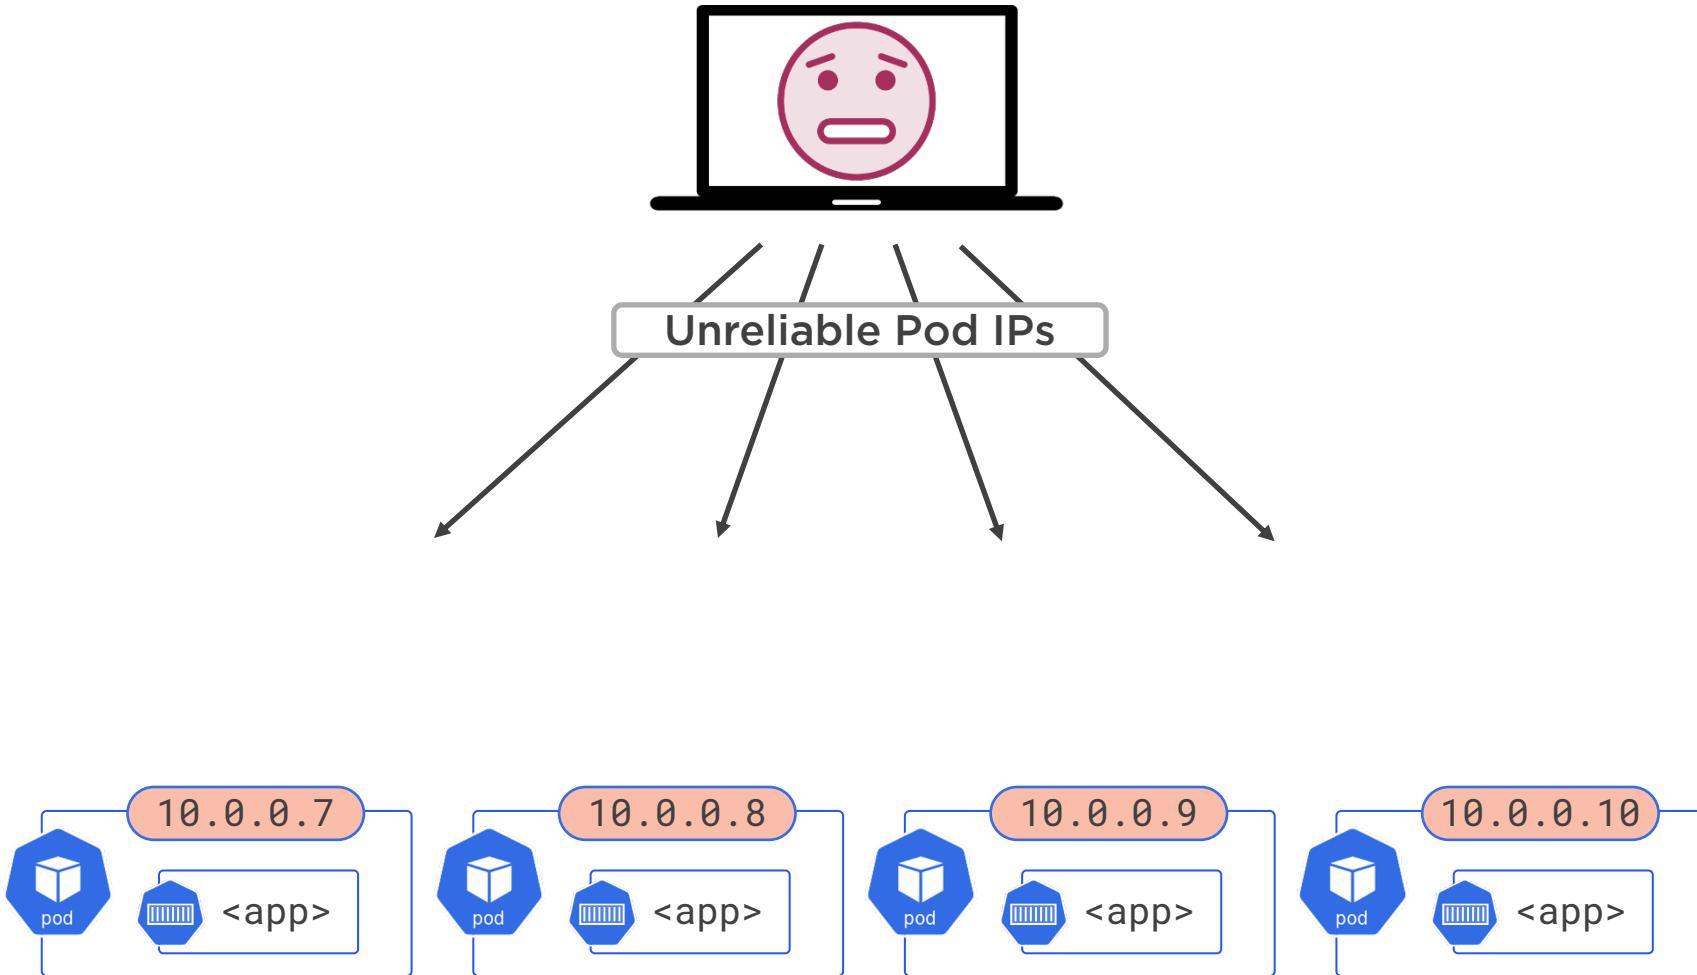


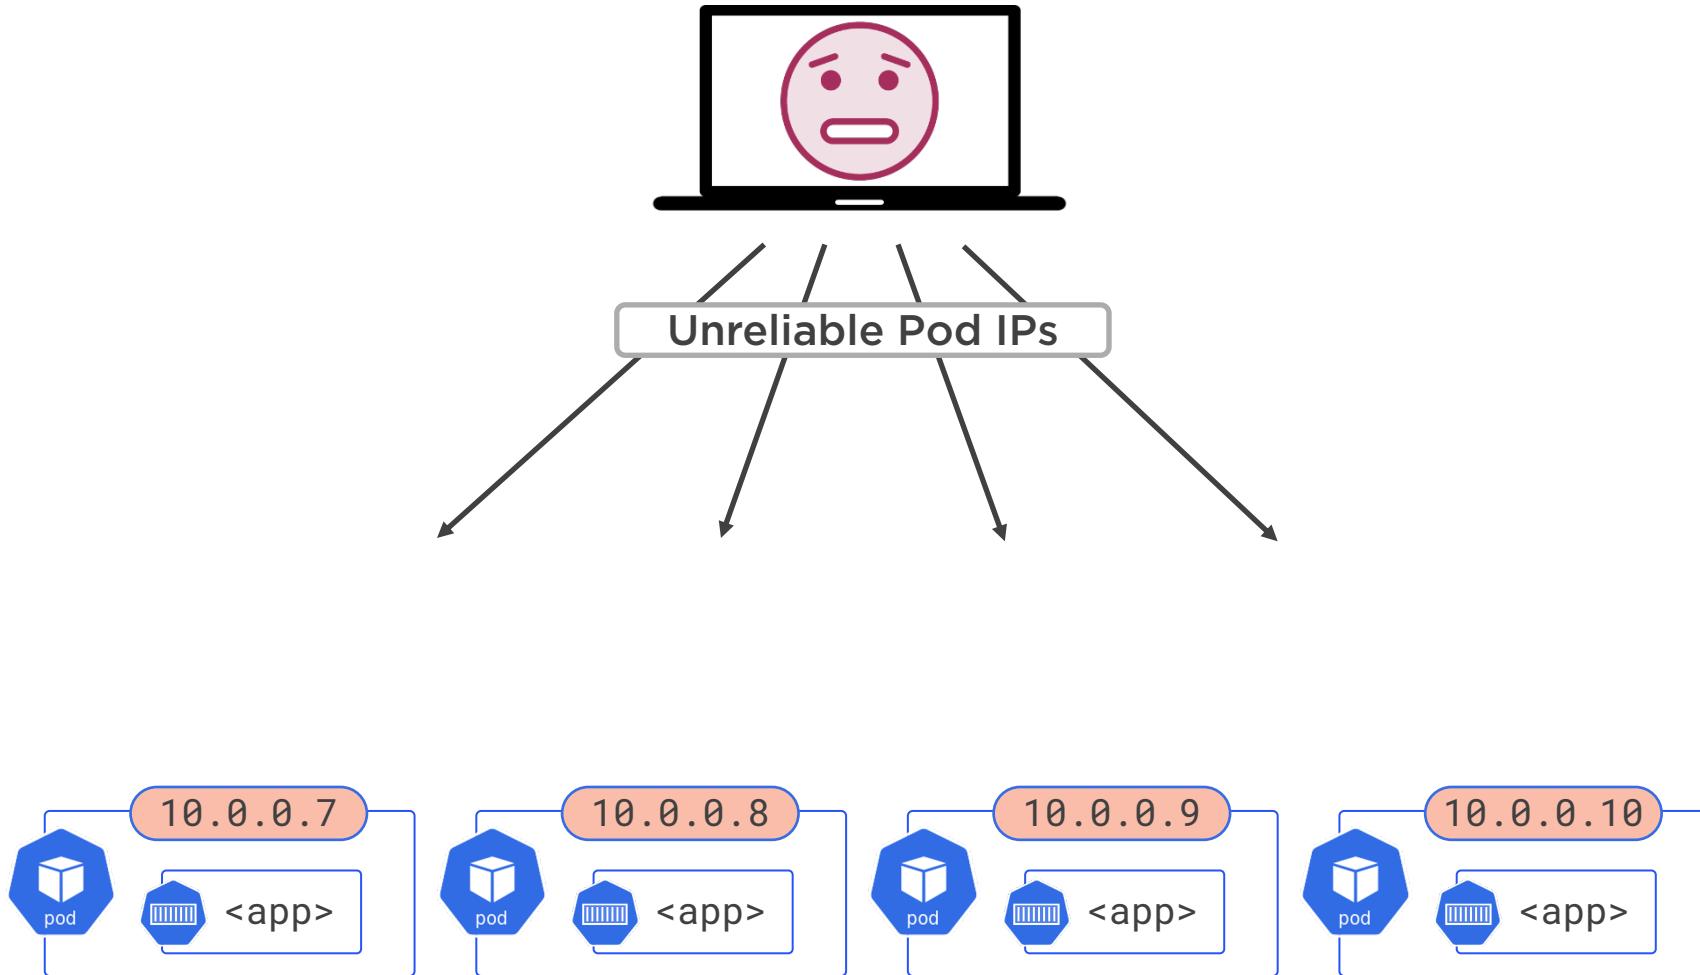


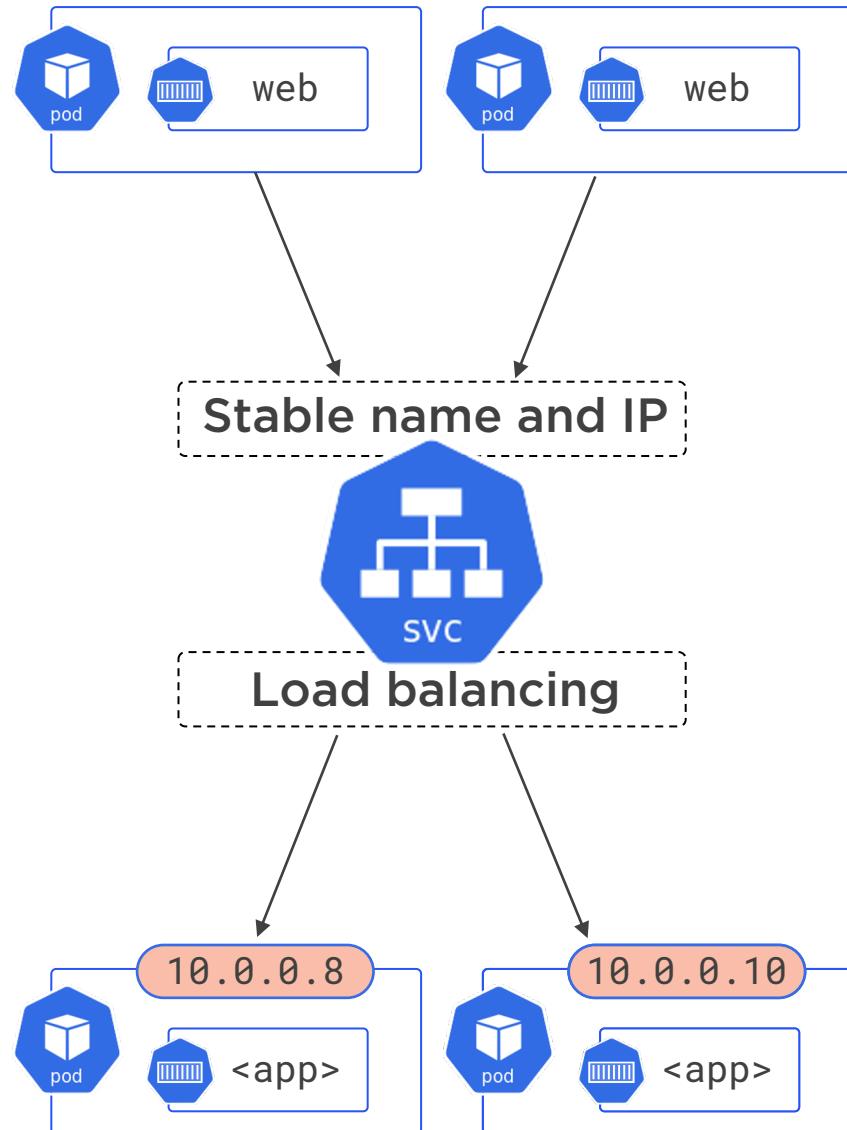


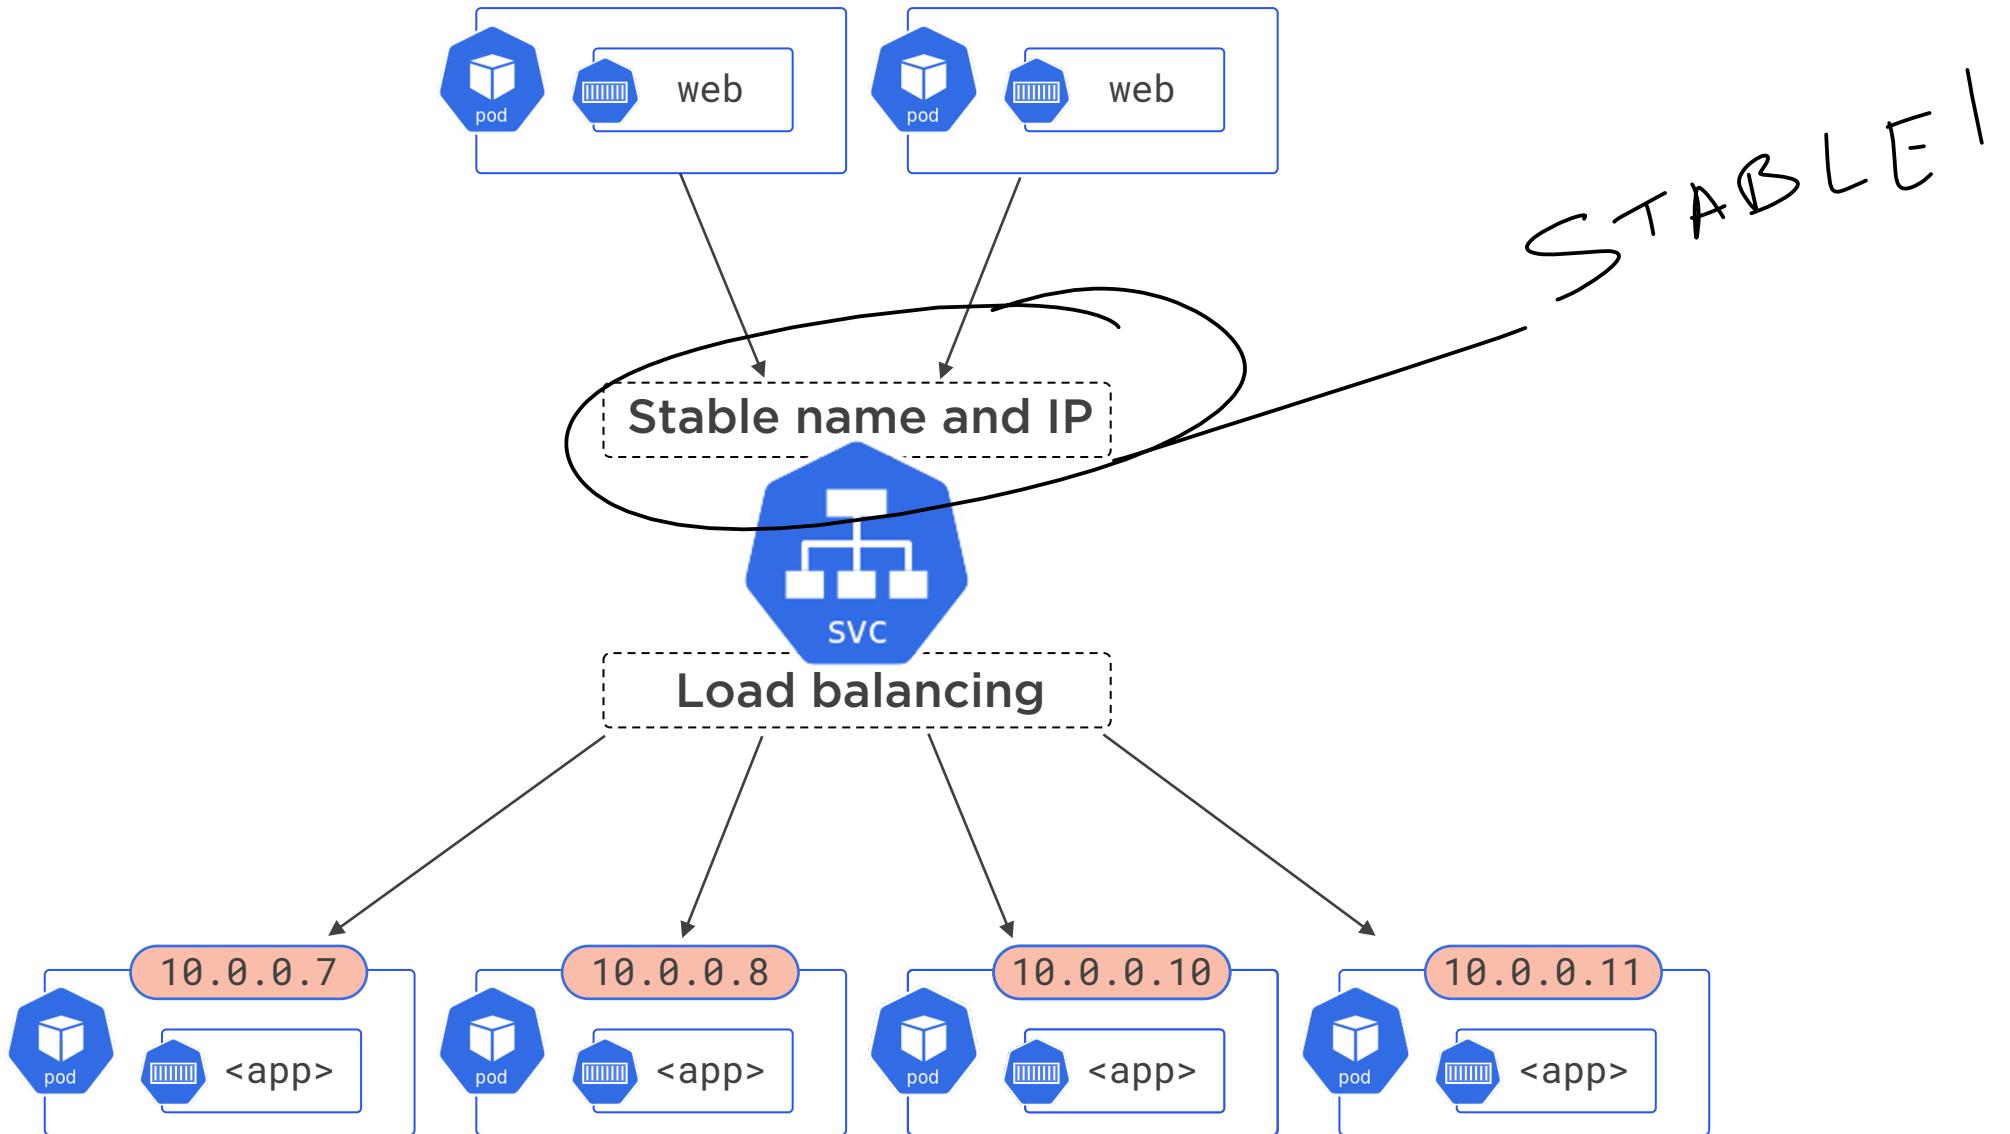


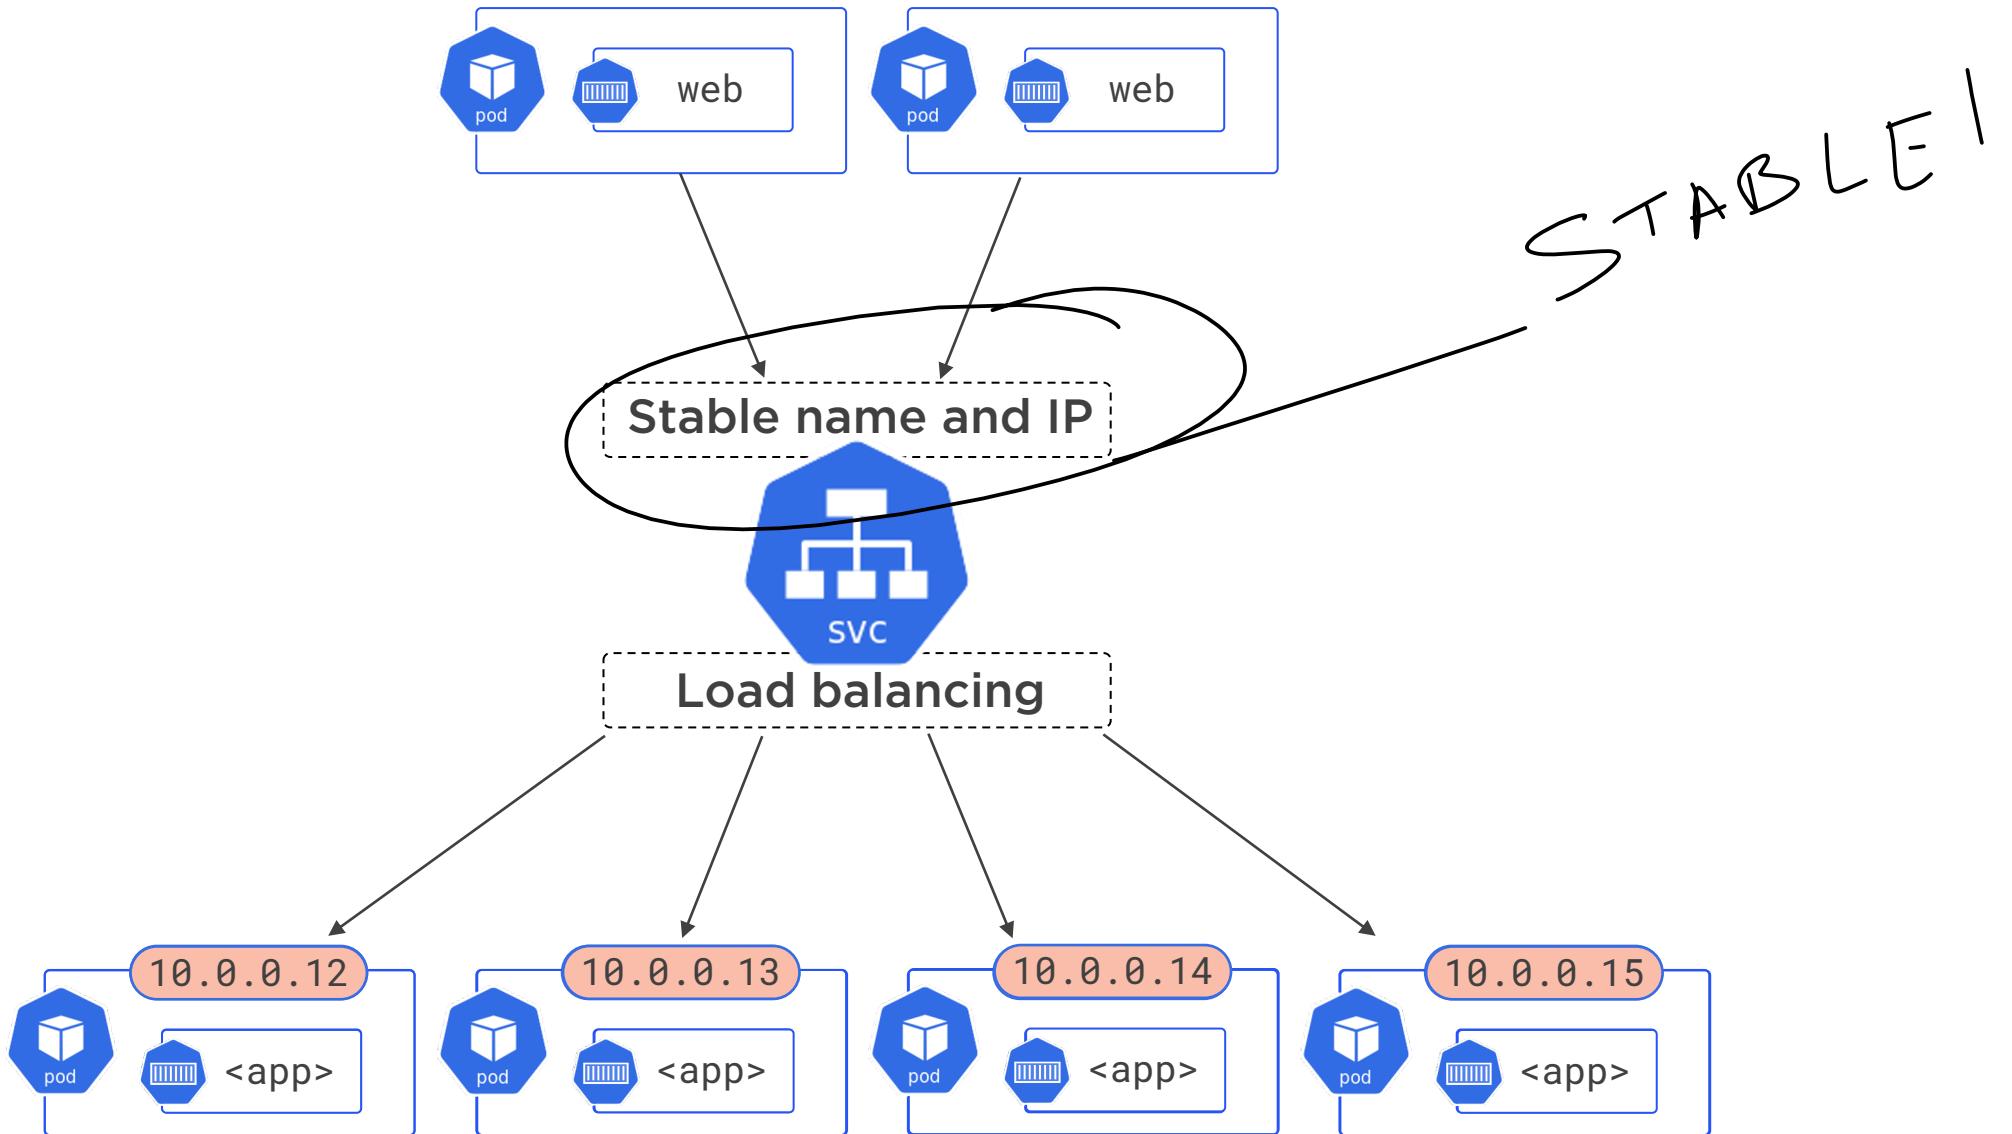










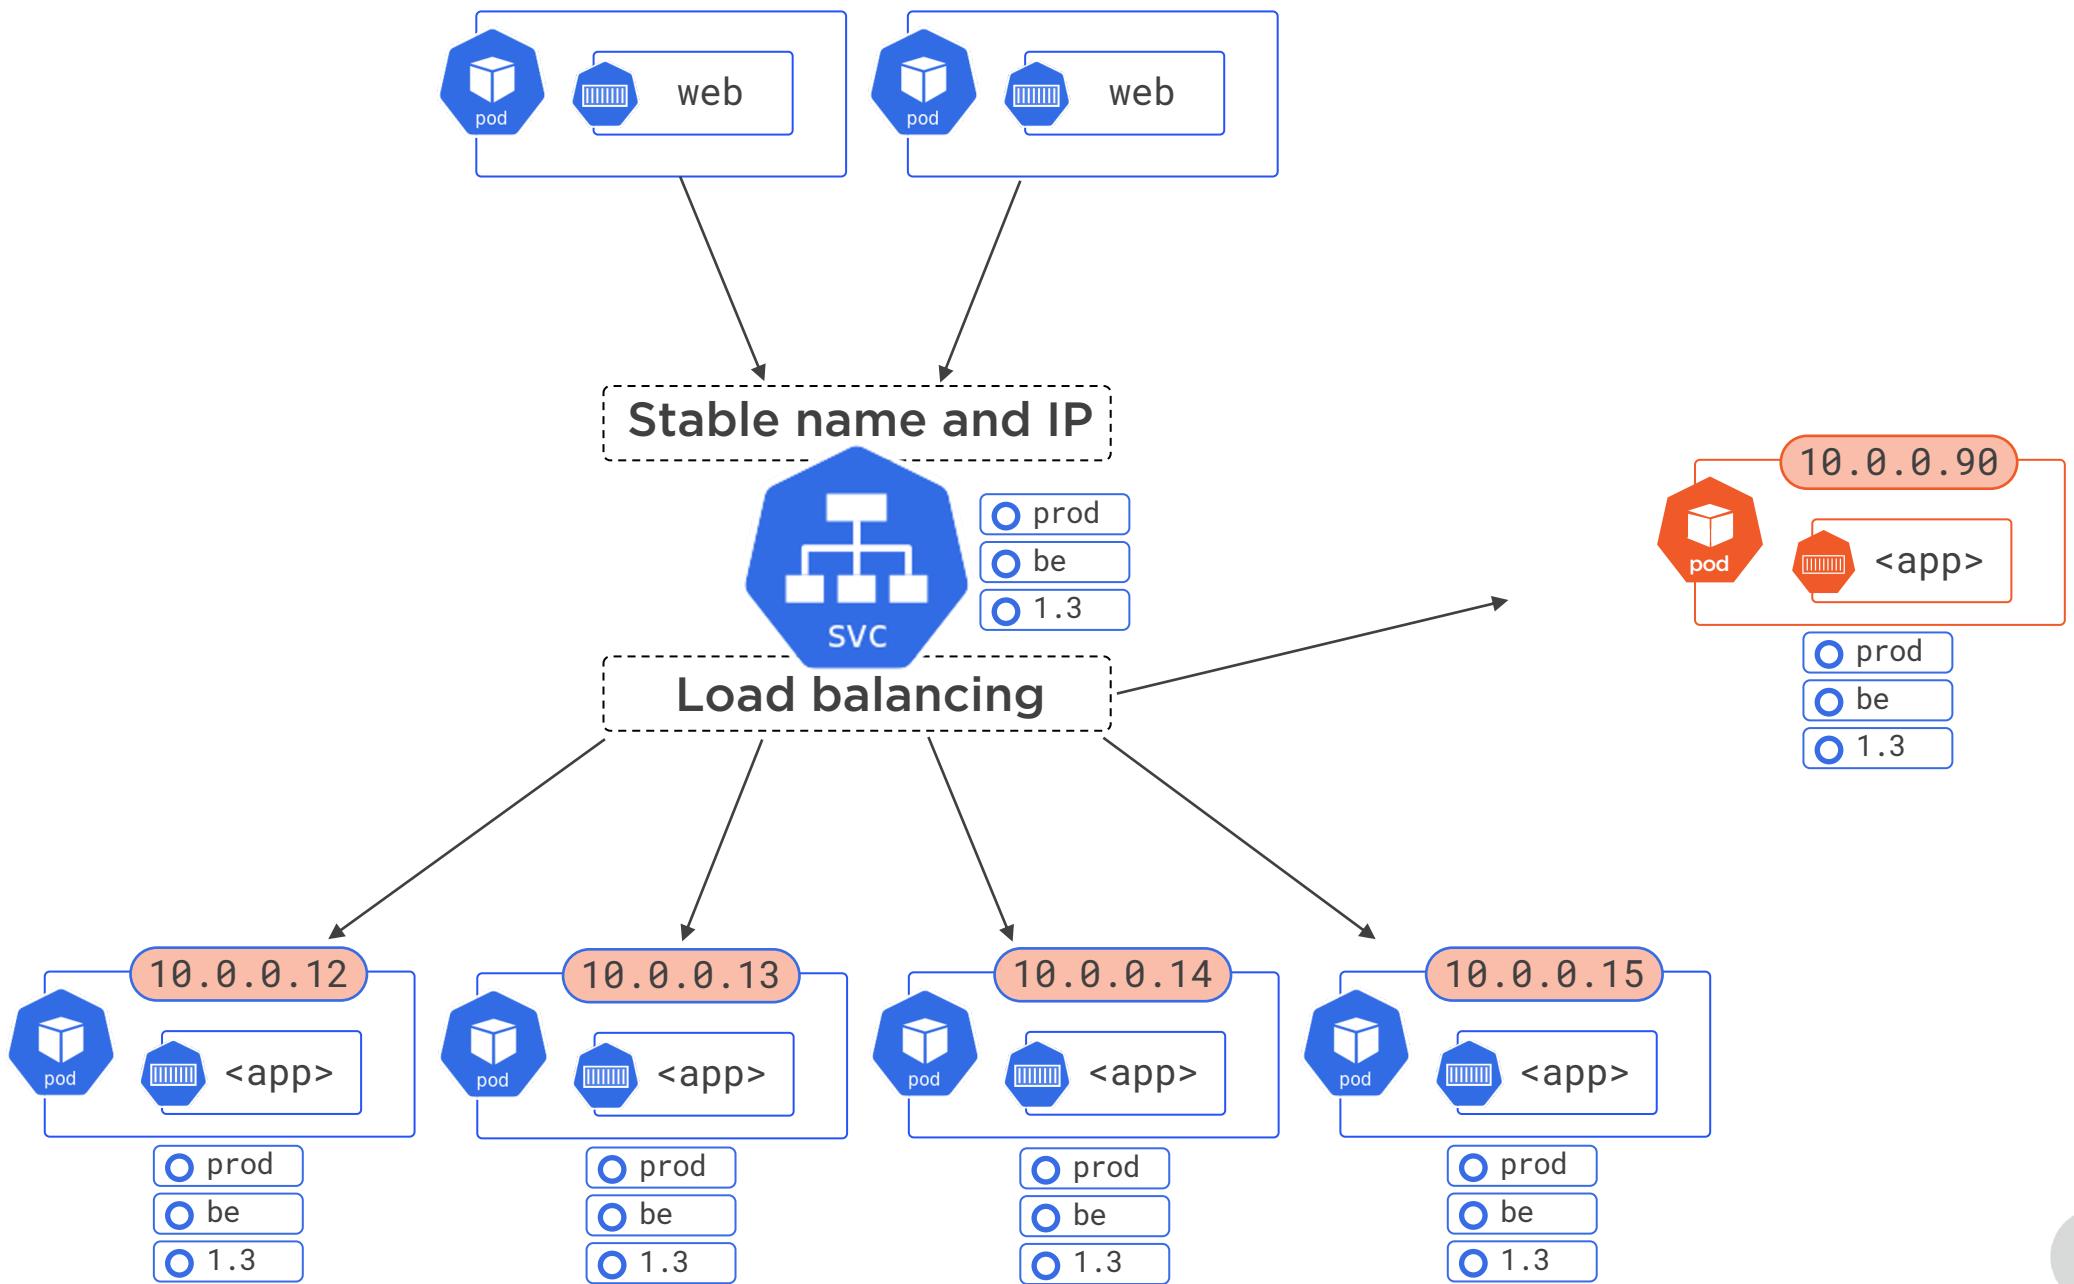


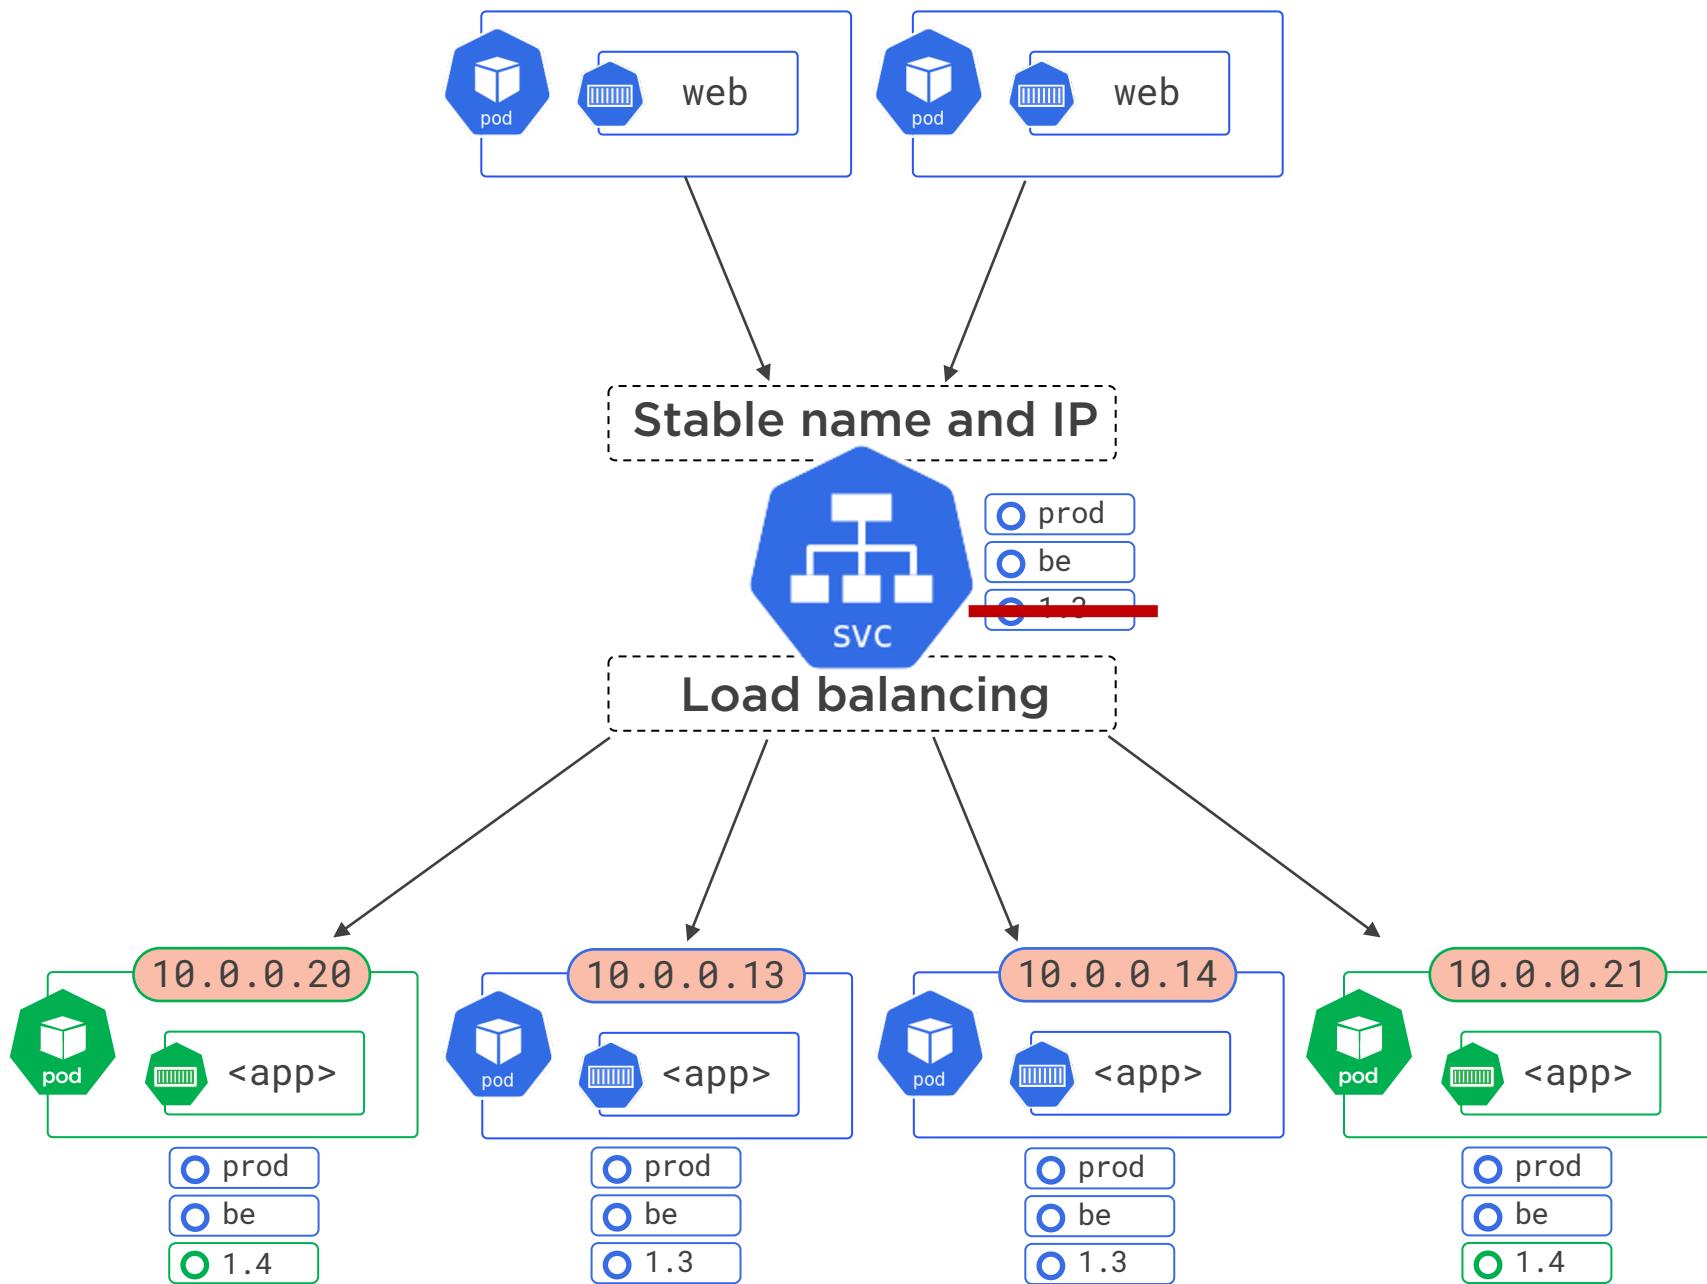


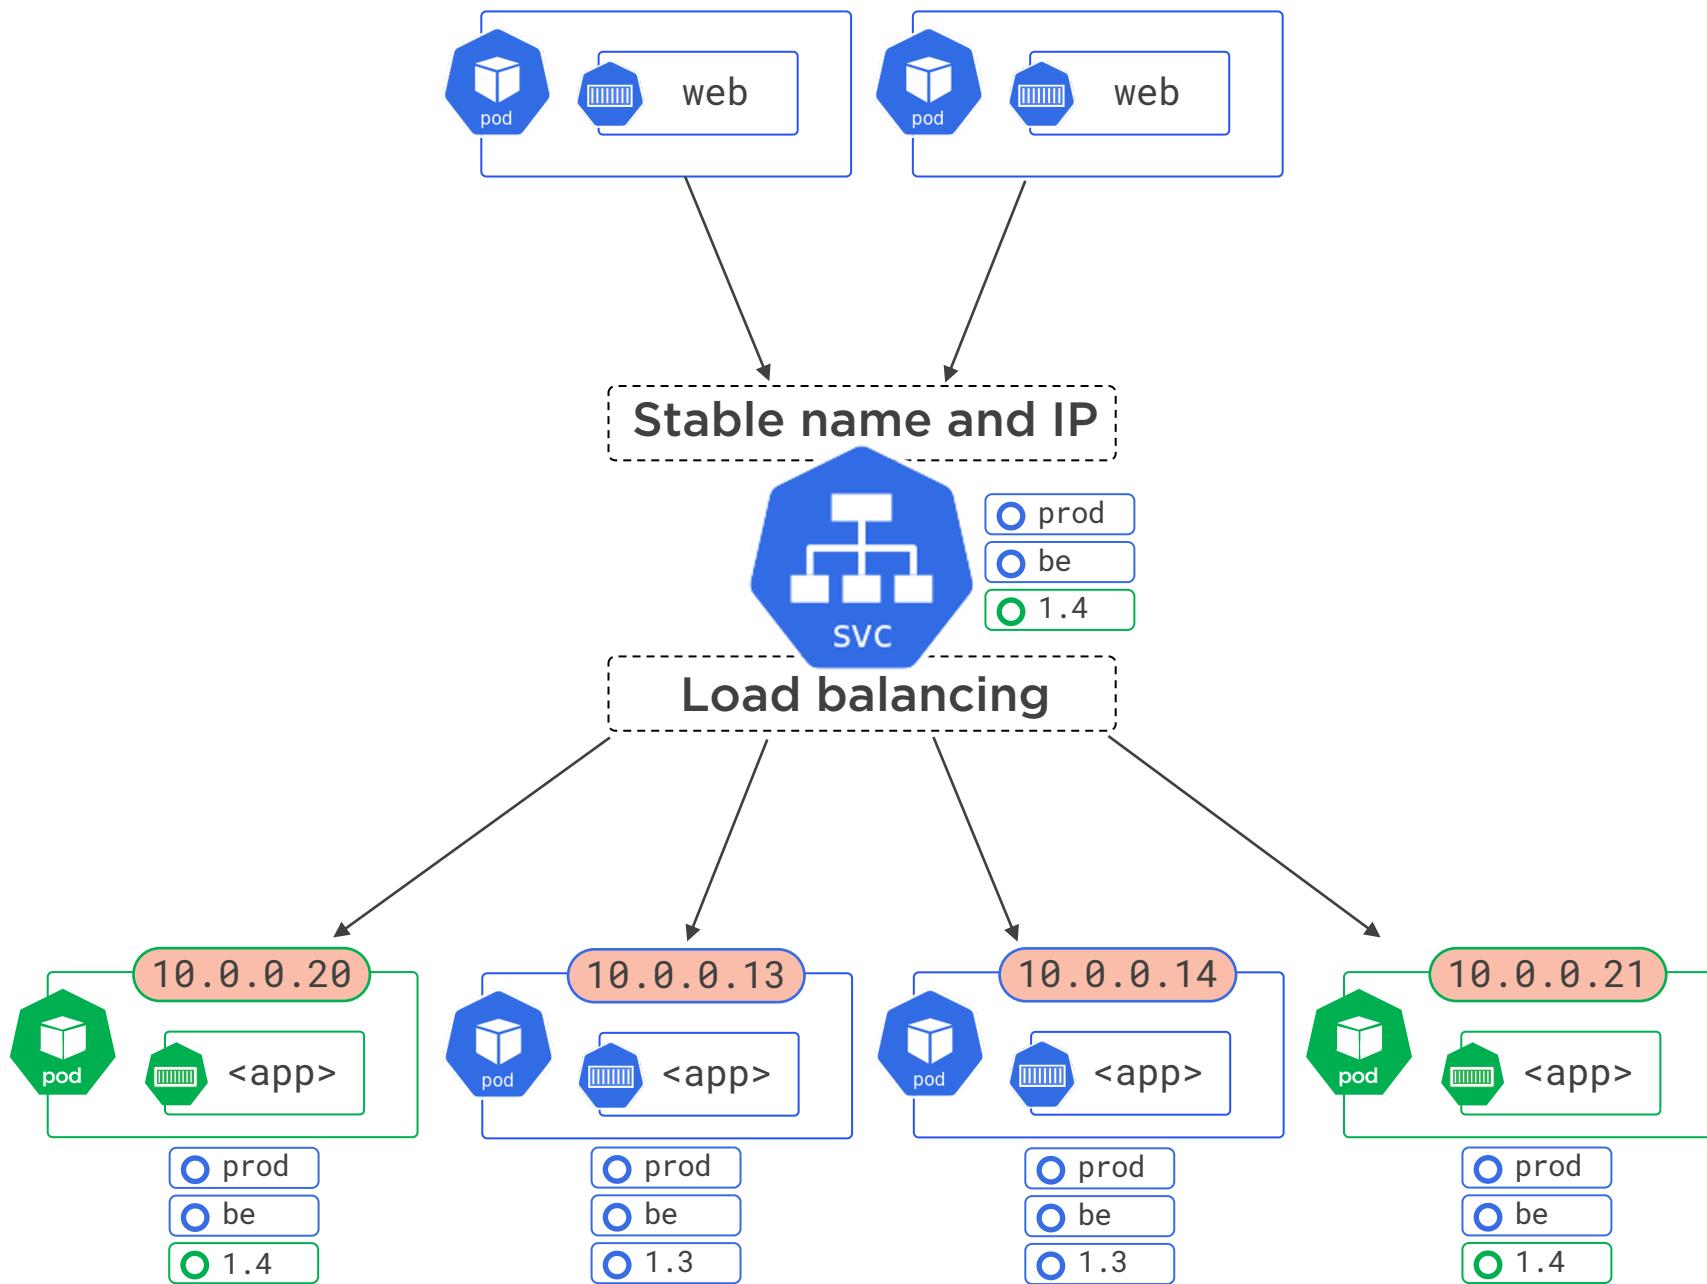
Labels

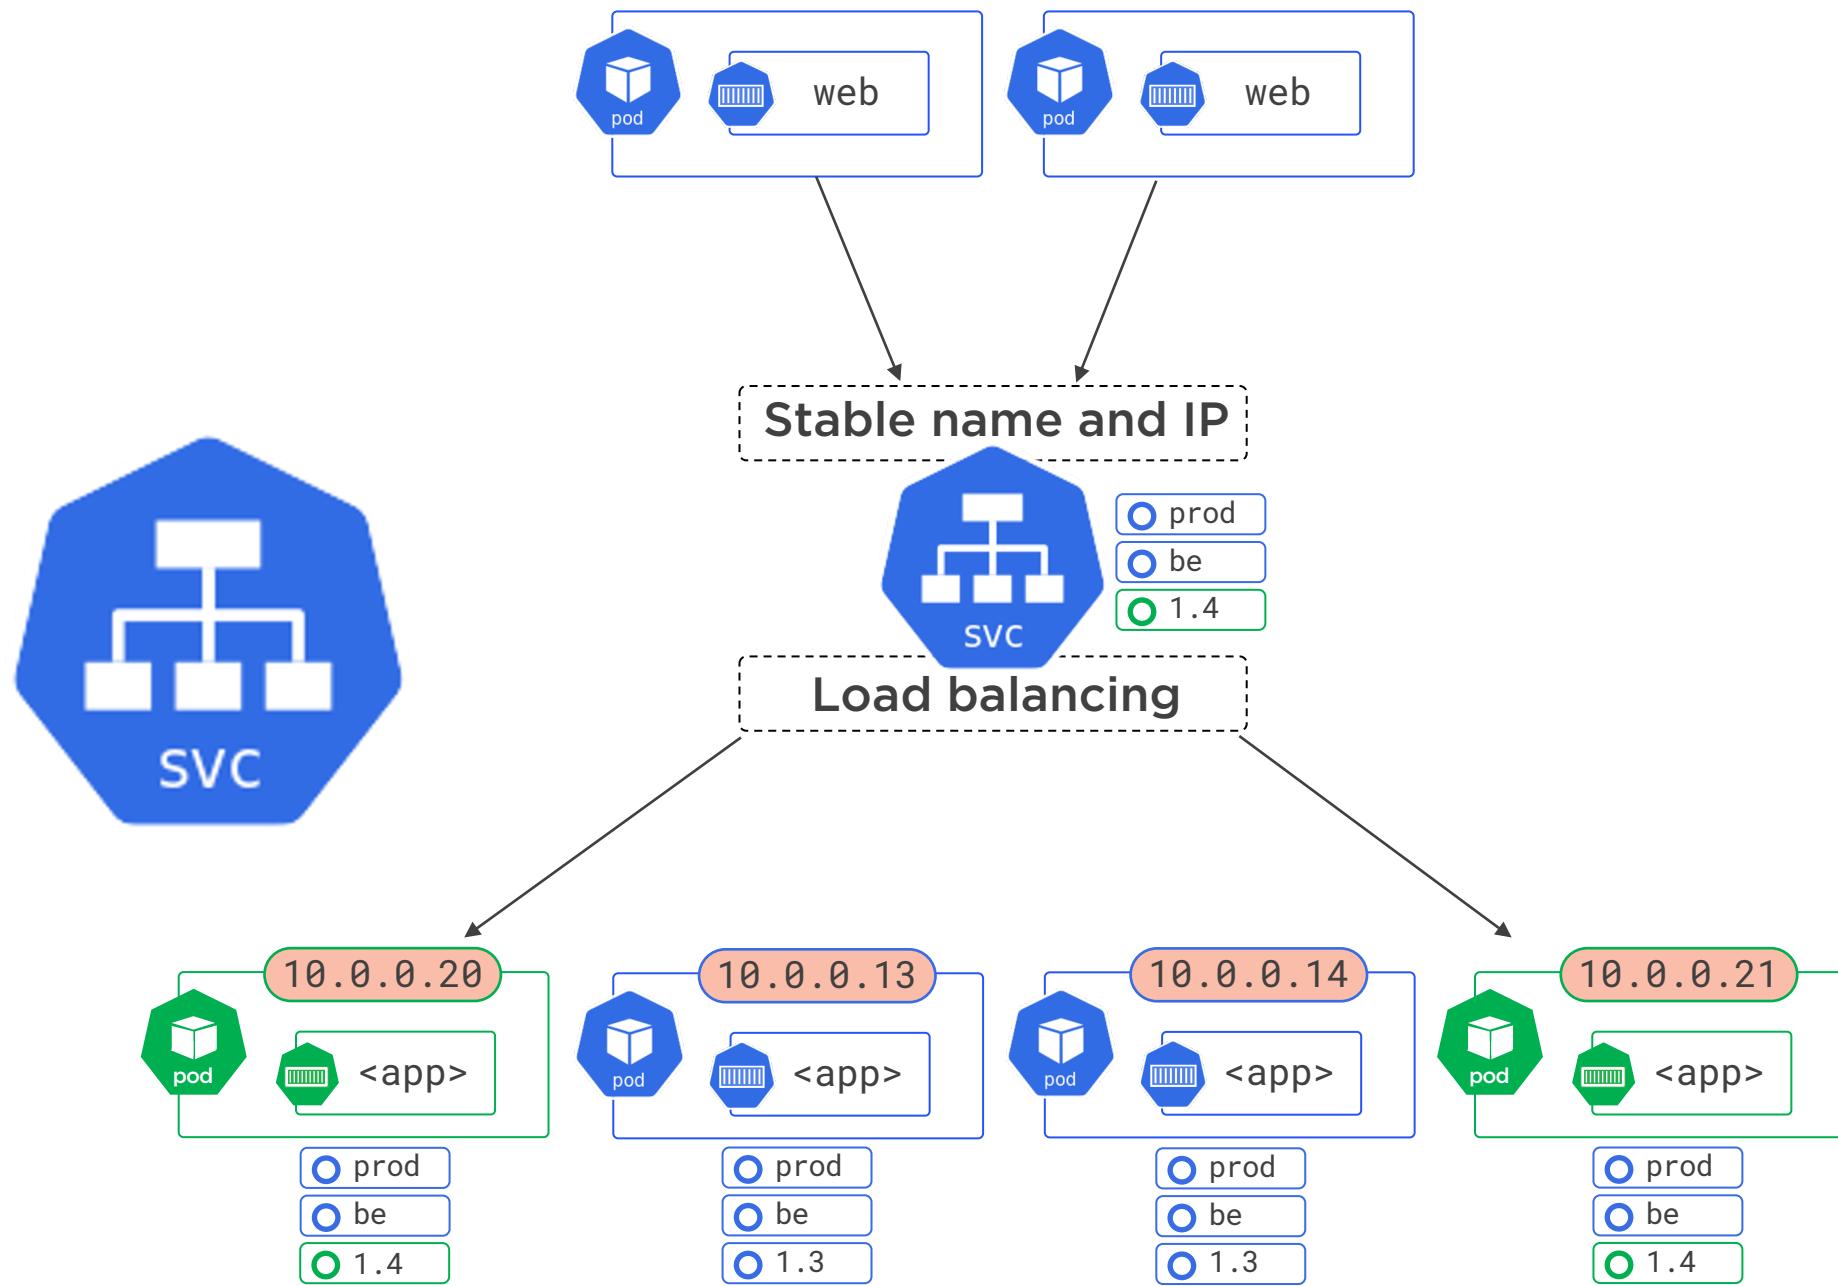
Labels are very simple and very powerful!













- Only sends traffic to healthy Pods**
- Can do session affinity**
- Can send traffic to endpoints outside the cluster**
- Can do TCP and UDP**



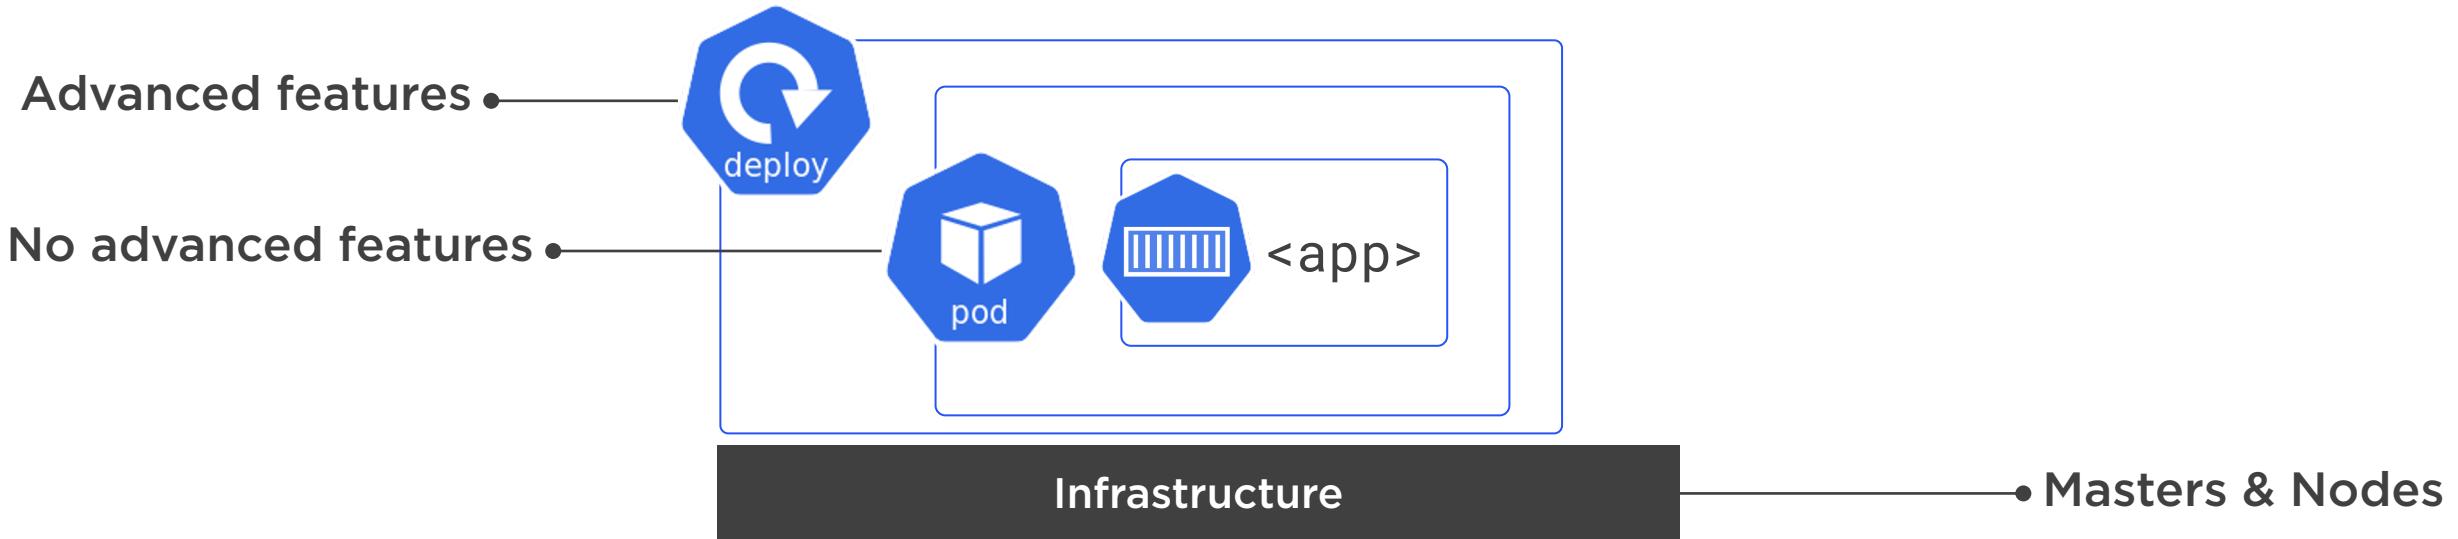
Up Next:

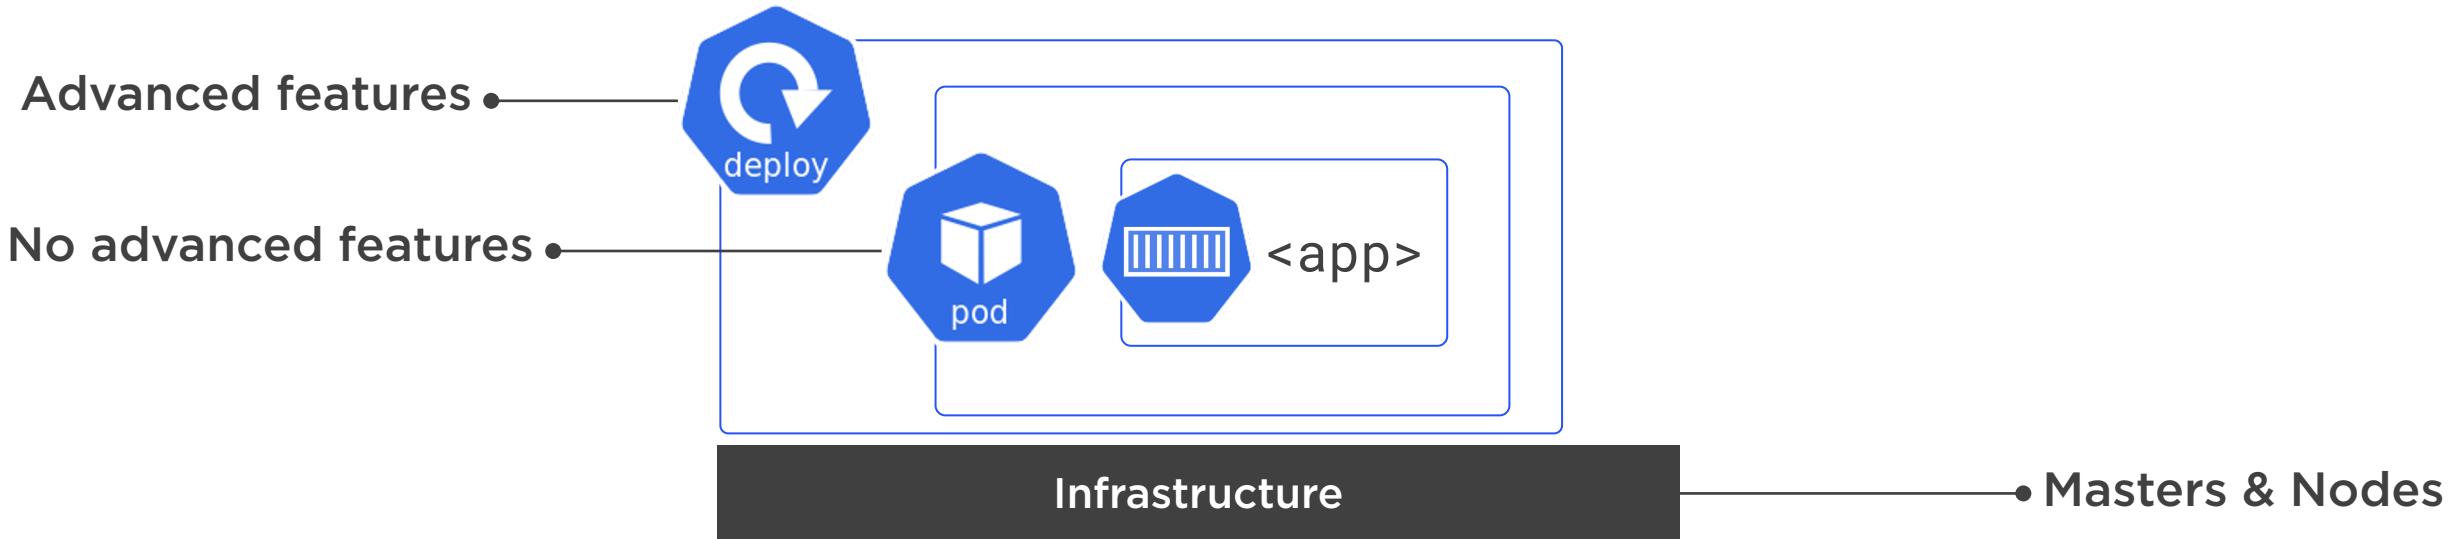
Game-changing Deployments



Game-changing Deployments









Stateless apps



One instance on every node



Stateful apps...



Time-based short-lived jobs





Stateless apps



One instance on every node

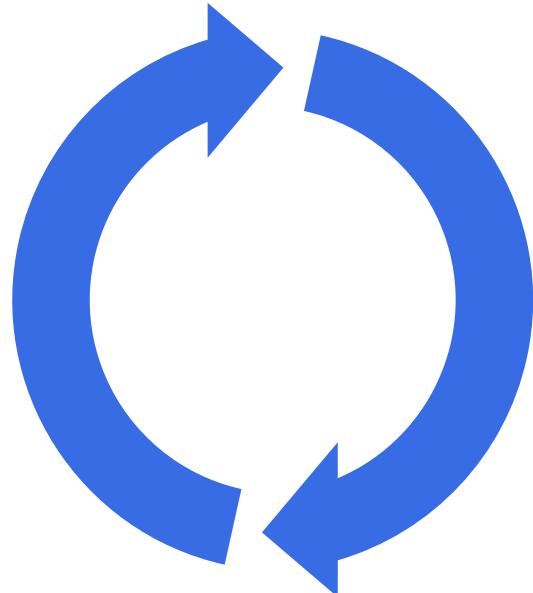


Stateful apps...



Time-based short-lived jobs





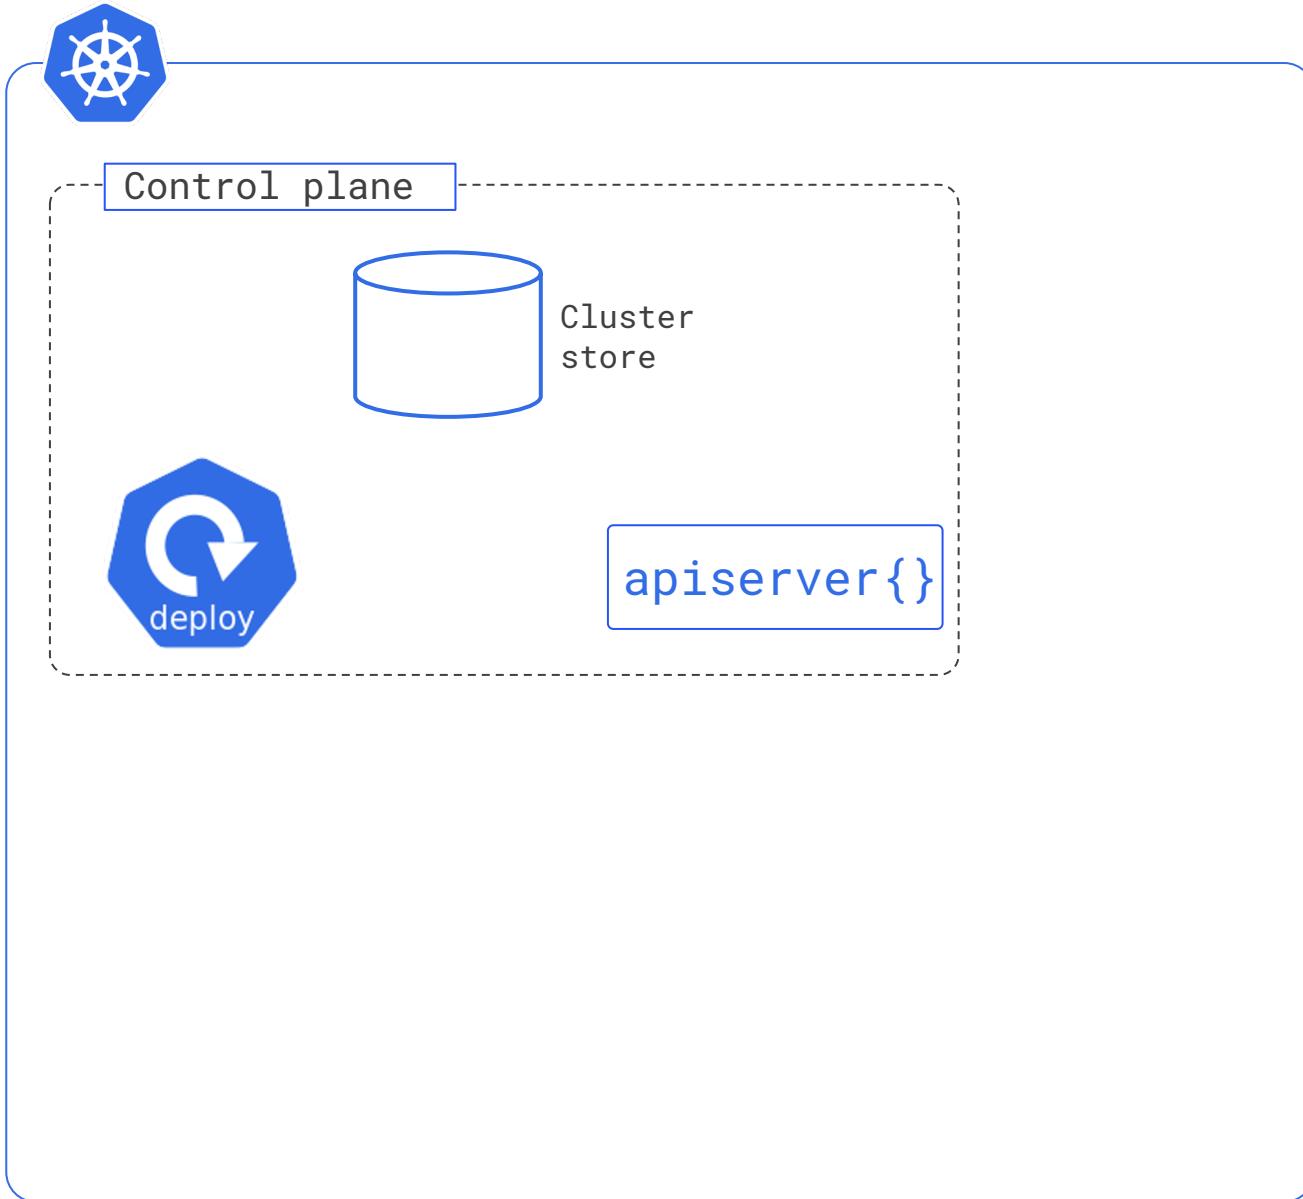
Deployment Controller/Reconciliation loop

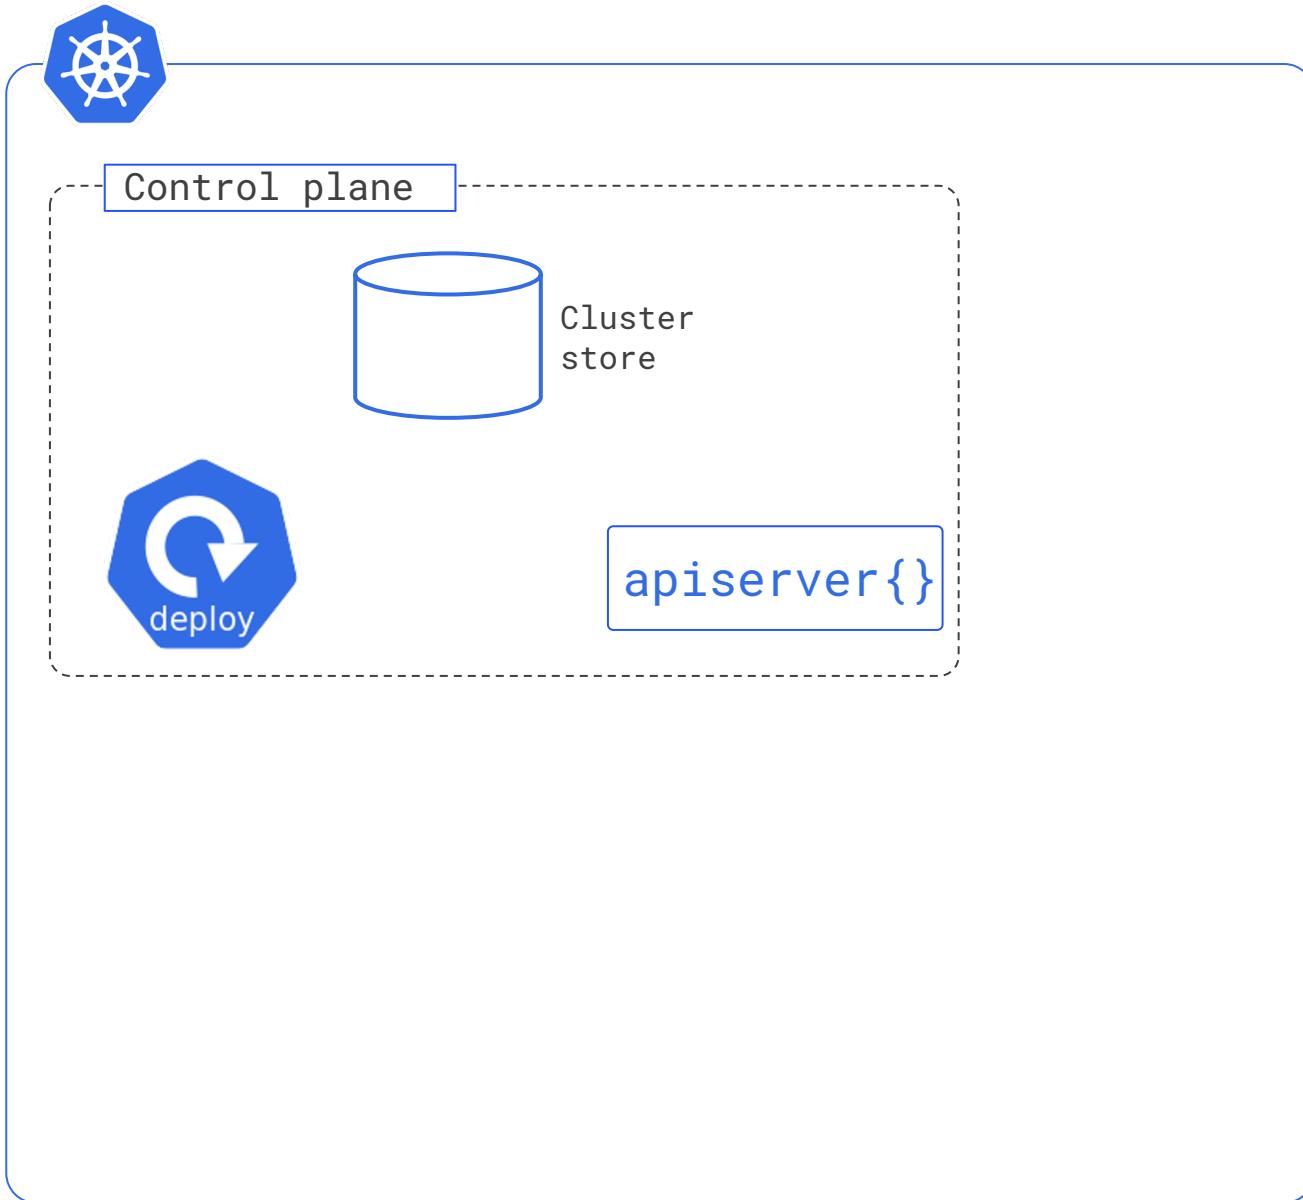
Watches API Server for new Deployments

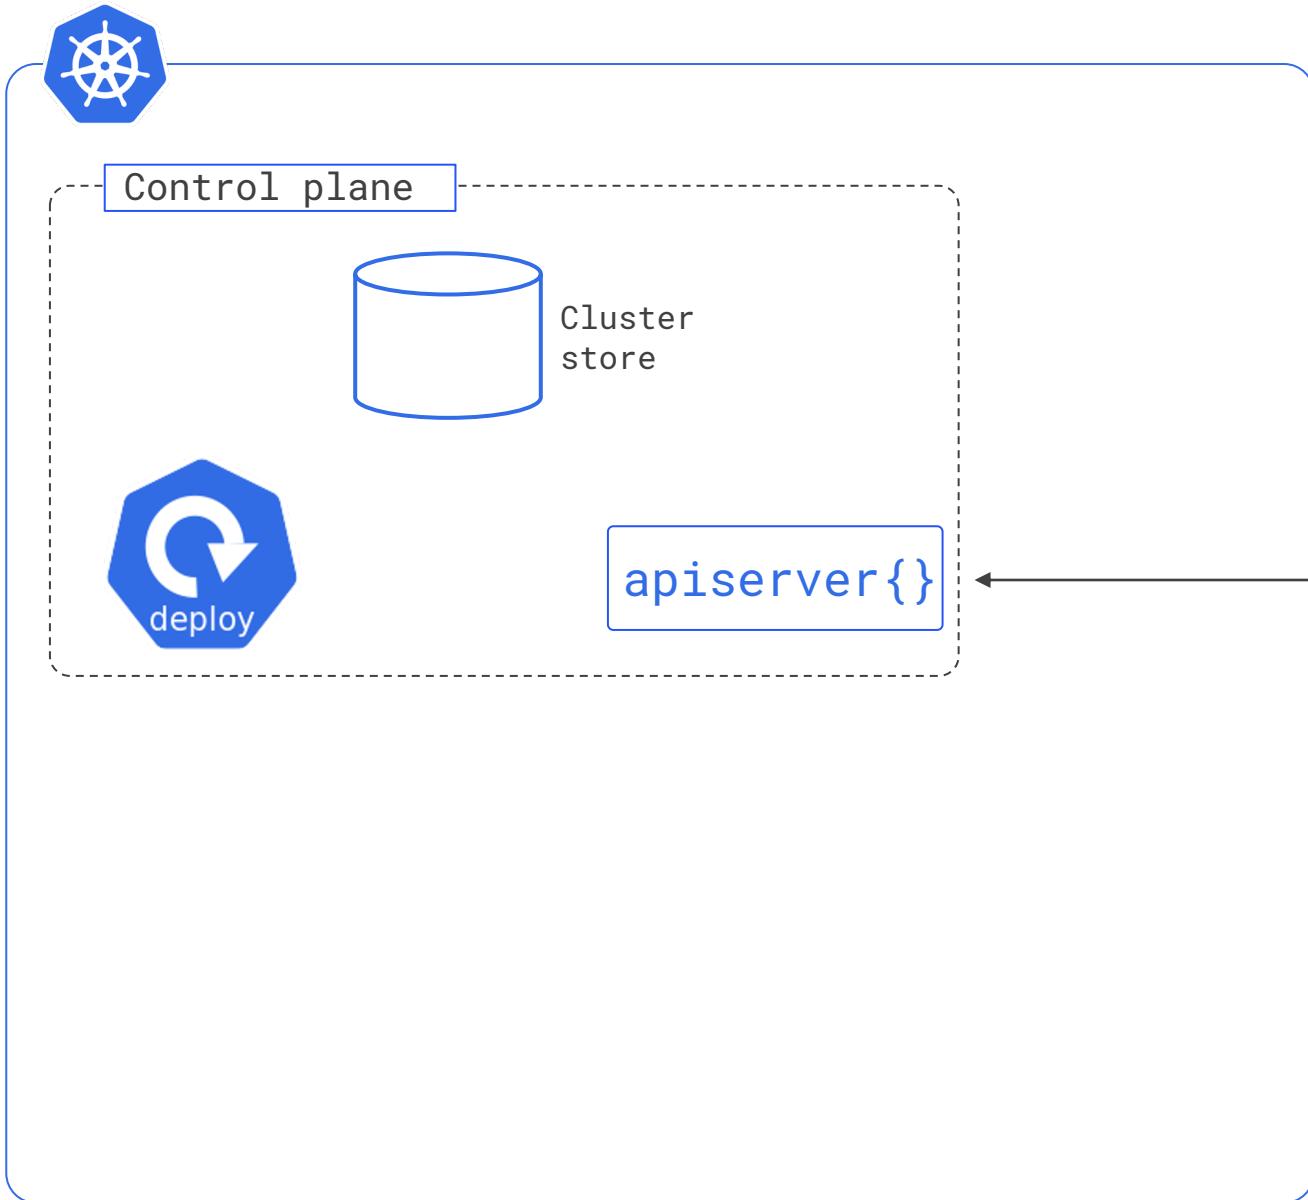
Implements them

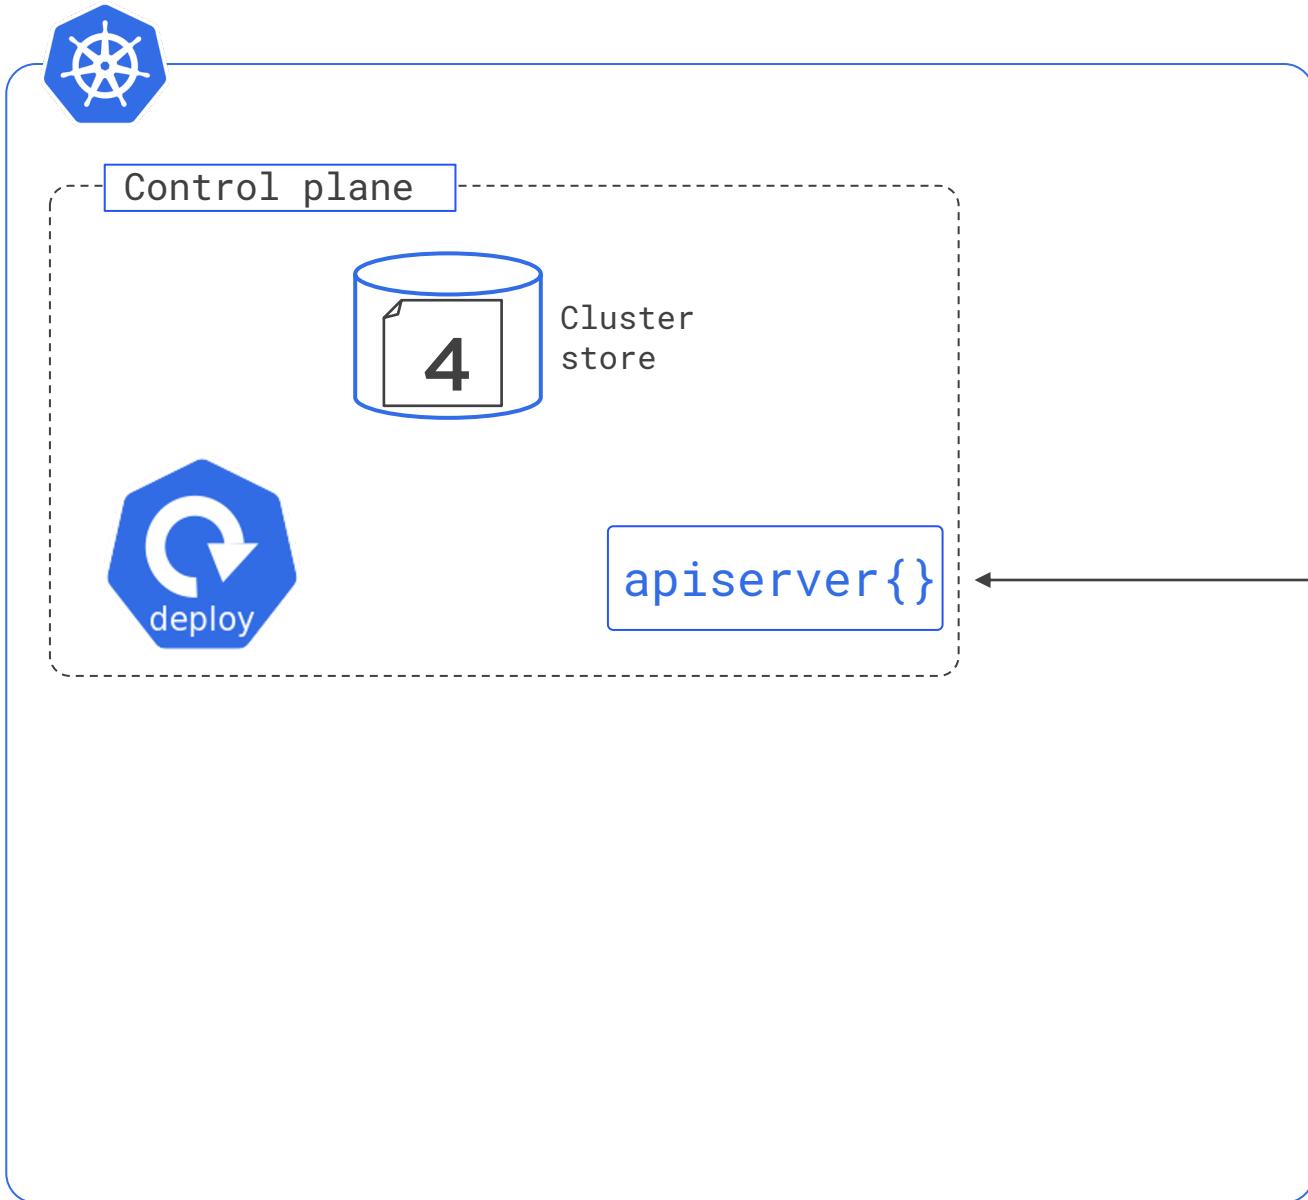
Constantly compares *observed state* with
desired state

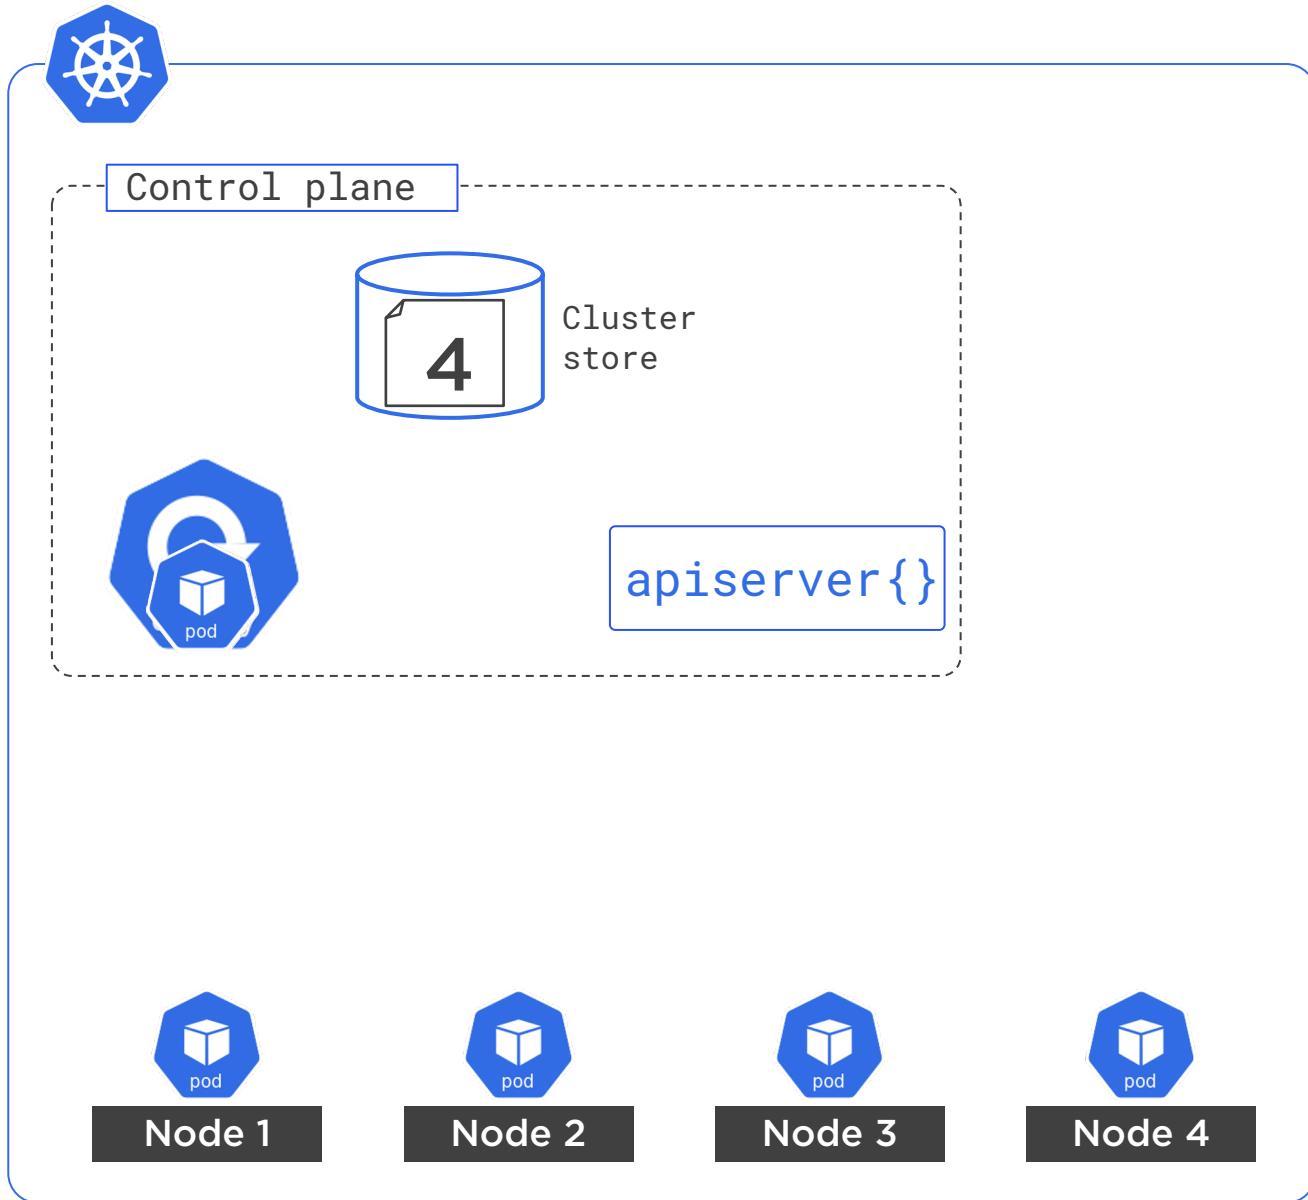


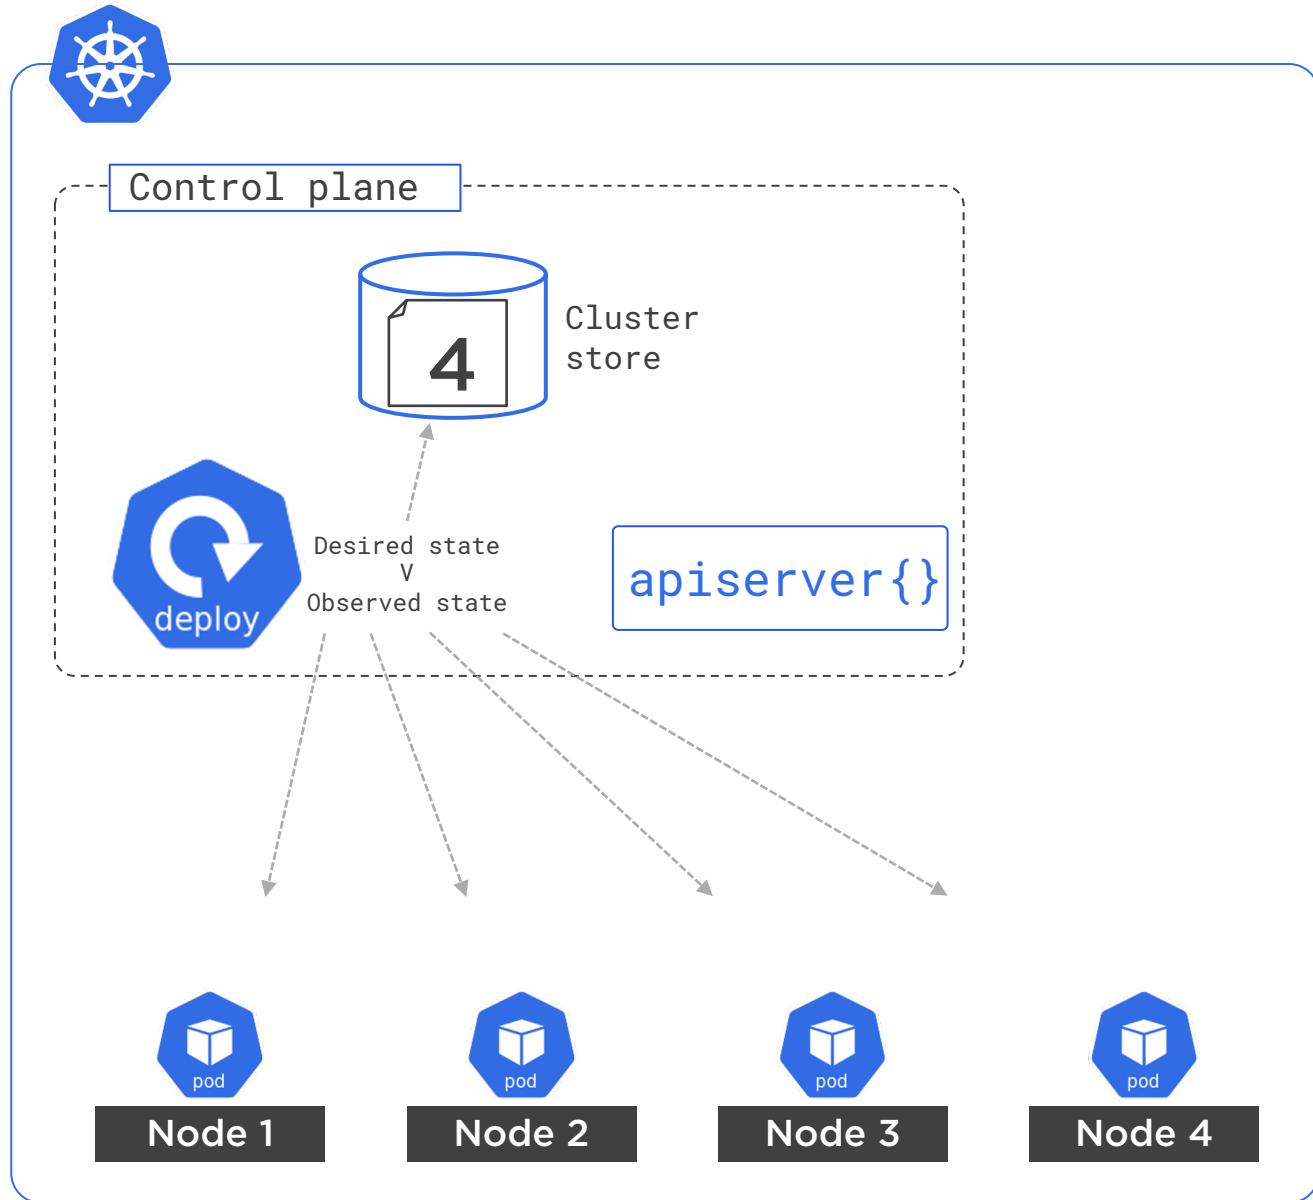


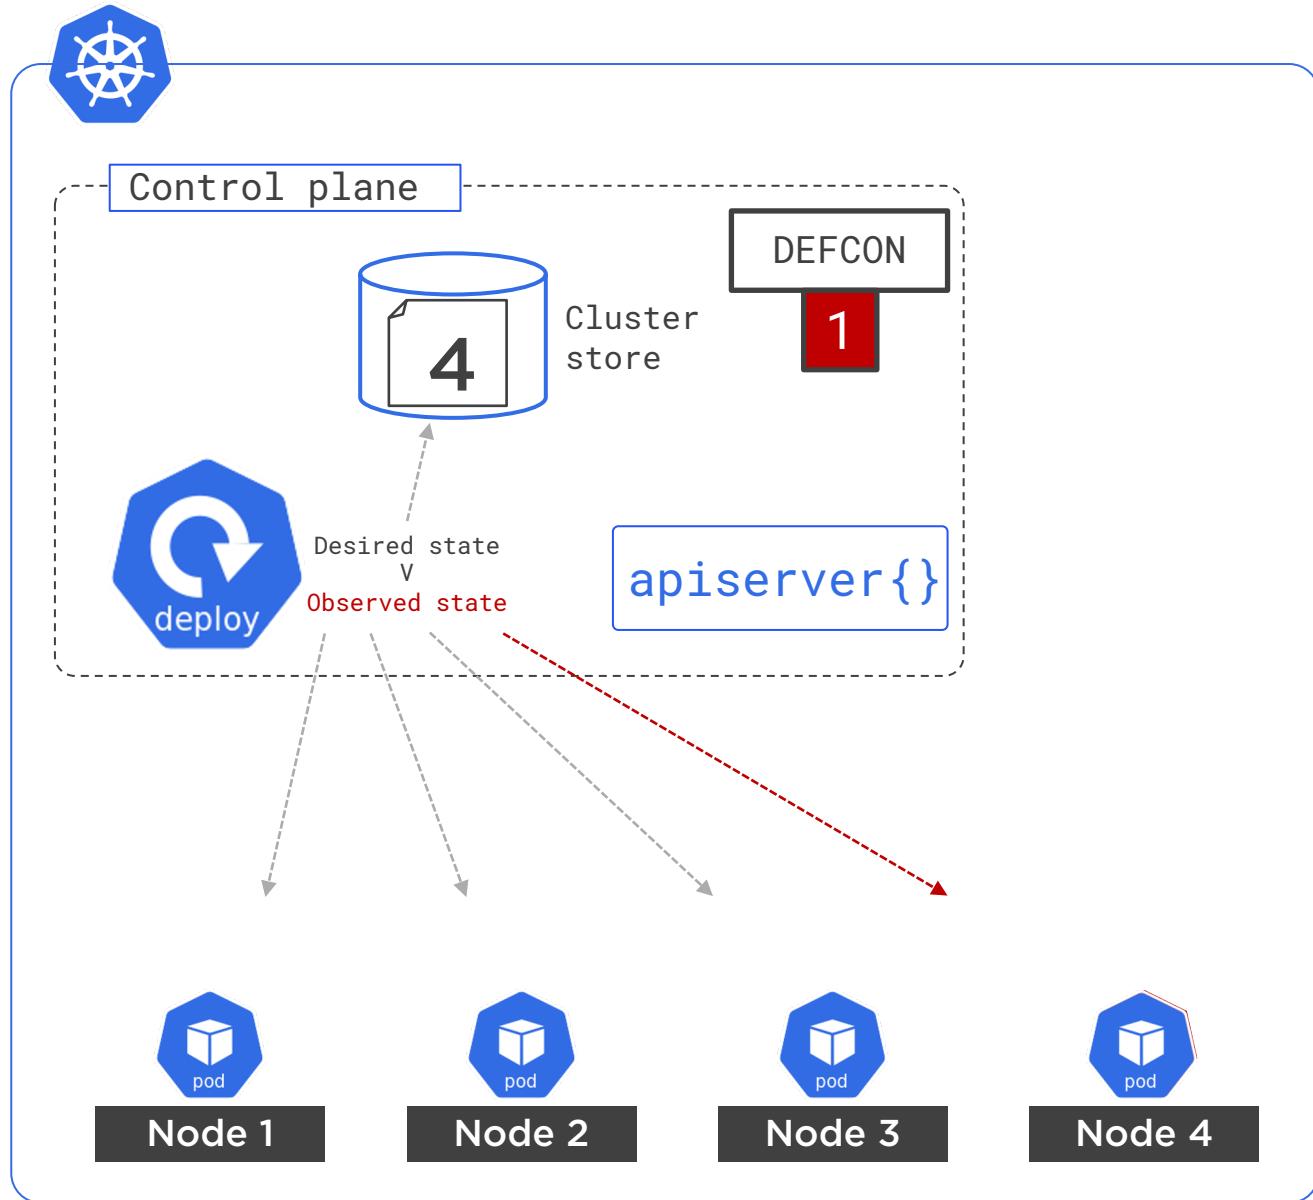


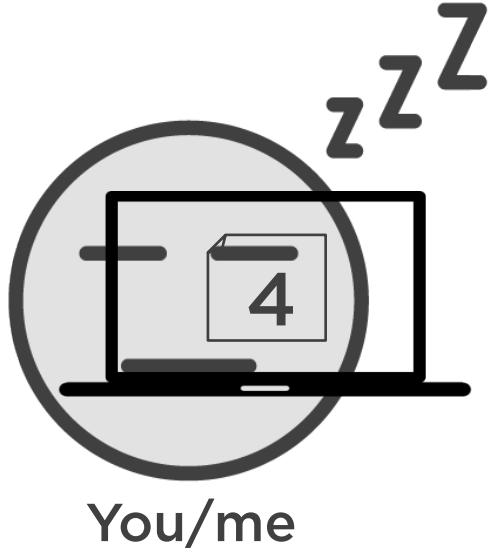
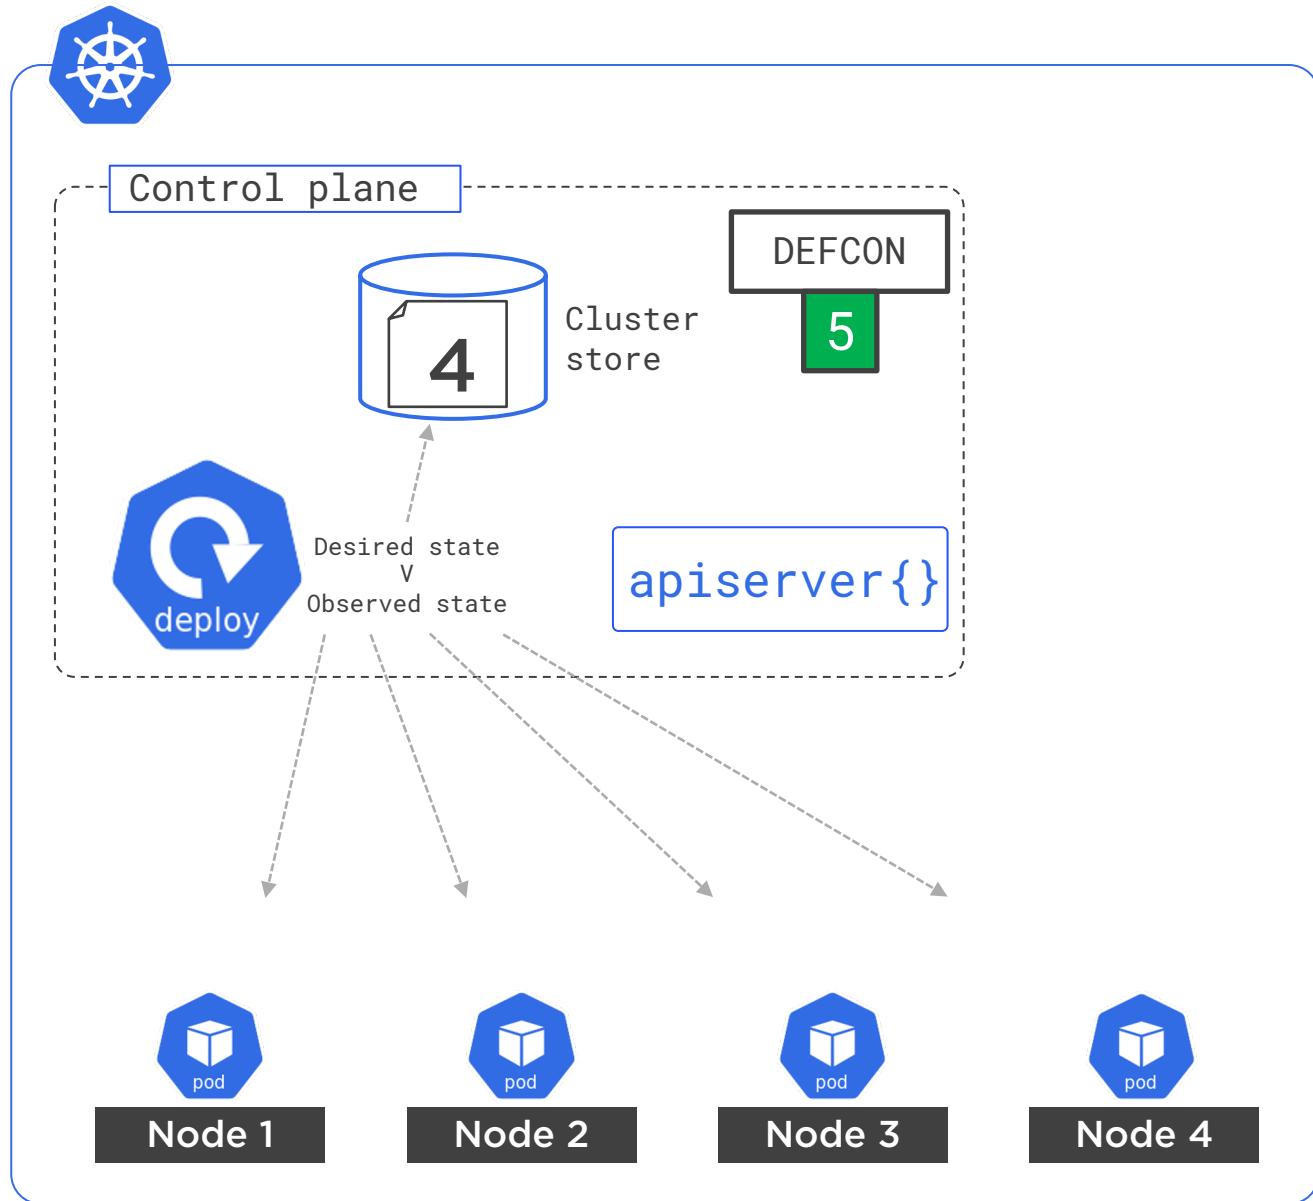


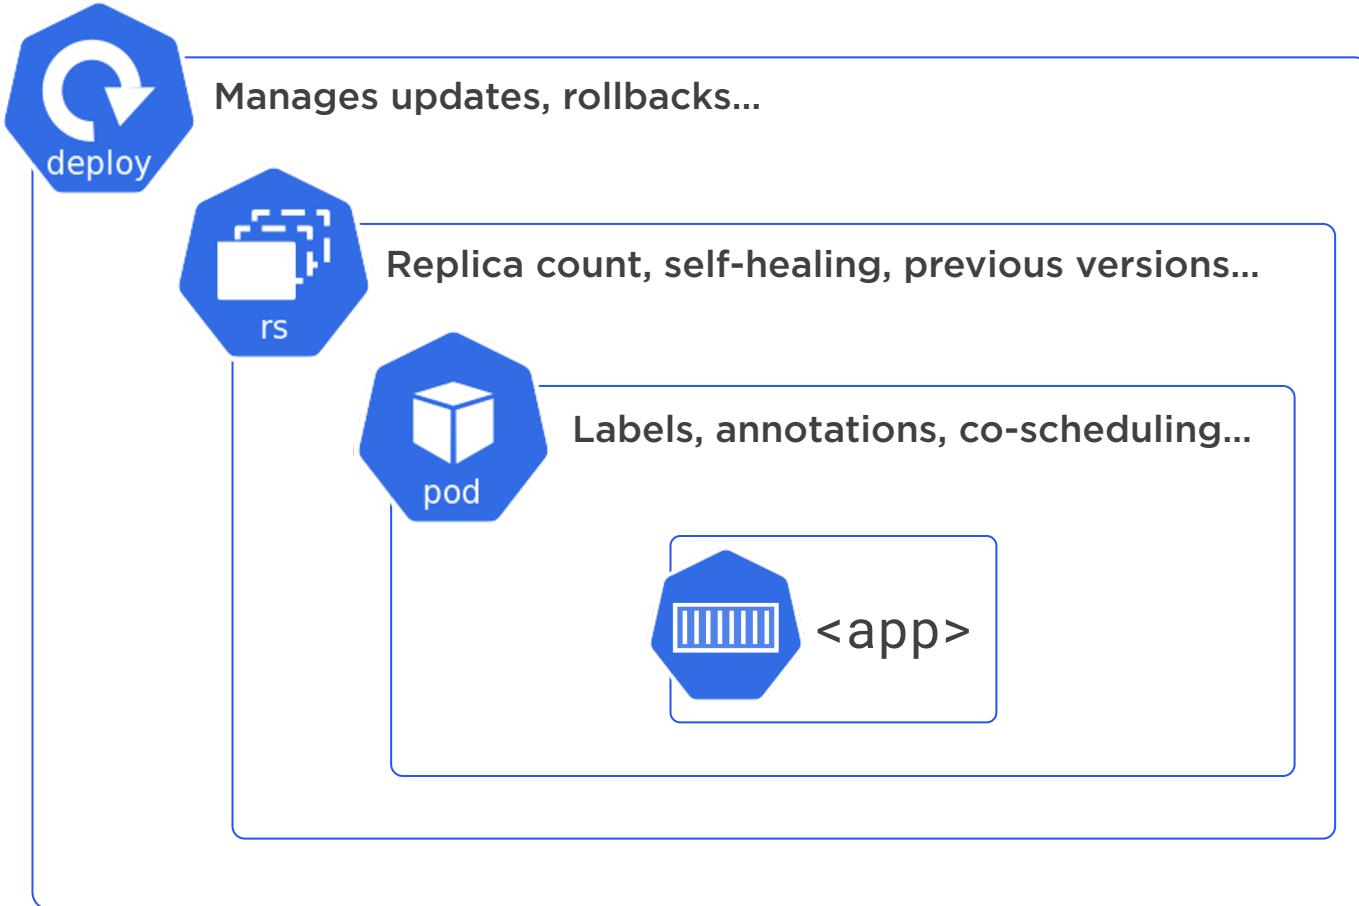


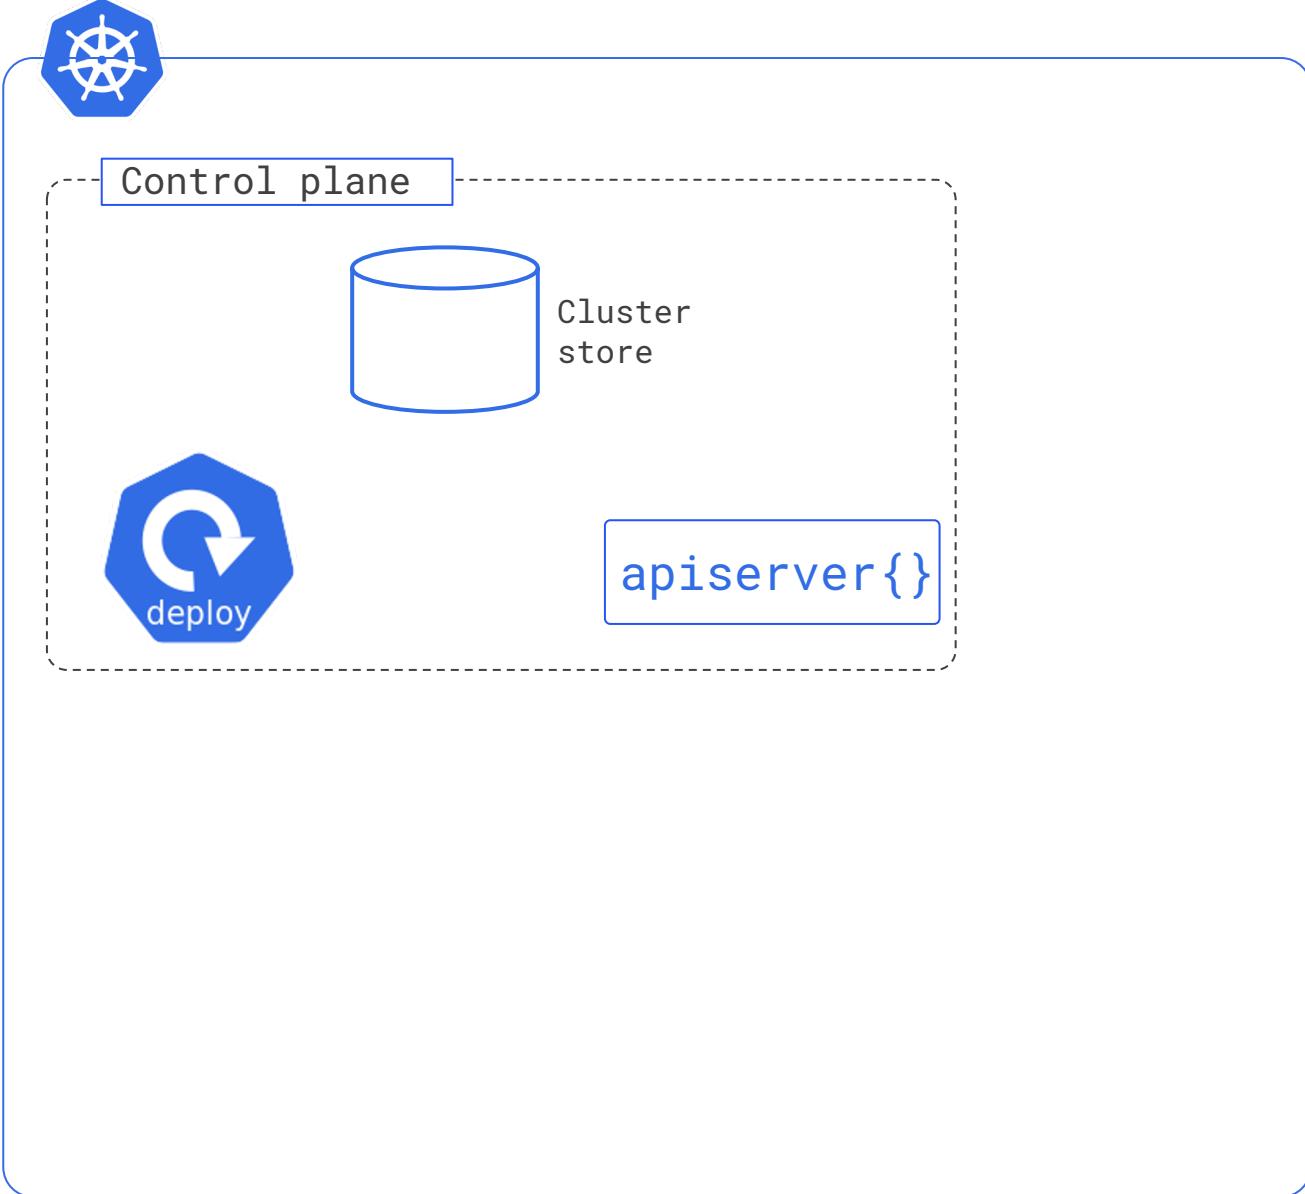


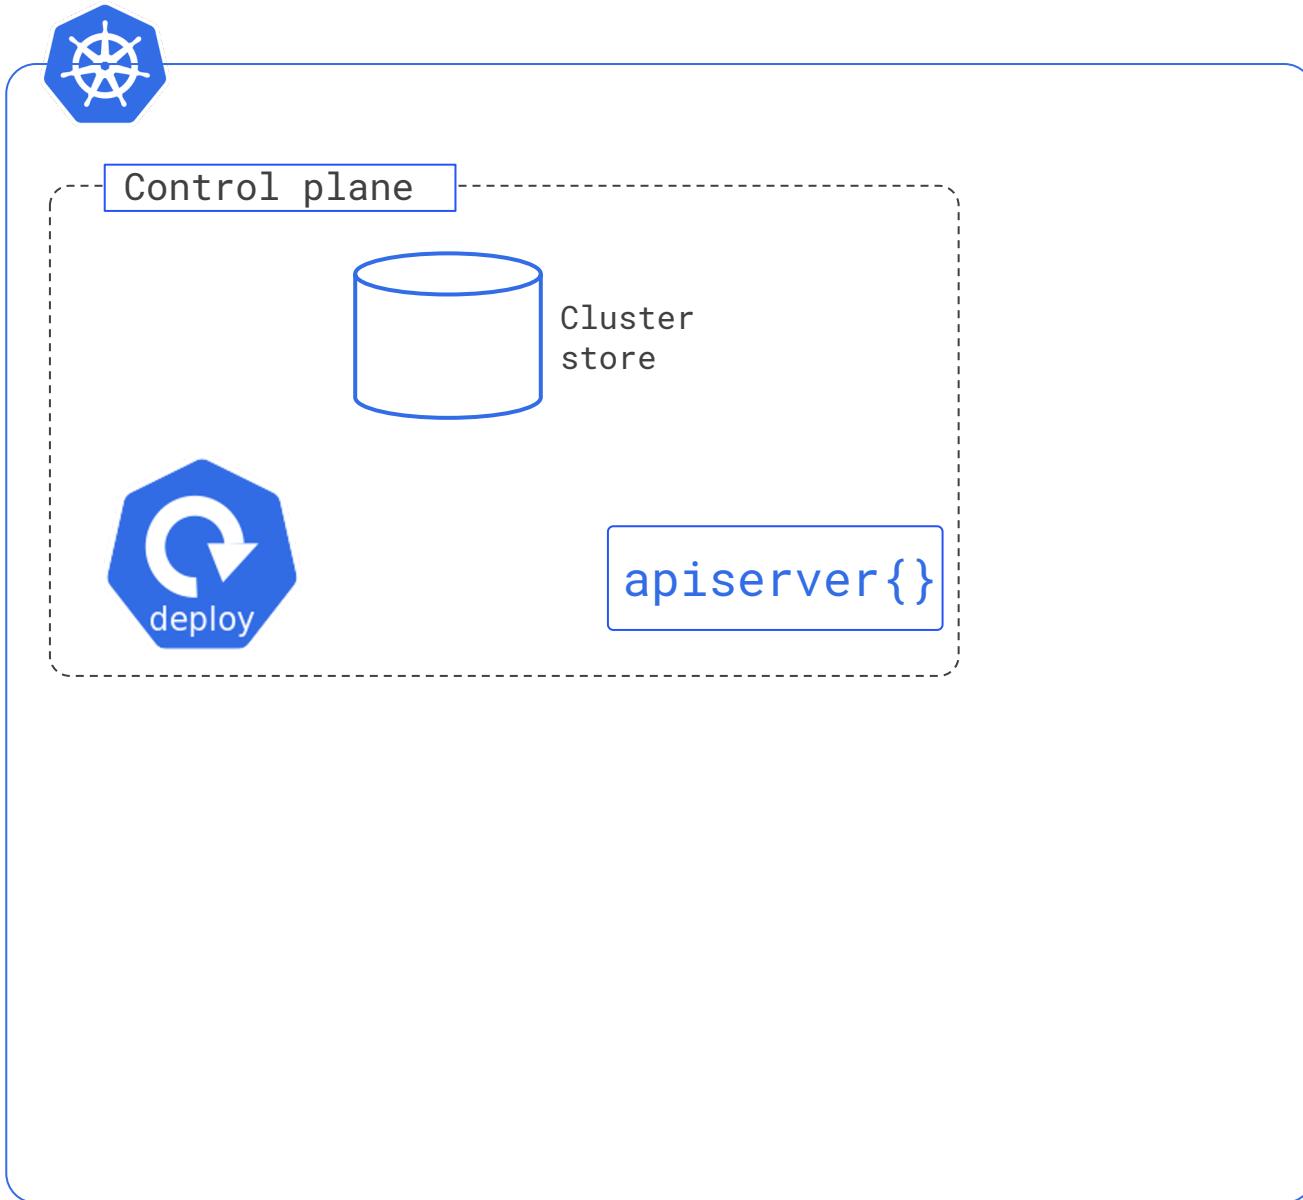


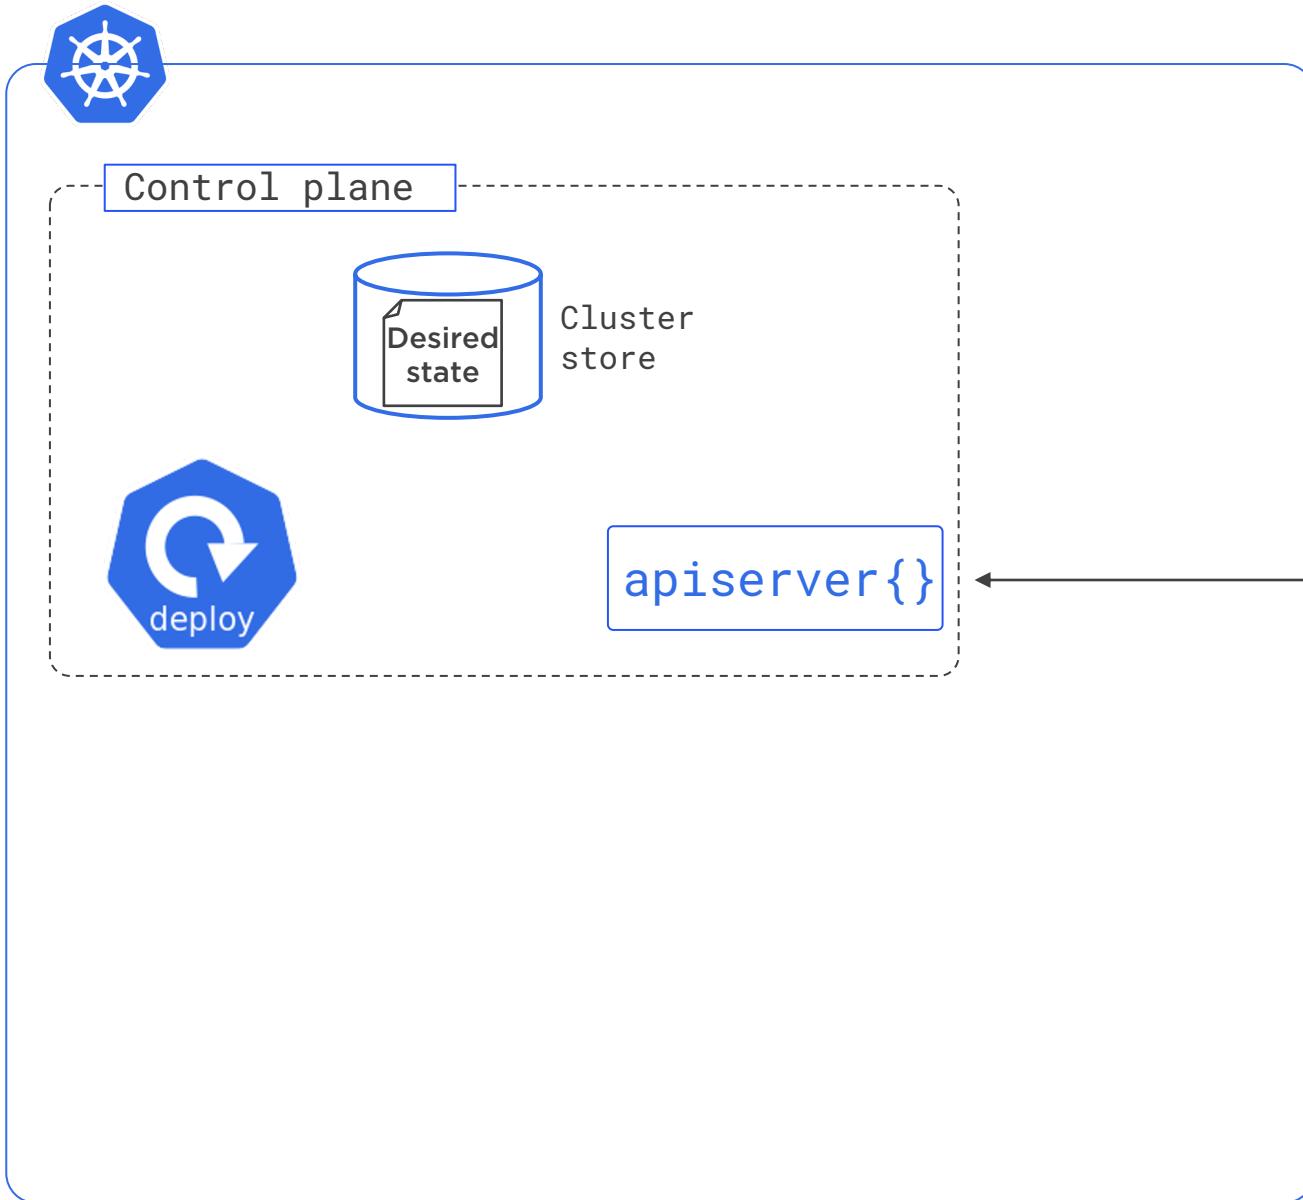


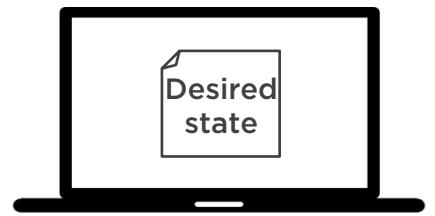
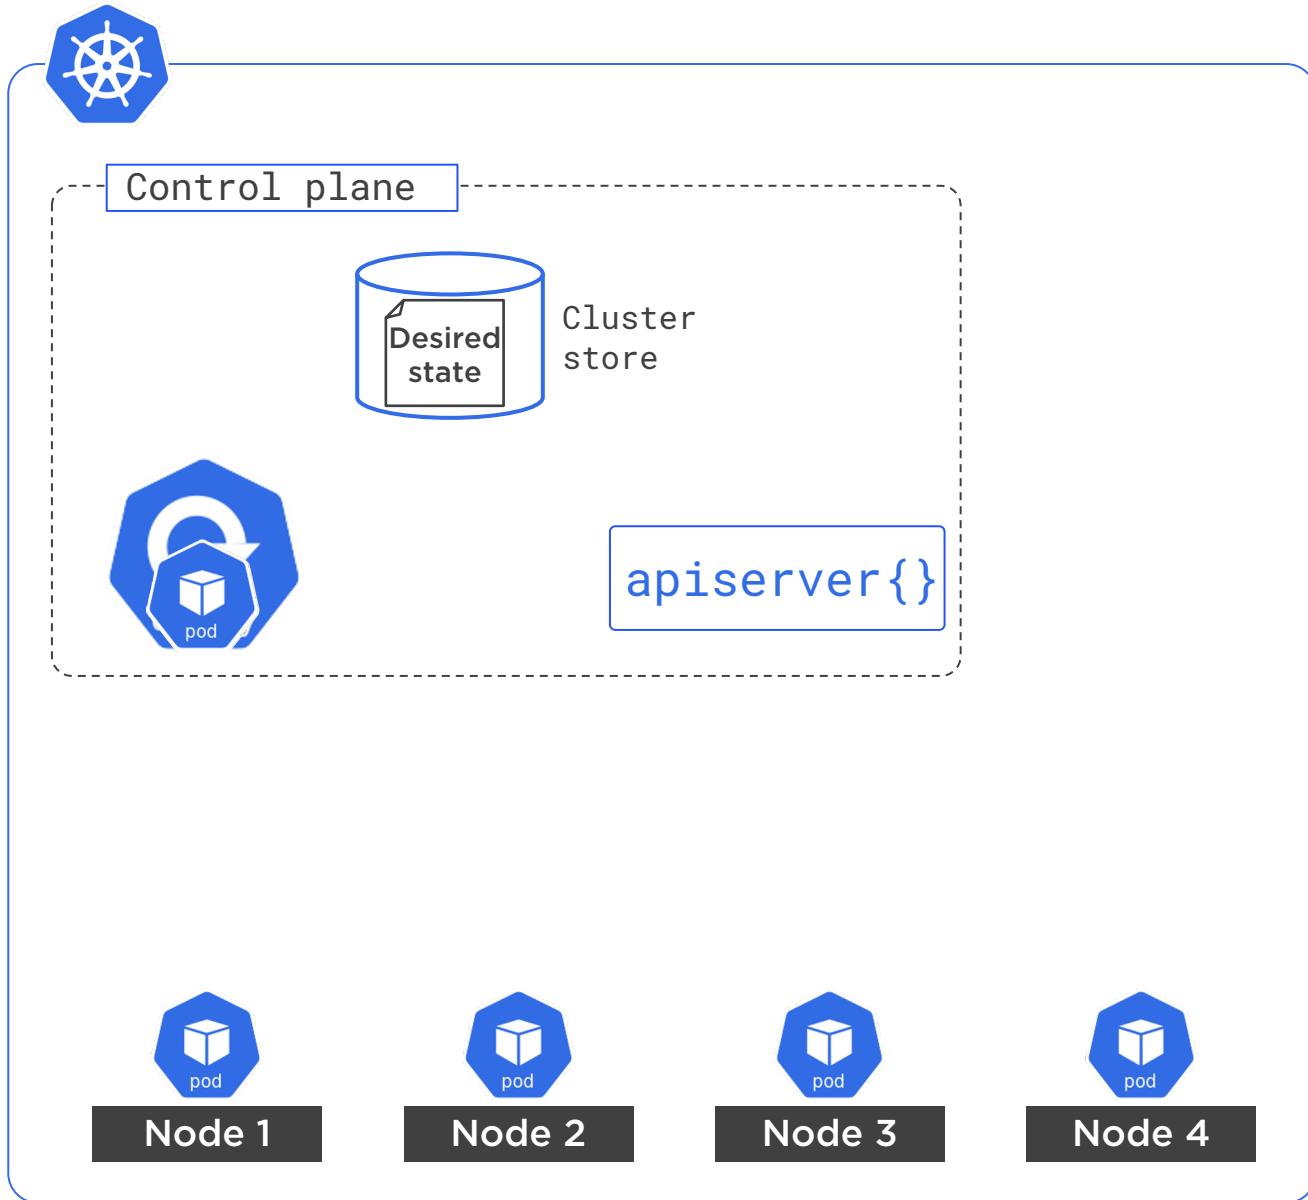


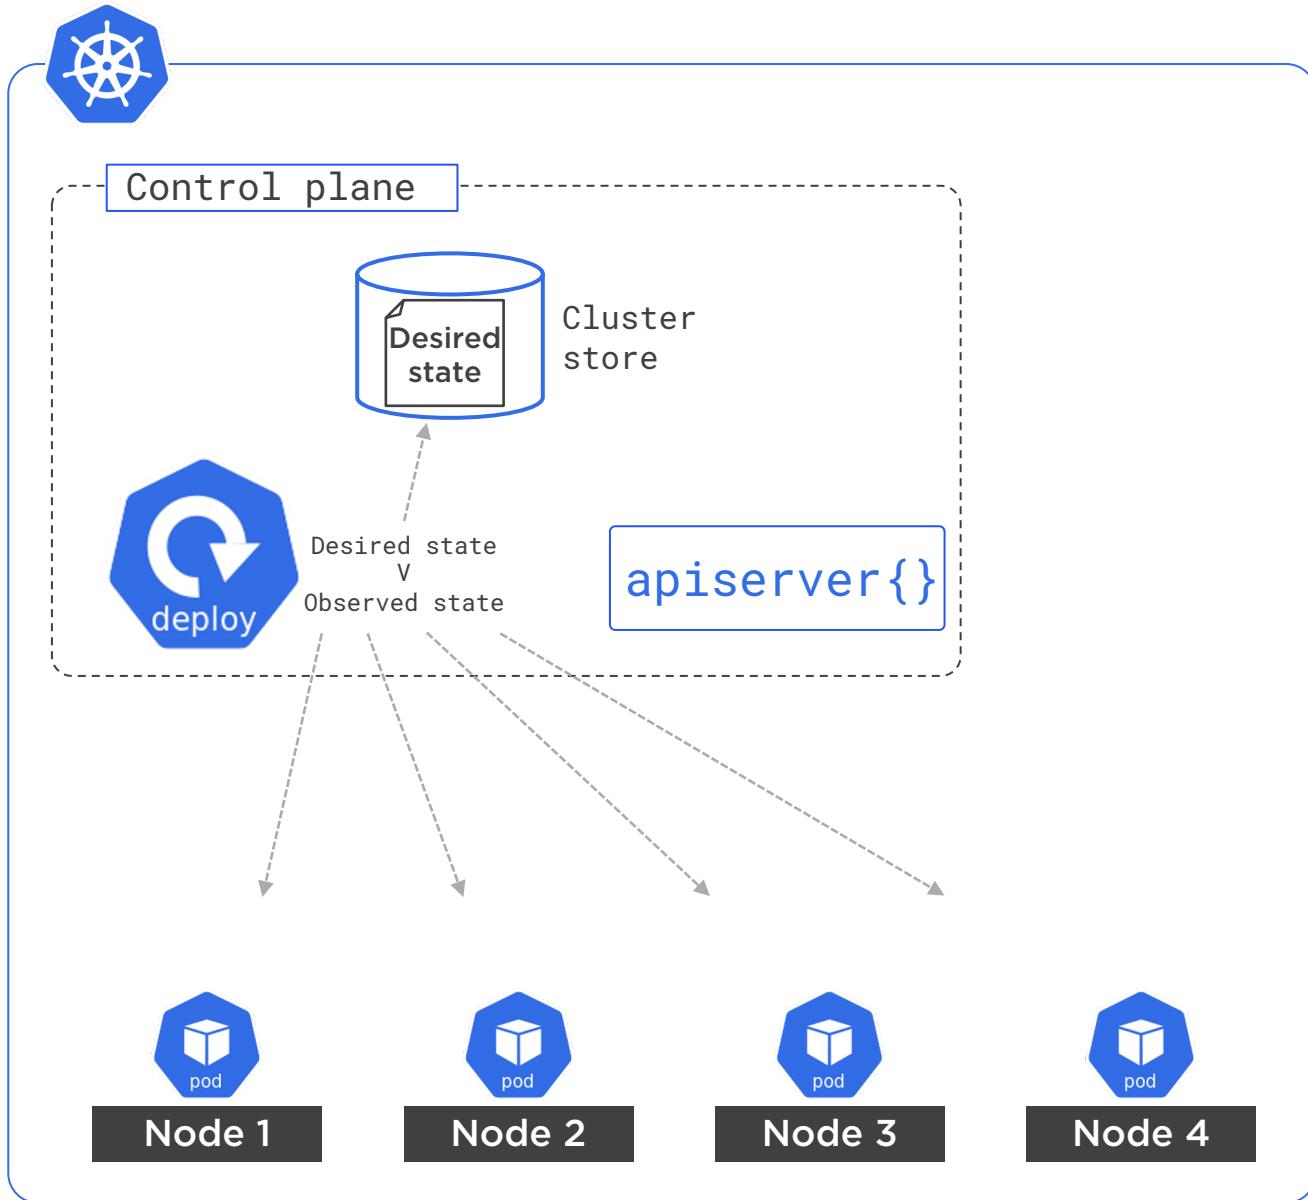












Up Next:
The K8s API and API Server



The K8s API and API Server





Atomic unit of scheduling



Replica count



Updates and rollbacks



Stable network abstraction





Atomic unit of scheduling



Replica count

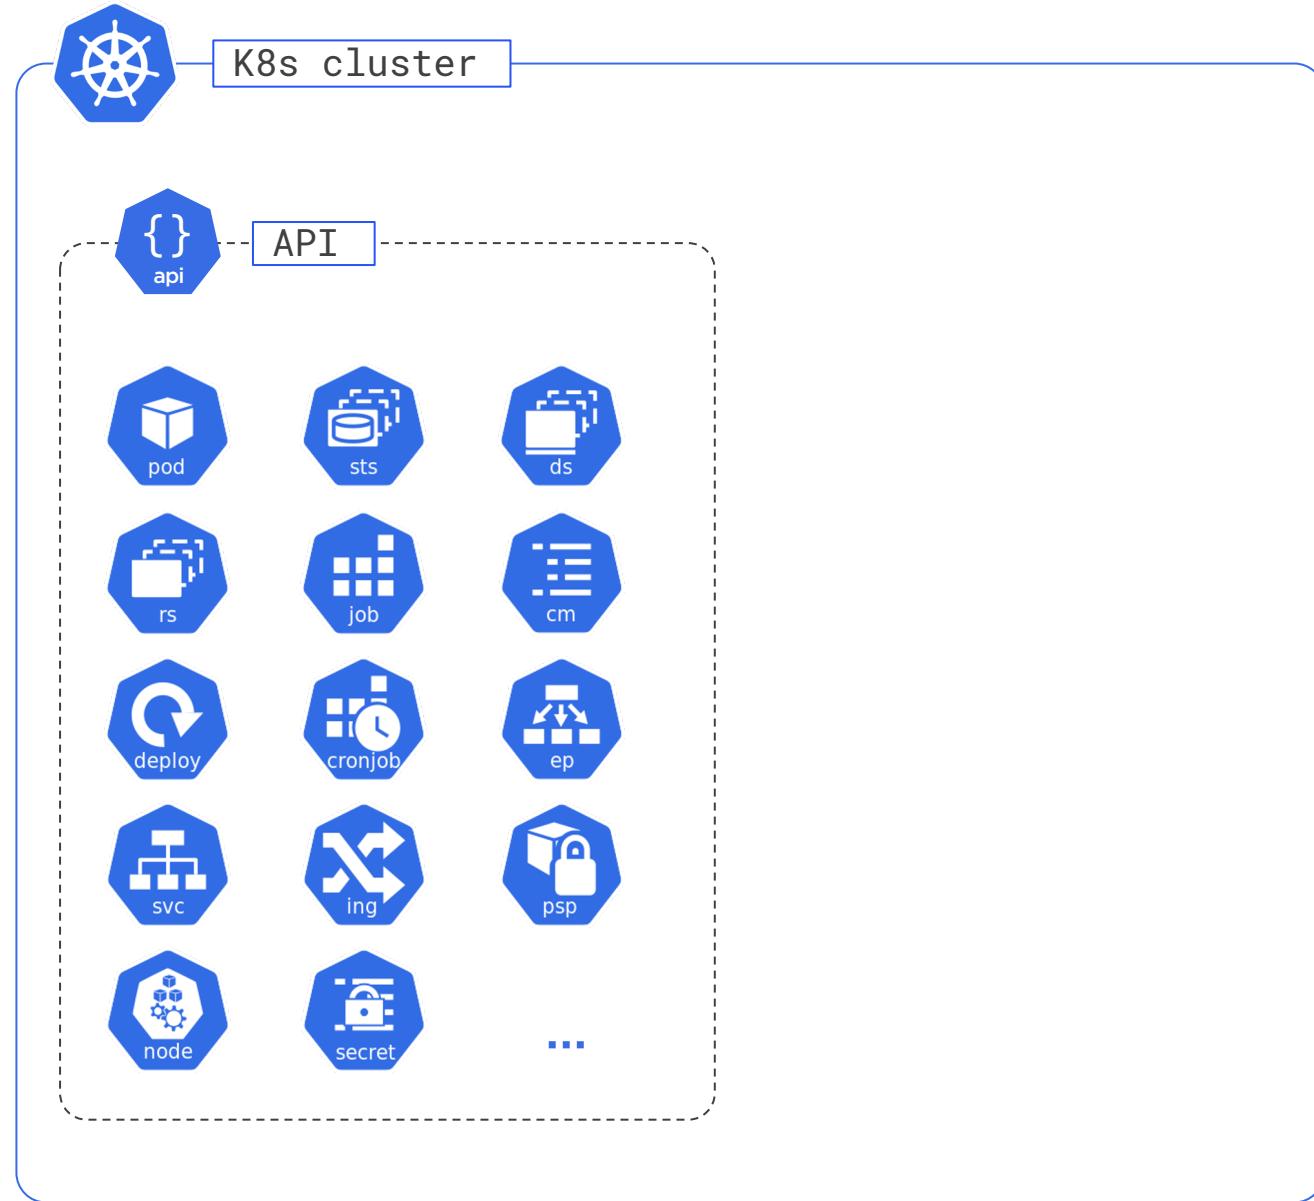


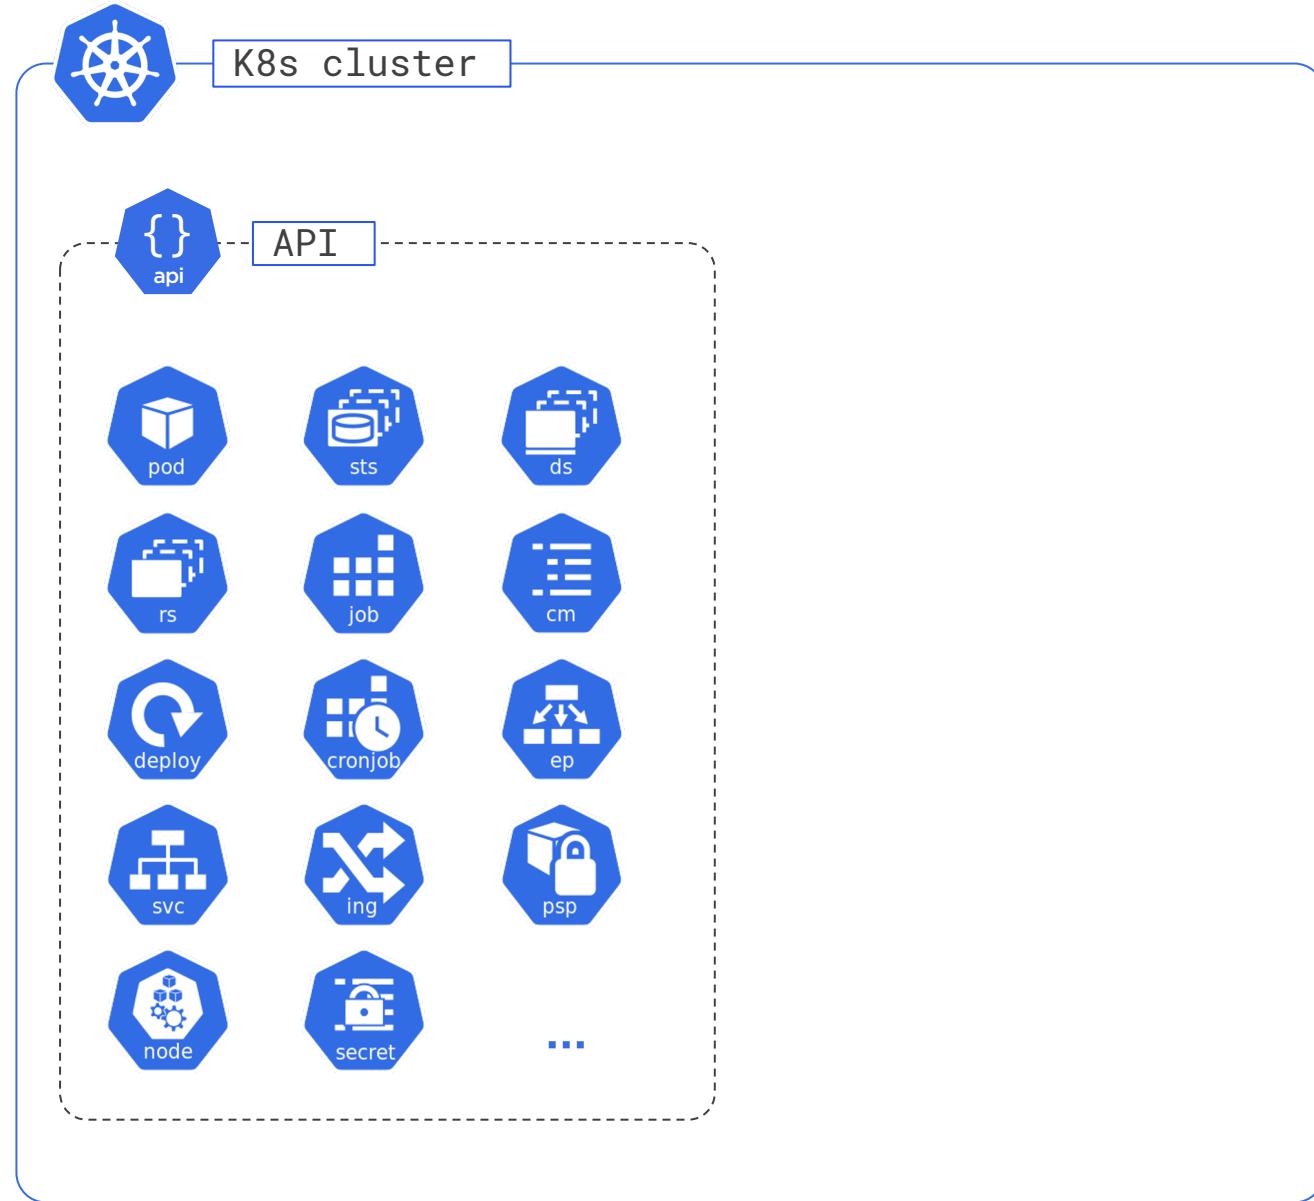
Updates and rollbacks

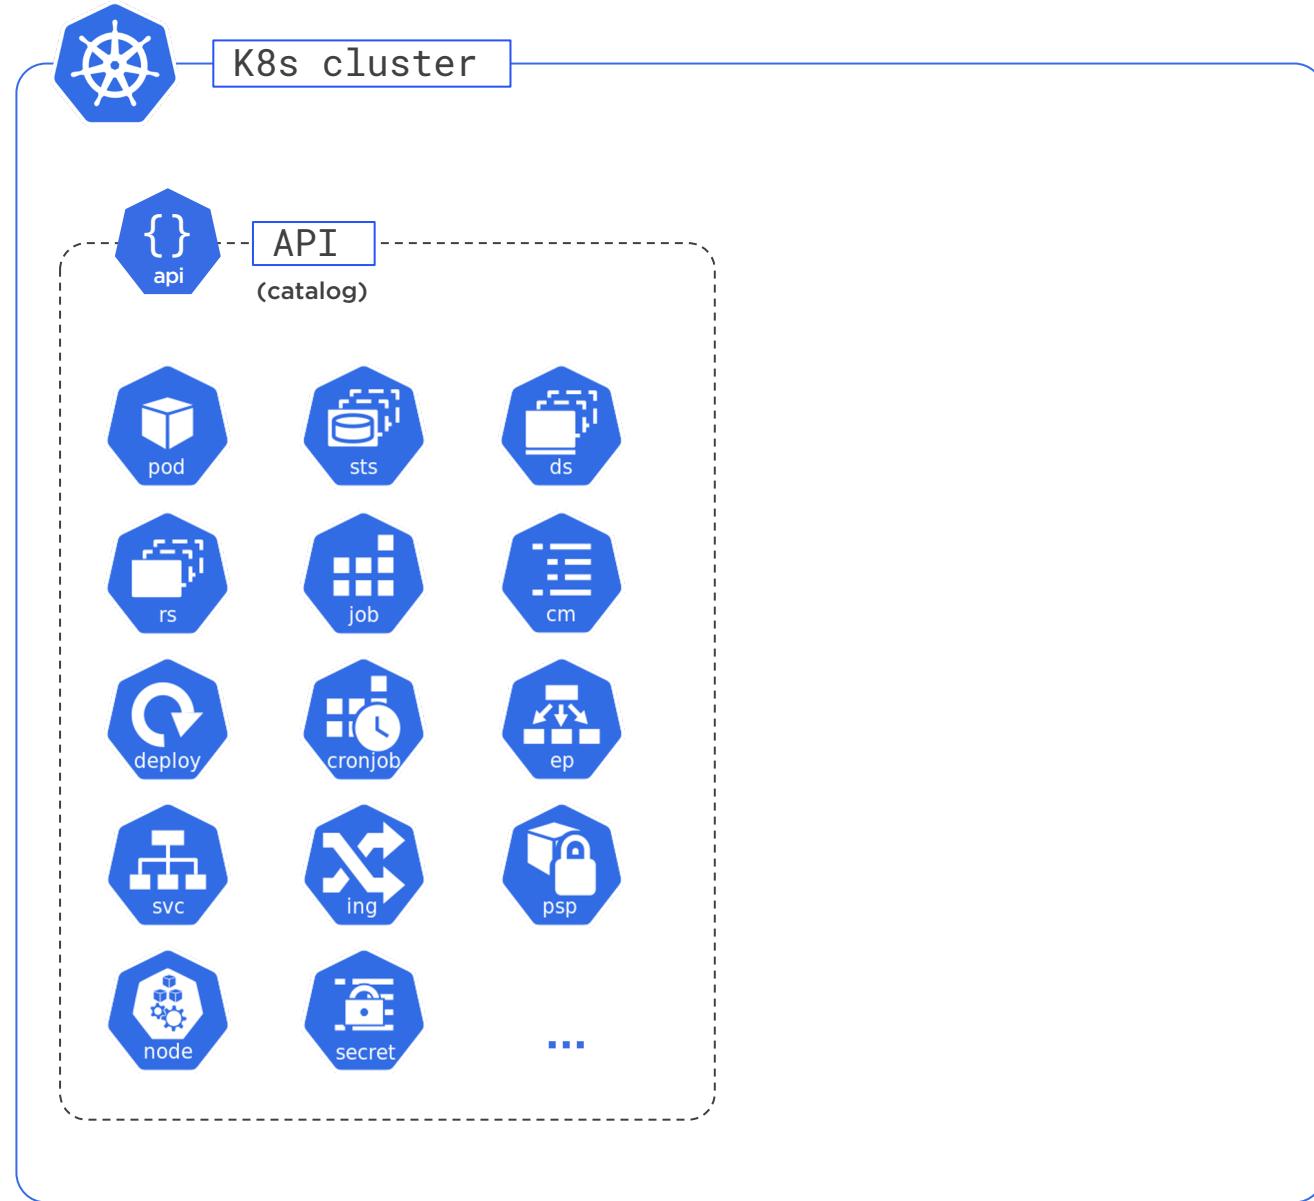


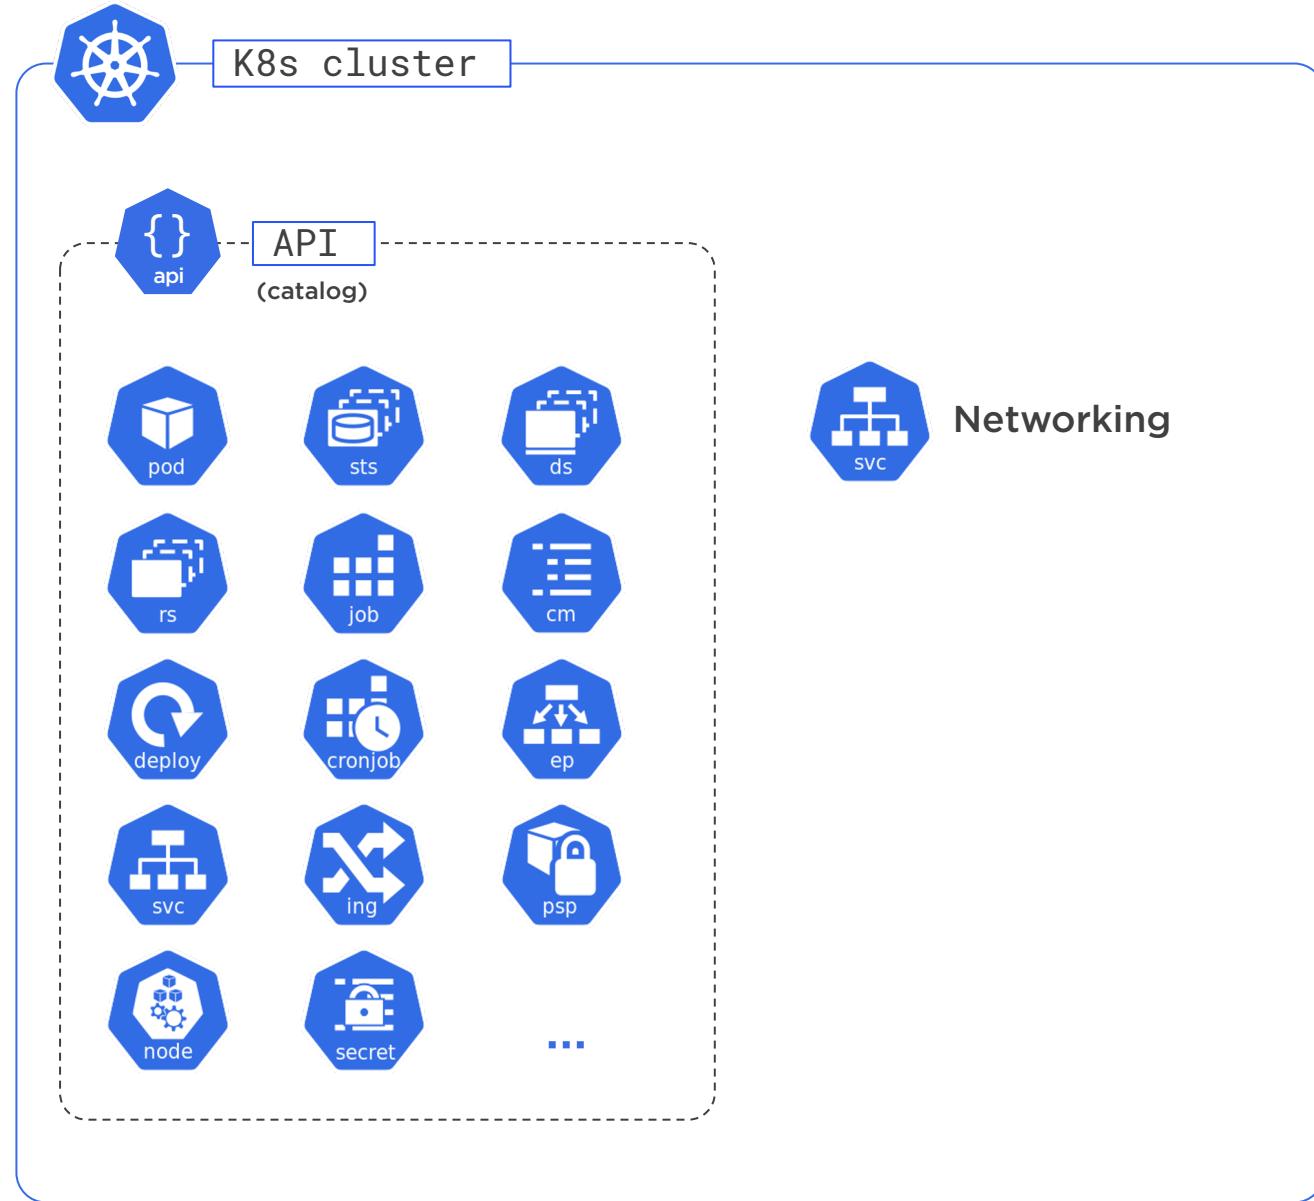
Stable network abstraction

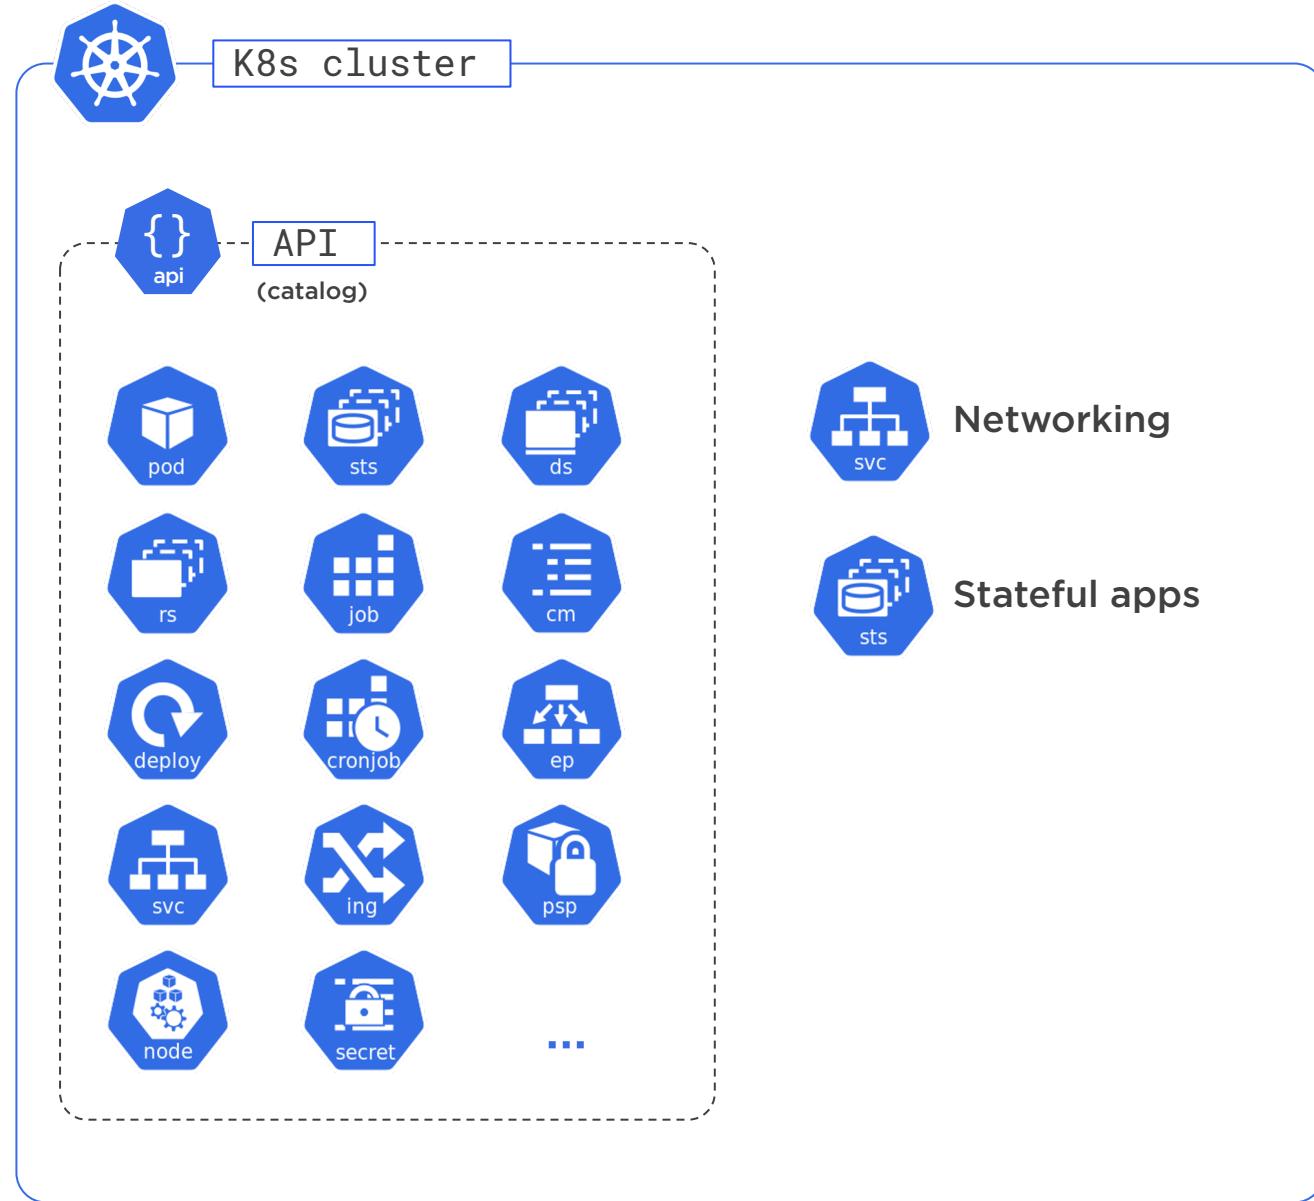


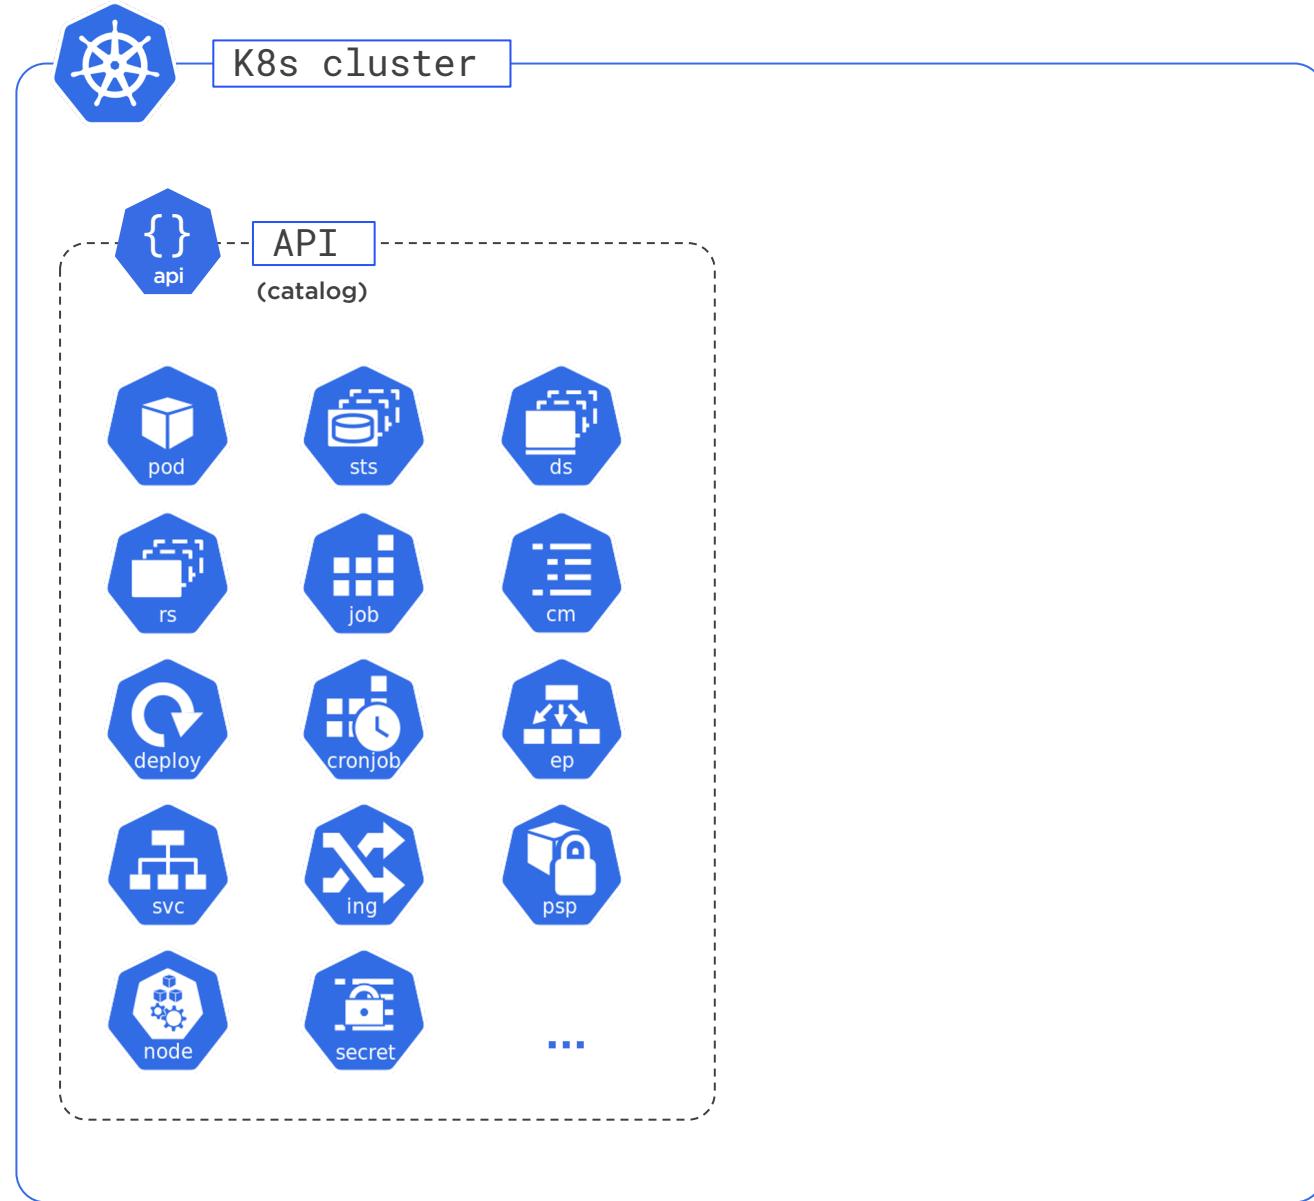


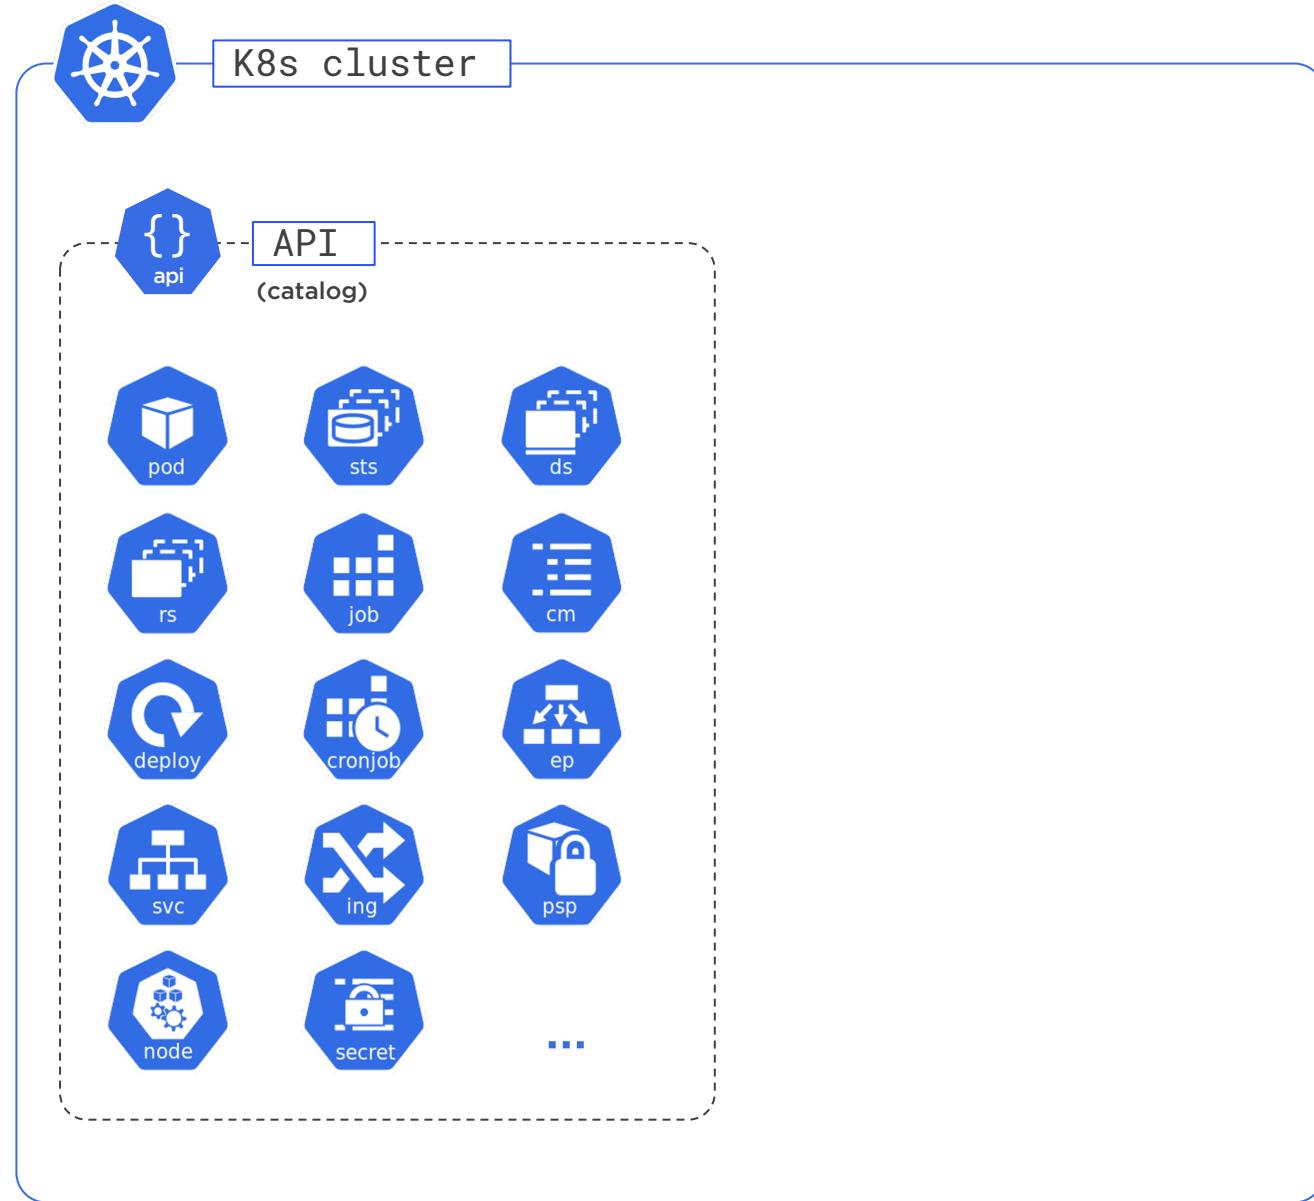


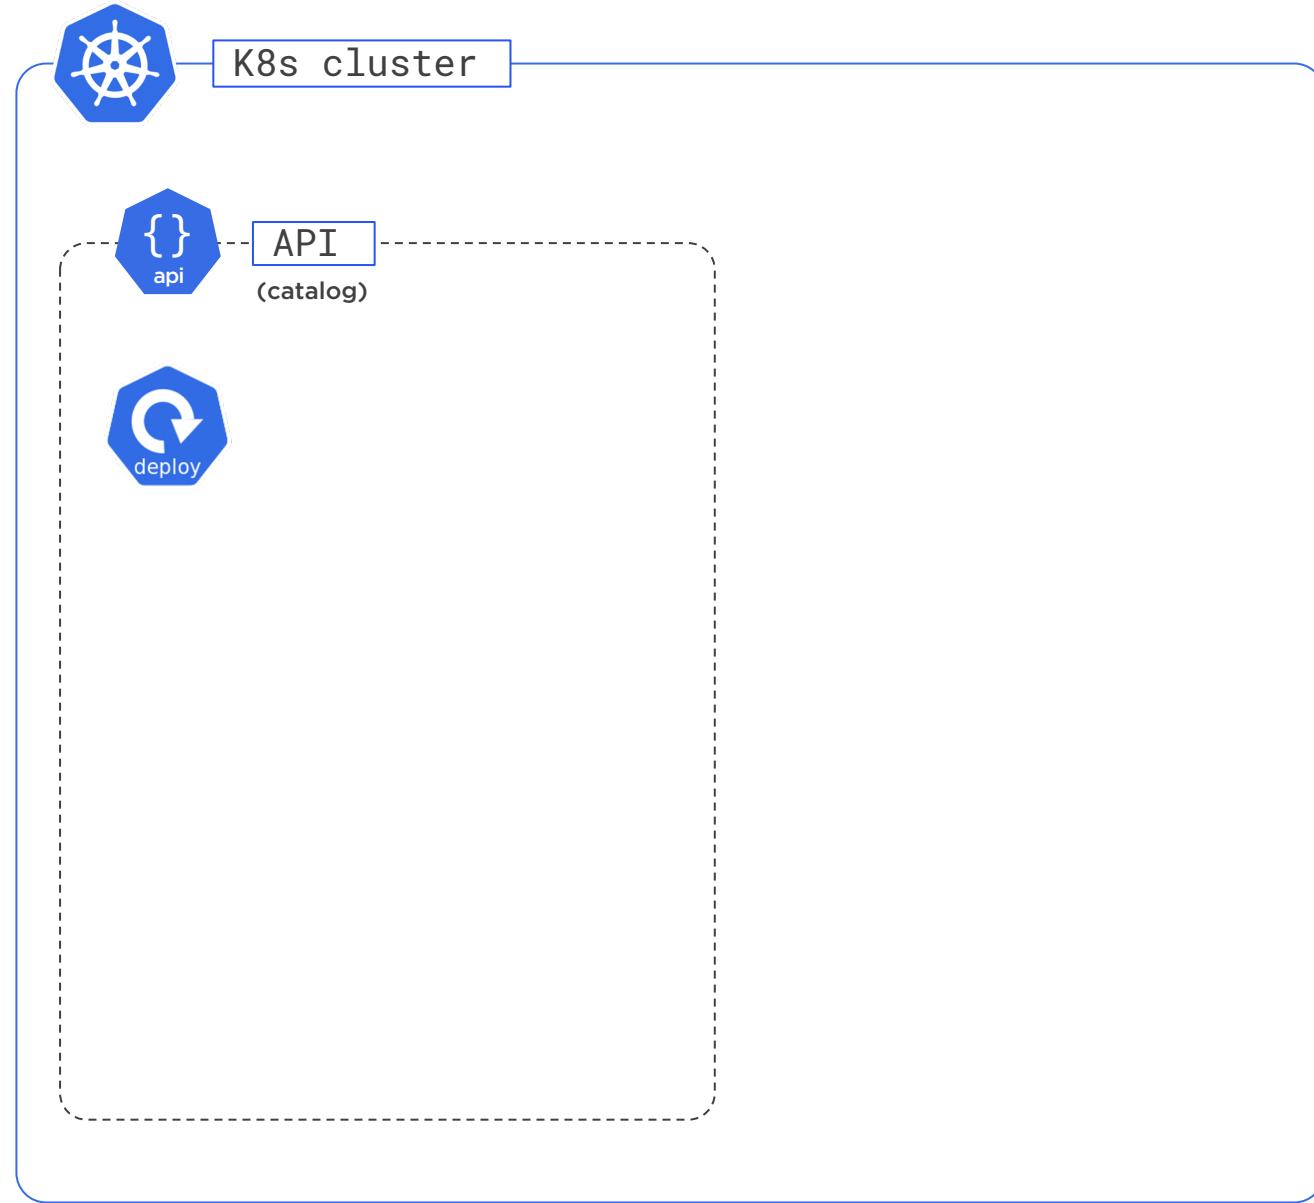


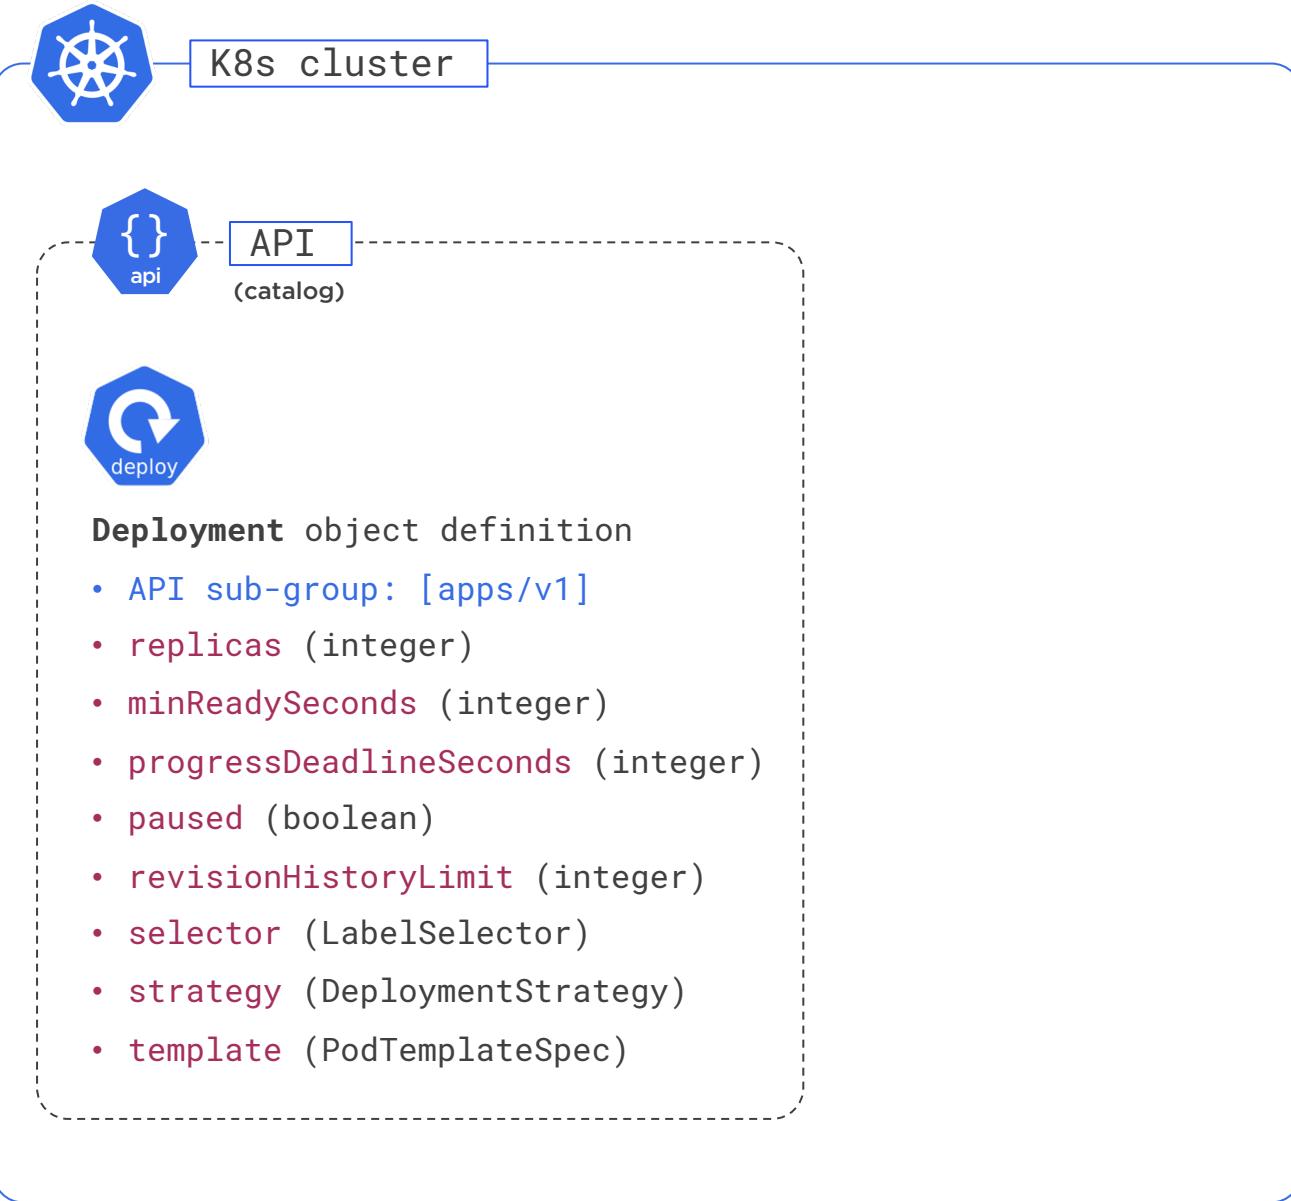


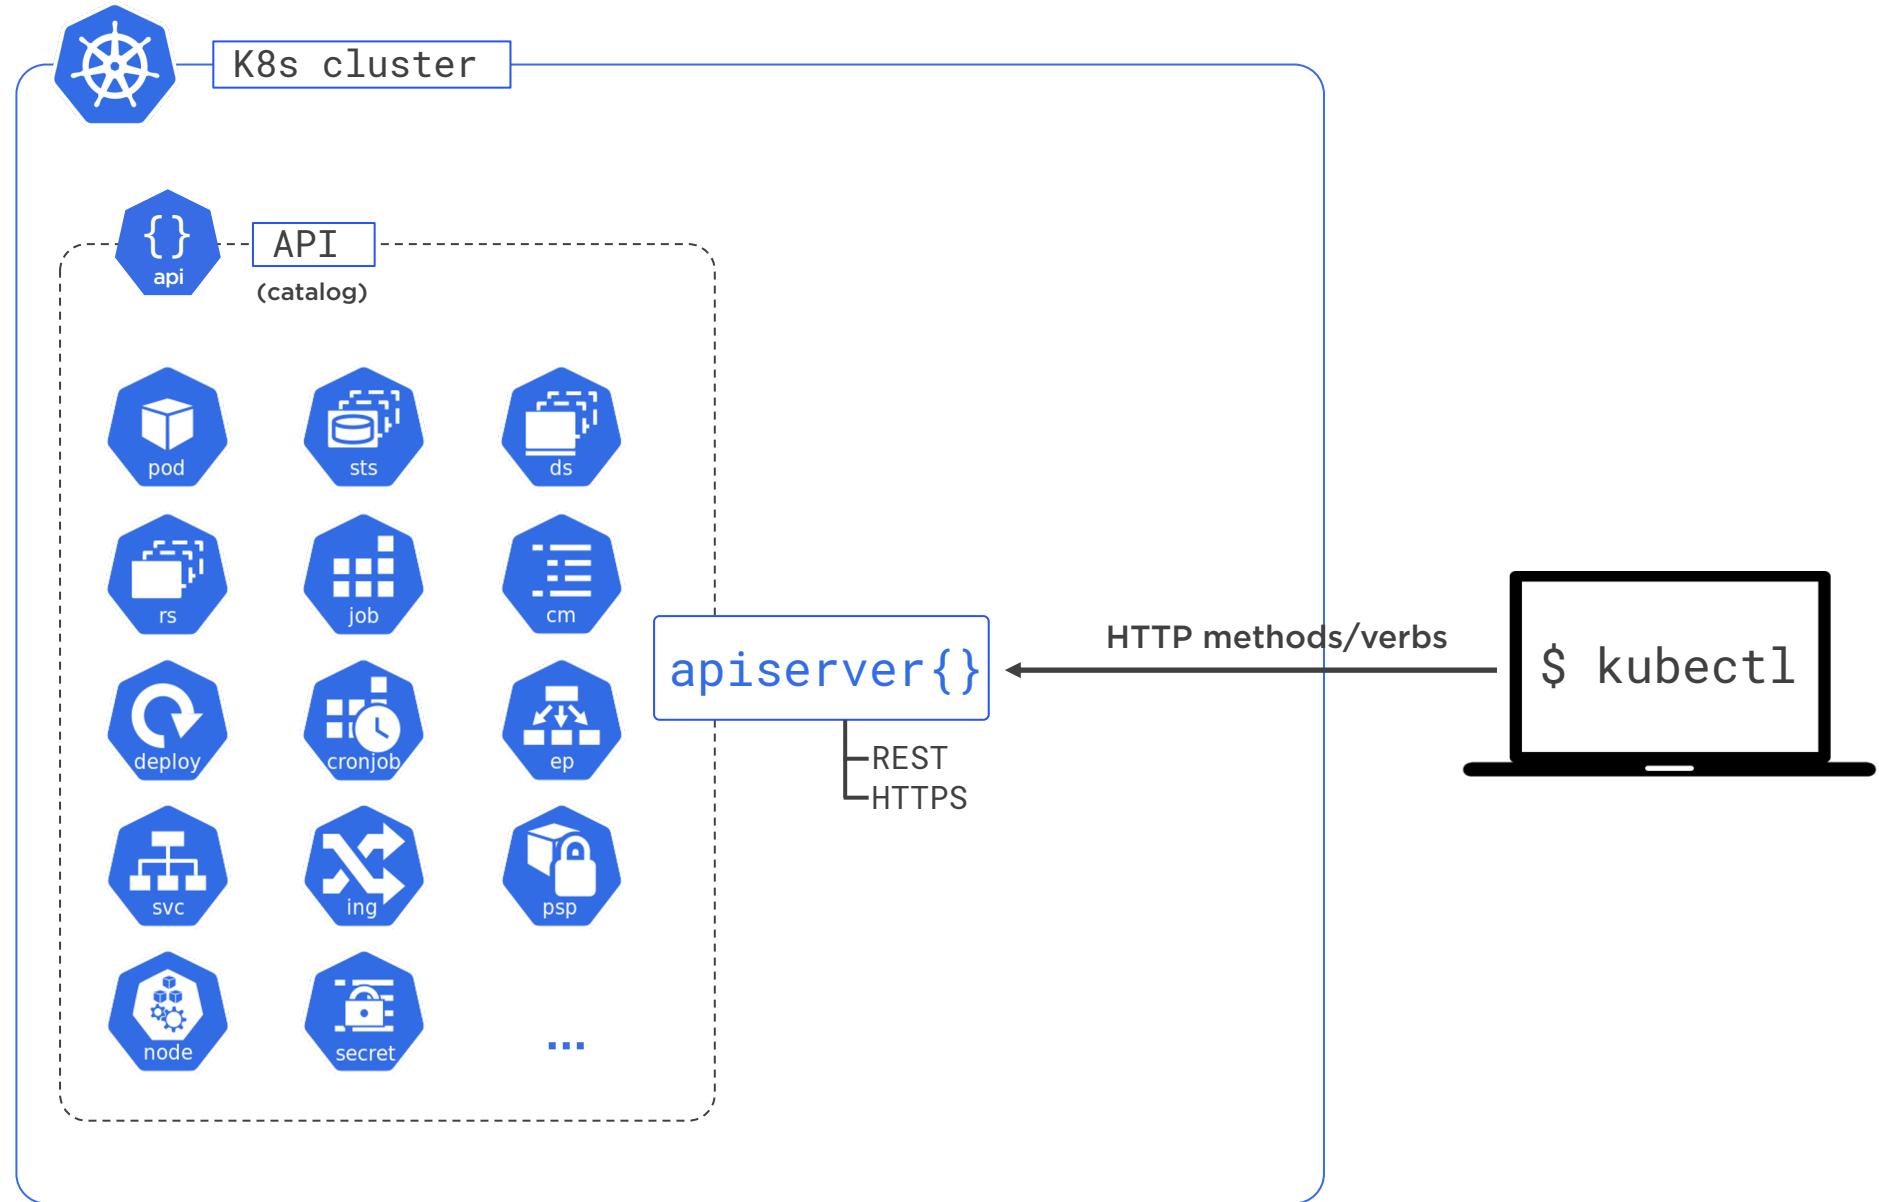


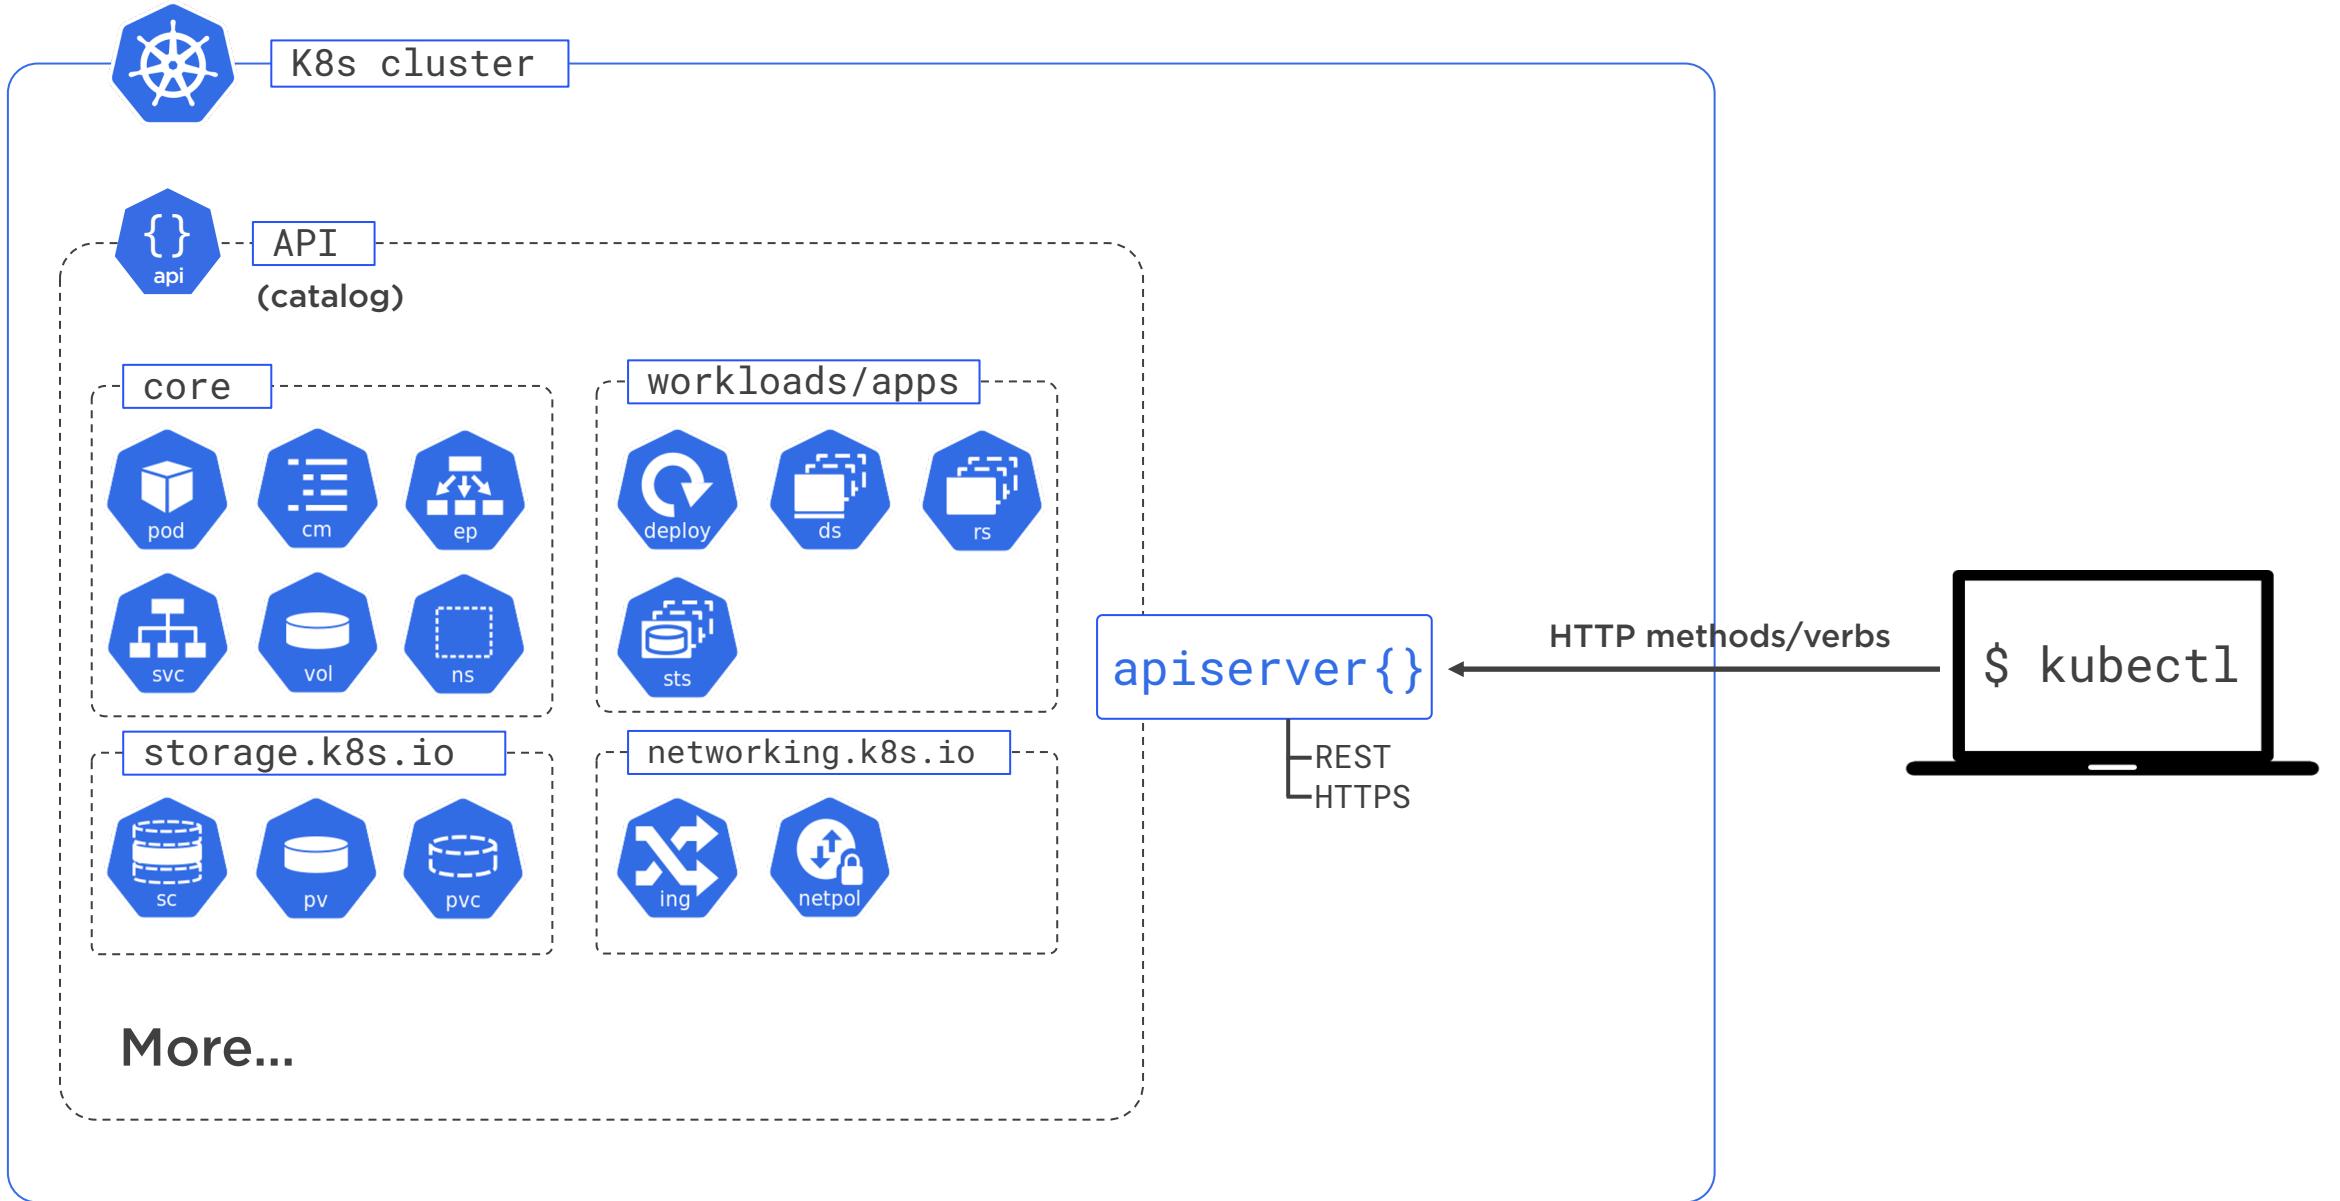










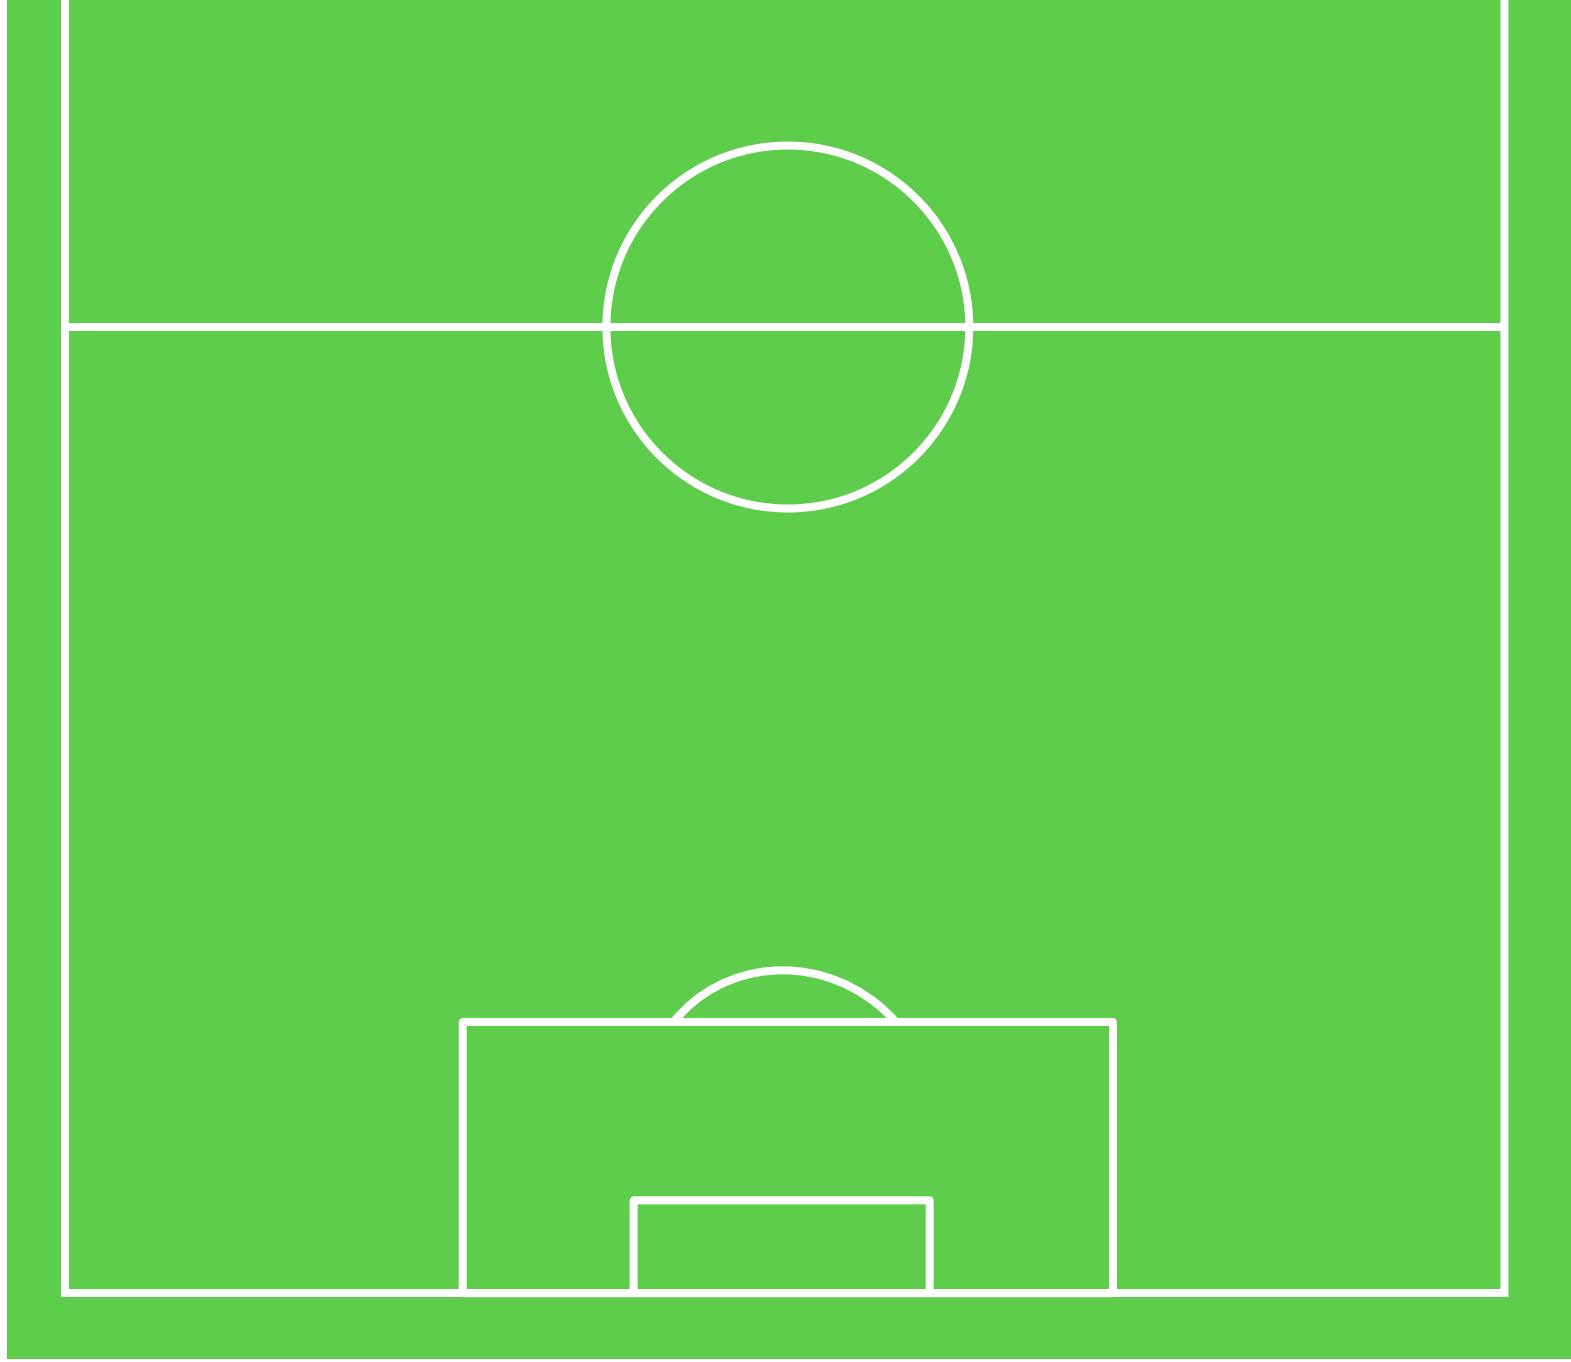


Up Next:
Epic Recap

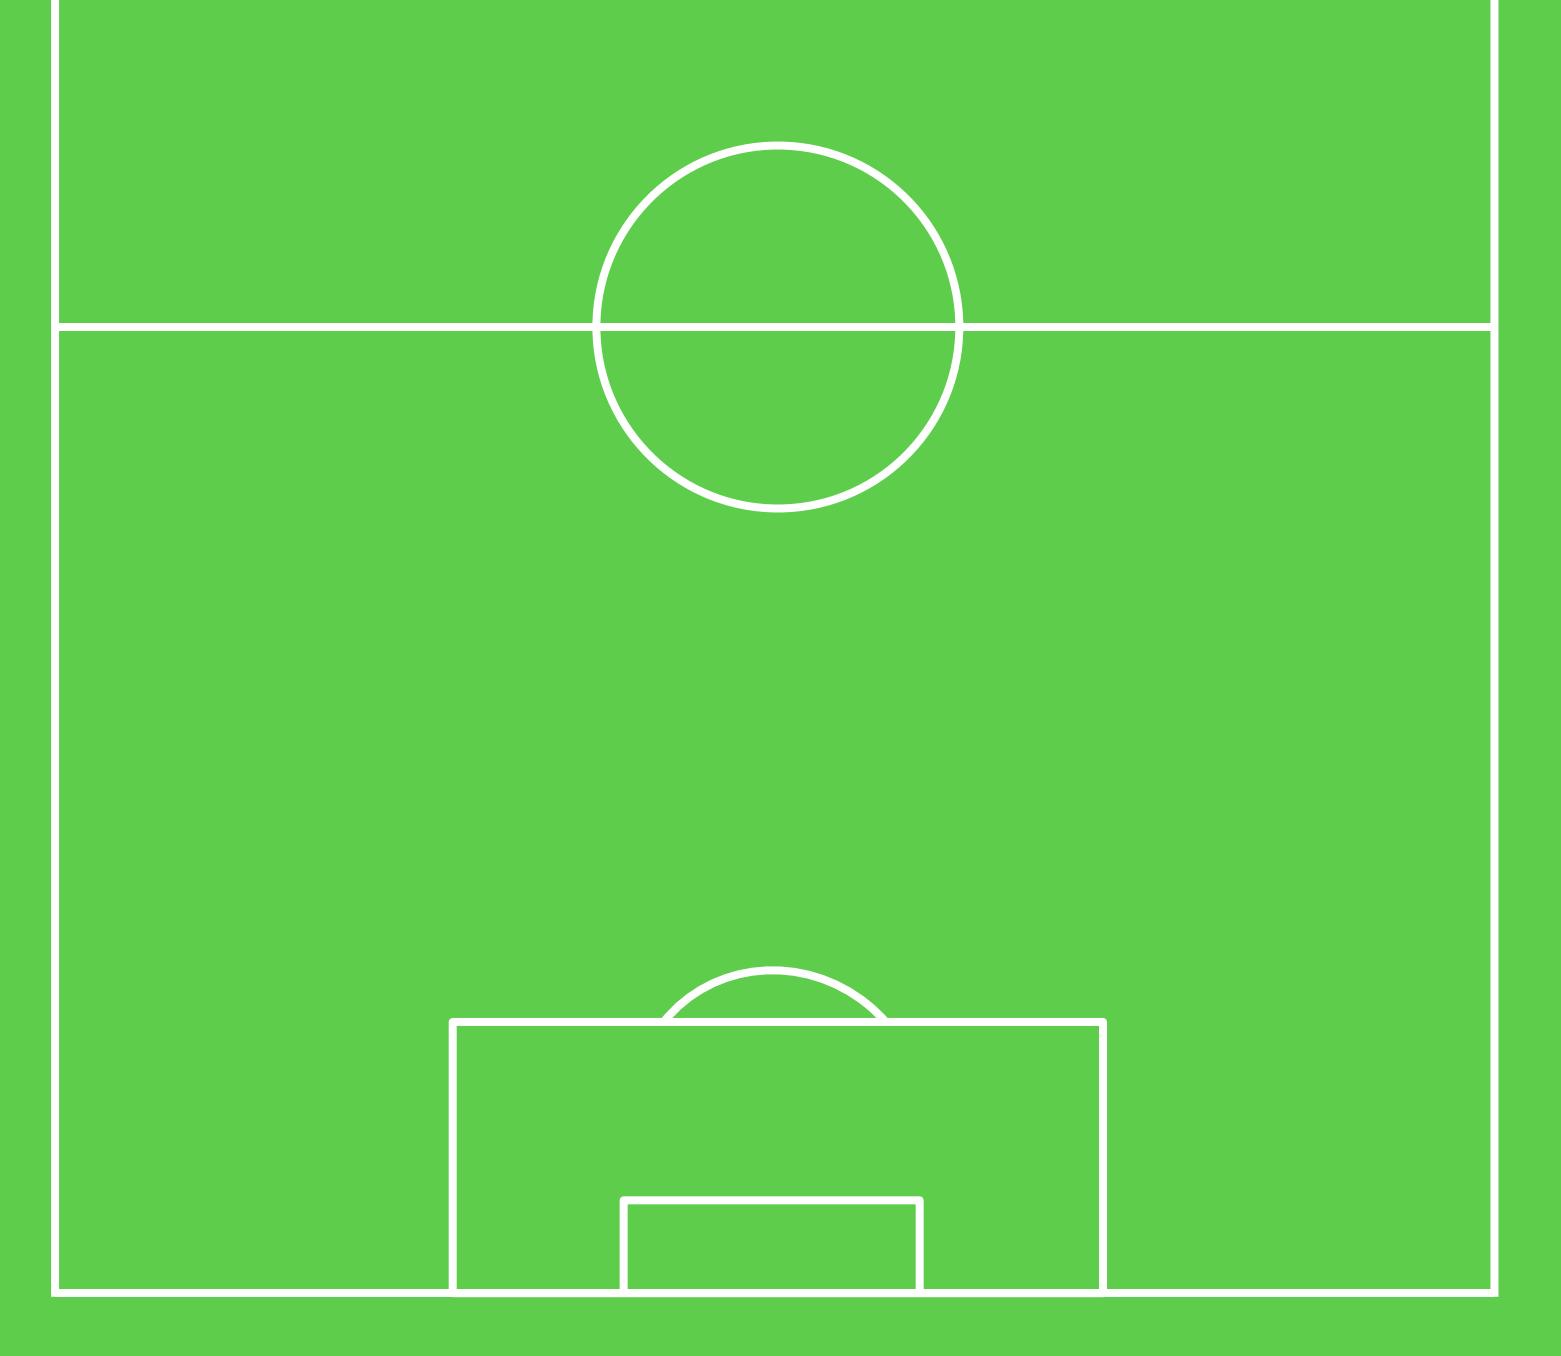
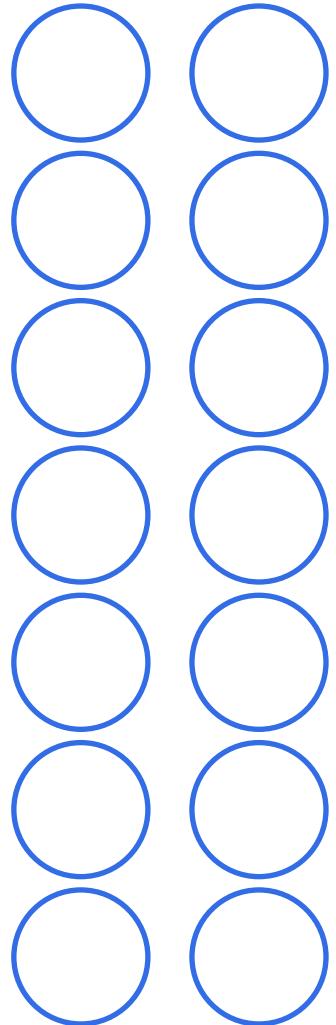


Epic Recap

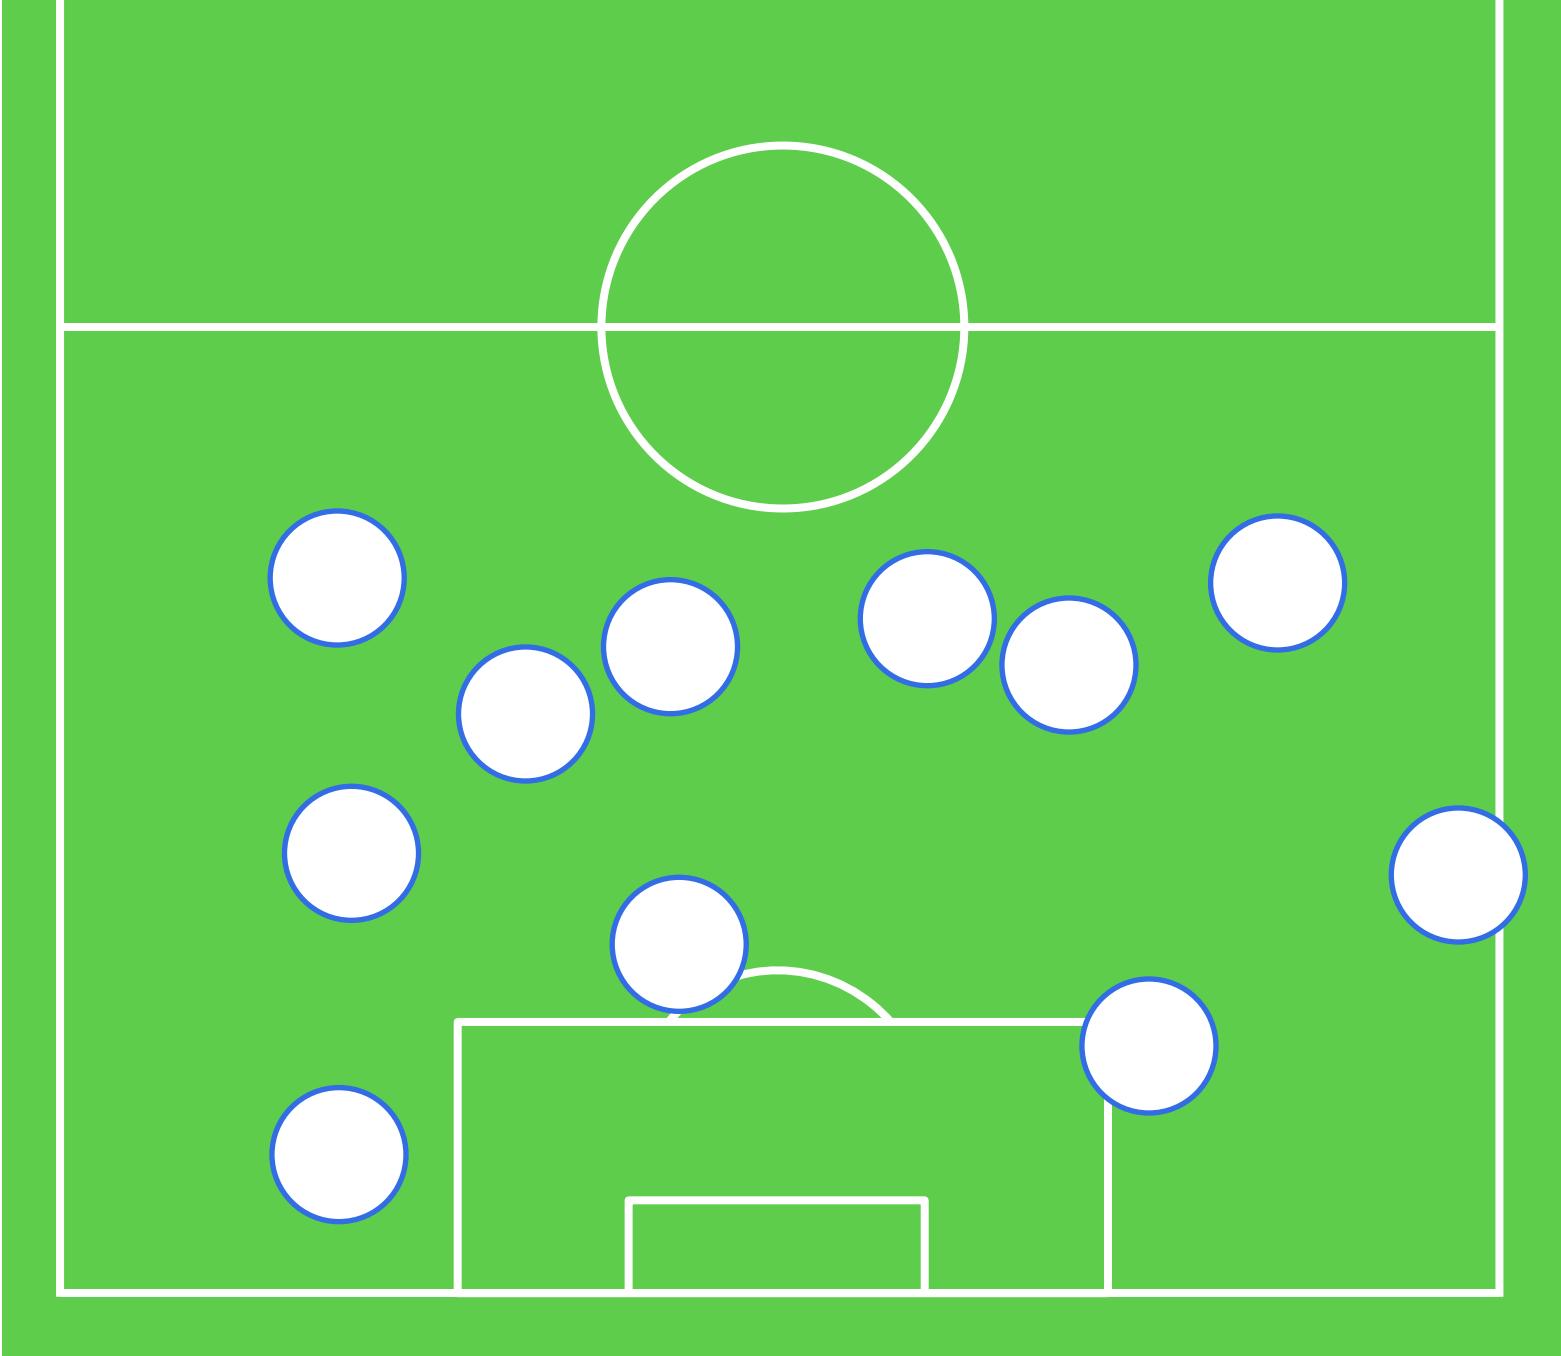
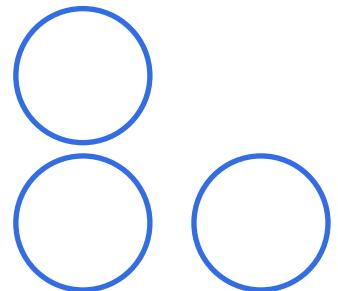




Team



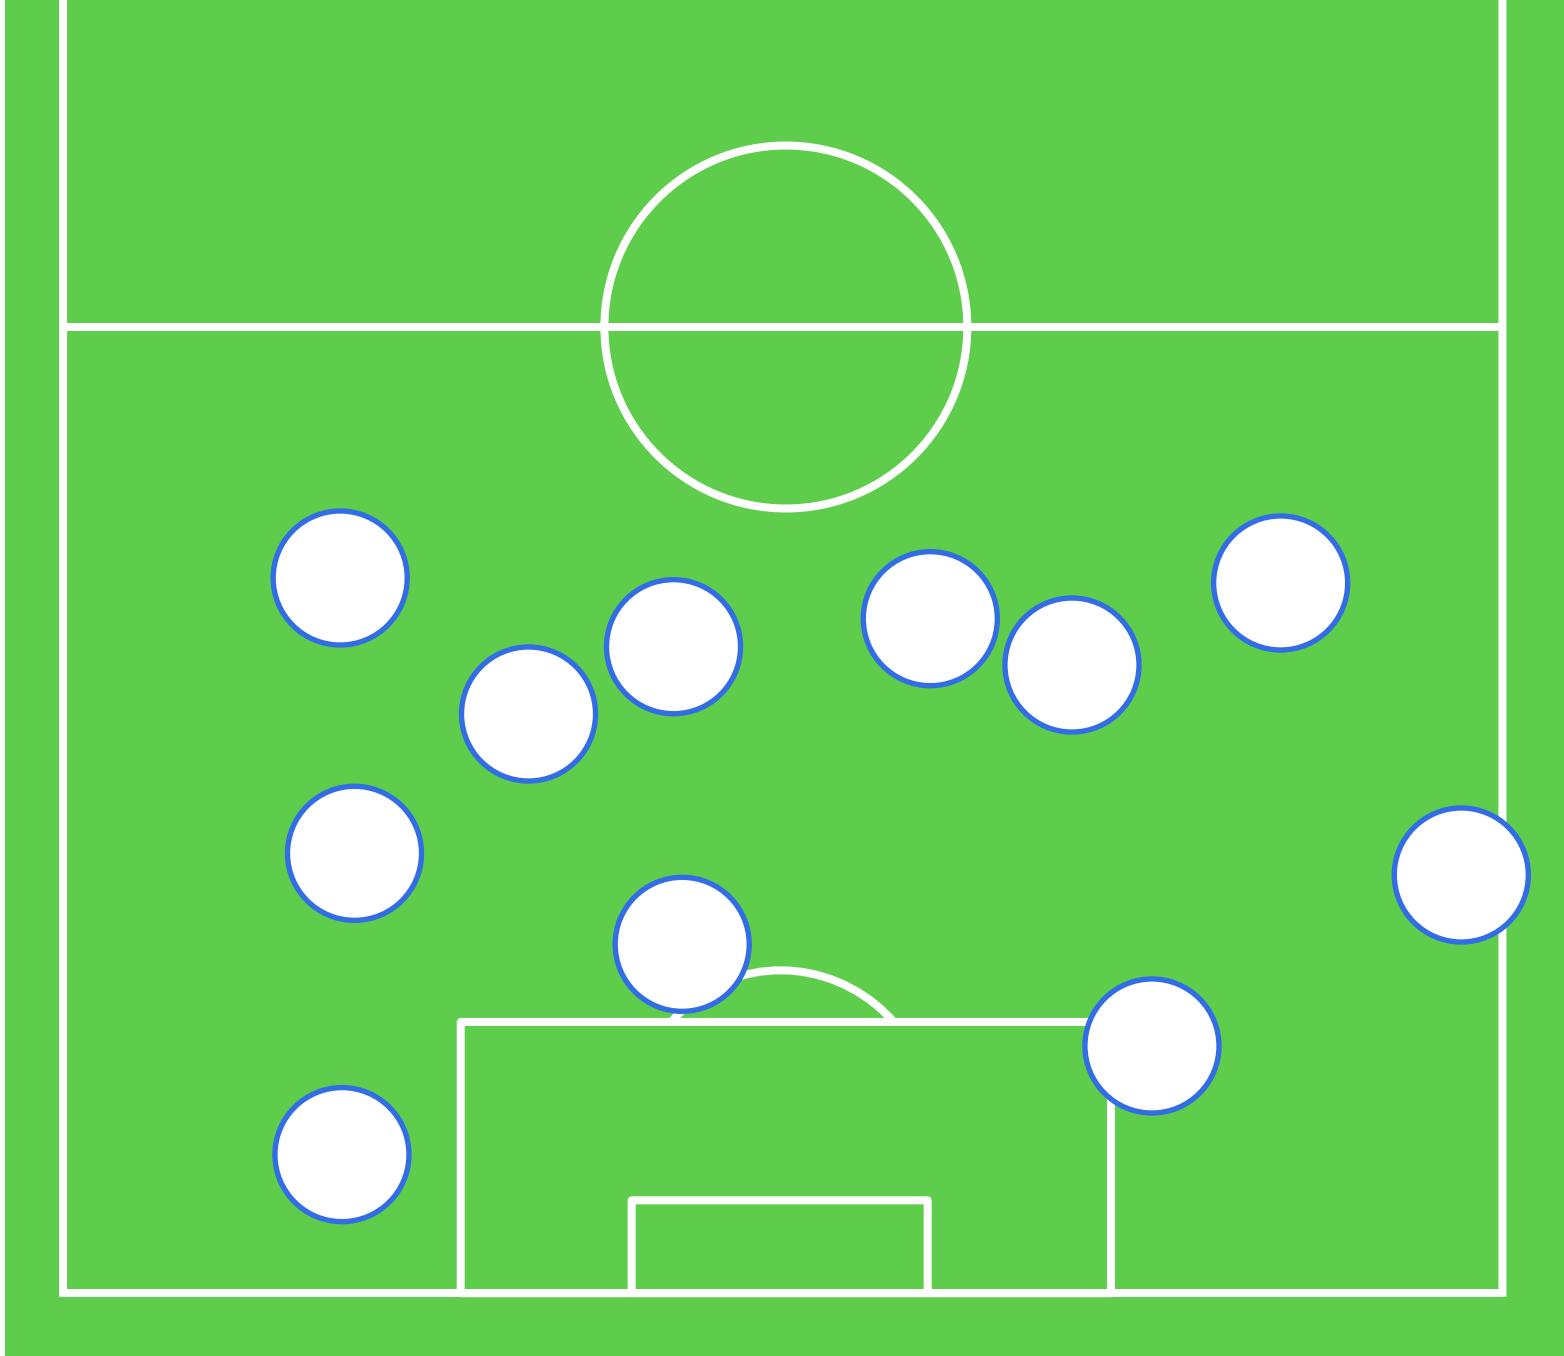
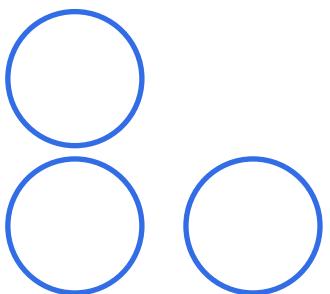
Team



Team



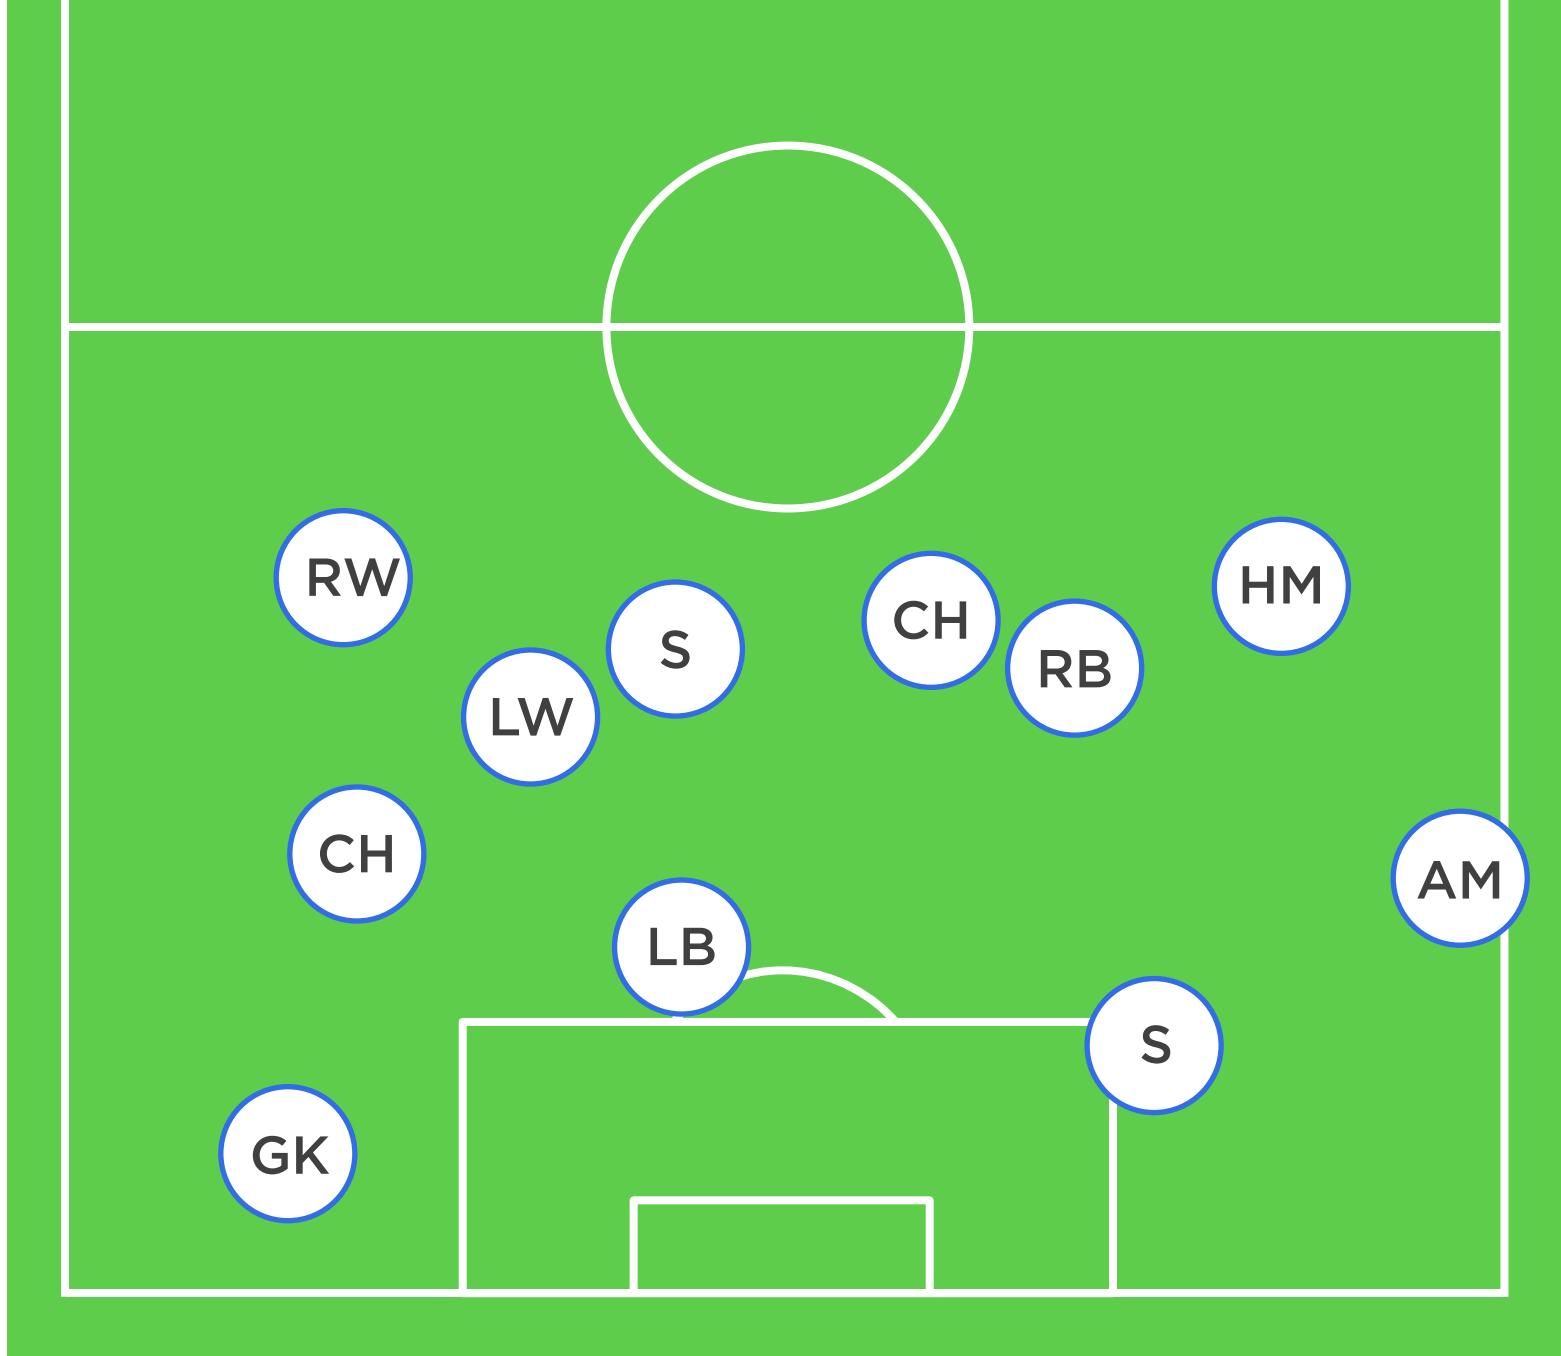
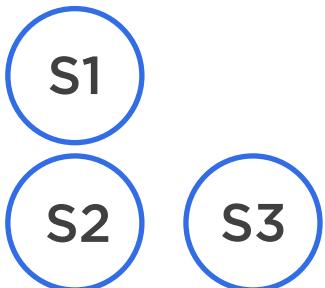
Manager
(coach)



Team



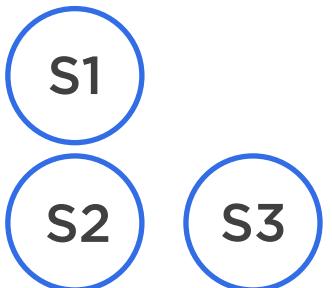
Manager
(coach)

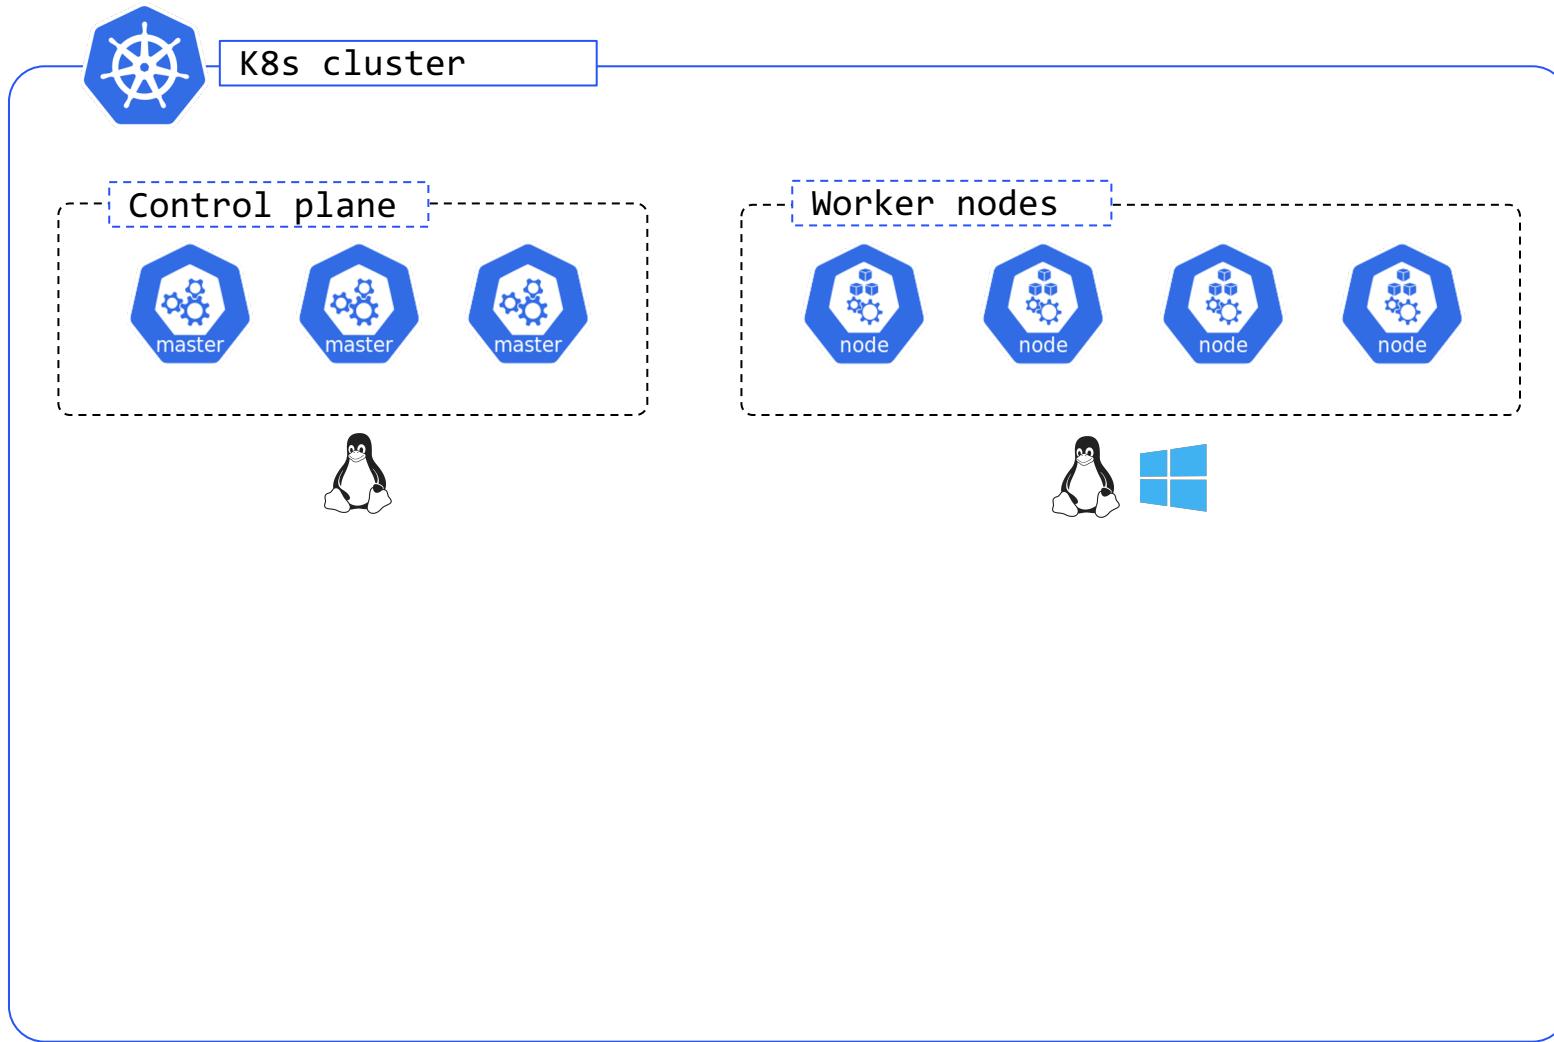


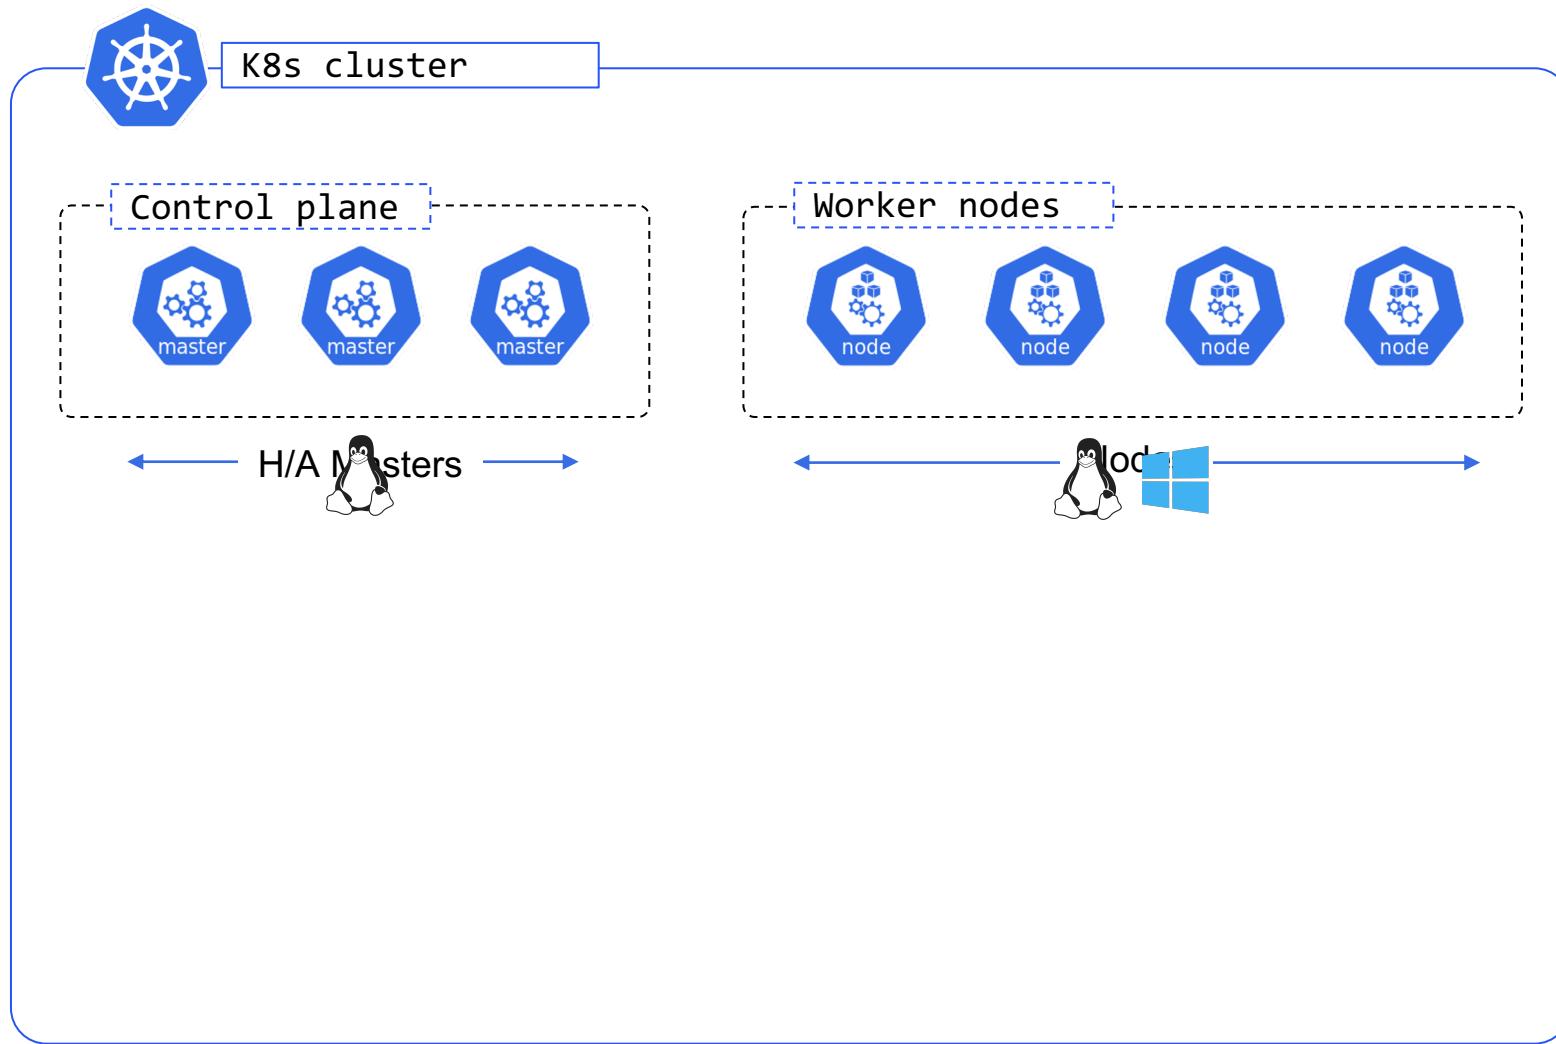
Team

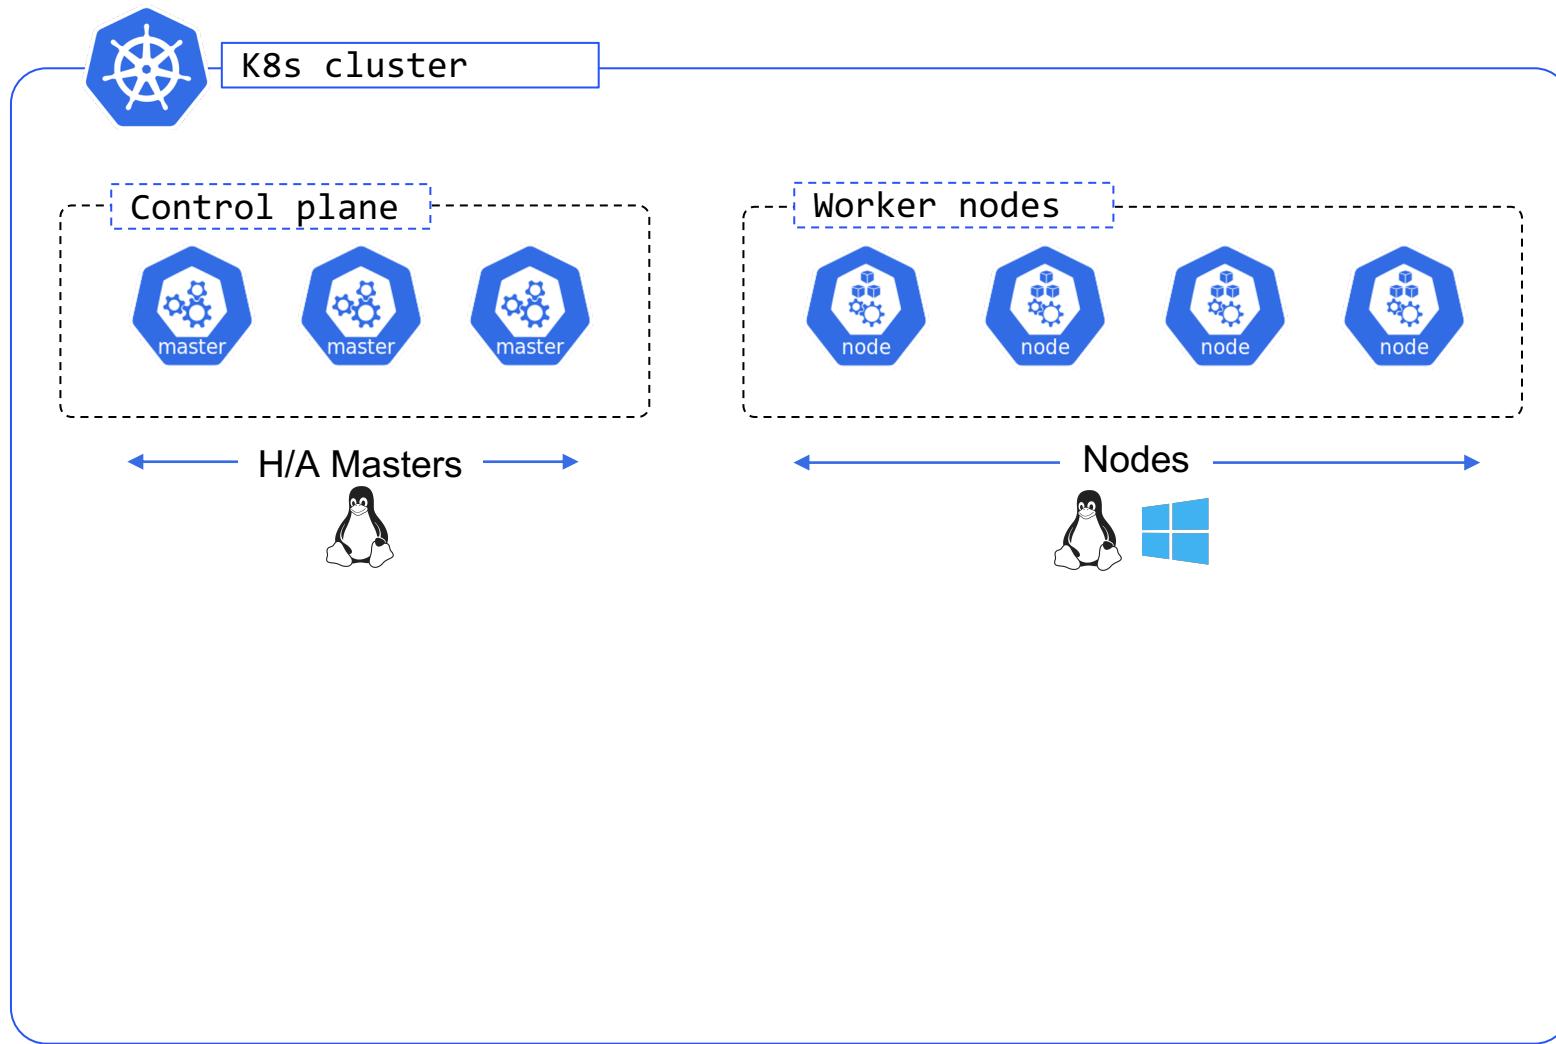


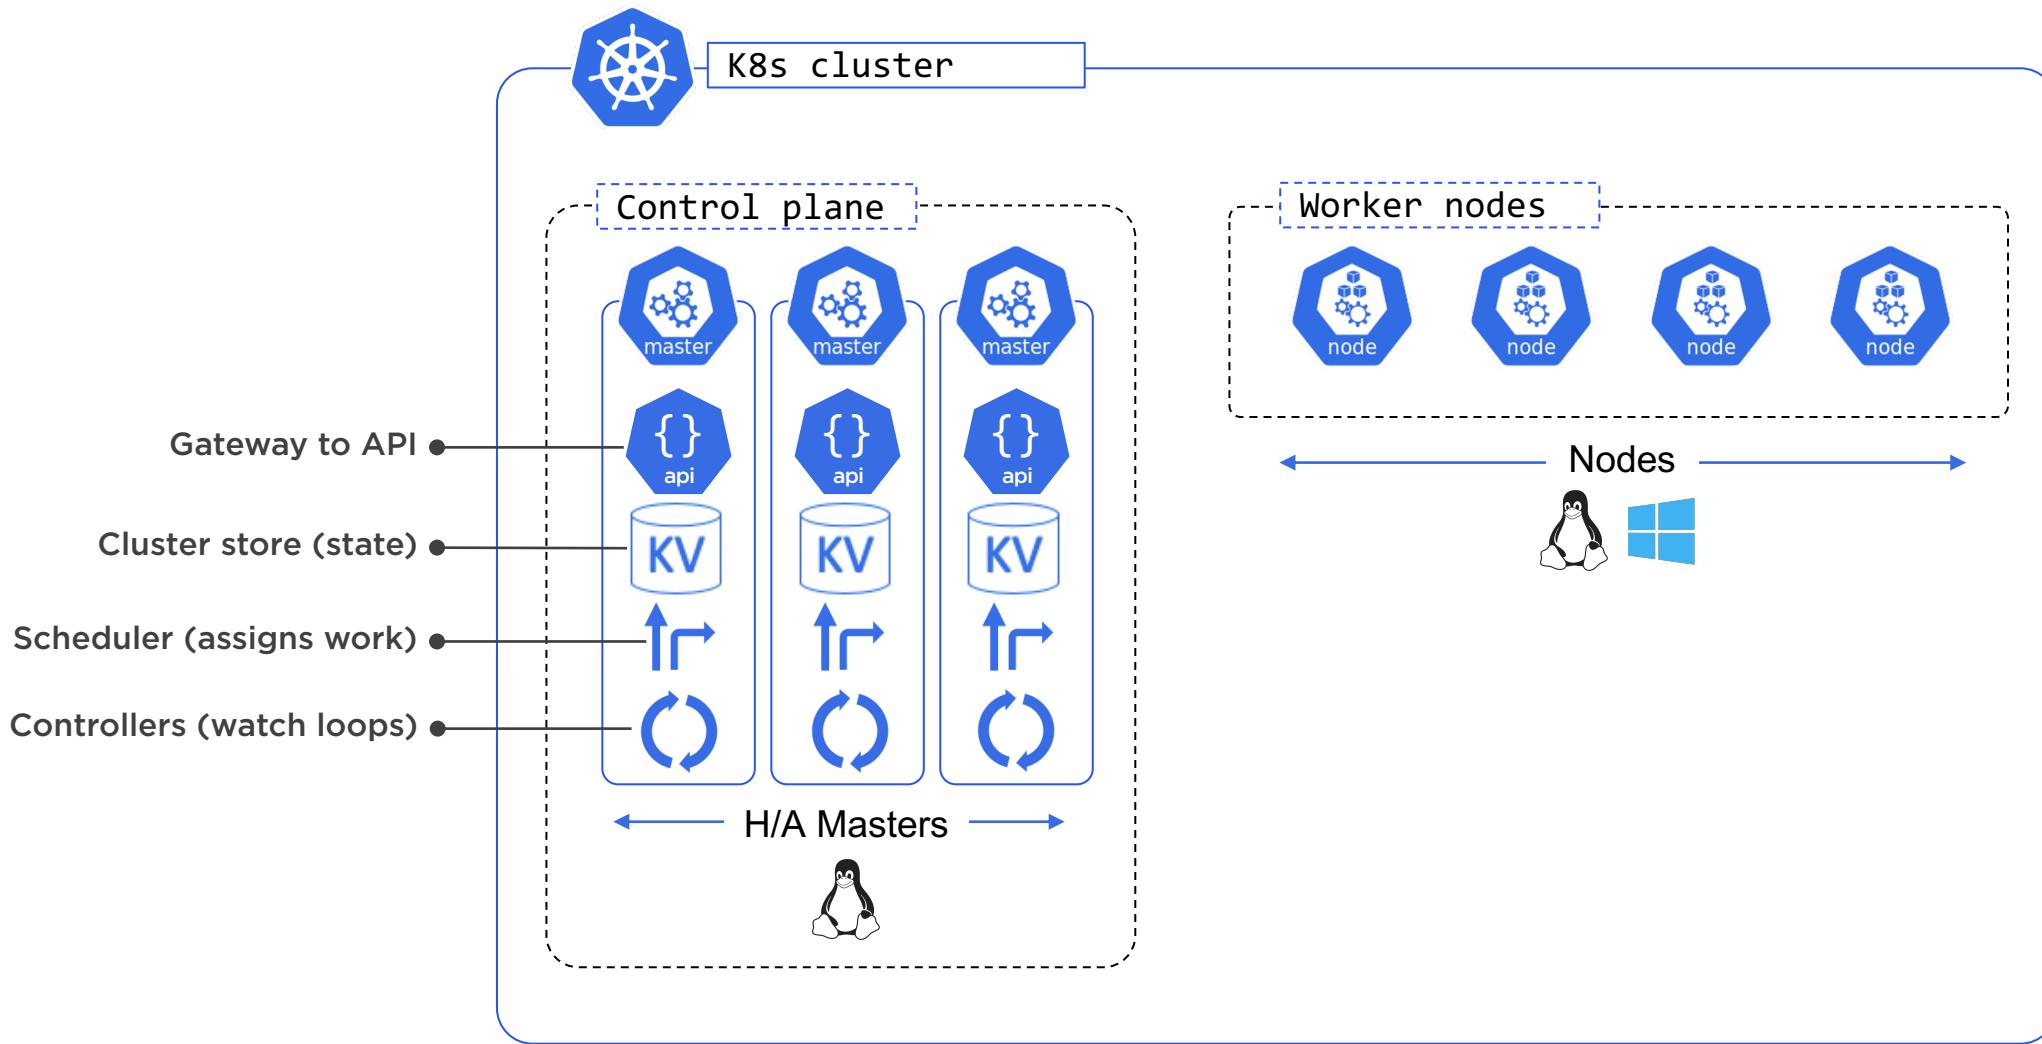
Manager
(coach)

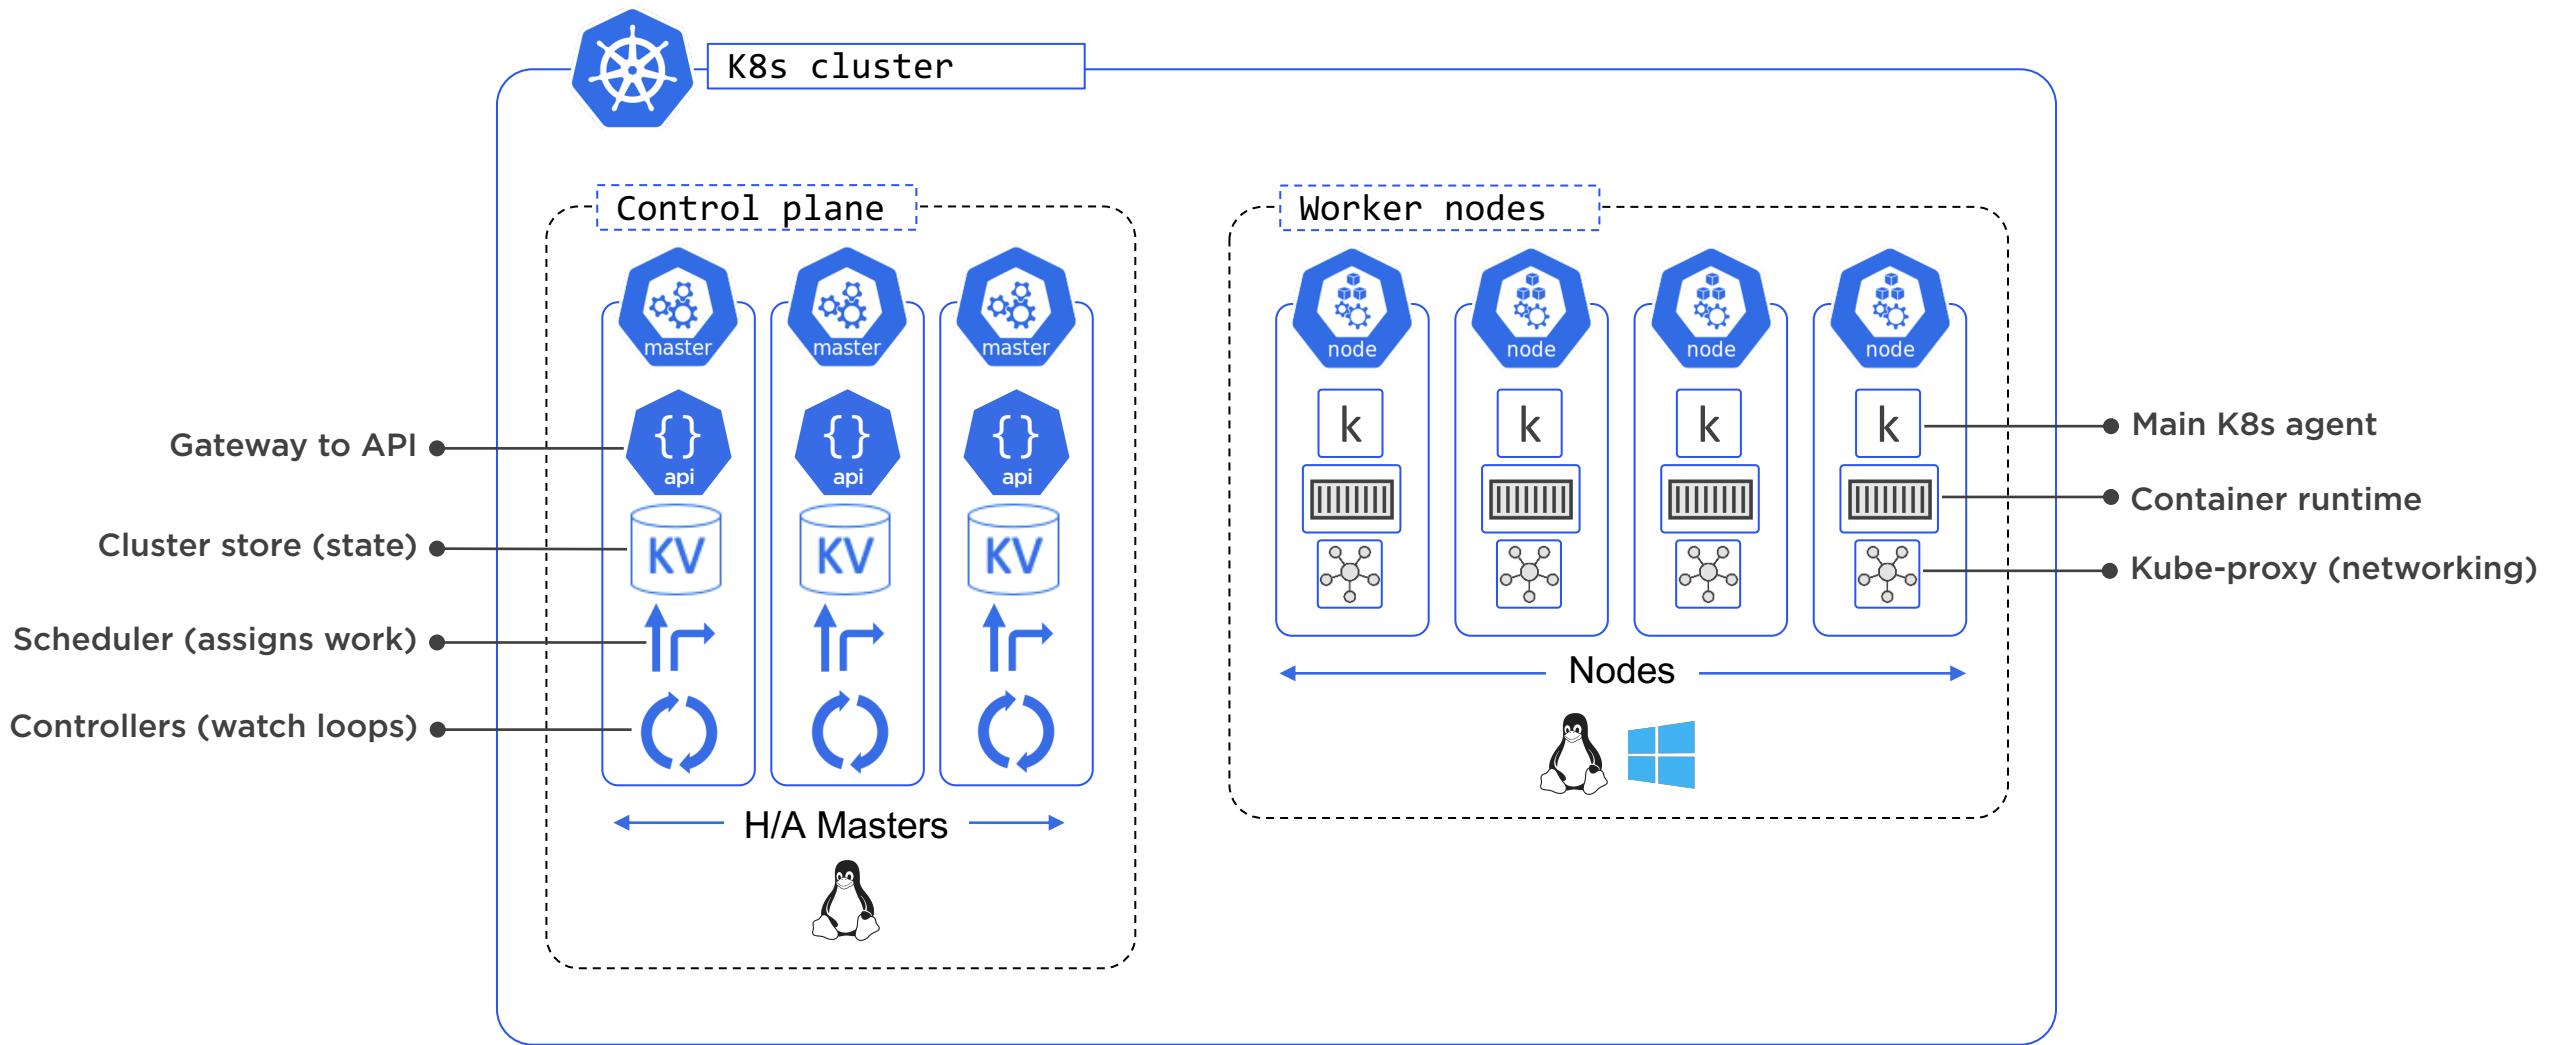


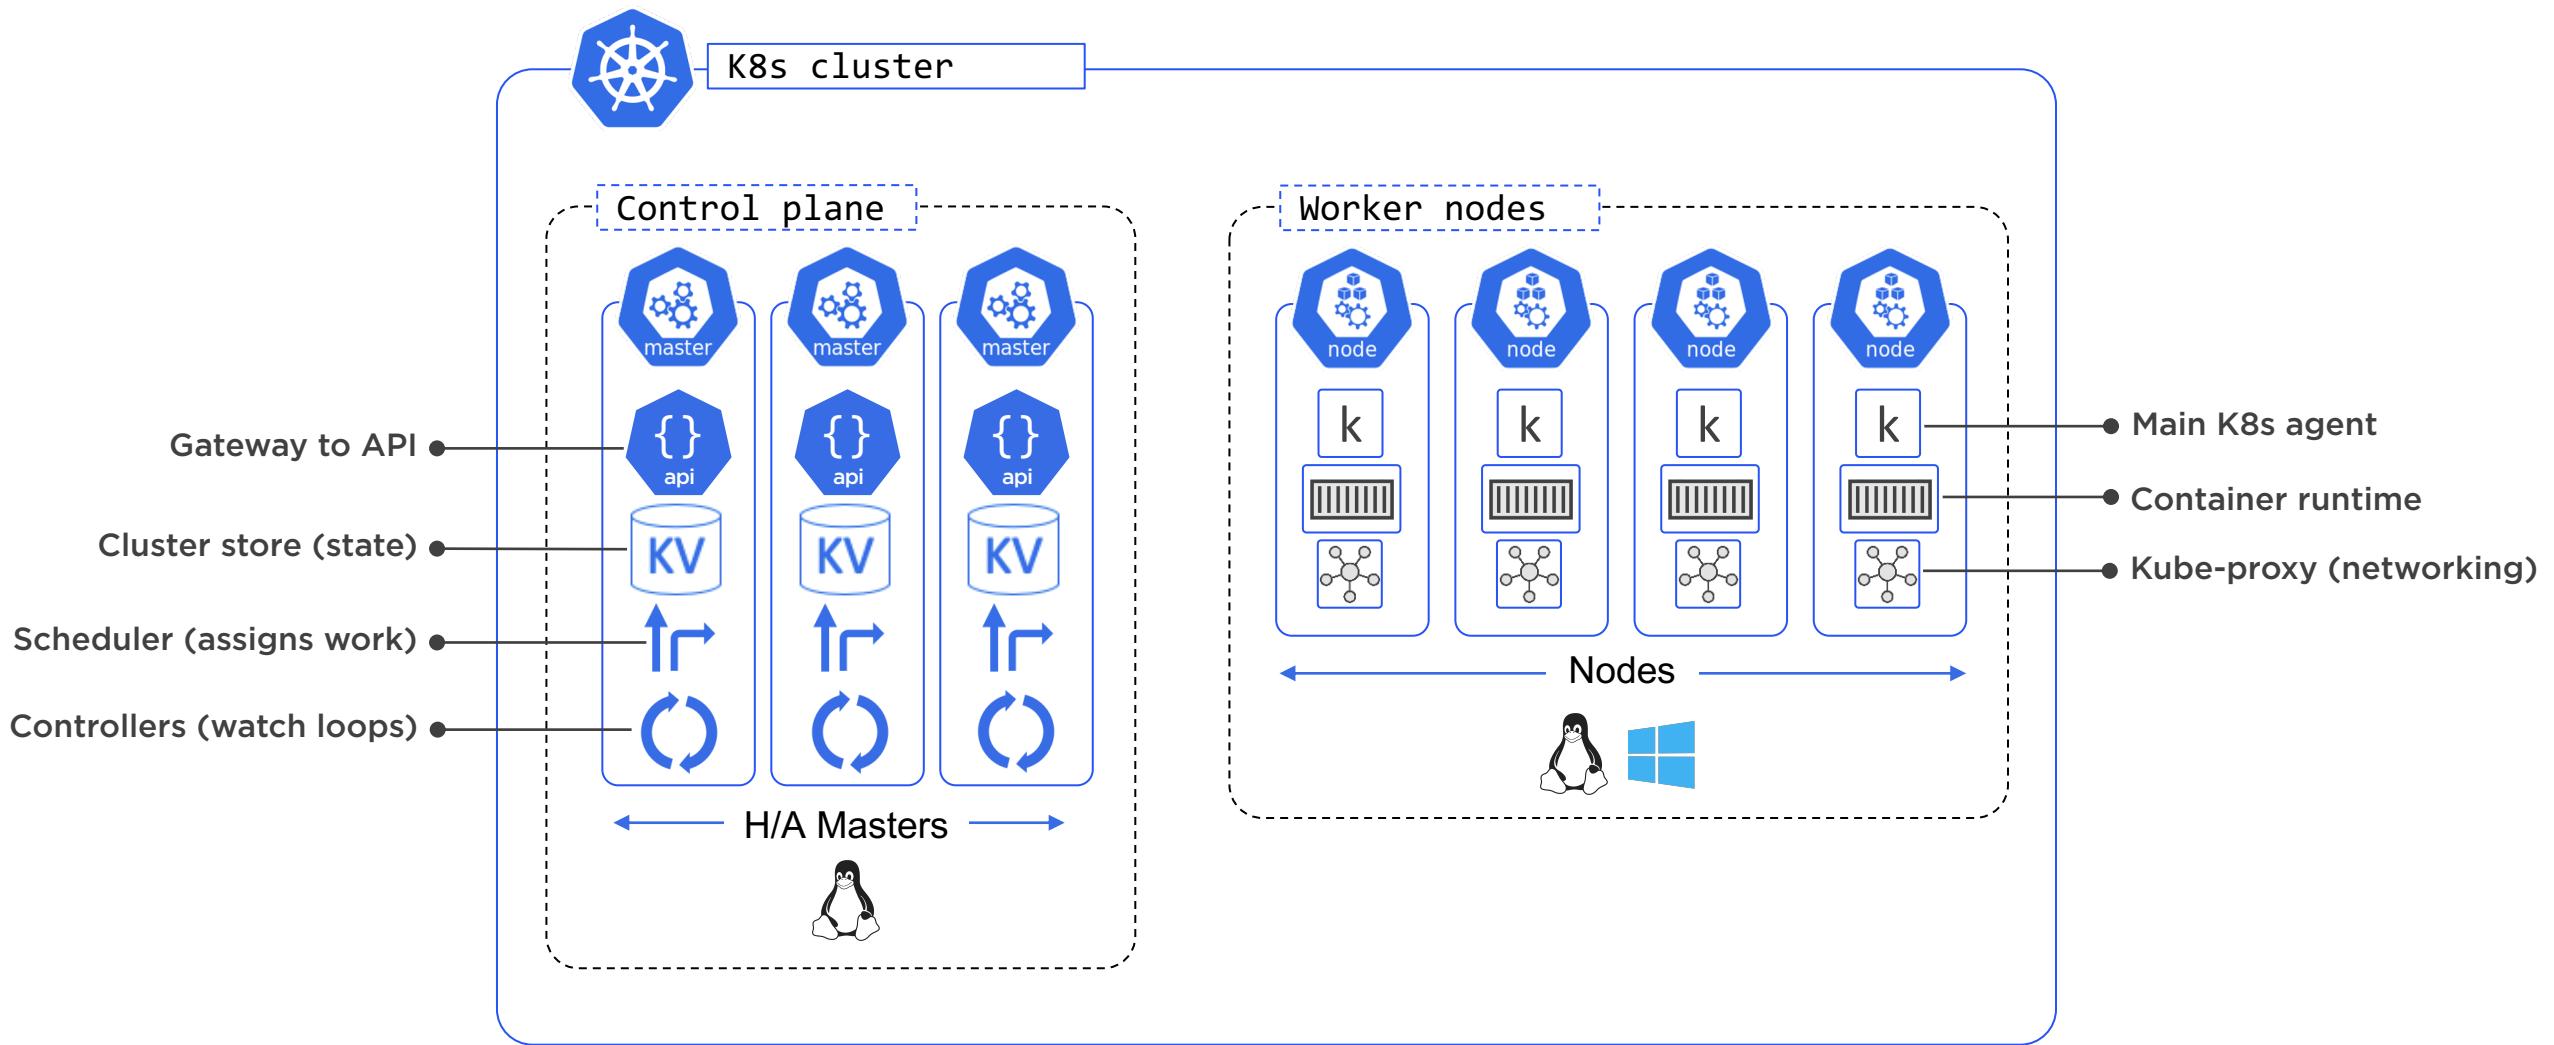


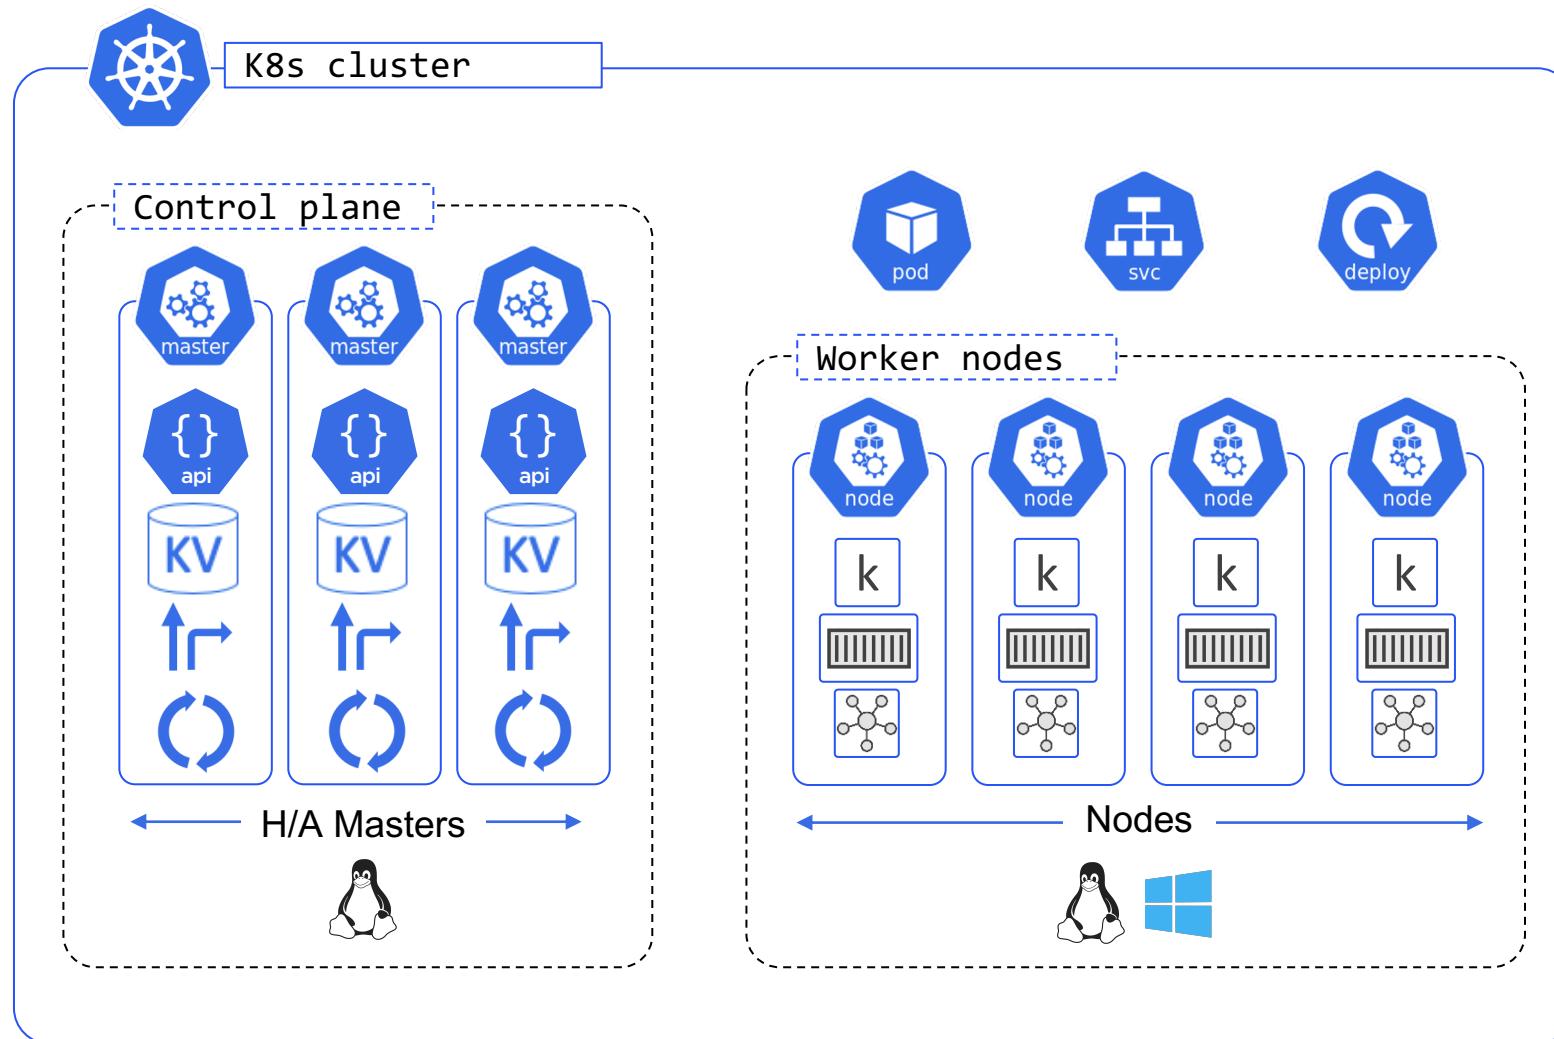












Up Next:
Getting Kubernetes

